

Creating DataFrames

We will create a `DataFrame` object from a dataset in a CSV file using the **`read_csv`** method on the pandas - **`pd`** - object. While the `read_csv` method can read dataset directly from a web URL, it's best to download to your computer first to ensure the data is persisted.

About the Data

We will use the dataset on nobel laureates available at:

<http://api.nobelprize.org/v1/laureates.csv>. Download the file to your computer using: **`curl -OL http://api.nobelprize.org/v1/laureates.csv`**

Imports

```
In [67]: import datetime
import numpy as np
import pandas as pd
```

Creating a Series

```
In [68]: np.random.seed(0) # set a seed for reproducibility
pd.Series(np.random.rand(5), name='random')
```

```
Out[68]: 0    0.548814
1    0.715189
2    0.602763
3    0.544883
4    0.423655
Name: random, dtype: float64
```

Creating a DataFrame from a Series

Use the `to_frame()` method:

```
In [69]: pd.Series(np.linspace(0, 10, num=5)).to_frame()
```

```
Out[69]:
```

	0
0	0.0
1	2.5
2	5.0
3	7.5
4	10.0

From a list of dictionaries

```
In [70]: pd.DataFrame([
    {'mag' : 5.2, 'place' : 'California'},
    {'mag' : 1.2, 'place' : 'Alaska'},
    {'mag' : 0.2, 'place' : 'California'},
])
```

```
Out[70]:
```

	mag	place
0	5.2	California
1	1.2	Alaska
2	0.2	California

From a NumPy array

```
In [71]: pd.DataFrame(
    np.array([
        [0, 0, 0],
        [1, 1, 1],
        [2, 4, 8],
        [3, 9, 27],
        [4, 16, 64]
    ]), columns=['n', 'n_squared', 'n_cubed']
)
```

```
Out[71]:
```

	n	n_squared	n_cubed
0	0	0	0
1	1	1	1
2	2	4	8
3	3	9	27
4	4	16	64

Creating a DataFrame by Reading in a CSV File

Finding information on the file before reading it in

Before attempting to read in a file, we can use the command line to see important information about the file that may determine how we read it in. We can run command line code from Jupyter Notebooks (thanks to IPython) by using `!` before the code.

Number of lines (row count)

For example, we can find out how many lines are in the file by using the `wc` utility (word count) and counting lines in the file (`-l`). Run the cell below to confirm the file has 1002 lines:

In [72]:

```
!wc -l laureates.csv # this will not work on Windows commandline
# On windows the !dir command will show directory contents
```

```
1002 laureates.csv
```

We can even capture the result of a command and use it in our Python code:

Reading in the file

Our file is small in size, has headers in the first row, and is comma-separated, so we don't need to provide any additional arguments to read in the file with `pd.read_csv()`, but be sure to check the [documentation](#) for possible arguments. To read data from file we can use

`pd.read_csv()` and for other delimited files, such as tab (`\t`), we can use the `read_csv()` function with the `sep` argument equal to the delimiter. We can use the `read_excel()` function for Excel files, the `read_json()` function for JSON (JavaScript Object Notation) files

In [73]:

```
import pandas as pd
df = pd.read_csv('laureates.csv')
```

Getting Documentation on Python elements

You can utilize the built-in `help()` function for documentation on Python elements. Simply run `help()`, passing in the package, module, class, object, method, or function. Assuming we aliased pandas as `pd` when we imported it, we can run `help(pd)` to see information on the pandas package; `help(pd.DataFrame)` for all the methods and attributes of a dataframe (note we can also pass in an already created DataFrame object instead); and `help(pd.read_csv)` to learn more about the pandas function for reading CSV files into Python and how to use it

Run the code cell below to see documentation for the DataFrame class.

In [74]:

```
import pandas as pd
help(pd.DataFrame)
```

Help on class DataFrame in module pandas.core.frame:

```
class DataFrame(pandas.core.generic.NDFrame, pandas.core.arraylike.OpsMixin)
| DataFrame(data=None, index: 'Optional[Axes]' = None, columns: 'Optional[Axes]' = None, dtype: 'Optional[Dtype]' = None, copy: 'bool' = False)
|
| Two-dimensional, size-mutable, potentially heterogeneous tabular data.
|
| Data structure also contains labeled axes (rows and columns).
| Arithmetic operations align on both row and column labels. Can be
| thought of as a dict-like container for Series objects. The primary
| pandas data structure.
|
| Parameters
| -----
| data : ndarray (structured or homogeneous), Iterable, dict, or DataFrame
|       Dict can contain Series, arrays, constants, dataclass or list-like objects
```

```

ts. If
    data is a dict, column order follows insertion-order.
    .. versionchanged:: 0.25.0
        If data is a list of dicts, column order follows insertion-order.

index : Index or array-like
    Index to use for resulting frame. Will default to RangeIndex if
    no indexing information part of input data and no index provided.
columns : Index or array-like
    Column labels to use for resulting frame. Will default to
    RangeIndex (0, 1, 2, ..., n) if no column labels are provided.
dtype : dtype, default None
    Data type to force. Only a single dtype is allowed. If None, infer.
copy : bool, default False
    Copy data from inputs. Only affects DataFrame / 2d ndarray input.

See Also
-----
DataFrame.from_records : Constructor from tuples, also record arrays.
DataFrame.from_dict : From dicts of Series, arrays, or dicts.
read_csv : Read a comma-separated values (csv) file into DataFrame.
read_table : Read general delimited file into DataFrame.
read_clipboard : Read text from clipboard into DataFrame.

Examples
-----
Constructing DataFrame from a dictionary.

>>> d = {'col1': [1, 2], 'col2': [3, 4]}
>>> df = pd.DataFrame(data=d)
>>> df
   col1  col2
0     1     3
1     2     4

Notice that the inferred dtype is int64.

>>> df.dtypes
col1    int64
col2    int64
dtype: object

To enforce a single dtype:

>>> df = pd.DataFrame(data=d, dtype=np.int8)
>>> df.dtypes
col1    int8
col2    int8
dtype: object

Constructing DataFrame from numpy ndarray:

>>> df2 = pd.DataFrame(np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]]),
...                     columns=['a', 'b', 'c'])
>>> df2
   a  b  c
0  1  2  3
1  4  5  6
2  7  8  9

Constructing DataFrame from dataclass:

>>> from dataclasses import make_dataclass
>>> Point = make_dataclass("Point", [("x", int), ("y", int)])

```

```

>>> pd.DataFrame([Point(0, 0), Point(0, 3), Point(2, 3)])
   x  y
0  0  0
1  0  3
2  2  3

Method resolution order:
  DataFrame
  pandas.core.generic.NDFrame
  pandas.core.base.PandasObject
  pandas.core.accessor.DirNamesMixin
  pandas.core.base.SelectionMixin
  pandas.core.indexing.IndexingMixin
  pandas.core.arraylike.OpsMixin
  builtins.object

Methods defined here:

  __divmod__(self, other) -> 'Tuple[DataFrame, DataFrame]'

  __getitem__(self, key)

  __init__(self, data=None, index: 'Optional[Axes]' = None, columns: 'Optional[Axes]' = None, dtype: 'Optional[Dtype]' = None, copy: 'bool' = False)
    Initialize self.  See help(type(self)) for accurate signature.

  __len__(self) -> 'int'
    Returns length of info axis, but here we use the index.

  __matmul__(self, other)
    Matrix multiplication using binary '@' operator in Python>=3.5.

  __rdivmod__(self, other) -> 'Tuple[DataFrame, DataFrame]'

  __repr__(self) -> 'str'
    Return a string representation for a particular DataFrame.

  __rmatmul__(self, other)
    Matrix multiplication using binary '@' operator in Python>=3.5.

  __setitem__(self, key, value)

  add(self, other, axis='columns', level=None, fill_value=None)
    Get Addition of dataframe and other, element-wise (binary operator `add`
    ).

    Equivalent to ``dataframe + other``, but with support to substitute a fi
    ll_value
    for missing data in one of the inputs. With reverse version, `radd`.

    Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
    arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

    Parameters
    -----
    other : scalar, sequence, Series, or DataFrame
        Any single or multiple element data structure, or list-like object.
    axis : {0 or 'index', 1 or 'columns'}
        Whether to compare by the index (0 or 'index') or columns
        (1 or 'columns'). For Series input, axis to match Series index on.
    level : int or label
        Broadcast across a level, matching Index values on the
        passed MultiIndex level.
    fill_value : float or None, default None
        Fill existing missing (NaN) values, and any new element needed for

```

successful DataFrame alignment, with this value before computation.
If data in both corresponding DataFrame locations is missing
the result will be missing.

Returns

DataFrame

Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df

	angles	degrees
circle	0	360
triangle	3	180
rectangle	4	360

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```

>>> df - [1, 2]
      angles  degrees
circle      -1     358
triangle     2     178
rectangle    3     358

>>> df.sub([1, 2], axis='columns')
      angles  degrees
circle      -1     358
triangle     2     178
rectangle    3     358

>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
...
      axis='index')
      angles  degrees
circle      -1     359
triangle     2     179
rectangle    3     359

Multiply a DataFrame of different shape with operator version.

>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...
      index=['circle', 'triangle', 'rectangle'])
>>> other
      angles
circle      0
triangle     3
rectangle    4

>>> df * other
      angles  degrees
circle      0      NaN
triangle     9      NaN
rectangle   16      NaN

>>> df.mul(other, fill_value=0)
      angles  degrees
circle      0      0.0
triangle     9      0.0
rectangle   16      0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...
      'degrees': [360, 180, 360, 360, 540, 720],
...
      index=[['A', 'A', 'A', 'B', 'B', 'B'],
...
      ['circle', 'triangle', 'rectangle', 'square', 'pentagon', 'hexagon']])
>>> df_multindex
      angles  degrees
A circle      0     360
  triangle     3     180
  rectangle    4     360
B square      4     360
  pentagon     5     540
  hexagon      6     720

>>> df.div(df_multindex, level=1, fill_value=0)
      angles  degrees
A circle      NaN     1.0

```

triangle	1.0	1.0
rectangle	1.0	1.0
B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

```
agg = aggregate(self, func=None, axis=0, *args, **kwargs)
```

```
aggregate(self, func=None, axis=0, *args, **kwargs)
```

Aggregate using one or more operations over the specified axis.

Parameters

func : function, str, list or dict

Function to use for aggregating the data. If a function, must either work when passed a DataFrame or when passed to DataFrame.apply.

Accepted combinations are:

- function
- string function name
- list of functions and/or function names, e.g. ``[np.sum, 'mean']``
- dict of axis labels -> functions, function names or list of such.

axis : {0 or 'index', 1 or 'columns'}, default 0

If 0 or 'index': apply function to each column.

If 1 or 'columns': apply function to each row.

***args**

Positional arguments to pass to `func`.

****kwargs**

Keyword arguments to pass to `func`.

Returns

scalar, Series or DataFrame

The return can be:

- * scalar : when Series.agg is called with single function
- * Series : when DataFrame.agg is called with a single function
- * DataFrame : when DataFrame.agg is called with several functions

Return scalar, Series or DataFrame.

The aggregation operations are always performed over an axis, either the index (default) or the column axis. This behavior is different from `numpy` aggregation functions (`mean`, `median`, `prod`, `sum`, `std`, `var`), where the default is to compute the aggregation of the flattened array, e.g., ``numpy.mean(arr_2d)`` as opposed to ``numpy.mean(arr_2d, axis=0)``.

`agg` is an alias for `aggregate`. Use the alias.

See Also

DataFrame.apply : Perform any type of operations.

DataFrame.transform : Perform transformation type operations.

core.groupby.GroupBy : Perform operations over groups.

core.resample.Resampler : Perform operations over resampled bins.

core.window.Rolling : Perform operations over rolling window.

core.window.Expanding : Perform operations over expanding window.

core.window.ExponentialMovingWindow : Perform operation over exponential

weighted

window.

Notes

`agg` is an alias for `aggregate`. Use the alias.

A passed user-defined-function will be passed a Series for evaluation.

Examples

```
>>> df = pd.DataFrame([[1, 2, 3],
...                     [4, 5, 6],
...                     [7, 8, 9],
...                     [np.nan, np.nan, np.nan]],
...                     columns=['A', 'B', 'C'])
```

Aggregate these functions over the rows.

```
>>> df.agg(['sum', 'min'])
      A      B      C
sum 12.0  15.0  18.0
min  1.0   2.0   3.0
```

Different aggregations per column.

```
>>> df.agg({'A' : ['sum', 'min'], 'B' : ['min', 'max']})
      A      B
sum 12.0  NaN
min  1.0  2.0
max  NaN  8.0
```

Aggregate different functions over the columns and rename the index of the resulting DataFrame.

```
>>> df.agg(x=('A', max), y=('B', 'min'), z=('C', np.mean))
      A      B      C
x  7.0  NaN  NaN
y  NaN  2.0  NaN
z  NaN  NaN  6.0
```

Aggregate over the columns.

```
>>> df.agg("mean", axis="columns")
0      2.0
1      5.0
2      8.0
3      NaN
dtype: float64
```

align(self, other, join='outer', axis=None, level=None, copy=True, fill_value=None, method=None, limit=None, fill_axis=0, broadcast_axis=None) -> 'DataFrame'

Align two objects on their axes with the specified join method.

Join method is specified for each axis Index.

Parameters

other : DataFrame or Series
join : {'outer', 'inner', 'left', 'right'}, default 'outer'
axis : allowed axis of the other object, default None
Align on index (0), columns (1), or both (None).
level : int or level name, default None
Broadcast across a level, matching Index values on the passed MultiIndex level.
copy : bool, default True
Always returns new objects. If copy=False and no reindexing is

```

        required then original objects are returned.
    fill_value : scalar, default np.NaN
        Value to use for missing values. Defaults to NaN, but can be any
        "compatible" value.
    method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None
        Method to use for filling holes in reindexed Series:

        - pad / ffill: propagate last valid observation forward to next valid
        - backfill / bfill: use NEXT valid observation to fill gap.

    limit : int, default None
        If method is specified, this is the maximum number of consecutive
        NaN values to forward/backward fill. In other words, if there is
        a gap with more than this number of consecutive NaNs, it will only
        be partially filled. If method is not specified, this is the
        maximum number of entries along the entire axis where NaNs will be
        filled. Must be greater than 0 if not None.
    fill_axis : {0 or 'index', 1 or 'columns'}, default 0
        Filling axis, method and limit.
    broadcast_axis : {0 or 'index', 1 or 'columns'}, default None
        Broadcast values along this axis, if aligning two objects of
        different dimensions.

    Returns
    -----
    (left, right) : (DataFrame, type of other)
        Aligned objects.

    all(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
        Return whether all elements are True, potentially over an axis.

    Returns True unless there at least one element within a series or
    along a Dataframe axis that is False or equivalent (e.g. zero or
    empty).

    Parameters
    -----
    axis : {0 or 'index', 1 or 'columns', None}, default 0
        Indicate which axis or axes should be reduced.

        * 0 / 'index' : reduce the index, return a Series whose index is the
          original column labels.
        * 1 / 'columns' : reduce the columns, return a Series whose index is
          the original index.
        * None : reduce all axes, return a scalar.

    bool_only : bool, default None
        Include only boolean columns. If None, will attempt to use everything,
        then use only boolean data. Not implemented for Series.
    skipna : bool, default True
        Exclude NA/null values. If the entire row/column is NA and skipna is
        True, then the result will be True, as for an empty row/column.
        If skipna is False, then NA are treated as True, because these are not
        equal to zero.
    level : int or level name, default None
        If the axis is a MultiIndex (hierarchical), count along a
        particular level, collapsing into a Series.
    **kwargs : any, default None
        Additional keywords have no effect but might be accepted for
        compatibility with NumPy.

```

Returns

Series or DataFrame

If level is specified, then, DataFrame is returned; otherwise, Series is returned.

See Also

Series.all : Return True if all elements are True.

DataFrame.any : Return True if one (or more) elements are True.

Examples

****Series****

```
>>> pd.Series([True, True]).all()
True
>>> pd.Series([True, False]).all()
False
>>> pd.Series([]).all()
True
>>> pd.Series([np.nan]).all()
True
>>> pd.Series([np.nan]).all(skipna=False)
True
```

****DataFrames****

Create a dataframe from a dictionary.

```
>>> df = pd.DataFrame({'col1': [True, True], 'col2': [True, False]})
>>> df
   col1  col2
0  True  True
1  True False
```

Default behaviour checks if column-wise values all return True.

```
>>> df.all()
col1    True
col2   False
dtype: bool
```

Specify ``axis='columns'`` to check if row-wise values all return True.

```
>>> df.all(axis='columns')
0    True
1   False
dtype: bool
```

Or ``axis=None`` for whether every value is True.

```
>>> df.all(axis=None)
False
```

any(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
Return whether any element is True, potentially over an axis.

Returns False unless there is at least one element within a series or along a Dataframe axis that is True or equivalent (e.g. non-zero or non-empty).

Parameters

```

axis : {0 or 'index', 1 or 'columns', None}, default 0
    Indicate which axis or axes should be reduced.

    * 0 / 'index' : reduce the index, return a Series whose index is the
      original column labels.
    * 1 / 'columns' : reduce the columns, return a Series whose index is
      original index.
    * None : reduce all axes, return a scalar.

bool_only : bool, default None
    Include only boolean columns. If None, will attempt to use everything,
    then use only boolean data. Not implemented for Series.

skipna : bool, default True
    Exclude NA/null values. If the entire row/column is NA and skipna is
    True, then the result will be False, as for an empty row/column.
    If skipna is False, then NA are treated as True, because these are not
    equal to zero.

level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.

**kwargs : any, default None
    Additional keywords have no effect but might be accepted for
    compatibility with NumPy.

```

Returns

```

-----
Series or DataFrame
    If level is specified, then, DataFrame is returned; otherwise, Series
    is returned.

```

See Also

```

-----
numpy.any : Numpy version of this method.
Series.any : Return whether any element is True.
Series.all : Return whether all elements are True.
DataFrame.any : Return whether any element is True over requested axis.
DataFrame.all : Return whether all elements are True over requested axis.

```

Examples

Series

For Series input, the output is a scalar indicating whether any element is True.

```

>>> pd.Series([False, False]).any()
False
>>> pd.Series([True, False]).any()
True
>>> pd.Series([]).any()
False
>>> pd.Series([np.nan]).any()
False
>>> pd.Series([np.nan]).any(skipna=False)
True

```

DataFrame

Whether each column contains at least one True element (the default).

```
>>> df = pd.DataFrame({"A": [1, 2], "B": [0, 2], "C": [0, 0]})
>>> df
   A  B  C
0  1  0  0
1  2  2  0

>>> df.any()
A      True
B      True
C     False
dtype: bool
```

Aggregating over the columns.

```
>>> df = pd.DataFrame({"A": [True, False], "B": [1, 2]})
>>> df
   A  B
0  True  1
1 False  2

>>> df.any(axis='columns')
0      True
1      True
dtype: bool
```

```
>>> df = pd.DataFrame({"A": [True, False], "B": [1, 0]})
>>> df
   A  B
0  True  1
1 False  0

>>> df.any(axis='columns')
0      True
1     False
dtype: bool
```

Aggregating over the entire DataFrame with ``axis=None``.

```
>>> df.any(axis=None)
True
```

``any`` for an empty DataFrame is an empty Series.

```
>>> pd.DataFrame([]).any()
Series([], dtype: bool)
```

-> `append(self, other, ignore_index=False, verify_integrity=False, sort=False)`
 'DataFrame'

Append rows of ``other`` to the end of caller, returning a new object.

Columns in ``other`` that are not in the caller are added as new columns.

Parameters

`other` : DataFrame or Series/dict-like object, or list of these
 The data to append.

`ignore_index` : bool, default False

If True, the resulting axis will be labeled 0, 1, ..., n - 1.

`verify_integrity` : bool, default False

If True, raise ValueError on creating index with duplicates.

`sort` : bool, default False

Sort columns if the columns of ``self`` and ``other`` are not aligned.

.. versionchanged:: 1.0.0

Changed to not sort by default.

Returns

DataFrame

See Also

`concat` : General function to concatenate DataFrame or Series objects.

Notes

If a list of dict/series is passed and the keys are all contained in the DataFrame's index, the order of the columns in the resulting DataFrame will be unchanged.

Iteratively appending rows to a DataFrame can be more computationally intensive than a single concatenate. A better solution is to append those rows to a list and then concatenate the list with the original DataFrame all at once.

Examples

```
>>> df = pd.DataFrame([[1, 2], [3, 4]], columns=list('AB'))
>>> df
   A  B
0  1  2
1  3  4
>>> df2 = pd.DataFrame([[5, 6], [7, 8]], columns=list('AB'))
>>> df.append(df2)
   A  B
0  1  2
1  3  4
0  5  6
1  7  8
```

With `'ignore_index'` set to True:

```
>>> df.append(df2, ignore_index=True)
   A  B
0  1  2
1  3  4
2  5  6
3  7  8
```

The following, while not recommended methods for generating DataFrames, show two ways to generate a DataFrame from multiple data sources.

Less efficient:

```
>>> df = pd.DataFrame(columns=['A'])
>>> for i in range(5):
...     df = df.append({'A': i}, ignore_index=True)
>>> df
   A
0  0
1  1
2  2
3  3
4  4
```

More efficient:

```
>>> pd.concat([pd.DataFrame([i], columns=['A']) for i in range(5)],
...           ignore_index=True)
```

```

      A
0    0
1    1
2    2
3    3
4    4

```

`apply(self, func, axis=0, raw=False, result_type=None, args=(), **kwargs)`
 Apply a function along an axis of the DataFrame.

Objects passed to the function are Series objects whose index is either the DataFrame's index (`axis=0`) or the DataFrame's columns (`axis=1`). By default (`result_type=None`), the final return type is inferred from the return type of the applied function. Otherwise, it depends on the `result_type` argument.

Parameters

`func` : function

Function to apply to each column or row.

`axis` : {0 or 'index', 1 or 'columns'}, default 0

Axis along which the function is applied:

- * 0 or 'index': apply function to each column.
- * 1 or 'columns': apply function to each row.

`raw` : bool, default False

Determines if row or column is passed as a Series or ndarray object:

- * `False` : passes each row or column as a Series to the function.
- * `True` : the passed function will receive ndarray objects instead.
If you are just applying a NumPy reduction function this will achieve much better performance.

`result_type` : {'expand', 'reduce', 'broadcast', None}, default None
 These only act when `axis=1` (columns):

- * 'expand' : list-like results will be turned into columns.
- * 'reduce' : returns a Series if possible rather than expanding list-like results. This is the opposite of 'expand'.
- * 'broadcast' : results will be broadcast to the original shape of the DataFrame, the original index and columns will be retained.

The default behaviour (None) depends on the return value of the applied function: list-like results will be returned as a Series of those. However if the apply function returns a Series these are expanded to columns.

`args` : tuple

Positional arguments to pass to `func` in addition to the array/series.

`**kwargs`

Additional keyword arguments to pass as keywords arguments to `func`.

Returns

Series or DataFrame

Result of applying `func` along the given axis of the DataFrame.

See Also

DataFrame.applymap: For elementwise operations.
 DataFrame.aggreate: Only perform aggregating type operations.
 DataFrame.transform: Only perform transforming type operations.

Examples

```
>>> df = pd.DataFrame([[4, 9]] * 3, columns=['A', 'B'])
>>> df
   A  B
0  4  9
1  4  9
2  4  9
```

Using a numpy universal function (in this case the same as ``np.sqrt(df)``):

```
>>> df.apply(np.sqrt)
   A  B
0  2.0  3.0
1  2.0  3.0
2  2.0  3.0
```

Using a reducing function on either axis

```
>>> df.apply(np.sum, axis=0)
A    12
B    27
dtype: int64
```

```
>>> df.apply(np.sum, axis=1)
0    13
1    13
2    13
dtype: int64
```

Returning a list-like will result in a Series

```
>>> df.apply(lambda x: [1, 2], axis=1)
0    [1, 2]
1    [1, 2]
2    [1, 2]
dtype: object
```

Passing ``result_type='expand'`` will expand list-like results to columns of a Dataframe

```
>>> df.apply(lambda x: [1, 2], axis=1, result_type='expand')
   0  1
0  1  2
1  1  2
2  1  2
```

Returning a Series inside the function is similar to passing ``result_type='expand'``. The resulting column names will be the Series index.

```
>>> df.apply(lambda x: pd.Series([1, 2], index=['foo', 'bar']), axis=1)
   foo  bar
0    1    2
1    1    2
2    1    2
```

Passing ``result_type='broadcast'`` will ensure the same shape result, whether list-like or scalar is returned by the function, and broadcast it along the axis. The resulting column names will

be the originals.

```
>>> df.apply(lambda x: [1, 2], axis=1, result_type='broadcast')
```

```
   A  B
0  1  2
1  1  2
2  1  2
```

`applymap(self, func, na_action: 'Optional[str]' = None) -> 'DataFrame'`
 Apply a function to a Dataframe elementwise.

This method applies a function that accepts and returns a scalar to every element of a DataFrame.

Parameters

`func` : callable

Python function, returns a single value from a single value.

`na_action` : {None, 'ignore'}, default None

If 'ignore', propagate NaN values, without passing them to func.

.. versionadded:: 1.2

Returns

DataFrame

Transformed DataFrame.

See Also

`DataFrame.apply` : Apply a function along input axis of DataFrame.

Examples

```
>>> df = pd.DataFrame([[1, 2.12], [3.356, 4.567]])
```

```
>>> df
   0      1
0  1.000  2.120
1  3.356  4.567
```

```
>>> df.applymap(lambda x: len(str(x)))
```

```
   0  1
0  3  4
1  5  5
```

Like `Series.map`, NA values can be ignored:

```
>>> df_copy = df.copy()
```

```
>>> df_copy.iloc[0, 0] = pd.NA
```

```
>>> df_copy.applymap(lambda x: len(str(x)), na_action='ignore')
```

```
   0  1
0  <NA>  4
1    5  5
```

Note that a vectorized version of ``func`` often exists, which will be much faster. You could square each number elementwise.

```
>>> df.applymap(lambda x: x**2)
```

```
   0      1
0  1.000000  4.494400
1 11.262736 20.857489
```

But it's better to avoid `applymap` in that case.

```
>>> df ** 2
```

	0	1
0	1.000000	4.494400
1	11.262736	20.857489

```
assign(self, **kwargs) -> 'DataFrame'
Assign new columns to a DataFrame.
```

Returns a new object with all original columns in addition to new ones. Existing columns that are re-assigned will be overwritten.

Parameters

****kwargs** : dict of {str: callable or Series}
 The column names are keywords. If the values are callable, they are computed on the DataFrame and assigned to the new columns. The callable must not change input DataFrame (though pandas doesn't check it). If the values are not callable, (e.g. a Series, scalar, or array), they are simply assigned.

Returns

DataFrame

A new DataFrame with the new columns in addition to all the existing columns.

Notes

Assigning multiple columns within the same ``assign`` is possible. Later items in ``***kwargs`` may refer to newly created or modified columns in 'df'; items are computed and assigned into 'df' in order.

Examples

```
>>> df = pd.DataFrame({'temp_c': [17.0, 25.0]},
...                    index=['Portland', 'Berkeley'])
>>> df
```

	temp_c
Portland	17.0
Berkeley	25.0

Where the value is a callable, evaluated on `df`:

```
>>> df.assign(temp_f=lambda x: x.temp_c * 9 / 5 + 32)
```

	temp_c	temp_f
Portland	17.0	62.6
Berkeley	25.0	77.0

Alternatively, the same behavior can be achieved by directly referencing an existing Series or sequence:

```
>>> df.assign(temp_f=df['temp_c'] * 9 / 5 + 32)
```

	temp_c	temp_f
Portland	17.0	62.6
Berkeley	25.0	77.0

You can create multiple columns within the same assign where one of the columns depends on another one defined within the same assign:

```
>>> df.assign(temp_f=lambda x: x['temp_c'] * 9 / 5 + 32,
...          temp_k=lambda x: (x['temp_f'] + 459.67) * 5 / 9)
```

	temp_c	temp_f	temp_k
Portland	17.0	62.6	290.15
Berkeley	25.0	77.0	298.15

```
| boxplot = boxplot_frame(self, column=None, by=None, ax=None, fontsize=None,
| rot=0, grid=True, figsize=None, layout=None, return_type=None, backend=None, **k
| wargs)
```

Make a box plot from DataFrame columns.

Make a box-and-whisker plot from DataFrame columns, optionally grouped by some other columns. A box plot is a method for graphically depicting groups of numerical data through their quartiles.

The box extends from the Q1 to Q3 quartile values of the data, with a line at the median (Q2). The whiskers extend from the edges of box to show the range of the data. By default, they extend no more th

an

thest

`1.5 * IQR (IQR = Q3 - Q1)` from the edges of the box, ending at the far data point within that interval. Outliers are plotted as separate dots.

For further details see

Wikipedia's entry for `boxplot` <https://en.wikipedia.org/wiki/Box_plot>

Parameters

column : str or list of str, optional

Column name or list of names, or vector.

Can be any valid input to :meth:`pandas.DataFrame.groupby`.

by : str or array-like, optional

Column in the DataFrame to :meth:`pandas.DataFrame.groupby`.

One box-plot will be done per value of columns in `by`.

ax : object of class matplotlib.axes.Axes, optional

The matplotlib axes to be used by boxplot.

fontsize : float or str

Tick label font size in points or as a string (e.g., `large`).

rot : int or float, default 0

The rotation angle of labels (in degrees)

with respect to the screen coordinate system.

grid : bool, default True

Setting this to True will show the grid.

figsize : A tuple (width, height) in inches

The size of the figure to create in matplotlib.

layout : tuple (rows, columns), optional

For example, (3, 5) will display the subplots

using 3 columns and 5 rows, starting from the top-left.

return_type : {'axes', 'dict', 'both'} or None, default 'axes'

The kind of object to return. The default is ``axes``.

* 'axes' returns the matplotlib axes the boxplot is drawn on.

* 'dict' returns a dictionary whose values are the matplotlib Lines of the boxplot.

* 'both' returns a namedtuple with the axes and dict.

* when grouping with ``by``, a Series mapping columns to ``return_type`` is returned.

If ``return_type`` is `None`, a NumPy array

of axes with the same shape as ``layout`` is returned.

backend : str, default None

Backend to use instead of the backend specified in the option

``plotting.backend``. For instance, 'matplotlib'. Alternatively, to specify the ``plotting.backend`` for the whole session, set

``pd.options.plotting.backend``.

.. versionadded:: 1.0.0

**kwargs

All other plotting keyword arguments to be passed to

:func:`matplotlib.pyplot.boxplot`.

Returns

result
See Notes.

See Also

Series.plot.hist: Make a histogram.
matplotlib.pyplot.boxplot : Matplotlib equivalent plot.

Notes

The return type depends on the `return_type` parameter:

- * 'axes' : object of class matplotlib.axes.Axes
- * 'dict' : dict of matplotlib.lines.Line2D objects
- * 'both' : a namedtuple with structure (ax, lines)

For data grouped with `by`, return a Series of the above or a numpy array:

- * :class:`pandas.Series`
- * :class:`numpy.array` (for `return_type = None`)

Use `return_type='dict'` when you want to tweak the appearance of the lines after plotting. In this case a dict containing the Lines making up the boxes, caps, fliers, medians, and whiskers is returned.

Examples

Boxplots can be created for every column in the dataframe by `df.boxplot()` or indicating the columns to be used:

```
.. plot::
    :context: close-figs

    >>> np.random.seed(1234)
    >>> df = pd.DataFrame(np.random.randn(10, 4),
    ...                   columns=['Col1', 'Col2', 'Col3', 'Col4'])
    >>> boxplot = df.boxplot(column=['Col1', 'Col2', 'Col3'])
```

Boxplots of variables distributions grouped by the values of a third variable can be created using the option `by`. For instance:

```
.. plot::
    :context: close-figs

    >>> df = pd.DataFrame(np.random.randn(10, 2),
    ...                   columns=['Col1', 'Col2'])
    >>> df['X'] = pd.Series(['A', 'A', 'A', 'A', 'A',
    ...                    'B', 'B', 'B', 'B', 'B'])
    >>> boxplot = df.boxplot(by='X')
```

A list of strings (i.e. `['X', 'Y']`) can be passed to `boxplot` in order to group the data by combination of the variables in the x-axis

```
.. plot::
    :context: close-figs

    >>> df = pd.DataFrame(np.random.randn(10, 3),
    ...                   columns=['Col1', 'Col2', 'Col3'])
    >>> df['X'] = pd.Series(['A', 'A', 'A', 'A', 'A',
    ...                    'A', 'A', 'A', 'A', 'A'])
```

s:

```

...         'B', 'B', 'B', 'B', 'B'])
>>> df['Y'] = pd.Series(['A', 'B', 'A', 'B', 'A',
...                     'B', 'A', 'B', 'A', 'B'])
>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by=['X', 'Y'])

```

The layout of boxplot can be adjusted giving a tuple to ``layout``:

```

.. plot::
   :context: close-figs

>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                       layout=(2, 1))

```

Additional formatting can be done to the boxplot, like suppressing the grid

(``grid=False``), rotating the labels in the x-axis (i.e. ``rot=45``) or changing the fontsize (i.e. ``fontsize=15``):

```

.. plot::
   :context: close-figs

>>> boxplot = df.boxplot(grid=False, rot=45, fontsize=15)

```

The parameter ``return_type`` can be used to select the type of element returned by ``boxplot``. When ``return_type='axes'`` is selected, the matplotlib axes on which the boxplot is drawn are returned:

```

>>> boxplot = df.boxplot(column=['Col1', 'Col2'], return_type='axes')
>>> type(boxplot)
<class 'matplotlib.axes._subplots.AxesSubplot'>

```

When grouping with ``by``, a Series mapping columns to ``return_type`` is returned:

```

>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                       return_type='axes')
>>> type(boxplot)
<class 'pandas.core.series.Series'>

```

If ``return_type`` is ``None``, a NumPy array of axes with the same shape as ``layout`` is returned:

```

>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                       return_type=None)
>>> type(boxplot)
<class 'numpy.ndarray'>

```

combine(self, other: 'DataFrame', func, fill_value=None, overwrite=True) -> 'DataFrame'

Perform column-wise combine with another DataFrame.

Combines a DataFrame with ``other`` DataFrame using ``func`` to element-wise combine columns. The row and column indexes of the resulting DataFrame will be the union of the two.

Parameters

other : DataFrame

The DataFrame to merge column-wise.

func : function

Function that takes two series as inputs and return a Series or a scalar. Used to merge the two dataframes column by columns.

fill_value : scalar value, default None

The value to fill NaNs with prior to passing any column to the

merge func.
 overwrite : bool, default True
 If True, columns in `self` that do not exist in `other` will be overwritten with NaNs.

Returns

 DataFrame

Combination of the provided DataFrames.

See Also

 DataFrame.combine_first : Combine two DataFrame objects and default to non-null values in frame calling the method.

Examples

 Combine using a simple function that chooses the smaller column.

```
>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [4, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> take_smaller = lambda s1, s2: s1 if s1.sum() < s2.sum() else s2
>>> df1.combine(df2, take_smaller)
```

	A	B
0	0	3
1	0	3

Example using a true element-wise combine function.

```
>>> df1 = pd.DataFrame({'A': [5, 0], 'B': [2, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> df1.combine(df2, np.minimum)
```

	A	B
0	1	2
1	0	3

Using `fill_value` fills Nones prior to passing the column to the merge function.

```
>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [None, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> df1.combine(df2, take_smaller, fill_value=-5)
```

	A	B
0	0	-5.0
1	0	4.0

However, if the same element in both dataframes is None, that None is preserved

```
>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [None, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [None, 3]})
>>> df1.combine(df2, take_smaller, fill_value=-5)
```

	A	B
0	0	-5.0
1	0	3.0

Example that demonstrates the use of `overwrite` and behavior when the axis differ between the dataframes.

```
>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [4, 4]})
>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [-10, 1], }, index=[1, 2])
>>> df1.combine(df2, take_smaller)
```

	A	B	C
0	NaN	NaN	NaN
1	NaN	3.0	-10.0

```
2 NaN 3.0 1.0
```

```
>>> df1.combine(df2, take_smaller, overwrite=False)
```

```
   A    B    C
0  0.0 NaN NaN
1  0.0 3.0 -10.0
2  NaN 3.0 1.0
```

Demonstrating the preference of the passed in dataframe.

```
>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [1, 1], }, index=[1, 2])
```

```
>>> df2.combine(df1, take_smaller)
```

```
   A    B    C
0  0.0 NaN NaN
1  0.0 3.0 NaN
2  NaN 3.0 NaN
```

```
>>> df2.combine(df1, take_smaller, overwrite=False)
```

```
   A    B    C
0  0.0 NaN NaN
1  0.0 3.0 1.0
2  NaN 3.0 1.0
```

`combine_first(self, other: 'DataFrame') -> 'DataFrame'`

Update null elements with value in the same location in `other`.

Combine two DataFrame objects by filling null values in one DataFrame with non-null values from other DataFrame. The row and column indexes of the resulting DataFrame will be the union of the two.

Parameters

`other : DataFrame`

Provided DataFrame to use to fill null values.

Returns

DataFrame

See Also

`DataFrame.combine` : Perform series-wise operation on two DataFrames using a given function.

Examples

```
>>> df1 = pd.DataFrame({'A': [None, 0], 'B': [None, 4]})
```

```
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
```

```
>>> df1.combine_first(df2)
```

```
   A    B
0  1.0  3.0
1  0.0  4.0
```

Null values still persist if the location of that null value does not exist in `other`

```
>>> df1 = pd.DataFrame({'A': [None, 0], 'B': [4, None]})
```

```
>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [1, 1]}, index=[1, 2])
```

```
>>> df1.combine_first(df2)
```

```
   A    B    C
0  NaN  4.0 NaN
1  0.0  3.0  1.0
2  NaN  3.0  1.0
```

`compare(self, other: 'DataFrame', align_axis: 'Axis' = 1, keep_shape: 'bool'`

```

= False, keep_equal: 'bool' = False) -> 'DataFrame'
    Compare to another DataFrame and show the differences.

    .. versionadded:: 1.1.0

Parameters
-----
other : DataFrame
    Object to compare with.

align_axis : {0 or 'index', 1 or 'columns'}, default 1
    Determine which axis to align the comparison on.

    * 0, or 'index' : Resulting differences are stacked vertically
      with rows drawn alternately from self and other.
    * 1, or 'columns' : Resulting differences are aligned horizontally
      with columns drawn alternately from self and other.

keep_shape : bool, default False
    If true, all rows and columns are kept.
    Otherwise, only the ones with different values are kept.

keep_equal : bool, default False
    If true, the result keeps values that are equal.
    Otherwise, equal values are shown as NaNs.

Returns
-----
DataFrame
    DataFrame that shows the differences stacked side by side.

    The resulting index will be a MultiIndex with 'self' and 'other'
    stacked alternately at the inner level.

Raises
-----
ValueError
    When the two DataFrames don't have identical labels or shape.

See Also
-----
Series.compare : Compare with another Series and show differences.
DataFrame.equals : Test whether two objects contain the same elements.

Notes
-----
Matching NaNs will not appear as a difference.

Can only compare identically-labeled
(i.e. same shape, identical row and column labels) DataFrames

Examples
-----
>>> df = pd.DataFrame(
...     {
...         "col1": ["a", "a", "b", "b", "a"],
...         "col2": [1.0, 2.0, 3.0, np.nan, 5.0],
...         "col3": [1.0, 2.0, 3.0, 4.0, 5.0]
...     },
...     columns=["col1", "col2", "col3"],
... )
>>> df
   col1  col2  col3
0    a   1.0   1.0
1    a   2.0   2.0

```



```

2    b    3.0    3.0
3    b   NaN    4.0
4    a    5.0    5.0

```

```

>>> df2 = df.copy()
>>> df2.loc[0, 'col1'] = 'c'
>>> df2.loc[2, 'col3'] = 4.0
>>> df2

```

```

   col1  col2  col3
0     c   1.0   1.0
1     a   2.0   2.0
2     b   3.0   4.0
3     b   NaN   4.0
4     a   5.0   5.0

```

Align the differences on columns

```

>>> df.compare(df2)
   col1      col3
   self other self other
0     a       c  NaN   NaN
2  NaN   NaN  3.0   4.0

```

Stack the differences on rows

```

>>> df.compare(df2, align_axis=0)
   col1  col3
0 self    a  NaN
  other   c  NaN
2 self  NaN  3.0
  other NaN  4.0

```

Keep the equal values

```

>>> df.compare(df2, keep_equal=True)
   col1      col3
   self other self other
0     a       c  1.0   1.0
2     b       b  3.0   4.0

```

Keep all original rows and columns

```

>>> df.compare(df2, keep_shape=True)
   col1      col2      col3
   self other self other self other
0     a       c  NaN   NaN  NaN   NaN
1  NaN   NaN  NaN   NaN  NaN   NaN
2  NaN   NaN  NaN   NaN  3.0   4.0
3  NaN   NaN  NaN   NaN  NaN   NaN
4  NaN   NaN  NaN   NaN  NaN   NaN

```

Keep all original rows and columns and also all original values

```

>>> df.compare(df2, keep_shape=True, keep_equal=True)
   col1      col2      col3
   self other self other self other
0     a       c  1.0   1.0  1.0   1.0
1     a       a  2.0   2.0  2.0   2.0
2     b       b  3.0   3.0  3.0   4.0
3     b       b  NaN   NaN  4.0   4.0
4     a       a  5.0   5.0  5.0   5.0

```

`corr(self, method='pearson', min_periods=1) -> 'DataFrame'`
 Compute pairwise correlation of columns, excluding NA/null values.

Parameters

method : {'pearson', 'kendall', 'spearman'} or callable

Method of correlation:

- * pearson : standard correlation coefficient
- * kendall : Kendall Tau correlation coefficient
- * spearman : Spearman rank correlation
- * callable: callable with input two 1d ndarrays and returning a float. Note that the returned matrix from corr will have 1 along the diagonals and will be symmetric regardless of the callable's behavior.

.. versionadded:: 0.24.0

min_periods : int, optional

Minimum number of observations required per pair of columns to have a valid result. Currently only available for Pearson and Spearman correlation.

Returns

DataFrame

Correlation matrix.

See Also

DataFrame.corrwith : Compute pairwise correlation with another DataFrame or Series.

Series.corr : Compute the correlation between two Series.

Examples

```
>>> def histogram_intersection(a, b):
...     v = np.minimum(a, b).sum().round(decimals=1)
...     return v
>>> df = pd.DataFrame([(0.2, 0.3), (0.0, 0.6), (0.6, 0.0), (0.2, 0.1)],
...                     columns=['dogs', 'cats'])
>>> df.corr(method=histogram_intersection)
      dogs  cats
dogs  1.0   0.3
cats  0.3   1.0
```

corrwith(self, other, axis=0, drop=False, method='pearson') -> 'Series'
Compute pairwise correlation.

Pairwise correlation is computed between rows or columns of DataFrame with rows or columns of Series or DataFrame. DataFrames are first aligned along both axes before computing the correlations.

Parameters

other : DataFrame, Series

Object with which to compute correlations.

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to use. 0 or 'index' to compute column-wise, 1 or 'columns'

row-wise.

drop : bool, default False

Drop missing indices from result.

method : {'pearson', 'kendall', 'spearman'} or callable

Method of correlation:

- * pearson : standard correlation coefficient

for

```
* kendall : Kendall Tau correlation coefficient
* spearman : Spearman rank correlation
* callable: callable with input two 1d ndarrays
            and returning a float.
```

```
.. versionadded:: 0.24.0
```

Returns

Series

Pairwise correlations.

See Also

`DataFrame.corr` : Compute pairwise correlation of columns.

```
count(self, axis=0, level=None, numeric_only=False)
Count non-NA cells for each column or row.
```

The values `None`, `NaN`, `NaT`, and optionally `numpy.inf` (depending on `pandas.options.mode.use_inf_as_na`) are considered NA.

Parameters

```
axis : {0 or 'index', 1 or 'columns'}, default 0
    If 0 or 'index' counts are generated for each column.
    If 1 or 'columns' counts are generated for each row.
level : int or str, optional
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a DataFrame.
    A str specifies the level name.
numeric_only : bool, default False
    Include only float, int or boolean data.
```

Returns

Series or DataFrame

For each column/row the number of non-NA/null entries.
If `level` is specified returns a `DataFrame`.

See Also

```
Series.count: Number of non-NA elements in a Series.
DataFrame.value_counts: Count unique combinations of columns.
DataFrame.shape: Number of DataFrame rows and columns (including NA
elements).
DataFrame.isna: Boolean same-sized DataFrame showing places of NA
elements.
```

Examples

Constructing DataFrame from a dictionary:

```
>>> df = pd.DataFrame({"Person":
...                     ["John", "Myla", "Lewis", "John", "Myla"],
...                     "Age": [24., np.nan, 21., 33, 26],
...                     "Single": [False, True, True, True, False]})
>>> df
   Person  Age  Single
0   John  24.0   False
1   Myla   NaN    True
2  Lewis  21.0    True
3   John  33.0    True
4   Myla  26.0   False
```

Notice the uncounted NA values:

```
>>> df.count()
Person      5
Age         4
Single      5
dtype: int64
```

Counts for each **row**:

```
>>> df.count(axis='columns')
0      3
1      2
2      3
3      3
4      3
dtype: int64
```

Counts for one level of a `MultiIndex`:

```
>>> df.set_index(["Person", "Single"]).count(level="Person")
Age
Person
John      2
Lewis     1
Myla      1
```

`cov(self, min_periods: 'Optional[int]' = None, ddof: 'Optional[int]' = 1) -> 'DataFrame'`

Compute pairwise covariance of columns, excluding NA/null values.

Compute the pairwise covariance among the series of a DataFrame. The returned data frame is the `covariance matrix` https://en.wikipedia.org/wiki/Covariance_matrix of the columns of the DataFrame.

Both NA and null values are automatically excluded from the calculation. (See the note below about bias from missing values.) A threshold can be set for the minimum number of observations for each value created. Comparisons with observations below this threshold will be returned as ``NaN``.

This method is generally used for the analysis of time series data to understand the relationship between different measures across time.

Parameters

`min_periods` : int, optional
Minimum number of observations required per pair of columns to have a valid result.

`ddof` : int, default 1
Delta degrees of freedom. The divisor used in calculations is ``N - ddof``, where ``N`` represents the number of elements.

.. versionadded:: 1.1.0

Returns

`DataFrame`
The covariance matrix of the series of the DataFrame.

See Also

Series.cov : Compute covariance with another Series.
 core.window.ExponentialMovingWindow.cov: Exponential weighted sample covariance.
 core.window.Expanding.cov : Expanding sample covariance.
 core.window.Rolling.cov : Rolling sample covariance.

Notes

 Returns the covariance matrix of the DataFrame's time series.
 The covariance is normalized by N-ddof.

For DataFrames that have Series that are missing data (assuming that data is `missing at random`
https://en.wikipedia.org/wiki/Missing_data#Missing_at_random>`__`)
 the returned covariance matrix will be an unbiased estimate
 of the variance and covariance between the member Series.

However, for many applications this estimate may not be acceptable because the estimate covariance matrix is not guaranteed to be positive semi-definite. This could lead to estimate correlations having absolute values which are greater than one, and/or a non-invertible covariance matrix. See `Estimation of covariance matrices`
https://en.wikipedia.org/w/index.php?title=Estimation_of_covariance_matrices>`__` for more details.

Examples

```
>>> df = pd.DataFrame([(1, 2), (0, 3), (2, 0), (1, 1)],
...                    columns=['dogs', 'cats'])
>>> df.cov()
           dogs      cats
dogs  0.666667 -1.000000
cats -1.000000  1.666667

>>> np.random.seed(42)
>>> df = pd.DataFrame(np.random.randn(1000, 5),
...                    columns=['a', 'b', 'c', 'd', 'e'])
>>> df.cov()
           a          b          c          d          e
a  0.998438 -0.020161  0.059277 -0.008943  0.014144
b -0.020161  1.059352 -0.008543 -0.024738  0.009826
c  0.059277 -0.008543  1.010670 -0.001486 -0.000271
d -0.008943 -0.024738 -0.001486  0.921297 -0.013692
e  0.014144  0.009826 -0.000271 -0.013692  0.977795
```

****Minimum number of periods****

This method also supports an optional ``min_periods`` keyword that specifies the required minimum number of non-NA observations for each column pair in order to have a valid result:

```
>>> np.random.seed(42)
>>> df = pd.DataFrame(np.random.randn(20, 3),
...                    columns=['a', 'b', 'c'])
>>> df.loc[df.index[:5], 'a'] = np.nan
>>> df.loc[df.index[5:10], 'b'] = np.nan
>>> df.cov(min_periods=12)
           a          b          c
a  0.316741      NaN -0.150812
b      NaN  1.248003  0.191417
c -0.150812  0.191417  0.895202
```

cummax(self, axis=None, skipna=True, *args, **kwargs)
 Return cumulative maximum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative maximum.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0
The index or the name of the axis. 0 is equivalent to None or 'index'.

skipna : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.

***args, **kwargs**
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

Series or DataFrame
Return cumulative maximum of Series or DataFrame.

See Also

core.window.Expanding.max : Similar functionality but ignores ``NaN`` values.

DataFrame.max : Return the maximum over DataFrame axis.

DataFrame.cummax : Return cumulative maximum over DataFrame axis.

DataFrame.cummin : Return cumulative minimum over DataFrame axis.

DataFrame.cumsum : Return cumulative sum over DataFrame axis.

DataFrame.cumprod : Return cumulative product over DataFrame axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cummax()
0    2.0
1    NaN
2    5.0
3    5.0
4    5.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cummax(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                     [3.0, np.nan],
...                     [1.0, 0.0]],
...                     columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the maximum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cummax()
   A    B
0  2.0  1.0
1  3.0  NaN
2  3.0  1.0
```

To iterate over columns and find the maximum in each row, use ``axis=1``

```
>>> df.cummax(axis=1)
   A    B
0  2.0  2.0
1  3.0  NaN
2  1.0  1.0
```

`cummin(self, axis=None, skipna=True, *args, **kwargs)`
Return cumulative minimum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative minimum.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0
The index or the name of the axis. 0 is equivalent to None or 'index'.

`skipna` : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.

`*args, **kwargs`
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

Series or DataFrame
Return cumulative minimum of Series or DataFrame.

See Also

`core.window.Expanding.min` : Similar functionality but ignores ``NaN`` values.

`DataFrame.min` : Return the minimum over DataFrame axis.

`DataFrame.cummax` : Return cumulative maximum over DataFrame axis.

`DataFrame.cummin` : Return cumulative minimum over DataFrame axis.

`DataFrame.cumsum` : Return cumulative sum over DataFrame axis.

`DataFrame.cumprod` : Return cumulative product over DataFrame axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cummin()
0    2.0
1    NaN
2    2.0
3   -1.0
4   -1.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cummin(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the minimum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cummin()
   A    B
0  2.0  1.0
1  2.0  NaN
2  1.0  0.0
```

To iterate over columns and find the minimum in each row, use ``axis=1``

```
>>> df.cummin(axis=1)
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

`cumprod(self, axis=None, skipna=True, *args, **kwargs)`
Return cumulative product over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative

product.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0
The index or the name of the axis. 0 is equivalent to None or 'index'.

skipna : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.

***args, **kwargs**
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

Series or DataFrame
Return cumulative product of Series or DataFrame.

See Also

core.window.Expanding.prod : Similar functionality but ignores ``NaN`` values.

DataFrame.prod : Return the product over DataFrame axis.

DataFrame.cummax : Return cumulative maximum over DataFrame axis.

DataFrame.cummin : Return cumulative minimum over DataFrame axis.

DataFrame.cumsum : Return cumulative sum over DataFrame axis.

DataFrame.cumprod : Return cumulative product over DataFrame axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cumprod()
0    2.0
1    NaN
2   10.0
3  -10.0
4   -0.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cumprod(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                     [3.0, np.nan],
...                     [1.0, 0.0]],
...                     columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the product in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cumprod()
   A    B
0  2.0  1.0
1  6.0  NaN
2  6.0  0.0
```

To iterate over columns and find the product in each row, use ``axis=1``

```
>>> df.cumprod(axis=1)
   A    B
0  2.0  2.0
1  3.0  NaN
2  1.0  0.0
```

`cumsum(self, axis=None, skipna=True, *args, **kwargs)`
Return cumulative sum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative sum.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0
The index or the name of the axis. 0 is equivalent to None or 'index'.

`skipna` : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.

`*args, **kwargs`
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

Series or DataFrame
Return cumulative sum of Series or DataFrame.

See Also

`core.window.Expanding.sum` : Similar functionality but ignores ``NaN`` values.

`DataFrame.sum` : Return the sum over DataFrame axis.

`DataFrame.cummax` : Return cumulative maximum over DataFrame axis.

`DataFrame.cummin` : Return cumulative minimum over DataFrame axis.

`DataFrame.cumsum` : Return cumulative sum over DataFrame axis.

`DataFrame.cumprod` : Return cumulative product over DataFrame axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cumsum()
0    2.0
1    NaN
2    7.0
3    6.0
4    6.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cumsum(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the sum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cumsum()
   A    B
0  2.0  1.0
1  5.0  NaN
2  6.0  1.0
```

To iterate over columns and find the sum in each row, use ``axis=1``

```
>>> df.cumsum(axis=1)
   A    B
0  2.0  3.0
1  3.0  NaN
2  1.0  1.0
```

`diff(self, periods: 'int' = 1, axis: 'Axis' = 0) -> 'DataFrame'`
First discrete difference of element.

Calculates the difference of a Dataframe element compared with another element in the Dataframe (default is element in previous row).

Parameters

`periods` : int, default 1
Periods to shift for calculating difference, accepts negative values.

`axis` : {0 or 'index', 1 or 'columns'}, default 0
Take difference over rows (0) or columns (1).

Returns

Dataframe
First differences of the Series.

See Also

`Dataframe.pct_change`: Percent change over given number of periods.
`Dataframe.shift`: Shift index by desired number of periods with an optional time freq.
`Series.diff`: First discrete difference of object.

Notes

For boolean dtypes, this uses `:meth:`operator.xor`` rather than `:meth:`operator.sub``.
The result is calculated according to current dtype in Dataframe, however dtype of the result is always float64.

Examples

Difference with previous row

```
>>> df = pd.DataFrame({'a': [1, 2, 3, 4, 5, 6],
...                     'b': [1, 1, 2, 3, 5, 8],
...                     'c': [1, 4, 9, 16, 25, 36]})
>>> df
   a  b   c
0  1  1   1
1  2  1   4
2  3  2   9
3  4  3  16
4  5  5  25
5  6  8  36
```

```
>>> df.diff()
   a    b    c
0 NaN NaN NaN
1 1.0 0.0 3.0
2 1.0 1.0 5.0
3 1.0 1.0 7.0
4 1.0 2.0 9.0
5 1.0 3.0 11.0
```

Difference with previous column

```
>>> df.diff(axis=1)
   a  b  c
0 NaN 0  0
1 NaN -1 3
2 NaN -1 7
3 NaN -1 13
4 NaN  0 20
5 NaN  2 28
```

Difference with 3rd previous row

```
>>> df.diff(periods=3)
      a      b      c
0  NaN  NaN  NaN
1  NaN  NaN  NaN
2  NaN  NaN  NaN
3  3.0  2.0  15.0
4  3.0  4.0  21.0
5  3.0  6.0  27.0
```

Difference with following row

```
>>> df.diff(periods=-1)
      a      b      c
0 -1.0  0.0 -3.0
1 -1.0 -1.0 -5.0
2 -1.0 -1.0 -7.0
3 -1.0 -2.0 -9.0
4 -1.0 -3.0 -11.0
5  NaN  NaN  NaN
```

Overflow in input dtype

```
>>> df = pd.DataFrame({'a': [1, 0]}, dtype=np.uint8)
>>> df.diff()
      a
0  NaN
1 255.0
```

`div = truediv(self, other, axis='columns', level=None, fill_value=None)`

`divide = truediv(self, other, axis='columns', level=None, fill_value=None)`

`dot(self, other)`

Compute the matrix multiplication between the DataFrame and other.

This method computes the matrix product between the DataFrame and the values of an other Series, DataFrame or a numpy array.

It can also be called using ``self @ other`` in Python >= 3.5.

Parameters

`other` : Series, DataFrame or array-like

The other object to compute the matrix product with.

Returns

Series or DataFrame

If other is a Series, return the matrix product between self and other as a Series. If other is a DataFrame or a numpy.array, return the matrix product of self and other in a DataFrame of a np.array.

See Also

`Series.dot`: Similar method for Series.

Notes

The dimensions of DataFrame and other must be compatible in order to compute the matrix multiplication. In addition, the column names of DataFrame and the index of other must contain the same values, as they will be aligned prior to the multiplication.

The dot method for Series computes the inner product, instead of the matrix product here.

Examples

Here we multiply a DataFrame with a Series.

```
>>> df = pd.DataFrame([[0, 1, -2, -1], [1, 1, 1, 1]])
>>> s = pd.Series([1, 1, 2, 1])
>>> df.dot(s)
0    -4
1     5
dtype: int64
```

Here we multiply a DataFrame with another DataFrame.

```
>>> other = pd.DataFrame([[0, 1], [1, 2], [-1, -1], [2, 0]])
>>> df.dot(other)
0    1
0    1    4
1    2    2
```

Note that the dot method give the same result as @

```
>>> df @ other
0    1
0    1    4
1    2    2
```

The dot method works also if other is an np.array.

```
>>> arr = np.array([[0, 1], [1, 2], [-1, -1], [2, 0]])
>>> df.dot(arr)
0    1
0    1    4
1    2    2
```

Note how shuffling of the objects does not change the result.

```
>>> s2 = s.reindex([1, 0, 2, 3])
>>> df.dot(s2)
0    -4
1     5
dtype: int64
```

```
drop(self, labels=None, axis=0, index=None, columns=None, level=None, inplace=False, errors='raise')
Drop specified labels from rows or columns.
```

Remove rows or columns by specifying label names and corresponding axis, or by specifying directly index or column names. When using a multi-index, labels on different levels can be removed by specifying the level.

Parameters

```
labels : single label or list-like
    Index or column labels to drop.
axis : {0 or 'index', 1 or 'columns'}, default 0
    Whether to drop labels from the index (0 or 'index') or columns (1 or 'columns').
index : single label or list-like
    Alternative to specifying axis (`labels, axis=0` is equivalent to `index=labels`).
columns : single label or list-like
```

Alternative to specifying axis (`axis=1` is equivalent to `columns=labels`).

level : int or level name, optional
For MultiIndex, level from which the labels will be removed.

inplace : bool, default False
If False, return a copy. Otherwise, do operation inplace and return None.

errors : {'ignore', 'raise'}, default 'raise'
If 'ignore', suppress error and only existing labels are dropped.

Returns

DataFrame or None

DataFrame without the removed index or column labels or None if `inplace=True`.

Raises

KeyError

If any of the labels is not found in the selected axis.

See Also

DataFrame.loc : Label-location based indexer for selection by label.
DataFrame.dropna : Return DataFrame with labels on given axis omitted where (all or any) data are missing.
DataFrame.drop_duplicates : Return DataFrame with duplicate rows removed, optionally only considering certain columns.
Series.drop : Return Series with specified index labels removed.

Examples

```
>>> df = pd.DataFrame(np.arange(12).reshape(3, 4),
...                    columns=['A', 'B', 'C', 'D'])
```

```
>>> df
   A  B  C  D
0  0  1  2  3
1  4  5  6  7
2  8  9 10 11
```

Drop columns

```
>>> df.drop(['B', 'C'], axis=1)
```

```
   A  D
0  0  3
1  4  7
2  8 11
```

```
>>> df.drop(columns=['B', 'C'])
```

```
   A  D
0  0  3
1  4  7
2  8 11
```

Drop a row by index

```
>>> df.drop([0, 1])
```

```
   A  B  C  D
2  8  9 10 11
```

Drop columns and/or rows of MultiIndex DataFrame

```
>>> midx = pd.MultiIndex(levels=[['lama', 'cow', 'falcon'],
...                               ['speed', 'weight', 'length']],
```

```

...         codes=[[0, 0, 0, 1, 1, 1, 2, 2, 2],
...                 [0, 1, 2, 0, 1, 2, 0, 1, 2]])
>>> df = pd.DataFrame(index=midx, columns=['big', 'small'],
...                     data=[[45, 30], [200, 100], [1.5, 1], [30, 20],
...                          [250, 150], [1.5, 0.8], [320, 250],
...                          [1, 0.8], [0.3, 0.2]])
>>> df

```

		big	small
lama	speed	45.0	30.0
	weight	200.0	100.0
	length	1.5	1.0
cow	speed	30.0	20.0
	weight	250.0	150.0
	length	1.5	0.8
falcon	speed	320.0	250.0
	weight	1.0	0.8
	length	0.3	0.2

```
>>> df.drop(index='cow', columns='small')
```

		big
lama	speed	45.0
	weight	200.0
	length	1.5
falcon	speed	320.0
	weight	1.0
	length	0.3

```
>>> df.drop(index='length', level=1)
```

		big	small
lama	speed	45.0	30.0
	weight	200.0	100.0
cow	speed	30.0	20.0
	weight	250.0	150.0
falcon	speed	320.0	250.0
	weight	1.0	0.8

`drop_duplicates(self, subset: 'Optional[Union[Hashable, Sequence[Hashable]]]' = None, keep: 'Union[str, bool]' = 'first', inplace: 'bool' = False, ignore_index: 'bool' = False) -> 'Optional[DataFrame]'`

Return DataFrame with duplicate rows removed.

Considering certain columns is optional. Indexes, including time indexes are ignored.

Parameters

`subset` : column label or sequence of labels, optional
Only consider certain columns for identifying duplicates, by default use all of the columns.

`keep` : {'first', 'last', False}, default 'first'
Determines which duplicates (if any) to keep.

- 'first' : Drop duplicates except for the first occurrence.
- 'last' : Drop duplicates except for the last occurrence.
- False : Drop all duplicates.

`inplace` : bool, default False
Whether to drop duplicates in place or to return a copy.

`ignore_index` : bool, default False
If True, the resulting axis will be labeled 0, 1, ..., n - 1.

.. versionadded:: 1.0.0

Returns

DataFrame or None

DataFrame with duplicates removed or None if 'inplace=True'.

See Also

`DataFrame.value_counts`: Count unique combinations of columns.

Examples

Consider dataset containing ramen rating.

```
>>> df = pd.DataFrame({
...     'brand': ['Yum Yum', 'Yum Yum', 'Indomie', 'Indomie', 'Indomi
e'],
...     'style': ['cup', 'cup', 'cup', 'pack', 'pack'],
...     'rating': [4, 4, 3.5, 15, 5]
... })
>>> df
   brand style  rating
0  Yum Yum   cup    4.0
1  Yum Yum   cup    4.0
2  Indomie   cup    3.5
3  Indomie  pack   15.0
4  Indomie  pack    5.0
```

By default, it removes duplicate rows based on all columns.

```
>>> df.drop_duplicates()
   brand style  rating
0  Yum Yum   cup    4.0
2  Indomie   cup    3.5
3  Indomie  pack   15.0
4  Indomie  pack    5.0
```

To remove duplicates on specific column(s), use ``subset``.

```
>>> df.drop_duplicates(subset=['brand'])
   brand style  rating
0  Yum Yum   cup    4.0
2  Indomie   cup    3.5
```

To remove duplicates and keep last occurrences, use ``keep``.

```
>>> df.drop_duplicates(subset=['brand', 'style'], keep='last')
   brand style  rating
1  Yum Yum   cup    4.0
2  Indomie   cup    3.5
4  Indomie  pack    5.0
```

`dropna(self, axis=0, how='any', thresh=None, subset=None, inplace=False)`
Remove missing values.

See the :ref:`User Guide <missing_data>` for more on which values are considered missing, and how to work with missing data.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0
Determine if rows or columns which contain missing values are removed.

- * 0, or 'index' : Drop rows which contain missing values.
- * 1, or 'columns' : Drop columns which contain missing value.

.. versionchanged:: 1.0.0

Pass tuple or list to drop on multiple axes.

Only a single axis is allowed.

how : {'any', 'all'}, default 'any'
Determine if row or column is removed from DataFrame, when we have at least one NA or all NA.

- * 'any' : If any NA values are present, drop that row or column.
- * 'all' : If all values are NA, drop that row or column.

thresh : int, optional
Require that many non-NA values.
subset : array-like, optional
Labels along other axis to consider, e.g. if you are dropping rows these would be a list of columns to include.
inplace : bool, default False
If True, do operation inplace and return None.

Returns

DataFrame or None
DataFrame with NA entries dropped from it or None if ``inplace=True`

See Also

DataFrame.isna: Indicate missing values.
DataFrame.notna : Indicate existing (non-missing) values.
DataFrame.fillna : Replace missing values.
Series.dropna : Drop missing values.
Index.dropna : Drop missing indices.

Examples

```
>>> df = pd.DataFrame({"name": ['Alfred', 'Batman', 'Catwoman'],
...                    "toy": [np.nan, 'Batmobile', 'Bullwhip'],
...                    "born": [pd.NaT, pd.Timestamp("1940-04-25"),
...                             pd.NaT]})
>>> df
```

	name	toy	born
0	Alfred	NaN	NaT
1	Batman	Batmobile	1940-04-25
2	Catwoman	Bullwhip	NaT

Drop the rows where at least one element is missing.

```
>>> df.dropna()
```

	name	toy	born
1	Batman	Batmobile	1940-04-25

Drop the columns where at least one element is missing.

```
>>> df.dropna(axis='columns')
```

	name
0	Alfred
1	Batman
2	Catwoman

Drop the rows where all elements are missing.

```
>>> df.dropna(how='all')
```

	name	toy	born
0	Alfred	NaN	NaT
1	Batman	Batmobile	1940-04-25
2	Catwoman	Bullwhip	NaT

Keep only the rows with at least 2 non-NA values.

```
>>> df.dropna(thresh=2)
      name      toy      born
1  Batman  Batmobile 1940-04-25
2  Catwoman  Bullwhip      NaT
```

Define in which columns to look for missing values.

```
>>> df.dropna(subset=['name', 'toy'])
      name      toy      born
1  Batman  Batmobile 1940-04-25
2  Catwoman  Bullwhip      NaT
```

Keep the DataFrame with valid entries in the same variable.

```
>>> df.dropna(inplace=True)
>>> df
      name      toy      born
1  Batman  Batmobile 1940-04-25
```

`one, keep: 'Union[str, bool]' = 'first') -> 'Series'`
 duplicated(self, subset: 'Optional[Union[Hashable, Sequence[Hashable]]]' = None, keep: 'Union[str, bool]' = 'first') -> 'Series'
 Return boolean Series denoting duplicate rows.

Considering certain columns is optional.

Parameters

`subset` : column label or sequence of labels, optional
 Only consider certain columns for identifying duplicates, by default use all of the columns.
`keep` : {'first', 'last', False}, default 'first'
 Determines which duplicates (if any) to mark.

– ‘first’ : Mark duplicates as ‘True’ except for the first occurrence.
 – ‘last’ : Mark duplicates as ‘True’ except for the last occurrence.
 – False : Mark all duplicates as ‘True’.

Returns

Series

Boolean series for each duplicated rows.

See Also

`Index.duplicated` : Equivalent method on index.
`Series.duplicated` : Equivalent method on Series.
`Series.drop_duplicates` : Remove duplicate values from Series.
`DataFrame.drop_duplicates` : Remove duplicate values from DataFrame.

Examples

Consider dataset containing ramen rating.

```
>>> df = pd.DataFrame({
...     'brand': ['Yum Yum', 'Yum Yum', 'Indomie', 'Indomie', 'Indomie'],
...     'style': ['cup', 'cup', 'cup', 'pack', 'pack'],
...     'rating': [4, 4, 3.5, 15, 5]
... })
>>> df
      brand style  rating
0  Yum Yum   cup      4.0
1  Yum Yum   cup      4.0
2  Indomie   cup      3.5
3  Indomie  pack     15.0
4  Indomie  pack      5.0
```

```

0 Yum Yum cup 4.0
1 Yum Yum cup 4.0
2 Indomie cup 3.5
3 Indomie pack 15.0
4 Indomie pack 5.0

```

By default, for each set of duplicated values, the first occurrence is set on False and all others on True.

```

>>> df.duplicated()
0 False
1 True
2 False
3 False
4 False
dtype: bool

```

By using 'last', the last occurrence of each set of duplicated values is set on False and all others on True.

```

>>> df.duplicated(keep='last')
0 True
1 False
2 False
3 False
4 False
dtype: bool

```

By setting ``keep`` on False, all duplicates are True.

```

>>> df.duplicated(keep=False)
0 True
1 True
2 False
3 False
4 False
dtype: bool

```

To find duplicates on specific column(s), use ``subset``.

```

>>> df.duplicated(subset=['brand'])
0 False
1 True
2 False
3 True
4 True
dtype: bool

```

`eq(self, other, axis='columns', level=None)`

Get Equal to of dataframe and other, element-wise (binary operator `eq`)

).

on

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

is

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters

`other` : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

`axis` : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns (1 or 'columns').
 level : int or label
 Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns

 DataFrame of bool
 Result of the comparison.

See Also

 DataFrame.eq : Compare DataFrames for equality elementwise.
 DataFrame.ne : Compare DataFrames for inequality elementwise.
 DataFrame.le : Compare DataFrames for less than inequality or equality elementwise.
 DataFrame.lt : Compare DataFrames for strictly less than inequality elementwise.
 DataFrame.ge : Compare DataFrames for greater than inequality or equality elementwise.
 DataFrame.gt : Compare DataFrames for strictly greater than inequality elementwise.

Notes

 Mismatched indices will be unioned together.
 `NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

 >>> df = pd.DataFrame({'cost': [250, 150, 100],
 ... 'revenue': [100, 250, 300]},
 ... index=['A', 'B', 'C'])
 >>> df
 cost revenue
 A 250 100
 B 150 250
 C 100 300

Comparison with a scalar, using either the operator or method:

>>> df == 100
 cost revenue
 A False True
 B False False
 C True False

 >>> df.eq(100)
 cost revenue
 A False True
 B False False
 C True False

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

>>> df != pd.Series([100, 250], index=["cost", "revenue"])
 cost revenue
 A True True
 B True False
 C False True

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
      cost  revenue
A   True    False
B   True     True
C   True     True
D   True     True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
      cost  revenue
A   True     True
B  False    False
C  False    False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
      cost  revenue
A   True    False
B  False     True
C   True    False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                        index=['A', 'B', 'C', 'D'])
```

```
>>> other
      revenue
A         300
B         250
C         100
D         150
```

```
>>> df.gt(other)
      cost  revenue
A  False    False
B  False    False
C  False     True
D  False    False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                       ['A', 'B', 'C', 'A', 'B', 'C']])
```

```
>>> df_multindex
      cost  revenue
Q1 A   250     100
   B   150     250
   C   100     300
Q2 A   150     200
   B   300     175
   C   220     225
```

```
>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A   True     True
   B   True     True
   C   True     True
```

Q2 A False True
 B True False
 C True False

```
eval(self, expr, inplace=False, **kwargs)
```

Evaluate a string describing operations on DataFrame columns.

Operates on columns only, not specific rows or elements. This allows ``eval`` to run arbitrary code, which can make you vulnerable to code injection if you pass user input to this function.

Parameters

`expr` : str

The expression string to evaluate.

`inplace` : bool, default False

If the expression contains an assignment, whether to perform the operation inplace and mutate the existing DataFrame. Otherwise, a new DataFrame is returned.

`**kwargs`

See the documentation for `:func:`eval`` for complete details on the keyword arguments accepted by `:meth:`~pandas.DataFrame.query``.

Returns

ndarray, scalar, pandas object, or None

The result of the evaluation or None if ``inplace=True``.

See Also

`DataFrame.query` : Evaluates a boolean expression to query the columns of a frame.

`DataFrame.assign` : Can evaluate an expression or function to create new values for a column.

`eval` : Evaluate a Python expression as a string using various backends.

Notes

For more details see the API documentation for `:func:`~eval``.

For detailed examples see `:ref:`enhancing performance with eval <enhancingperf.eval>``.

Examples

```
>>> df = pd.DataFrame({'A': range(1, 6), 'B': range(10, 0, -2)})
```

```
>>> df
```

```
   A  B
0  1 10
1  2  8
2  3  6
3  4  4
4  5  2
```

```
>>> df.eval('A + B')
```

```
0    11
1    10
2     9
3     8
4     7
```

```
dtype: int64
```

Assignment is allowed though by default the original DataFrame is not modified.

```
>>> df.eval('C = A + B')
   A  B  C
0  1 10 11
1  2  8 10
2  3  6  9
3  4  4  8
4  5  2  7
>>> df
   A  B
0  1 10
1  2  8
2  3  6
3  4  4
4  5  2
```

Use ``inplace=True`` to modify the original DataFrame.

```
>>> df.eval('C = A + B', inplace=True)
>>> df
   A  B  C
0  1 10 11
1  2  8 10
2  3  6  9
3  4  4  8
4  5  2  7
```

Multiple columns can be assigned to using multi-line expressions:

```
>>> df.eval(
...     '''
...     C = A + B
...     D = A - B
...     '''
... )
   A  B  C  D
0  1 10 11 -9
1  2  8 10 -6
2  3  6  9 -3
3  4  4  8  0
4  5  2  7  3
```

`explode(self, column: 'Union[str, Tuple]', ignore_index: 'bool' = False) -> 'DataFrame'`

Transform each element of a list-like to a row, replicating index values.

.. versionadded:: 0.25.0

Parameters

`column` : str or tuple

Column to explode.

`ignore_index` : bool, default False

If True, the resulting index will be labeled 0, 1, ..., n - 1.

.. versionadded:: 1.1.0

Returns

DataFrame

Exploded lists to rows of the subset columns;
index will be duplicated for these rows.

Raises

ValueError :
if columns of the frame are not unique.

See Also

DataFrame.unstack : Pivot a level of the (necessarily hierarchical) index labels.
DataFrame.melt : Unpivot a DataFrame from wide format to long format.
Series.explode : Explode a DataFrame from list-like columns to long form

at.

Notes

This routine will explode list-likes including lists, tuples, sets, Series, and np.ndarray. The result dtype of the subset rows will be object. Scalars will be returned unchanged, and empty list-likes will result in a np.nan for that row. In addition, the ordering of rows in the

e

output will be non-deterministic when exploding sets.

Examples

>>> df = pd.DataFrame({'A': [[1, 2, 3], 'foo', [], [3, 4]], 'B': 1})
>>> df

	A	B
0	[1, 2, 3]	1
1	foo	1
2	[]	1
3	[3, 4]	1

>>> df.explode('A')

	A	B
0	1	1
0	2	1
0	3	1
1	foo	1
2	NaN	1
3	3	1
3	4	1

fillna(self, value=None, method=None, axis=None, inplace=False, limit=None, downcast=None) -> 'Optional[DataFrame]'

Fill NA/NaN values using the specified method.

Parameters

value : scalar, dict, Series, or DataFrame

Value to use to fill holes (e.g. 0), alternately a dict/Series/DataFrame of values specifying which value to use for each index (for a Series) or column (for a DataFrame). Values not in the dict/Series/DataFrame will not be filled. This value cannot be a list.

method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None

Method to use for filling holes in reindexed Series
pad / ffill: propagate last valid observation forward to next valid
backfill / bfill: use next valid observation to fill gap.

axis : {0 or 'index', 1 or 'columns'}

Axis along which to fill missing values.

inplace : bool, default False

If True, fill in-place. Note: this will modify any other views on this object (e.g., a no-copy slice for a column in a DataFrame).

limit : int, default None

If method is specified, this is the maximum number of consecutive NaN values to forward/backward fill. In other words, if there is

a gap with more than this number of consecutive NaNs, it will only be partially filled. If method is not specified, this is the maximum number of entries along the entire axis where NaNs will be filled. Must be greater than 0 if not None.

downcast : dict, default is None

A dict of item->dtype of what to downcast if possible, or the string 'infer' which will try to downcast to an appropriate equal type (e.g. float64 to int64 if possible).

Returns

DataFrame or None

Object with missing values filled or None if ``inplace=True``.

See Also

interpolate : Fill NaN values using interpolation.

reindex : Conform object to new index.

asfreq : Convert TimeSeries to specified frequency.

Examples

```
>>> df = pd.DataFrame([[np.nan, 2, np.nan, 0],
...                     [3, 4, np.nan, 1],
...                     [np.nan, np.nan, np.nan, 5],
...                     [np.nan, 3, np.nan, 4]],
...                     columns=list('ABCD'))
>>> df
   A    B    C    D
0 NaN  2.0 NaN  0
1 3.0  4.0 NaN  1
2 NaN  NaN NaN  5
3 NaN  3.0 NaN  4
```

Replace all NaN elements with 0s.

```
>>> df.fillna(0)
   A    B    C    D
0  0.0  2.0  0.0  0
1  3.0  4.0  0.0  1
2  0.0  0.0  0.0  5
3  0.0  3.0  0.0  4
```

We can also propagate non-null values forward or backward.

```
>>> df.fillna(method='ffill')
   A    B    C    D
0 NaN  2.0 NaN  0
1  3.0  4.0 NaN  1
2  3.0  4.0 NaN  5
3  3.0  3.0 NaN  4
```

Replace all NaN elements in column 'A', 'B', 'C', and 'D', with 0, 1, 2, and 3 respectively.

```
>>> values = {'A': 0, 'B': 1, 'C': 2, 'D': 3}
>>> df.fillna(value=values)
   A    B    C    D
0  0.0  2.0  2.0  0
1  3.0  4.0  2.0  1
2  0.0  1.0  2.0  5
3  0.0  3.0  2.0  4
```

Only replace the first NaN element.

```

>>> df.fillna(value=values, limit=1)
   A    B    C    D
0  0.0  2.0  2.0  0
1  3.0  4.0  NaN  1
2  NaN  1.0  NaN  5
3  NaN  3.0  NaN  4

```

`floordiv(self, other, axis='columns', level=None, fill_value=None)`
 Get Integer division of dataframe and other, element-wise (binary operation or ``floordiv``).

Equivalent to ``dataframe // other``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, ``rfloordiv``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

Parameters

`other` : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.

`axis` : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

`level` : int or label
 Broadcast across a level, matching Index values on the passed MultiIndex level.

`fill_value` : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

`DataFrame.add` : Add DataFrames.
`DataFrame.sub` : Subtract DataFrames.
`DataFrame.mul` : Multiply DataFrames.
`DataFrame.div` : Divide DataFrames (float division).
`DataFrame.truediv` : Divide DataFrames (float division).
`DataFrame.floordiv` : Divide DataFrames (integer division).
`DataFrame.mod` : Calculate modulo (remainder after division).
`DataFrame.pow` : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

```

>>> df = pd.DataFrame({'angles': [0, 3, 4],
...                     'degrees': [360, 180, 360]},
...                     index=['circle', 'triangle', 'rectangle'])
>>> df

```

	angles	degrees
circle	0	360
triangle	3	180
rectangle	4	360

Add a scalar with operator version which return the same results.

```
>>> df + 1
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

```
>>> df.add(1)
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

Divide by constant with reverse version.

```
>>> df.div(10)
      angles  degrees
circle    0.0     36.0
triangle  0.3     18.0
rectangle 0.4     36.0
```

```
>>> df.rdiv(10)
      angles  degrees
circle      inf  0.027778
triangle  3.333333  0.055556
rectangle 2.500000  0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
      angles  degrees
circle     -1     358
triangle    2     178
rectangle    3     358
```

```
>>> df.sub([1, 2], axis='columns')
      angles  degrees
circle     -1     358
triangle    2     178
rectangle    3     358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
...
...         axis='index'))
      angles  degrees
circle     -1     359
triangle    2     179
rectangle    3     359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                        index=['circle', 'triangle', 'rectangle'])
>>> other
      angles
circle      0
triangle     3
rectangle     4

>>> df * other
      angles  degrees
circle      0      NaN
```

```

triangle      9      NaN
rectangle     16      NaN

>>> df.mul(other, fill_value=0)
      angles  degrees
circle      0      0.0
triangle     9      0.0
rectangle    16      0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                                'degrees': [360, 180, 360, 360, 540, 720],
...                                index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                       ['circle', 'triangle', 'rectangle',
...                                       'square', 'pentagon', 'hexagon']]})

>>> df_multindex
      angles  degrees
A circle      0      360
  triangle     3      180
  rectangle     4      360
B square      4      360
  pentagon     5      540
  hexagon      6      720

>>> df.div(df_multindex, level=1, fill_value=0)
      angles  degrees
A circle      NaN      1.0
  triangle     1.0      1.0
  rectangle     1.0      1.0
B square      0.0      0.0
  pentagon     0.0      0.0
  hexagon      0.0      0.0

ge(self, other, axis='columns', level=None)
Get Greater than or equal to of dataframe and other, element-wise (binary operator `ge`).

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters
-----
other : scalar, sequence, Series, or DataFrame
    Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}, default 'columns'
    Whether to compare by the index (0 or 'index') or columns (1 or 'columns').
level : int or label
    Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns
-----
DataFrame of bool
    Result of the comparison.

```

See Also

DataFrame.eq : Compare DataFrames for equality elementwise.
 DataFrame.ne : Compare DataFrames for inequality elementwise.
 DataFrame.le : Compare DataFrames for less than inequality
 or equality elementwise.
 DataFrame.lt : Compare DataFrames for strictly less than
 inequality elementwise.
 DataFrame.ge : Compare DataFrames for greater than inequality
 or equality elementwise.
 DataFrame.gt : Compare DataFrames for strictly greater than
 inequality elementwise.

Notes

Mismatched indices will be unioned together.
 `NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                     'revenue': [100, 250, 300]},
...                     index=['A', 'B', 'C'])
>>> df
```

	cost	revenue
A	250	100
B	150	250
C	100	300

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
```

	cost	revenue
A	False	True
B	False	False
C	True	False

```
>>> df.eq(100)
```

	cost	revenue
A	False	True
B	False	False
C	True	False

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
```

	cost	revenue
A	True	True
B	True	False
C	False	True

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
```

	cost	revenue
A	True	False
B	True	True
C	True	True
D	True	True

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
      cost  revenue
A   True    True
B  False   False
C  False   False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
      cost  revenue
A   True    False
B  False    True
C   True    False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                        index=['A', 'B', 'C', 'D'])
>>> other
      revenue
A         300
B         250
C         100
D         150
```

```
>>> df.gt(other)
      cost  revenue
A  False   False
B  False   False
C  False    True
D  False   False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                       ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost  revenue
Q1 A    250     100
   B    150     250
   C    100     300
Q2 A    150     200
   B    300     175
   C    220     225

>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A   True    True
   B   True    True
   C   True    True
Q2 A  False    True
   B   True   False
   C   True   False
```

```
groupby(self, by=None, axis=0, level=None, as_index: 'bool' = True, sort: 'bool' = True, group_keys: 'bool' = True, squeeze: 'bool' = <object object at 0x152b0a4c52f0>, observed: 'bool' = False, dropna: 'bool' = True) -> 'DataFrameGroupBy'
```

Group DataFrame using a mapper or by a Series of columns.

A groupby operation involves some combination of splitting the

object, applying a function, and combining the results. This can be used to group large amounts of data and compute operations on these groups.

Parameters

`by` : mapping, function, label, or list of labels
Used to determine the groups for the groupby.
If `by` is a function, it's called on each value of the object's index. If a dict or Series is passed, the Series or dict VALUES will be used to determine the groups (the Series' values are first aligned; see `align()` method). If an ndarray is passed, the values are used as-is to determine the groups. A label or list of labels may be passed to group by the columns in `self`. Notice that a tuple is interpreted as a (single) key.

`axis` : {0 or 'index', 1 or 'columns'}, default 0
Split along rows (0) or columns (1).

`level` : int, level name, or sequence of such, default None
If the axis is a MultiIndex (hierarchical), group by a particular level or levels.

`as_index` : bool, default True
For aggregated output, return object with group labels as the index. Only relevant for DataFrame input. `as_index=False` is effectively "SQL-style" grouped output.

`sort` : bool, default True
Sort group keys. Get better performance by turning this off. Note this does not influence the order of observations within each group. Groupby preserves the order of rows within each group.

`group_keys` : bool, default True
When calling `apply`, add group keys to index to identify pieces.

`squeeze` : bool, default False
Reduce the dimensionality of the return type if possible, otherwise return a consistent type.

.. deprecated:: 1.1.0

`observed` : bool, default False
This only applies if any of the groupers are Categoricals.
If True: only show observed values for categorical groupers.
If False: show all values for categorical groupers.

`dropna` : bool, default True
If True, and if group keys contain NA values, NA values together with row/column will be dropped.
If False, NA values will also be treated as the key in groups

.. versionadded:: 1.1.0

Returns

DataFrameGroupBy

Returns a groupby object that contains information about the groups.

See Also

`resample` : Convenience method for frequency conversion and resampling of time series.

Notes

See the `user guide` <https://pandas.pydata.org/pandas-docs/stable/groupby.html> for more.

Examples

```
>>> df = pd.DataFrame({'Animal': ['Falcon', 'Falcon',
```



```
...           'Parrot', 'Parrot'],
...           'Max Speed': [380., 370., 24., 26.]})
```

```
>>> df
   Animal  Max Speed
0  Falcon    380.0
1  Falcon    370.0
2  Parrot     24.0
3  Parrot     26.0
>>> df.groupby(['Animal']).mean()
           Max Speed
Animal
Falcon    375.0
Parrot    25.0
```

****Hierarchical Indexes****

We can groupby different levels of a hierarchical index using the `level` parameter:

```
>>> arrays = [['Falcon', 'Falcon', 'Parrot', 'Parrot'],
...           ['Captive', 'Wild', 'Captive', 'Wild']]
>>> index = pd.MultiIndex.from_arrays(arrays, names=('Animal', 'Type'))
>>> df = pd.DataFrame({'Max Speed': [390., 350., 30., 20.]},
...                   index=index)
>>> df
           Max Speed
Animal Type
Falcon Captive    390.0
        Wild     350.0
Parrot Captive     30.0
        Wild     20.0
>>> df.groupby(level=0).mean()
           Max Speed
Animal
Falcon    370.0
Parrot     25.0
>>> df.groupby(level="Type").mean()
           Max Speed
Type
Captive    210.0
Wild      185.0
```

We can also choose to include NA in group keys or not by setting `dropna` parameter, the default setting is `True`:

```
>>> l = [[1, 2, 3], [1, None, 4], [2, 1, 3], [1, 2, 2]]
>>> df = pd.DataFrame(l, columns=["a", "b", "c"])
```

```
>>> df.groupby(by=["b"]).sum()
   a  c
b
1.0 2  3
2.0 2  5
```

```
>>> df.groupby(by=["b"], dropna=False).sum()
   a  c
b
1.0 2  3
2.0 2  5
NaN 1  4
```

```
>>> l = [["a", 12, 12], [None, 12.3, 33.], ["b", 12.3, 123], ["a", 1,
1]]
>>> df = pd.DataFrame(l, columns=["a", "b", "c"])
```

```
>>> df.groupby(by="a").sum()
      b      c
a
a  13.0  13.0
b  12.3  123.0

>>> df.groupby(by="a", dropna=False).sum()
      b      c
a
a  13.0  13.0
b  12.3  123.0
NaN 12.3  33.0
```

`gt(self, other, axis='columns', level=None)`
 Get Greater than of dataframe and other, element-wise (binary operator ``gt``).

Among flexible wrappers (``eq``, ``ne``, ``le``, ``lt``, ``ge``, ``gt``) to comparison operators.

Equivalent to ``==``, ``!=``, ``<=``, ``<``, ``>=``, ``>`` with support to choose axis (rows or columns) and level for comparison.

Parameters

`other` : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
`axis` : {0 or 'index', 1 or 'columns'}, default 'columns'
 Whether to compare by the index (0 or 'index') or columns (1 or 'columns').
`level` : int or label
 Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns

DataFrame of bool
 Result of the comparison.

See Also

`DataFrame.eq` : Compare DataFrames for equality elementwise.
`DataFrame.ne` : Compare DataFrames for inequality elementwise.
`DataFrame.le` : Compare DataFrames for less than inequality or equality elementwise.
`DataFrame.lt` : Compare DataFrames for strictly less than inequality elementwise.
`DataFrame.ge` : Compare DataFrames for greater than inequality or equality elementwise.
`DataFrame.gt` : Compare DataFrames for strictly greater than inequality elementwise.

Notes

Mismatched indices will be unioned together.
``NaN`` values are considered different (i.e. ``NaN` != `NaN``).

Examples

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                     'revenue': [100, 250, 300]},
...                     index=['A', 'B', 'C'])
>>> df
```

	cost	revenue
A	250	100
B	150	250
C	100	300

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
      cost  revenue
A  False    True
B  False    False
C   True    False

>>> df.eq(100)
      cost  revenue
A  False    True
B  False    False
C   True    False
```

When `other` is a `:class:`Series``, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
      cost  revenue
A   True    True
B   True    False
C  False    True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
      cost  revenue
A   True    False
B   True    True
C   True    True
D   True    True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
      cost  revenue
A   True    True
B  False    False
C  False    False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
      cost  revenue
A   True    False
B  False    True
C   True    False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                        index=['A', 'B', 'C', 'D'])
>>> other
      revenue
A         300
B         250
C         100
D         150
```

```
>>> df.gt(other)
      cost  revenue
A  False   False
B  False   False
C  False    True
D  False   False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                               ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost  revenue
Q1 A    250     100
   B    150     250
   C    100     300
Q2 A    150     200
   B    300     175
   C    220     225

>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A    True     True
   B    True     True
   C    True     True
Q2 A   False     True
   B    True    False
   C    True    False
```

```
hist = hist_frame(data: 'DataFrame', column: Union[Hashable, NoneType, Sequence[Union[Hashable, NoneType]]] = None, by=None, grid: bool = True, xlabelsize: Union[int, NoneType] = None, xrot: Union[float, NoneType] = None, ylabelsize: Union[int, NoneType] = None, yrot: Union[float, NoneType] = None, ax=None, sharex: bool = False, sharey: bool = False, figsize: Union[Tuple[int, int], NoneType] = None, layout: Union[Tuple[int, int], NoneType] = None, bins: Union[int, Sequence[int]] = 10, backend: Union[str, NoneType] = None, legend: bool = False, **kwargs)
```

Make a histogram of the DataFrame's.

A `histogram` is a representation of the distribution of data. This function calls `matplotlib.pyplot.hist`, on each series in the DataFrame, resulting in one histogram per column.

.. `_histogram`: <https://en.wikipedia.org/wiki/Histogram>

Parameters

```
data : DataFrame
    The pandas object holding the data.
column : str or sequence
    If passed, will be used to limit data to a subset of columns.
by : object, optional
    If passed, then used to form histograms for separate groups.
grid : bool, default True
    Whether to show axis grid lines.
xlabelsize : int, default None
    If specified changes the x-axis label size.
xrot : float, default None
    Rotation of x axis labels. For example, a value of 90 displays the x labels rotated 90 degrees clockwise.
```

```

ylabelsize : int, default None
    If specified changes the y-axis label size.
yrot : float, default None
    Rotation of y axis labels. For example, a value of 90 displays the
    y labels rotated 90 degrees clockwise.
ax : Matplotlib axes object, default None
    The axes to plot the histogram on.
sharex : bool, default True if ax is None else False
    In case subplots=True, share x axis and set some x axis labels to
    invisible; defaults to True if ax is None otherwise False if an ax
    is passed in.
    Note that passing in both an ax and sharex=True will alter all x axis
    labels for all subplots in a figure.
sharey : bool, default False
    In case subplots=True, share y axis and set some y axis labels to
    invisible.
figsize : tuple
    The size in inches of the figure to create. Uses the value in
    `matplotlib.rcParams` by default.
layout : tuple, optional
    Tuple of (rows, columns) for the layout of the histograms.
bins : int or sequence, default 10
    Number of histogram bins to be used. If an integer is given, bins +
    bin edges are calculated and returned. If bins is a sequence, gives
    bin edges, including left edge of first bin and right edge of last
    bin. In this case, bins is returned unmodified.

backend : str, default None
    Backend to use instead of the backend specified in the option
    ``plotting.backend``. For instance, 'matplotlib'. Alternatively, to
    specify the ``plotting.backend`` for the whole session, set
    ``pd.options.plotting.backend``.

.. versionadded:: 1.0.0

legend : bool, default False
    Whether to show the legend.

.. versionadded:: 1.1.0

**kwargs
    All other plotting keyword arguments to be passed to
    :meth:`matplotlib.pyplot.hist`.

Returns
-----
matplotlib.AxesSubplot or numpy.ndarray of them

See Also
-----
matplotlib.pyplot.hist : Plot a histogram using matplotlib.

Examples
-----
This example draws a histogram based on the length and width of
some animals, displayed in three bins

.. plot::
    :context: close-figs

    >>> df = pd.DataFrame({
    ...     'length': [1.5, 0.5, 1.2, 0.9, 3],
    ...     'width': [0.7, 0.2, 0.15, 0.2, 1.1]

```

```

...     }, index=['pig', 'rabbit', 'duck', 'chicken', 'horse'])
>>> hist = df.hist(bins=3)

idxmax(self, axis=0, skipna=True) -> 'Series'
Return index of first occurrence of maximum over requested axis.

NA/null values are excluded.

Parameters
-----
axis : {0 or 'index', 1 or 'columns'}, default 0
    The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for column-wise.
skipna : bool, default True
    Exclude NA/null values. If an entire row/column is NA, the result will be NA.

Returns
-----
Series
    Indexes of maxima along the specified axis.

Raises
-----
ValueError
    * If the row/column is empty

See Also
-----
Series.idxmax : Return index of the maximum element.

Notes
-----
This method is the DataFrame version of ``ndarray.argmax``.

Examples
-----
Consider a dataset containing food consumption in Argentina.

>>> df = pd.DataFrame({'consumption': [10.51, 103.11, 55.48],
...                     'co2_emissions': [37.2, 19.66, 1712]},
...                     index=['Pork', 'Wheat Products', 'Beef'])

>>> df
           consumption  co2_emissions
Pork                10.51           37.20
Wheat Products      103.11           19.66
Beef                 55.48          1712.00

By default, it returns the index for the maximum value in each column.

>>> df.idxmax()
consumption    Wheat Products
co2_emissions                Beef
dtype: object

To return the index for the maximum value in each row, use ``axis="columns"``.

>>> df.idxmax(axis="columns")
Pork                co2_emissions
Wheat Products      consumption
Beef                 co2_emissions
dtype: object

```

```

idxmin(self, axis=0, skipna=True) -> 'Series'
    Return index of first occurrence of minimum over requested axis.

    NA/null values are excluded.

    Parameters
    -----
    axis : {0 or 'index', 1 or 'columns'}, default 0
        The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for column-wise.
    skipna : bool, default True
        Exclude NA/null values. If an entire row/column is NA, the result will be NA.

    Returns
    -----
    Series
        Indexes of minima along the specified axis.

    Raises
    -----
    ValueError
        * If the row/column is empty

    See Also
    -----
    Series.idxmin : Return index of the minimum element.

    Notes
    -----
    This method is the DataFrame version of ``ndarray.argmin``.

    Examples
    -----
    Consider a dataset containing food consumption in Argentina.

    >>> df = pd.DataFrame({'consumption': [10.51, 103.11, 55.48],
    ...                     'co2_emissions': [37.2, 19.66, 1712]},
    ...                     index=['Pork', 'Wheat Products', 'Beef'])

    >>> df
               consumption  co2_emissions
    Pork                10.51             37.20
    Wheat Products      103.11             19.66
    Beef                55.48            1712.00

    By default, it returns the index for the minimum value in each column.

    >>> df.idxmin()
    consumption                Pork
    co2_emissions      Wheat Products
    dtype: object

    To return the index for the minimum value in each row, use ``axis="columns"``.

    >>> df.idxmin(axis="columns")
    Pork                consumption
    Wheat Products      co2_emissions
    Beef                consumption
    dtype: object

    info(self, verbose: 'Optional[bool]' = None, buf: 'Optional[IO[str]]' = None,
    max_cols: 'Optional[int]' = None, memory_usage: 'Optional[Union[bool, str]]'
    = None, show_counts: 'Optional[bool]' = None, null_counts: 'Optional[bool]' = No

```

```
ne) -> 'None'
```

```
Print a concise summary of a DataFrame.
```

This method prints information about a DataFrame including the index dtype and columns, non-null values and memory usage.

Parameters

data : DataFrame

DataFrame to print information about.

verbose : bool, optional

Whether to print the full summary. By default, the setting in ``pandas.options.display.max_info_columns`` is followed.

buf : writable buffer, defaults to sys.stdout

Where to send the output. By default, the output is printed to sys.stdout. Pass a writable buffer if you need to further process the output.

max_cols : int, optional

When to switch from the verbose to the truncated output. If the DataFrame has more than `max_cols` columns, the truncated output is used. By default, the setting in ``pandas.options.display.max_info_columns`` is used.

memory_usage : bool, str, optional

Specifies whether total memory usage of the DataFrame elements (including the index) should be displayed. By default, this follows the ``pandas.options.display.memory_usage`` setting.

True always show memory usage. False never shows memory usage. A value of 'deep' is equivalent to "True with deep introspection". Memory usage is shown in human-readable units (base-2 representation). Without deep introspection a memory estimation is made based in column dtype and number of rows assuming values consume the same memory amount for corresponding dtypes. With deep memory introspection, a real memory usage calculation is performed at the cost of computational resources.

show_counts : bool, optional

Whether to show the non-null counts. By default, this is shown only if the DataFrame is smaller than ``pandas.options.display.max_info_rows`` and ``pandas.options.display.max_info_columns``. A value of True always shows the counts, and False never shows the counts.

null_counts : bool, optional

.. deprecated:: 1.2.0
Use show_counts instead.

Returns

None

This method prints a summary of a DataFrame and returns None.

See Also

DataFrame.describe: Generate descriptive statistics of DataFrame columns.

DataFrame.memory_usage: Memory usage of DataFrame columns.

Examples

```
>>> int_values = [1, 2, 3, 4, 5]
>>> text_values = ['alpha', 'beta', 'gamma', 'delta', 'epsilon']
>>> float_values = [0.0, 0.25, 0.5, 0.75, 1.0]
>>> df = pd.DataFrame({"int_col": int_values, "text_col": text_values,
...                    "float_col": float_values})
>>> df
   int_col text_col  float_col
```


0	1	alpha	0.00
1	2	beta	0.25
2	3	gamma	0.50
3	4	delta	0.75
4	5	epsilon	1.00

Prints information of all columns:

```
>>> df.info(verbose=True)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   int_col     5 non-null      int64
1   text_col    5 non-null      object
2   float_col   5 non-null      float64
dtypes: float64(1), int64(1), object(1)
memory usage: 248.0+ bytes
```

Prints a summary of columns count and its dtypes but not per column information:

```
>>> df.info(verbose=False)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Columns: 3 entries, int_col to float_col
dtypes: float64(1), int64(1), object(1)
memory usage: 248.0+ bytes
```

Pipe output of DataFrame.info to buffer instead of sys.stdout, get buffer content and writes to a text file:

```
>>> import io
>>> buffer = io.StringIO()
>>> df.info(buf=buffer)
>>> s = buffer.getvalue()
>>> with open("df_info.txt", "w",
...         encoding="utf-8") as f: # doctest: +SKIP
...     f.write(s)
260
```

The `memory_usage` parameter allows deep introspection mode, specially useful for big DataFrames and fine-tune memory optimization:

```
>>> random_strings_array = np.random.choice(['a', 'b', 'c'], 10 ** 6)
>>> df = pd.DataFrame({
...     'column_1': np.random.choice(['a', 'b', 'c'], 10 ** 6),
...     'column_2': np.random.choice(['a', 'b', 'c'], 10 ** 6),
...     'column_3': np.random.choice(['a', 'b', 'c'], 10 ** 6)
... })
>>> df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   column_1    1000000 non-null  object
1   column_2    1000000 non-null  object
2   column_3    1000000 non-null  object
dtypes: object(3)
memory usage: 22.9+ MB

>>> df.info(memory_usage='deep')
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   column_1    1000000 non-null  object
1   column_2    1000000 non-null  object
2   column_3    1000000 non-null  object
dtypes: object(3)
memory usage: 165.9 MB

```

```

insert(self, loc, column, value, allow_duplicates=False) -> 'None'
Insert column into DataFrame at specified location.

```

Raises a `ValueError` if `'column'` is already contained in the `DataFrame`, unless `'allow_duplicates'` is set to `True`.

Parameters

```

loc : int
    Insertion index. Must verify 0 <= loc <= len(columns).
column : str, number, or hashable object
    Label of the inserted column.
value : int, Series, or array-like
allow_duplicates : bool, optional

```

```

isin(self, values) -> 'DataFrame'
Whether each element in the DataFrame is contained in values.

```

Parameters

```

values : iterable, Series, DataFrame or dict
    The result will only be true at a location if all the
    labels match. If 'values' is a Series, that's the index. If
    'values' is a dict, the keys must be the column names,
    which must match. If 'values' is a DataFrame,
    then both the index and column labels must match.

```

Returns

```

DataFrame
    DataFrame of booleans showing whether each element in the DataFrame
    is contained in values.

```

See Also

```

DataFrame.eq: Equality test for DataFrame.
Series.isin: Equivalent method on Series.
Series.str.contains: Test if pattern or regex is contained within a
    string of a Series or Index.

```

Examples

```

>>> df = pd.DataFrame({'num_legs': [2, 4], 'num_wings': [2, 0]},
...                     index=['falcon', 'dog'])
>>> df
      num_legs  num_wings
falcon         2         2
dog            4         0

```

When `'values'` is a list check whether every value in the `DataFrame` is present in the list (which animals have 0 or 2 legs or wings)

```

>>> df.isin([0, 2])
      num_legs  num_wings
falcon      True      True

```

```
dog          False      True
```

When ``values`` is a dict, we can pass values to check for each column separately:

```
>>> df.isin({'num_wings': [0, 3]})
      num_legs  num_wings
falcon      False      False
dog          False       True
```

When ``values`` is a Series or DataFrame the index and column must match. Note that 'falcon' does not match based on the number of legs in df2.

```
>>> other = pd.DataFrame({'num_legs': [8, 2], 'num_wings': [0, 2]},
...                       index=['spider', 'falcon'])
>>> df.isin(other)
      num_legs  num_wings
falcon      True       True
dog          False      False
```

isna(self) -> 'DataFrame'
Detect missing values.

Return a boolean same-sized object indicating if the values are NA. NA values, such as None or :attr:`numpy.NaN`, gets mapped to True values.

Everything else gets mapped to False values. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``).

Returns

DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is an NA value.

See Also

DataFrame.isnull : Alias of isna.

DataFrame.notna : Boolean inverse of isna.

DataFrame.dropna : Omit axes labels with missing values.

isna : Top-level isna.

Examples

Show which entries in a DataFrame are NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
...                        born=[pd.NaT, pd.Timestamp('1939-05-27'),
...                               pd.Timestamp('1940-04-25')],
...                        name=['Alfred', 'Batman', ''],
...                        toy=[None, 'Batmobile', 'Joker']))
>>> df
   age      born  name      toy
0  5.0      NaT  Alfred    None
1  6.0  1939-05-27  Batman  Batmobile
2  NaN  1940-04-25      Joker

>>> df.isna()
   age      born  name      toy
0  False    True  False    True
1  False    False  False    False
2    True    False  False    False
```

Show which entries in a Series are NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64

>>> ser.isna()
0    False
1    False
2     True
dtype: bool
```

`isnull(self) -> 'DataFrame'`
Detect missing values.

Return a boolean same-sized object indicating if the values are NA. NA values, such as None or :attr:`numpy.NaN`, gets mapped to True values. Everything else gets mapped to False values. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``).

Returns

DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is an NA value.

See Also

DataFrame.isnull : Alias of isna.
DataFrame.notna : Boolean inverse of isna.
DataFrame.dropna : Omit axes labels with missing values.
isna : Top-level isna.

Examples

Show which entries in a DataFrame are NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
...                        born=[pd.NaT, pd.Timestamp('1939-05-27'),
...                             pd.Timestamp('1940-04-25')],
...                        name=['Alfred', 'Batman', ''],
...                        toy=[None, 'Batmobile', 'Joker']))
>>> df
   age    born  name    toy
0  5.0    NaT  Alfred  None
1  6.0  1939-05-27  Batman  Batmobile
2  NaN  1940-04-25      Joker

>>> df.isna()
   age  born  name  toy
0  False  True  False  True
1  False  False  False  False
2   True  False  False  False
```

Show which entries in a Series are NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
```

```
2    NaN
dtype: float64
```

```
>>> ser.isna()
0    False
1    False
2     True
dtype: bool
```

```
items(self) -> 'Iterable[Tuple[Label, Series]]'
Iterate over (column name, Series) pairs.
```

Iterates over the DataFrame columns, returning a tuple with the column name and the content as a Series.

Yields

```
label : object
    The column names for the DataFrame being iterated over.
content : Series
    The column entries belonging to each label, as a Series.
```

See Also

```
DataFrame.iterrows : Iterate over DataFrame rows as
    (index, Series) pairs.
DataFrame.itertuples : Iterate over DataFrame rows as namedtuples
    of the values.
```

Examples

```
>>> df = pd.DataFrame({'species': ['bear', 'bear', 'marsupial'],
...                     'population': [1864, 22000, 80000]},
...                     index=['panda', 'polar', 'koala'])
>>> df
      species  population
panda   bear         1864
polar   bear        22000
koala  marsupial    80000
>>> for label, content in df.items():
...     print(f'label: {label}')
...     print(f'content: {content}', sep='\n')
...
label: species
content:
panda      bear
polar      bear
koala  marsupial
Name: species, dtype: object
label: population
content:
panda      1864
polar     22000
koala     80000
Name: population, dtype: int64
```

```
iteritems(self) -> 'Iterable[Tuple[Label, Series]]'
Iterate over (column name, Series) pairs.
```

Iterates over the DataFrame columns, returning a tuple with the column name and the content as a Series.

Yields

```
label : object
```

The column names for the DataFrame being iterated over.
 content : Series
 The column entries belonging to each label, as a Series.

See Also

 DataFrame.iterrows : Iterate over DataFrame rows as
 (index, Series) pairs.
 DataFrame.itertuples : Iterate over DataFrame rows as namedtuples
 of the values.

Examples

```
>>> df = pd.DataFrame({'species': ['bear', 'bear', 'marsupial'],
...                     'population': [1864, 22000, 80000]},
...                     index=['panda', 'polar', 'koala'])
>>> df
```

```
   species  population
panda   bear      1864
polar   bear     22000
koala  marsupial  80000
```

```
>>> for label, content in df.items():
...     print(f'label: {label}')
...     print(f'content: {content}', sep='\n')
... 
```

```
label: species
content:
panda      bear
polar      bear
koala  marsupial
Name: species, dtype: object
label: population
content:
panda      1864
polar     22000
koala     80000
Name: population, dtype: int64
```

iterrows(self) -> 'Iterable[Tuple[Label, Series]]'
 Iterate over DataFrame rows as (index, Series) pairs.

Yields

 index : label or tuple of label
 The index of the row. A tuple for a `MultiIndex`.
 data : Series
 The data of the row as a Series.

See Also

 DataFrame.itertuples : Iterate over DataFrame rows as namedtuples of the
 values.

DataFrame.items : Iterate over (column name, Series) pairs.

Notes

1. Because ``iterrows`` returns a Series for each row, it does **not** preserve dtypes across the rows (dtypes are preserved across columns for DataFrames). For example,

```
>>> df = pd.DataFrame([[1, 1.5]], columns=['int', 'float'])
>>> row = next(df.iterrows())[1]
>>> row
int      1.0
float    1.5
```

```
Name: 0, dtype: float64
>>> print(row['int'].dtype)
float64
>>> print(df['int'].dtype)
int64
```

To preserve dtypes while iterating over the rows, it is better to use `:meth:`itertuples`` which returns namedtuples of the values and which is generally faster than ``iterrows``.

2. You should **never modify** something you are iterating over. This is not guaranteed to work in all cases. Depending on the data types, the iterator returns a copy and not a view, and writing to it will have no effect.

```
itertuples(self, index: 'bool' = True, name: 'Optional[str]' = 'Pandas')
Iterate over DataFrame rows as namedtuples.
```

Parameters

```
index : bool, default True
    If True, return the index as the first element of the tuple.
name : str or None, default "Pandas"
    The name of the returned namedtuples or None to return regular
    tuples.
```

Returns

```
iterator
    An object to iterate over namedtuples for each row in the
    DataFrame with the first field possibly being the index and
    following fields being the column values.
```

See Also

```
DataFrame.iterrows : Iterate over DataFrame rows as (index, Series)
    pairs.
DataFrame.items : Iterate over (column name, Series) pairs.
```

Notes

The column names will be renamed to positional names if they are invalid Python identifiers, repeated, or start with an underscore. On python versions < 3.7 regular tuples are returned for DataFrames with a large number of columns (>254).

Examples

```
>>> df = pd.DataFrame({'num_legs': [4, 2], 'num_wings': [0, 2]},
...                    index=['dog', 'hawk'])
>>> df
   num_legs  num_wings
dog         4         0
hawk        2         2
>>> for row in df.itertuples():
...     print(row)
...
Pandas(Index='dog', num_legs=4, num_wings=0)
Pandas(Index='hawk', num_legs=2, num_wings=2)
```

By setting the ``index`` parameter to `False` we can remove the index as the first element of the tuple:

```
>>> for row in df.itertuples(index=False):
...     print(row)
```

```
...
Pandas(num_legs=4, num_wings=0)
Pandas(num_legs=2, num_wings=2)
```

With the `name` parameter set we set a custom name for the yielded namedtuples:

```
>>> for row in df.itertuples(name='Animal'):
...     print(row)
...
Animal(Index='dog', num_legs=4, num_wings=0)
Animal(Index='hawk', num_legs=2, num_wings=2)
```

```
join(self, other, on=None, how='left', lsuffix='', rsuffix='', sort=False) -
> 'DataFrame'
```

Join columns of another DataFrame.

Join columns with `other` DataFrame either on index or on a key column. Efficiently join multiple DataFrame objects by index at once by passing a list.

Parameters

other : DataFrame, Series, or list of DataFrame

Index should be similar to one of the columns in this one. If a Series is passed, its name attribute must be set, and that will be used as the column name in the resulting joined DataFrame.

on : str, list of str, or array-like, optional

Column or index level name(s) in the caller to join on the index in `other`, otherwise joins index-on-index. If multiple values given, the `other` DataFrame must have a MultiIndex. Can pass an array as the join key if it is not already contained in the calling DataFrame. Like an Excel VLOOKUP operation.

how : {'left', 'right', 'outer', 'inner'}, default 'left'

How to handle the operation of the two objects.

* left: use calling frame's index (or column if on is specified)

* right: use `other`'s index.

* outer: form union of calling frame's index (or column if on is specified) with `other`'s index, and sort it. lexicographically.

* inner: form intersection of calling frame's index (or column if on is specified) with `other`'s index, preserving the order of the calling's one.

lsuffix : str, default ''

Suffix to use from left frame's overlapping columns.

rsuffix : str, default ''

Suffix to use from right frame's overlapping columns.

sort : bool, default False

Order result DataFrame lexicographically by the join key. If False, the order of the join key depends on the join type (how keyword).

Returns

DataFrame

A dataframe containing columns from both the caller and `other`.

See Also

DataFrame.merge : For column(s)-on-column(s) operations.

Notes

Parameters `on`, `lsuffix`, and `rsuffix` are not supported when passing a list of `DataFrame` objects.

Support for specifying index levels as the `on` parameter was added in version 0.23.0.

Examples

```
>>> df = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K3', 'K4', 'K5'],
...                     'A': ['A0', 'A1', 'A2', 'A3', 'A4', 'A5']})
```

```
>>> df
   key  A
0  K0  A0
1  K1  A1
2  K2  A2
3  K3  A3
4  K4  A4
5  K5  A5
```

```
>>> other = pd.DataFrame({'key': ['K0', 'K1', 'K2'],
...                       'B': ['B0', 'B1', 'B2']})
```

```
>>> other
   key  B
0  K0  B0
1  K1  B1
2  K2  B2
```

Join DataFrames using their indexes.

```
>>> df.join(other, lsuffix='_caller', rsuffix='_other')
   key_caller  A key_other  B
0           K0  A0         K0  B0
1           K1  A1         K1  B1
2           K2  A2         K2  B2
3           K3  A3        NaN NaN
4           K4  A4        NaN NaN
5           K5  A5        NaN NaN
```

If we want to join using the key columns, we need to set key to be the index in both `df` and `other`. The joined DataFrame will have key as its index.

```
>>> df.set_index('key').join(other.set_index('key'))
   A  B
key
K0  A0  B0
K1  A1  B1
K2  A2  B2
K3  A3  NaN
K4  A4  NaN
K5  A5  NaN
```

Another option to join using the key columns is to use the `on` parameter. DataFrame.join always uses `other`'s index but we can use any column in `df`. This method preserves the original DataFrame's index in the result.

```
>>> df.join(other.set_index('key'), on='key')
   key  A  B
0  K0  A0  B0
1  K1  A1  B1
2  K2  A2  B2
3  K3  A3  NaN
4  K4  A4  NaN
5  K5  A5  NaN
```

```

kurt(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
    Return unbiased kurtosis over requested axis.

    Kurtosis obtained using Fisher's definition of
    kurtosis (kurtosis of normal == 0.0). Normalized by N-1.

    Parameters
    -----
    axis : {index (0), columns (1)}
        Axis for the function to be applied on.
    skipna : bool, default True
        Exclude NA/null values when computing the result.
    level : int or level name, default None
        If the axis is a MultiIndex (hierarchical), count along a
        particular level, collapsing into a Series.
    numeric_only : bool, default None
        Include only float, int, boolean columns. If None, will attempt to u
        se
            everything, then use only numeric data. Not implemented for Series.
    **kwargs
        Additional keyword arguments to be passed to the function.

    Returns
    -----
    Series or DataFrame (if level specified)

    kurtosis = kurt(self, axis=None, skipna=None, level=None, numeric_only=None,
    **kwargs)

    le(self, other, axis='columns', level=None)
    Get Less than or equal to of dataframe and other, element-wise (binary o
    perator `le`).

    Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparis
    on
        operators.

    Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose ax
    is
        (rows or columns) and level for comparison.

    Parameters
    -----
    other : scalar, sequence, Series, or DataFrame
        Any single or multiple element data structure, or list-like object.
    axis : {0 or 'index', 1 or 'columns'}, default 'columns'
        Whether to compare by the index (0 or 'index') or columns
        (1 or 'columns').
    level : int or label
        Broadcast across a level, matching Index values on the passed
        MultiIndex level.

    Returns
    -----
    DataFrame of bool
        Result of the comparison.

    See Also
    -----
    DataFrame.eq : Compare DataFrames for equality elementwise.
    DataFrame.ne : Compare DataFrames for inequality elementwise.
    DataFrame.le : Compare DataFrames for less than inequality
        or equality elementwise.
    DataFrame.lt : Compare DataFrames for strictly less than

```

inequality elementwise.
 DataFrame.ge : Compare DataFrames for greater than inequality or equality elementwise.
 DataFrame.gt : Compare DataFrames for strictly greater than inequality elementwise.

Notes

Mismatched indices will be unioned together.
 `NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                     'revenue': [100, 250, 300]},
...                     index=['A', 'B', 'C'])
>>> df
```

	cost	revenue
A	250	100
B	150	250
C	100	300

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
```

	cost	revenue
A	False	True
B	False	False
C	True	False

```
>>> df.eq(100)
```

	cost	revenue
A	False	True
B	False	False
C	True	False

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
```

	cost	revenue
A	True	True
B	True	False
C	False	True

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
```

	cost	revenue
A	True	False
B	True	True
C	True	True
D	True	True

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
```

	cost	revenue
A	True	True
B	False	False
C	False	False

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
      cost  revenue
A   True    False
B   False    True
C   True    False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                       index=['A', 'B', 'C', 'D'])
>>> other
      revenue
A         300
B         250
C         100
D         150
```

```
>>> df.gt(other)
      cost  revenue
A   False    False
B   False    False
C   False     True
D   False    False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                              index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                     ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost  revenue
Q1 A    250     100
   B    150     250
   C    100     300
Q2 A    150     200
   B    300     175
   C    220     225
```

```
>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A   True     True
   B   True     True
   C   True     True
Q2 A  False     True
   B   True    False
   C   True    False
```

lookup(self, row_labels, col_labels) -> 'np.ndarray'
 Label-based "fancy indexing" function for DataFrame.
 Given equal-length arrays of row and column labels, return an
 array of the values corresponding to each (row, col) pair.

.. deprecated:: 1.2.0
 DataFrame.lookup is deprecated,
 use DataFrame.melt and DataFrame.loc instead.
 For further details see
 :ref:`Looking up values by index/column labels <indexing.lookup>`.

Parameters

row_labels : sequence

The row labels to use for lookup.
`col_labels` : sequence
 The column labels to use for lookup.

Returns

`numpy.ndarray`
 The found values.

`lt(self, other, axis='columns', level=None)`
 Get Less than of dataframe and other, element-wise (binary operator ``lt``

```).

Among flexible wrappers (``eq``, ``ne``, ``le``, ``lt``, ``ge``, ``gt``) to comparison operators.

Equivalent to ``==``, ``!=``, ``<=``, ``<``, ``>=``, ``>`` with support to choose axis (rows or columns) and level for comparison.

#### Parameters

-----  
`other` : scalar, sequence, Series, or DataFrame  
 Any single or multiple element data structure, or list-like object.  
`axis` : {0 or 'index', 1 or 'columns'}, default 'columns'  
 Whether to compare by the index (0 or 'index') or columns (1 or 'columns').  
`level` : int or label  
 Broadcast across a level, matching Index values on the passed MultiIndex level.

#### Returns

-----  
`DataFrame of bool`  
 Result of the comparison.

#### See Also

-----  
`DataFrame.eq` : Compare DataFrames for equality elementwise.  
`DataFrame.ne` : Compare DataFrames for inequality elementwise.  
`DataFrame.le` : Compare DataFrames for less than inequality or equality elementwise.  
`DataFrame.lt` : Compare DataFrames for strictly less than inequality elementwise.  
`DataFrame.ge` : Compare DataFrames for greater than inequality or equality elementwise.  
`DataFrame.gt` : Compare DataFrames for strictly greater than inequality elementwise.

#### Notes

-----  
 Mismatched indices will be unioned together.  
``NaN`` values are considered different (i.e. ``NaN` != `NaN``).

#### Examples

-----  

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
... 'revenue': [100, 250, 300]},
... index=['A', 'B', 'C'])
>>> df
```

|   | cost | revenue |
|---|------|---------|
| A | 250  | 100     |
| B | 150  | 250     |
| C | 100  | 300     |

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
 cost revenue
A False True
B False False
C True False
```

```
>>> df.eq(100)
 cost revenue
A False True
B False False
C True False
```

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
 cost revenue
A True True
B True False
C False True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
 cost revenue
A True False
B True True
C True True
D True True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
 cost revenue
A True True
B False False
C False False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
 cost revenue
A True False
B False True
C True False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
... index=['A', 'B', 'C', 'D'])
>>> other
 revenue
A 300
B 250
C 100
D 150

>>> df.gt(other)
 cost revenue
A False False
```

```

B False False
C False True
D False False

```

Compare to a MultiIndex by level.

```

>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
... 'revenue': [100, 250, 300, 200, 175, 225]},
... index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
... ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
 cost revenue
Q1 A 250 100
 B 150 250
 C 100 300
Q2 A 150 200
 B 300 175
 C 220 225

>>> df.le(df_multindex, level=1)
 cost revenue
Q1 A True True
 B True True
 C True True
Q2 A False True
 B True False
 C True False

```

`mad(self, axis=None, skipna=None, level=None)`

Return the mean absolute deviation of the values over the requested axis.

#### Parameters

`axis : {index (0), columns (1)}`

Axis for the function to be applied on.

`skipna : bool, default None`

Exclude NA/null values when computing the result.

`level : int or level name, default None`

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

#### Returns

Series or DataFrame (if level specified)

`max(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)`

Return the maximum of the values over the requested axis.

If you want the `*index*` of the maximum, use ```idxmax```. This is the equivalent of the ```numpy.ndarray``` method ```argmax```.

#### Parameters

`axis : {index (0), columns (1)}`

Axis for the function to be applied on.

`skipna : bool, default True`

Exclude NA/null values when computing the result.

`level : int or level name, default None`

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

`numeric_only : bool, default None`

Include only float, int, boolean columns. If None, will attempt to u

se

everything, then use only numeric data. Not implemented for Series.

**\*\*kwargs**

Additional keyword arguments to be passed to the function.

**Returns**

Series or DataFrame (if level specified)

**See Also**

Series.sum : Return the sum.

Series.min : Return the minimum.

Series.max : Return the maximum.

Series.idxmin : Return the index of the minimum.

Series.idxmax : Return the index of the maximum.

DataFrame.sum : Return the sum over the requested axis.

DataFrame.min : Return the minimum over the requested axis.

DataFrame.max : Return the maximum over the requested axis.

DataFrame.idxmin : Return the index of the minimum over the requested axis.

DataFrame.idxmax : Return the index of the maximum over the requested axis.

is.

is.

**Examples**

```
>>> idx = pd.MultiIndex.from_arrays([
... ['warm', 'warm', 'cold', 'cold'],
... ['dog', 'falcon', 'fish', 'spider']],
... names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
>>> s
```

```
blooded animal
warm dog 4
 falcon 2
cold fish 0
 spider 8
Name: legs, dtype: int64
```

```
>>> s.max()
8
```

Max using level names, as well as indices.

```
>>> s.max(level='blooded')
blooded
warm 4
cold 8
Name: legs, dtype: int64
```

```
>>> s.max(level=0)
blooded
warm 4
cold 8
Name: legs, dtype: int64
```

mean(self, axis=None, skipna=None, level=None, numeric\_only=None, \*\*kwargs)  
Return the mean of the values over the requested axis.

**Parameters**

axis : {index (0), columns (1)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.



```

level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a Series.
numeric_only : bool, default None
 Include only float, int, boolean columns. If None, will attempt to u
se
 everything, then use only numeric data. Not implemented for Series.
**kwargs
 Additional keyword arguments to be passed to the function.

Returns

Series or DataFrame (if level specified)

median(self, axis=None, skipna=None, level=None, numeric_only=None, **kwarg
s)
 Return the median of the values over the requested axis.

Parameters

axis : {index (0), columns (1)}
 Axis for the function to be applied on.
skipna : bool, default True
 Exclude NA/null values when computing the result.
level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a Series.
numeric_only : bool, default None
 Include only float, int, boolean columns. If None, will attempt to u
se
 everything, then use only numeric data. Not implemented for Series.
**kwargs
 Additional keyword arguments to be passed to the function.

Returns

Series or DataFrame (if level specified)

melt(self, id_vars=None, value_vars=None, var_name=None, value_name='value',
col_level=None, ignore_index=True) -> 'DataFrame'
 Unpivot a DataFrame from wide to long format, optionally leaving identif
iers set.

 This function is useful to massage a DataFrame into a format where one
 or more columns are identifier variables (`id_vars`), while all other
 columns, considered measured variables (`value_vars`), are "unpivoted" t
o
 the row axis, leaving just two non-identifier columns, 'variable' and
 'value'.

Parameters

id_vars : tuple, list, or ndarray, optional
 Column(s) to use as identifier variables.
value_vars : tuple, list, or ndarray, optional
 Column(s) to unpivot. If not specified, uses all columns that
 are not set as `id_vars`.
var_name : scalar
 Name to use for the 'variable' column. If None it uses
 ``frame.columns.name`` or 'variable'.
value_name : scalar, default 'value'
 Name to use for the 'value' column.
col_level : int or str, optional
 If columns are a MultiIndex then use this level to melt.
ignore_index : bool, default True

```

If True, original index is ignored. If False, the original index is retained.

Index labels will be repeated as necessary.

.. versionadded:: 1.1.0

Returns

-----  
DataFrame

Unpivoted DataFrame.

See Also

-----  
melt : Identical method.

pivot\_table : Create a spreadsheet-style pivot table as a DataFrame.

DataFrame.pivot : Return reshaped DataFrame organized by given index / column values.

DataFrame.explode : Explode a DataFrame from list-like columns to long format.

Examples

```
>>> df = pd.DataFrame({'A': {0: 'a', 1: 'b', 2: 'c'},
... 'B': {0: 1, 1: 3, 2: 5},
... 'C': {0: 2, 1: 4, 2: 6}})
```

```
>>> df
 A B C
0 a 1 2
1 b 3 4
2 c 5 6
```

```
>>> df.melt(id_vars=['A'], value_vars=['B'])
```

```
 A variable value
0 a B 1
1 b B 3
2 c B 5
```

```
>>> df.melt(id_vars=['A'], value_vars=['B', 'C'])
```

```
 A variable value
0 a B 1
1 b B 3
2 c B 5
3 a C 2
4 b C 4
5 c C 6
```

The names of 'variable' and 'value' columns can be customized:

```
>>> df.melt(id_vars=['A'], value_vars=['B'],
... var_name='myVarname', value_name='myValname')
```

```
 A myVarname myValname
0 a B 1
1 b B 3
2 c B 5
```

Original index values can be kept around:

```
>>> df.melt(id_vars=['A'], value_vars=['B', 'C'], ignore_index=False)
```

```
 A variable value
0 a B 1
1 b B 3
2 c B 5
0 a C 2
1 b C 4
2 c C 6
```

If you have multi-index columns:

```
>>> df.columns = [list('ABC'), list('DEF')]
>>> df
 A B C
 D E F
0 a 1 2
1 b 3 4
2 c 5 6

>>> df.melt(col_level=0, id_vars=['A'], value_vars=['B'])
 A variable value
0 a B 1
1 b B 3
2 c B 5

>>> df.melt(id_vars=[('A', 'D')], value_vars=[('B', 'E')])
 (A, D) variable_0 variable_1 value
0 a B E 1
1 b B E 3
2 c B E 5
```

`memory_usage(self, index=True, deep=False) -> 'Series'`

Return the memory usage of each column in bytes.

The memory usage can optionally include the contribution of the index and elements of `'object'` dtype.

This value is displayed in `'DataFrame.info'` by default. This can be suppressed by setting `'pandas.options.display.memory_usage'` to `False`.

#### Parameters

`index` : bool, default True

Specifies whether to include the memory usage of the DataFrame's index in returned Series. If `'index=True'`, the memory usage of the index is the first item in the output.

`deep` : bool, default False

If True, introspect the data deeply by interrogating `'object'` dtypes for system-level memory consumption, and include it in the returned values.

#### Returns

##### Series

A Series whose index is the original column names and whose values is the memory usage of each column in bytes.

#### See Also

`numpy.ndarray.nbytes` : Total bytes consumed by the elements of an ndarray.

`Series.memory_usage` : Bytes consumed by a Series.

`Categorical` : Memory-efficient array for string values with many repeated values.

`DataFrame.info` : Concise summary of a DataFrame.

#### Examples

```
>>> dtypes = ['int64', 'float64', 'complex128', 'object', 'bool']
>>> data = dict([(t, np.ones(shape=5000, dtype=int).astype(t))
... for t in dtypes])
>>> df = pd.DataFrame(data)
>>> df.head()
```

|   | int64 | float64 | complex128 | object | bool |
|---|-------|---------|------------|--------|------|
| 0 | 1     | 1.0     | 1.0+0.0j   | 1      | True |
| 1 | 1     | 1.0     | 1.0+0.0j   | 1      | True |
| 2 | 1     | 1.0     | 1.0+0.0j   | 1      | True |
| 3 | 1     | 1.0     | 1.0+0.0j   | 1      | True |
| 4 | 1     | 1.0     | 1.0+0.0j   | 1      | True |

```
>>> df.memory_usage()
```

```
Index 128
int64 40000
float64 40000
complex128 80000
object 40000
bool 5000
dtype: int64
```

```
>>> df.memory_usage(index=False)
```

```
int64 40000
float64 40000
complex128 80000
object 40000
bool 5000
dtype: int64
```

The memory footprint of `object` dtype columns is ignored by default:

```
>>> df.memory_usage(deep=True)
```

```
Index 128
int64 40000
float64 40000
complex128 80000
object 180000
bool 5000
dtype: int64
```

Use a Categorical for efficient storage of an object-dtype column with many repeated values.

```
>>> df['object'].astype('category').memory_usage(deep=True)
```

```
5244
```

```
merge(self, right, how='inner', on=None, left_on=None, right_on=None, left_index=False, right_index=False, sort=False, suffixes=('_x', '_y'), copy=True, indicator=False, validate=None) -> 'DataFrame'
```

Merge DataFrame or named Series objects with a database-style join.

The join is done on columns or indexes. If joining columns on columns, the DataFrame indexes \*will be ignored\*. Otherwise if joining indexes on indexes or indexes on a column or columns, the index will be passed on.

When performing a cross merge, no column specifications to merge on are allowed.

Parameters

right : DataFrame or named Series

Object to merge with.

how : {'left', 'right', 'outer', 'inner', 'cross'}, default 'inner'

Type of merge to be performed.

\* left: use only keys from left frame, similar to a SQL left outer join;

preserve key order.

\* right: use only keys from right frame, similar to a SQL right outer join;

```

r join;
| preserve key order.
| * outer: use union of keys from both frames, similar to a SQL full o
uter
| join; sort keys lexicographically.
| * inner: use intersection of keys from both frames, similar to a SQL
inner
| join; preserve the order of the left keys.
| * cross: creates the cartesian product from both frames, preserves t
he order
| of the left keys.
| .. versionadded:: 1.2.0
|
| on : label or list
| Column or index level names to join on. These must be found in both
| DataFrames. If `on` is None and not merging on indexes then this def
aults
| to the intersection of the columns in both DataFrames.
left_on : label or list, or array-like
| Column or index level names to join on in the left DataFrame. Can al
so
| be an array or list of arrays of the length of the left DataFrame.
| These arrays are treated as if they are columns.
right_on : label or list, or array-like
| Column or index level names to join on in the right DataFrame. Can a
lso
| be an array or list of arrays of the length of the right DataFrame.
| These arrays are treated as if they are columns.
left_index : bool, default False
| Use the index from the left DataFrame as the join key(s). If it is a
| MultiIndex, the number of keys in the other DataFrame (either the in
dex
| or a number of columns) must match the number of levels.
right_index : bool, default False
| Use the index from the right DataFrame as the join key. Same caveats
as
| left_index.
sort : bool, default False
| Sort the join keys lexicographically in the result DataFrame. If Fal
se,
| the order of the join keys depends on the join type (how keyword).
suffixes : list-like, default is ("_x", "_y")
| A length-2 sequence where each element is optionally a string
| indicating the suffix to add to overlapping column names in
| `left` and `right` respectively. Pass a value of `None` instead
| of a string to indicate that the column name from `left` or
| `right` should be left as-is, with no suffix. At least one of the
| values must not be None.
copy : bool, default True
| If False, avoid copy if possible.
indicator : bool or str, default False
| If True, adds a column to the output DataFrame called "_merge" with
| information on the source of each row. The column can be given a dif
ferent
| name by providing a string argument. The column will have a Categori
cal
| type with the value of "left_only" for observations whose merge key
only
| appears in the left DataFrame, "right_only" for observations
| whose merge key only appears in the right DataFrame, and "both"
| if the observation's merge key is found in both DataFrames.
|
| validate : str, optional
| If specified, checks if merge is of specified type.

```

- \* "one\_to\_one" or "1:1": check if merge keys are unique in both left and right datasets.
- \* "one\_to\_many" or "1:m": check if merge keys are unique in left dataset.
- \* "many\_to\_one" or "m:1": check if merge keys are unique in right dataset.
- \* "many\_to\_many" or "m:m": allowed, but does not result in checks.

## Returns

### DataFrame

A DataFrame of the two merged objects.

## See Also

`merge_ordered` : Merge with optional filling/interpolation.

`merge_asof` : Merge on nearest keys.

`DataFrame.join` : Similar method using indices.

## Notes

Support for specifying index levels as the ``on``, ``left_on``, and ``right_on`` parameters was added in version 0.23.0

Support for merging named Series objects was added in version 0.24.0

## Examples

```
>>> df1 = pd.DataFrame({'lkey': ['foo', 'bar', 'baz', 'foo'],
... 'value': [1, 2, 3, 5]})
>>> df2 = pd.DataFrame({'rkey': ['foo', 'bar', 'baz', 'foo'],
... 'value': [5, 6, 7, 8]})
>>> df1
 lkey value
0 foo 1
1 bar 2
2 baz 3
3 foo 5
>>> df2
 rkey value
0 foo 5
1 bar 6
2 baz 7
3 foo 8
```

Merge df1 and df2 on the lkey and rkey columns. The value columns have the default suffixes, `_x` and `_y`, appended.

```
>>> df1.merge(df2, left_on='lkey', right_on='rkey')
 lkey value_x rkey value_y
0 foo 1 foo 5
1 foo 1 foo 8
2 foo 5 foo 5
3 foo 5 foo 8
4 bar 2 bar 6
5 baz 3 baz 7
```

Merge DataFrames df1 and df2 with specified left and right suffixes appended to any overlapping columns.

```
>>> df1.merge(df2, left_on='lkey', right_on='rkey',
... suffixes=('_left', '_right'))
 lkey value_left rkey value_right
0 foo 1 foo 5
1 foo 1 foo 8
```

|   |     |   |     |   |
|---|-----|---|-----|---|
| 2 | foo | 5 | foo | 5 |
| 3 | foo | 5 | foo | 8 |
| 4 | bar | 2 | bar | 6 |
| 5 | baz | 3 | baz | 7 |

ave

Merge DataFrames df1 and df2, but raise an exception if the DataFrames have any overlapping columns.

se))

```
>>> df1.merge(df2, left_on='lkey', right_on='rkey', suffixes=(False, False))
```

Traceback (most recent call last):

```
...
ValueError: columns overlap but no suffix specified:
Index(['value'], dtype='object')
```

```
>>> df1 = pd.DataFrame({'a': ['foo', 'bar'], 'b': [1, 2]})
```

```
>>> df2 = pd.DataFrame({'a': ['foo', 'baz'], 'c': [3, 4]})
```

```
>>> df1
```

```
 a b
0 foo 1
1 bar 2
```

```
>>> df2
```

```
 a c
0 foo 3
1 baz 4
```

```
>>> df1.merge(df2, how='inner', on='a')
```

```
 a b c
0 foo 1 3
```

```
>>> df1.merge(df2, how='left', on='a')
```

```
 a b c
0 foo 1 3.0
1 bar 2 NaN
```

```
>>> df1 = pd.DataFrame({'left': ['foo', 'bar']})
```

```
>>> df2 = pd.DataFrame({'right': [7, 8]})
```

```
>>> df1
```

```
 left
0 foo
1 bar
```

```
>>> df2
```

```
 right
0 7
1 8
```

```
>>> df1.merge(df2, how='cross')
```

```
 left right
0 foo 7
1 foo 8
2 bar 7
3 bar 8
```

`min(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)`  
Return the minimum of the values over the requested axis.

If you want the `*index*` of the minimum, use ```idxmin```. This is the equivalent of the ```numpy.ndarray``` method ```argmin```.

Parameters

`axis` : {index (0), columns (1)}

Axis for the function to be applied on.

`skipna` : bool, default True

```

Exclude NA/null values when computing the result.
level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a Series.
numeric_only : bool, default None
 Include only float, int, boolean columns. If None, will attempt to u
 everything, then use only numeric data. Not implemented for Series.
**kwargs
 Additional keyword arguments to be passed to the function.

Returns

Series or DataFrame (if level specified)

See Also

Series.sum : Return the sum.
Series.min : Return the minimum.
Series.max : Return the maximum.
Series.idxmin : Return the index of the minimum.
Series.idxmax : Return the index of the maximum.
DataFrame.sum : Return the sum over the requested axis.
DataFrame.min : Return the minimum over the requested axis.
DataFrame.max : Return the maximum over the requested axis.
DataFrame.idxmin : Return the index of the minimum over the requested ax
DataFrame.idxmax : Return the index of the maximum over the requested ax

Examples

>>> idx = pd.MultiIndex.from_arrays([
... ['warm', 'warm', 'cold', 'cold'],
... ['dog', 'falcon', 'fish', 'spider']],
... names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
>>> s
blooded animal
warm dog 4
 falcon 2
cold fish 0
 spider 8
Name: legs, dtype: int64

>>> s.min()
0

Min using level names, as well as indices.

>>> s.min(level='blooded')
blooded
warm 2
cold 0
Name: legs, dtype: int64

>>> s.min(level=0)
blooded
warm 2
cold 0
Name: legs, dtype: int64

mod(self, other, axis='columns', level=None, fill_value=None)
 Get Modulo of dataframe and other, element-wise (binary operator `mod`).

```



Equivalent to ``dataframe % other``, but with support to substitute a fill\_value for missing data in one of the inputs. With reverse version, `rmod`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

#### Parameters

**other** : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.

**axis** : {0 or 'index', 1 or 'columns'}  
Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

**level** : int or label  
Broadcast across a level, matching Index values on the passed MultiIndex level.

**fill\_value** : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

**DataFrame**  
Result of the arithmetic operation.

#### See Also

**DataFrame.add** : Add DataFrames.  
**DataFrame.sub** : Subtract DataFrames.  
**DataFrame.mul** : Multiply DataFrames.  
**DataFrame.div** : Divide DataFrames (float division).  
**DataFrame.truediv** : Divide DataFrames (float division).  
**DataFrame.floordiv** : Divide DataFrames (integer division).  
**DataFrame.mod** : Calculate modulo (remainder after division).  
**DataFrame.pow** : Calculate exponential power.

#### Notes

Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

|           |   |     |
|-----------|---|-----|
| circle    | 1 | 361 |
| triangle  | 4 | 181 |
| rectangle | 5 | 361 |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub([1, 2], axis='columns')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 359     |
| triangle  | 2      | 179     |
| rectangle | 3      | 359     |

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

|           | angles |
|-----------|--------|
| circle    | 0      |
| triangle  | 3      |
| rectangle | 4      |

```
>>> df * other
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | NaN     |
| triangle  | 9      | NaN     |
| rectangle | 16     | NaN     |

```
>>> df.mul(other, fill_value=0)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 0.0     |
| triangle  | 9      | 0.0     |
| rectangle | 16     | 0.0     |

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
```

```

0]},
e',
n']]
...
'degrees': [360, 180, 360, 360, 540, 72
index=[['A', 'A', 'A', 'B', 'B', 'B'],
['circle', 'triangle', 'rectangl
'square', 'pentagon', 'hexago

```

```

>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

```

```

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

```

mode(self, axis=0, numeric\_only=False, dropna=True) -> 'DataFrame'  
 Get the mode(s) of each element along the selected axis.

The mode of a set of values is the value that appears most often.  
 It can be multiple values.

#### Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0  
 The axis to iterate over while searching for the mode:

- \* 0 or 'index' : get mode of each column
- \* 1 or 'columns' : get mode of each row.

numeric\_only : bool, default False  
 If True, only apply to numeric columns.  
 dropna : bool, default True  
 Don't consider counts of NaN/NaT.

.. versionadded:: 0.24.0

#### Returns

DataFrame  
 The modes of each column or row.

#### See Also

Series.mode : Return the highest frequency value in a Series.  
 Series.value\_counts : Return the counts of values in a Series.

#### Examples

```

>>> df = pd.DataFrame([('bird', 2, 2),
... ('mammal', 4, np.nan),
... ('arthropod', 8, 0),
... ('bird', 2, np.nan)],
... index=('falcon', 'horse', 'spider', 'ostrich'),
... columns=('species', 'legs', 'wings'))
>>> df

```

|         | species   | legs | wings |
|---------|-----------|------|-------|
| falcon  | bird      | 2    | 2.0   |
| horse   | mammal    | 4    | NaN   |
| spider  | arthropod | 8    | 0.0   |
| ostrich | bird      | 2    | NaN   |

By default, missing values are not considered, and the mode of wings are both 0 and 2. Because the resulting DataFrame has two rows, the second row of ``species`` and ``legs`` contains ``NaN``.

```
>>> df.mode()
 species legs wings
0 bird 2.0 0.0
1 NaN NaN 2.0
```

Setting ``dropna=False`` ``NaN`` values are considered and they can be the mode (like for wings).

```
>>> df.mode(dropna=False)
 species legs wings
0 bird 2 NaN
```

Setting ``numeric\_only=True``, only the mode of numeric columns is computed, and columns of other types are ignored.

```
>>> df.mode(numeric_only=True)
 legs wings
0 2.0 0.0
1 NaN 2.0
```

To compute the mode over columns and not rows, use the axis parameter:

```
>>> df.mode(axis='columns', numeric_only=True)
 0 1
falcon 2.0 NaN
horse 4.0 NaN
spider 0.0 8.0
ostrich 2.0 NaN
```

```
mul(self, other, axis='columns', level=None, fill_value=None)
```

Get Multiplication of dataframe and other, element-wise (binary operator `mul`).

Equivalent to ``dataframe \* other``, but with support to substitute a fill\_value for missing data in one of the inputs. With reverse version, `rmul`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

#### Parameters

**other** : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

**axis** : {0 or 'index', 1 or 'columns'}

Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

**level** : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

**fill\_value** : float or None, default None

Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

## Returns

-----  
DataFrame

Result of the arithmetic operation.

## See Also

-----  
DataFrame.add : Add DataFrames.  
DataFrame.sub : Subtract DataFrames.  
DataFrame.mul : Multiply DataFrames.  
DataFrame.div : Divide DataFrames (float division).  
DataFrame.truediv : Divide DataFrames (float division).  
DataFrame.floordiv : Divide DataFrames (integer division).  
DataFrame.mod : Calculate modulo (remainder after division).  
DataFrame.pow : Calculate exponential power.

## Notes

-----  
Mismatched indices will be unioned together.

## Examples

-----  

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

  

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|  | angles | degrees |
|--|--------|---------|
|--|--------|---------|

```

circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
e')),
 axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359

Multiply a DataFrame of different shape with operator version.

>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4

>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 72
0]],
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
e',
... ['circle', 'triangle', 'rectangle',
n']]
... 'square', 'pentagon', 'hexagon'])

>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0

```

```

pentagon 0.0 0.0
hexagon 0.0 0.0

```

```
multiply = mul(self, other, axis='columns', level=None, fill_value=None)
```

```
ne(self, other, axis='columns', level=None)
```

Get Not equal to of dataframe and other, element-wise (binary operator ``ne``).

Among flexible wrappers (``eq``, ``ne``, ``le``, ``lt``, ``ge``, ``gt``) to comparison operators.

Equivalent to ``==``, ``!=``, ``<=``, ``<``, ``>=``, ``>`` with support to choose axis (rows or columns) and level for comparison.

#### Parameters

`other` : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

`axis` : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns (1 or 'columns').

`level` : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

#### Returns

DataFrame of bool

Result of the comparison.

#### See Also

`DataFrame.eq` : Compare DataFrames for equality elementwise.

`DataFrame.ne` : Compare DataFrames for inequality elementwise.

`DataFrame.le` : Compare DataFrames for less than inequality or equality elementwise.

`DataFrame.lt` : Compare DataFrames for strictly less than inequality elementwise.

`DataFrame.ge` : Compare DataFrames for greater than inequality or equality elementwise.

`DataFrame.gt` : Compare DataFrames for strictly greater than inequality elementwise.

#### Notes

Mismatched indices will be unioned together.

``NaN`` values are considered different (i.e. ``NaN` != `NaN``).

#### Examples

```

>>> df = pd.DataFrame({'cost': [250, 150, 100],
... 'revenue': [100, 250, 300]},
... index=['A', 'B', 'C'])
>>> df
 cost revenue
A 250 100
B 150 250
C 100 300

```

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
```

```

 cost revenue
A False True
B False False
C True False

```

```

>>> df.eq(100)
 cost revenue
A False True
B False False
C True False

```

When `other` is a `:class:`Series``, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```

>>> df != pd.Series([100, 250], index=["cost", "revenue"])
 cost revenue
A True True
B True False
C False True

```

Use the method to control the broadcast axis:

```

>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
 cost revenue
A True False
B True True
C True True
D True True

```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```

>>> df == [250, 100]
 cost revenue
A True True
B False False
C False False

```

Use the method to control the axis:

```

>>> df.eq([250, 250, 100], axis='index')
 cost revenue
A True False
B False True
C True False

```

Compare to a DataFrame of different shape.

```

>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
... index=['A', 'B', 'C', 'D'])

```

```

>>> other
 revenue
A 300
B 250
C 100
D 150

```

```

>>> df.gt(other)
 cost revenue
A False False
B False False
C False True
D False False

```



```

Compare to a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
... 'revenue': [100, 250, 300, 200, 175, 225]},
... index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
... ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex

```

|      | cost | revenue |
|------|------|---------|
| Q1 A | 250  | 100     |
| B    | 150  | 250     |
| C    | 100  | 300     |
| Q2 A | 150  | 200     |
| B    | 300  | 175     |
| C    | 220  | 225     |

```

>>> df.le(df_multindex, level=1)

```

|      | cost  | revenue |
|------|-------|---------|
| Q1 A | True  | True    |
| B    | True  | True    |
| C    | True  | True    |
| Q2 A | False | True    |
| B    | True  | False   |
| C    | True  | False   |

`nlargest(self, n, columns, keep='first') -> 'DataFrame'`

Return the first `n` rows ordered by `columns` in descending order.

Return the first `n` rows with the largest values in `columns`, in descending order. The columns that are not specified are returned as well, but not used for ordering.

This method is equivalent to `df.sort_values(columns, ascending=False).head(n)`, but more performant.

#### Parameters

`n` : int

Number of rows to return.

`columns` : label or list of labels

Column label(s) to order by.

`keep` : {'first', 'last', 'all'}, default 'first'

Where there are duplicate values:

- 'first' : prioritize the first occurrence(s)
- 'last' : prioritize the last occurrence(s)
- 'all' : do not drop any duplicates, even it means selecting more than `n` items.

.. versionadded:: 0.24.0

#### Returns

`DataFrame`

The first `n` rows ordered by the given columns in descending order.

#### See Also

`DataFrame.nsmallest` : Return the first `n` rows ordered by `columns` in ascending order.

`DataFrame.sort_values` : Sort DataFrame by the values.

`DataFrame.head` : Return the first `n` rows without re-ordering.

## Notes

This function cannot be used with all column types. For example, when specifying columns with `object` or `category` dtypes, ``TypeError`` is raised.

## Examples

```
>>> df = pd.DataFrame({'population': [59000000, 65000000, 434000,
... 434000, 434000, 337000, 11300,
... 11300, 11300],
... 'GDP': [1937894, 2583560, 12011, 4520, 12128,
... 17036, 182, 38, 311],
... 'alpha-2': ["IT", "FR", "MT", "MV", "BN",
... "IS", "NR", "TV", "AI"]},
... index=["Italy", "France", "Malta",
... "Maldives", "Brunei", "Iceland",
... "Nauru", "Tuvalu", "Anguilla"])
```

```
>>> df
```

|          | population | GDP     | alpha-2 |
|----------|------------|---------|---------|
| Italy    | 59000000   | 1937894 | IT      |
| France   | 65000000   | 2583560 | FR      |
| Malta    | 434000     | 12011   | MT      |
| Maldives | 434000     | 4520    | MV      |
| Brunei   | 434000     | 12128   | BN      |
| Iceland  | 337000     | 17036   | IS      |
| Nauru    | 11300      | 182     | NR      |
| Tuvalu   | 11300      | 38      | TV      |
| Anguilla | 11300      | 311     | AI      |

In the following example, we will use ``nlargest`` to select the three rows having the largest values in column "population".

```
>>> df.nlargest(3, 'population')
```

|        | population | GDP     | alpha-2 |
|--------|------------|---------|---------|
| France | 65000000   | 2583560 | FR      |
| Italy  | 59000000   | 1937894 | IT      |
| Malta  | 434000     | 12011   | MT      |

When using ``keep='last'``, ties are resolved in reverse order:

```
>>> df.nlargest(3, 'population', keep='last')
```

|        | population | GDP     | alpha-2 |
|--------|------------|---------|---------|
| France | 65000000   | 2583560 | FR      |
| Italy  | 59000000   | 1937894 | IT      |
| Brunei | 434000     | 12128   | BN      |

When using ``keep='all'``, all duplicate items are maintained:

```
>>> df.nlargest(3, 'population', keep='all')
```

|          | population | GDP     | alpha-2 |
|----------|------------|---------|---------|
| France   | 65000000   | 2583560 | FR      |
| Italy    | 59000000   | 1937894 | IT      |
| Malta    | 434000     | 12011   | MT      |
| Maldives | 434000     | 4520    | MV      |
| Brunei   | 434000     | 12128   | BN      |

To order by the largest values in column "population" and then "GDP", we can specify multiple columns like in the next example.

```
>>> df.nlargest(3, ['population', 'GDP'])
```

|        | population | GDP     | alpha-2 |
|--------|------------|---------|---------|
| France | 65000000   | 2583560 | FR      |
| Italy  | 59000000   | 1937894 | IT      |

Brunei 434000 12128 BN

notna(self) -> 'DataFrame'

Detect existing (non-missing) values.

Return a boolean same-sized object indicating if the values are not NA. Non-missing values get mapped to True. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use\_inf\_as\_na = True``). NA values, such as None or :attr:`numpy.NaN`, get mapped to False values.

Returns

-----  
DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is not an NA value.

See Also

-----  
DataFrame.notnull : Alias of notna.

DataFrame.isna : Boolean inverse of notna.

DataFrame.dropna : Omit axes labels with missing values.

notna : Top-level notna.

Examples

-----  
Show which entries in a DataFrame are not NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
... born=[pd.NaT, pd.Timestamp('1939-05-27'),
... pd.Timestamp('1940-04-25')],
... name=['Alfred', 'Batman', ''],
... toy=[None, 'Batmobile', 'Joker']))
```

```
>>> df
 age born name toy
0 5.0 NaT Alfred None
1 6.0 1939-05-27 Batman Batmobile
2 NaN 1940-04-25 Joker
```

```
>>> df.notna()
 age born name toy
0 True False True False
1 True True True True
2 False True True True
```

Show which entries in a Series are not NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
```

```
>>> ser
0 5.0
1 6.0
2 NaN
dtype: float64
```

```
>>> ser.notna()
0 True
1 True
2 False
dtype: bool
```

notnull(self) -> 'DataFrame'

Detect existing (non-missing) values.

Return a boolean same-sized object indicating if the values are not NA.

Non-missing values get mapped to True. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use\_inf\_as\_na = True``). NA values, such as None or :attr:`numpy.NaN`, get mapped to False values.

## Returns

### DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is not an NA value.

## See Also

DataFrame.notnull : Alias of notna.

DataFrame.isna : Boolean inverse of notna.

DataFrame.dropna : Omit axes labels with missing values.

notna : Top-level notna.

## Examples

Show which entries in a DataFrame are not NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
... born=[pd.NaT, pd.Timestamp('1939-05-27'),
... pd.Timestamp('1940-04-25')],
... name=['Alfred', 'Batman', ''],
... toy=[None, 'Batmobile', 'Joker']))
```

```
>>> df
 age born name toy
0 5.0 NaT Alfred None
1 6.0 1939-05-27 Batman Batmobile
2 NaN 1940-04-25 Joker
```

```
>>> df.notna()
 age born name toy
0 True False True False
1 True True True True
2 False True True True
```

Show which entries in a Series are not NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
```

```
>>> ser
0 5.0
1 6.0
2 NaN
dtype: float64
```

```
>>> ser.notna()
0 True
1 True
2 False
dtype: bool
```

nsmallest(self, n, columns, keep='first') -> 'DataFrame'

Return the first `n` rows ordered by `columns` in ascending order.

Return the first `n` rows with the smallest values in `columns`, in ascending order. The columns that are not specified are returned as well, but not used for ordering.

This method is equivalent to  
 ``df.sort\_values(columns, ascending=True).head(n)`` , but more performant.

## Parameters

```

n : int
 Number of items to retrieve.
columns : list or str
 Column name or names to order by.
keep : {'first', 'last', 'all'}, default 'first'
 Where there are duplicate values:

 - ``first`` : take the first occurrence.
 - ``last`` : take the last occurrence.
 - ``all`` : do not drop any duplicates, even it means
 selecting more than `n` items.

.. versionadded:: 0.24.0

```

## Returns

```
DataFrame
```

## See Also

```

DataFrame.nlargest : Return the first `n` rows ordered by `columns` in
 descending order.
DataFrame.sort_values : Sort DataFrame by the values.
DataFrame.head : Return the first `n` rows without re-ordering.

```

## Examples

```

>>> df = pd.DataFrame({'population': [59000000, 65000000, 434000,
... 434000, 434000, 337000, 337000,
... 11300, 11300],
... 'GDP': [1937894, 2583560, 12011, 4520, 12128,
... 17036, 182, 38, 311],
... 'alpha-2': ["IT", "FR", "MT", "MV", "BN",
... "IS", "NR", "TV", "AI"]},
... index=["Italy", "France", "Malta",
... "Maldives", "Brunei", "Iceland",
... "Nauru", "Tuvalu", "Anguilla"])
>>> df

```

|          | population | GDP     | alpha-2 |
|----------|------------|---------|---------|
| Italy    | 59000000   | 1937894 | IT      |
| France   | 65000000   | 2583560 | FR      |
| Malta    | 434000     | 12011   | MT      |
| Maldives | 434000     | 4520    | MV      |
| Brunei   | 434000     | 12128   | BN      |
| Iceland  | 337000     | 17036   | IS      |
| Nauru    | 337000     | 182     | NR      |
| Tuvalu   | 11300      | 38      | TV      |
| Anguilla | 11300      | 311     | AI      |

In the following example, we will use ``nsmallest`` to select the three rows having the smallest values in column "population".

```

>>> df.nsmallest(3, 'population')

```

|          | population | GDP   | alpha-2 |
|----------|------------|-------|---------|
| Tuvalu   | 11300      | 38    | TV      |
| Anguilla | 11300      | 311   | AI      |
| Iceland  | 337000     | 17036 | IS      |

When using ``keep='last'``, ties are resolved in reverse order:

```

>>> df.nsmallest(3, 'population', keep='last')

```

|  | population | GDP | alpha-2 |
|--|------------|-----|---------|
|--|------------|-----|---------|

|          |        |     |    |
|----------|--------|-----|----|
| Anguilla | 11300  | 311 | AI |
| Tuvalu   | 11300  | 38  | TV |
| Nauru    | 337000 | 182 | NR |

When using ``keep='all'``, all duplicate items are maintained:

```
>>> df.nsmallest(3, 'population', keep='all')
 population GDP alpha-2
Tuvalu 11300 38 TV
Anguilla 11300 311 AI
Iceland 337000 17036 IS
Nauru 337000 182 NR
```

e can To order by the smallest values in column "population" and then "GDP", w  
specify multiple columns like in the next example.

```
>>> df.nsmallest(3, ['population', 'GDP'])
 population GDP alpha-2
Tuvalu 11300 38 TV
Anguilla 11300 311 AI
Nauru 337000 182 NR
```

nunique(self, axis=0, dropna=True) -> 'Series'  
Count distinct observations over requested axis.

Return Series with number of distinct observations. Can ignore NaN values.

#### Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0  
The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for column-wise.  
dropna : bool, default True  
Don't include NaN in the counts.

#### Returns

Series

#### See Also

Series.nunique: Method nunique for Series.  
DataFrame.count: Count non-NA cells for each column or row.

#### Examples

```
>>> df = pd.DataFrame({'A': [1, 2, 3], 'B': [1, 1, 1]})
>>> df.nunique()
A 3
B 1
dtype: int64

>>> df.nunique(axis=1)
0 1
1 2
2 2
dtype: int64
```

pivot(self, index=None, columns=None, values=None) -> 'DataFrame'  
Return reshaped DataFrame organized by given index / column values.

Reshape data (produce a "pivot" table) based on column values. Uses unique values from specified `index` / `columns` to form axes of the

resulting DataFrame. This function does not support data aggregation, multiple values will result in a MultiIndex in the columns. See the :ref:`User Guide <reshaping>` for more on reshaping.

#### Parameters

`index` : str or object or a list of str, optional  
Column to use to make new frame's index. If None, uses existing index.

.. versionchanged:: 1.1.0  
Also accept list of index names.

`columns` : str or object or a list of str  
Column to use to make new frame's columns.

.. versionchanged:: 1.1.0  
Also accept list of columns names.

`values` : str, object or a list of the previous, optional  
Column(s) to use for populating new frame's values. If not specified, all remaining columns will be used and the result will have hierarchically indexed columns.

#### Returns

##### DataFrame

Returns reshaped DataFrame.

#### Raises

##### ValueError:

When there are any `index`, `columns` combinations with multiple values. `DataFrame.pivot\_table` when you need to aggregate.

#### See Also

`DataFrame.pivot_table` : Generalization of pivot that can handle duplicate values for one index/column pair.

`DataFrame.unstack` : Pivot based on the index values instead of a column.

`wide_to_long` : Wide panel to long format. Less flexible but more user-friendly than melt.

#### Notes

For finer-tuned control, see hierarchical indexing documentation along with the related stack/unstack methods.

#### Examples

```
>>> df = pd.DataFrame({'foo': ['one', 'one', 'one', 'two', 'two',
... 'two'],
... 'bar': ['A', 'B', 'C', 'A', 'B', 'C'],
... 'baz': [1, 2, 3, 4, 5, 6],
... 'zoo': ['x', 'y', 'z', 'q', 'w', 't']})
>>> df
 foo bar baz zoo
0 one A 1 x
1 one B 2 y
2 one C 3 z
3 two A 4 q
4 two B 5 w
5 two C 6 t
```

```
>>> df.pivot(index='foo', columns='bar', values='baz')
bar A B C
foo
one 1 2 3
two 4 5 6

>>> df.pivot(index='foo', columns='bar')['baz']
bar A B C
foo
one 1 2 3
two 4 5 6

>>> df.pivot(index='foo', columns='bar', values=['baz', 'zoo'])
 baz zoo
bar A B C A B C
foo
one 1 2 3 x y z
two 4 5 6 q w t
```

You could also assign a list of column names or a list of index names.

```
>>> df = pd.DataFrame({
... "lev1": [1, 1, 1, 2, 2, 2],
... "lev2": [1, 1, 2, 1, 1, 2],
... "lev3": [1, 2, 1, 2, 1, 2],
... "lev4": [1, 2, 3, 4, 5, 6],
... "values": [0, 1, 2, 3, 4, 5]})
>>> df
 lev1 lev2 lev3 lev4 values
0 1 1 1 1 0
1 1 1 2 2 1
2 1 2 1 3 2
3 2 1 2 4 3
4 2 1 1 5 4
5 2 2 2 6 5

>>> df.pivot(index="lev1", columns=["lev2", "lev3"], values="values")
lev2 1 2
lev3 1 2
lev1
1 0.0 1.0 2.0 NaN
2 4.0 3.0 NaN 5.0

>>> df.pivot(index=["lev1", "lev2"], columns=["lev3"], values="values")
 lev3 1 2
lev1 lev2
1 1 0.0 1.0
 2 2.0 NaN
2 1 4.0 3.0
 2 NaN 5.0
```

A `ValueError` is raised if there are any duplicates.

```
>>> df = pd.DataFrame({"foo": ['one', 'one', 'two', 'two'],
... "bar": ['A', 'A', 'B', 'C'],
... "baz": [1, 2, 3, 4]})
>>> df
 foo bar baz
0 one A 1
1 one A 2
2 two B 3
3 two C 4
```

Notice that the first two rows are the same for our ``index`` and ``columns`` arguments.



```
>>> df.pivot(index='foo', columns='bar', values='baz')
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: Index contains duplicate entries, cannot reshape
```

```
pivot_table(self, values=None, index=None, columns=None, aggfunc='mean', fill_value=None, margins=False, dropna=True, margins_name='All', observed=False) -> DataFrame'
```

Create a spreadsheet-style pivot table as a DataFrame.

The levels in the pivot table will be stored in MultiIndex objects (hierarchical indexes) on the index and columns of the result DataFrame.

#### Parameters

**values** : column to aggregate, optional

**index** : column, Grouper, array, or list of the previous

If an array is passed, it must be the same length as the data. The list can contain any of the other types (except list).

Keys to group by on the pivot table index. If an array is passed, it is being used as the same manner as column values.

**columns** : column, Grouper, array, or list of the previous

If an array is passed, it must be the same length as the data. The list can contain any of the other types (except list).

Keys to group by on the pivot table column. If an array is passed, it is being used as the same manner as column values.

**aggfunc** : function, list of functions, dict, default numpy.mean

If list of functions passed, the resulting pivot table will have hierarchical columns whose top level are the function names (inferred from the function objects themselves)

If dict is passed, the key is column to aggregate and value is function or list of functions.

**fill\_value** : scalar, default None

Value to replace missing values with (in the resulting pivot table, after aggregation).

**margins** : bool, default False

Add all row / columns (e.g. for subtotal / grand totals).

**dropna** : bool, default True

Do not include columns whose entries are all NaN.

**margins\_name** : str, default 'All'

Name of the row / column that will contain the totals when margins is True.

**observed** : bool, default False

This only applies if any of the groupers are Categoricals.

If True: only show observed values for categorical groupers.

If False: show all values for categorical groupers.

.. versionchanged:: 0.25.0

#### Returns

DataFrame

An Excel style pivot table.

#### See Also

**DataFrame.pivot** : Pivot without aggregation that can handle non-numeric data.

**DataFrame.melt**: Unpivot a DataFrame from wide to long format, optionally leaving identifiers set.

**wide\_to\_long** : Wide panel to long format. Less flexible but more user-friendly than melt.

#### Examples

```

>>> df = pd.DataFrame({"A": ["foo", "foo", "foo", "foo", "foo",
... "bar", "bar", "bar", "bar"],
... "B": ["one", "one", "one", "two", "two",
... "one", "one", "two", "two"],
... "C": ["small", "large", "large", "small",
... "small", "large", "small", "small",
... "large"],
... "D": [1, 2, 2, 3, 3, 4, 5, 6, 7],
... "E": [2, 4, 5, 5, 6, 6, 8, 9, 9]})
>>> df
 A B C D E
0 foo one small 1 2
1 foo one large 2 4
2 foo one large 2 5
3 foo two small 3 5
4 foo two small 3 6
5 bar one large 4 6
6 bar one small 5 8
7 bar two small 6 9
8 bar two large 7 9

```

This first example aggregates values by taking the sum.

```

>>> table = pd.pivot_table(df, values='D', index=['A', 'B'],
... columns=['C'], aggfunc=np.sum)
>>> table
C large small
A B
bar one 4.0 5.0
 two 7.0 6.0
foo one 4.0 1.0
 two NaN 6.0

```

We can also fill missing values using the `fill\_value` parameter.

```

>>> table = pd.pivot_table(df, values='D', index=['A', 'B'],
... columns=['C'], aggfunc=np.sum, fill_value=0)
>>> table
C large small
A B
bar one 4 5
 two 7 6
foo one 4 1
 two 0 6

```

The next example aggregates by taking the mean across multiple columns.

```

>>> table = pd.pivot_table(df, values=['D', 'E'], index=['A', 'C'],
... aggfunc={'D': np.mean,
... 'E': np.mean})
>>> table
 D E
A C
bar large 5.500000 7.500000
 small 5.500000 8.500000
foo large 2.000000 4.500000
 small 2.333333 4.333333

```

We can also calculate multiple types of aggregations for any given value column.

```

>>> table = pd.pivot_table(df, values=['D', 'E'], index=['A', 'C'],
... aggfunc={'D': np.mean,
... 'E': [min, max, np.mean]})
>>>

```

```
>>> table
 D E
 mean max mean min
A C
bar large 5.500000 9.0 7.500000 6.0
 small 5.500000 9.0 8.500000 8.0
foo large 2.000000 5.0 4.500000 4.0
 small 2.333333 6.0 4.333333 2.0
```

pop(self, item: 'Label') -> 'Series'  
Return item and drop from frame. Raise KeyError if not found.

Parameters  
-----  
item : label  
Label of column to be popped.

Returns  
-----  
Series

Examples  
-----

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0),
... ('parrot', 'bird', 24.0),
... ('lion', 'mammal', 80.5),
... ('monkey', 'mammal', np.nan)],
... columns=('name', 'class', 'max_speed'))
>>> df
 name class max_speed
0 falcon bird 389.0
1 parrot bird 24.0
2 lion mammal 80.5
3 monkey mammal NaN

>>> df.pop('class')
0 bird
1 bird
2 mammal
3 mammal
Name: class, dtype: object

>>> df
 name max_speed
0 falcon 389.0
1 parrot 24.0
2 lion 80.5
3 monkey NaN
```

pow(self, other, axis='columns', level=None, fill\_value=None)  
Get Exponential power of dataframe and other, element-wise (binary operator `pow`).

Equivalent to ``dataframe \*\* other``, but with support to substitute a fill\_value for missing data in one of the inputs. With reverse version, `rpow`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

Parameters  
-----  
other : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.  
axis : {0 or 'index', 1 or 'columns'}

Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

level : int or label  
Broadcast across a level, matching Index values on the passed MultiIndex level.

fill\_value : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

-----  
DataFrame

Result of the arithmetic operation.

#### See Also

-----  
DataFrame.add : Add DataFrames.  
DataFrame.sub : Subtract DataFrames.  
DataFrame.mul : Multiply DataFrames.  
DataFrame.div : Divide DataFrames (float division).  
DataFrame.truediv : Divide DataFrames (float division).  
DataFrame.floordiv : Divide DataFrames (integer division).  
DataFrame.mod : Calculate modulo (remainder after division).  
DataFrame.pow : Calculate exponential power.

#### Notes

-----  
Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
 angles degrees
circle inf 0.027778
triangle 3.333333 0.055556
rectangle 2.500000 0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
...
 axis='index'))
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...
 index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4
```

```
>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN
```

```
>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...
 'degrees': [360, 180, 360, 360, 540, 720]},
...
 index=[['A', 'A', 'A', 'B', 'B', 'B'],
...
 ['circle', 'triangle', 'rectangle',
...
 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
```

|          |   |     |
|----------|---|-----|
| B square | 4 | 360 |
| pentagon | 5 | 540 |
| hexagon  | 6 | 720 |

```
>>> df.div(df_multindex, level=1, fill_value=0)
```

|           | angles | degrees |
|-----------|--------|---------|
| A circle  | NaN    | 1.0     |
| triangle  | 1.0    | 1.0     |
| rectangle | 1.0    | 1.0     |
| B square  | 0.0    | 0.0     |
| pentagon  | 0.0    | 0.0     |
| hexagon   | 0.0    | 0.0     |

```
prod(self, axis=None, skipna=None, level=None, numeric_only=None, min_count=
0, **kwargs)
```

Return the product of the values over the requested axis.

Parameters

`axis` : {index (0), columns (1)}

Axis for the function to be applied on.

`skipna` : bool, default True

Exclude NA/null values when computing the result.

`level` : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

`numeric_only` : bool, default None

Include only float, int, boolean columns. If None, will attempt to u

se

everything, then use only numeric data. Not implemented for Series.

`min_count` : int, default 0

er than

The required number of valid values to perform the operation. If few

``min\_count`` non-NA values are present the result will be NA.

`**kwargs`

Additional keyword arguments to be passed to the function.

Returns

Series or DataFrame (if level specified)

See Also

`Series.sum` : Return the sum.

`Series.min` : Return the minimum.

`Series.max` : Return the maximum.

`Series.idxmin` : Return the index of the minimum.

`Series.idxmax` : Return the index of the maximum.

`DataFrame.sum` : Return the sum over the requested axis.

`DataFrame.min` : Return the minimum over the requested axis.

`DataFrame.max` : Return the maximum over the requested axis.

`DataFrame.idxmin` : Return the index of the minimum over the requested ax

is.

`DataFrame.idxmax` : Return the index of the maximum over the requested ax

is.

Examples

By default, the product of an empty or all-NA Series is ``1``

```
>>> pd.Series([]).prod()
1.0
```

This can be controlled with the ``min\_count`` parameter

```
>>> pd.Series([]).prod(min_count=1)
nan
```

Thanks to the ``skipna`` parameter, ``min\_count`` handles all-NA and empty series identically.

```
>>> pd.Series([np.nan]).prod()
1.0
```

```
>>> pd.Series([np.nan]).prod(min_count=1)
nan
```

```
product = prod(self, axis=None, skipna=None, level=None, numeric_only=None,
min_count=0, **kwargs)
```

```
quantile(self, q=0.5, axis=0, numeric_only=True, interpolation='linear')
Return values at the given quantile over requested axis.
```

#### Parameters

```
q : float or array-like, default 0.5 (50% quantile)
 Value between 0 <= q <= 1, the quantile(s) to compute.
axis : {0, 1, 'index', 'columns'}, default 0
 Equals 0 or 'index' for row-wise, 1 or 'columns' for column-wise.
numeric_only : bool, default True
 If False, the quantile of datetime and timedelta data will be
 computed as well.
interpolation : {'linear', 'lower', 'higher', 'midpoint', 'nearest'}
 This optional parameter specifies the interpolation method to use,
 when the desired quantile lies between two data points `i` and `j`:

 * linear: `i + (j - i) * fraction`, where `fraction` is the
 fractional part of the index surrounded by `i` and `j`.
 * lower: `i`.
 * higher: `j`.
 * nearest: `i` or `j` whichever is nearest.
 * midpoint: `(i + j) / 2.
```

#### Returns

##### Series or DataFrame

```
If ``q`` is an array, a DataFrame will be returned where the
index is ``q``, the columns are the columns of self, and the
values are the quantiles.
If ``q`` is a float, a Series will be returned where the
index is the columns of self and the values are the quantiles.
```

#### See Also

```
core.window.Rolling.quantile: Rolling quantile.
numpy.percentile: Numpy function to compute the percentile.
```

#### Examples

```
>>> df = pd.DataFrame(np.array([[1, 1], [2, 10], [3, 100], [4, 100]]),
... columns=['a', 'b'])
>>> df.quantile(.1)
a 1.3
b 3.7
Name: 0.1, dtype: float64
>>> df.quantile([.1, .5])
 a b
0.1 1.3 3.7
0.5 2.5 55.0
```

Specifying `numeric_only=False` will also compute the quantile of datetime and timedelta data.

```
>>> df = pd.DataFrame({'A': [1, 2],
... 'B': [pd.Timestamp('2010'),
... pd.Timestamp('2011')],
... 'C': [pd.Timedelta('1 days'),
... pd.Timedelta('2 days')]}))
>>> df.quantile(0.5, numeric_only=False)
A 1.5
B 2010-07-02 12:00:00
C 1 days 12:00:00
Name: 0.5, dtype: object
```

`query(self, expr, inplace=False, **kwargs)`  
Query the columns of a DataFrame with a boolean expression.

#### Parameters

`expr : str`

The query string to evaluate.

You can refer to variables in the environment by prefixing them with an '@' character like ```@a + b```.

You can refer to column names that are not valid Python variable names by surrounding them in backticks. Thus, column names containing spaces or punctuations (besides underscores) or starting with digits must be surrounded by backticks. (For example, a column named "Area (cm^2)" would be referenced as ```Area (cm^2)```). Column names which are Python keywords (like "list", "for", "import", etc) cannot be used.

For example, if one of your columns is called ```a a``` and you want to sum it with ```b```, your query should be ```a a + b```.

.. versionadded:: 0.25.0  
Backtick quoting introduced.

.. versionadded:: 1.0.0  
Expanding functionality of backtick quoting for more than only spaces.

`inplace : bool`

Whether the query should modify the data in place or return a modified copy.

`**kwargs`

See the documentation for `:func:`eval`` for complete details on the keyword arguments accepted by `:meth:`DataFrame.query``.

#### Returns

DataFrame or None

DataFrame resulting from the provided query expression or None if ```inplace=True```.

#### See Also

`eval` : Evaluate a string describing operations on



DataFrame columns.  
 DataFrame.eval : Evaluate a string describing operations on  
 DataFrame columns.

## Notes

The result of the evaluation of this expression is first passed to :attr:`DataFrame.loc` and if that fails because of a multidimensional key (e.g., a DataFrame) then the result will be passed to :meth:`DataFrame.\_\_getitem\_\_`.

This method uses the top-level :func:`eval` function to evaluate the passed query.

The :meth:`~pandas.DataFrame.query` method uses a slightly modified Python syntax by default. For example, the ``&`` and ``|`` (bitwise) operators have the precedence of their boolean cousins, :keyword:`and` and :keyword:`or`. This *is* syntactically valid Python, however the semantics are different.

You can change the semantics of the expression by passing the keyword argument ``parser='python'``. This enforces the same semantics as evaluation in Python space. Likewise, you can pass ``engine='python'`` to evaluate an expression using Python itself as a backend. This is not recommended as it is inefficient compared to using ``numexpr`` as the engine.

The :attr:`DataFrame.index` and :attr:`DataFrame.columns` attributes of the :class:`~pandas.DataFrame` instance are placed in the query namespace by default, which allows you to treat both the index and columns of the frame as a column in the frame.

The identifier ``index`` is used for the frame index; you can also use the name of the index to identify it in a query. Please note that Python keywords may not be used as identifiers.

For further details and examples see the ``query`` documentation in :ref:`indexing <indexing.query>`.

### \*Backtick quoted variables\*

Backtick quoted variables are parsed as literal Python code and are converted internally to a Python valid identifier. This can lead to the following problems.

During parsing a number of disallowed characters inside the backtick quoted string are replaced by strings that are allowed as a Python identifier.

These characters include all operators in Python, the space character, the

question mark, the exclamation mark, the dollar sign, and the euro sign. For other characters that fall outside the ASCII range (U+0001..U+007F) and those that are not further specified in PEP 3131, the query parser will raise an error.

This excludes whitespace different than the space character, but also the hashtag (as it is used for comments) and the backtick itself (backtick can also not be escaped).

In a special case, quotes that make a pair around a backtick can confuse the parser.

For example, ``'it's' > 'that's'`` will raise an error, as it forms a quoted string (``'s > 'that'``) with a backtick inside.

See also the Python documentation about lexical analysis ([https://docs.python.org/3/reference/lexical\\_analysis.html](https://docs.python.org/3/reference/lexical_analysis.html))

in combination with the source code in :mod:`pandas.core.computation.par  
sing`.

### Examples

```
>>> df = pd.DataFrame({'A': range(1, 6),
... 'B': range(10, 0, -2),
... 'C C': range(10, 5, -1)})
>>> df
 A B C C
0 1 10 10
1 2 8 9
2 3 6 8
3 4 4 7
4 5 2 6
>>> df.query('A > B')
 A B C C
4 5 2 6
```

The previous expression is equivalent to

```
>>> df[df.A > df.B]
 A B C C
4 5 2 6
```

For columns with spaces in their name, you can use backtick quoting.

```
>>> df.query('B == `C C`')
 A B C C
0 1 10 10
```

The previous expression is equivalent to

```
>>> df[df.B == df['C C']]
 A B C C
0 1 10 10
```

radd(self, other, axis='columns', level=None, fill\_value=None)  
Get Addition of dataframe and other, element-wise (binary operator `radd`  
)

Equivalent to ``other + dataframe``, but with support to substitute a fi  
ll\_value

for missing data in one of the inputs. With reverse version, `add`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to  
arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

### Parameters

other : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.  
axis : {0 or 'index', 1 or 'columns'}  
Whether to compare by the index (0 or 'index') or columns  
(1 or 'columns'). For Series input, axis to match Series index on.  
level : int or label  
Broadcast across a level, matching Index values on the  
passed MultiIndex level.  
fill\_value : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for  
successful DataFrame alignment, with this value before computation.  
If data in both corresponding DataFrame locations is missing  
the result will be missing.

### Returns

-----  
 DataFrame  
 Result of the arithmetic operation.

See Also  
 -----

DataFrame.add : Add DataFrames.  
 DataFrame.sub : Subtract DataFrames.  
 DataFrame.mul : Multiply DataFrames.  
 DataFrame.div : Divide DataFrames (float division).  
 DataFrame.truediv : Divide DataFrames (float division).  
 DataFrame.floordiv : Divide DataFrames (integer division).  
 DataFrame.mod : Calculate modulo (remainder after division).  
 DataFrame.pow : Calculate exponential power.

Notes  
 -----

Mismatched indices will be unioned together.

Examples  
 -----

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|          | angles | degrees |
|----------|--------|---------|
| circle   | -1     | 358     |
| triangle | 2      | 178     |

```

rectangle 3 358

>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
e')),
... axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359

Multiply a DataFrame of different shape with operator version.

>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4

>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 72
0]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
e',
... 'square', 'pentagon', 'hexagon
n']]])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

```

```

 rdiv = rtruediv(self, other, axis='columns', level=None, fill_value=None)
 reindex(self, labels=None, index=None, columns=None, axis=None, method=None,
copy=True, level=None, fill_value=nan, limit=None, tolerance=None)
 Conform Series/DataFrame to new index with optional filling logic.

```

Places NA/NaN in locations having no value in the previous index. A new object is produced unless the new index is equivalent to the current one and ``copy=False``.

#### Parameters

keywords for axes : array-like, optional

New labels / index to conform to, should be specified using keywords. Preferably an Index object to avoid duplicating data.

method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}

Method to use for filling holes in reindexed DataFrame.

Please note: this is only applicable to DataFrames/Series with a monotonically increasing/decreasing index.

\* None (default): don't fill gaps

\* pad / ffill: Propagate last valid observation forward to next valid.

\* backfill / bfill: Use next valid observation to fill gap.

\* nearest: Use nearest valid observations to fill gap.

copy : bool, default True

Return a new object, even if the passed indexes are the same.

level : int or name

Broadcast across a level, matching Index values on the passed MultiIndex level.

fill\_value : scalar, default np.NaN

Value to use for missing values. Defaults to NaN, but can be any "compatible" value.

limit : int, default None

Maximum number of consecutive elements to forward or backward fill.

tolerance : optional

Maximum distance between original and new labels for inexact matches. The values of the index at the matching locations must satisfy the equation ``abs(index[indexer] - target) <= tolerance``.

Tolerance may be a scalar value, which applies the same tolerance to all values, or list-like, which applies variable tolerance per element. List-like includes list, tuple, array, Series, and must be the same size as the index and its dtype must exactly match the index's type.

#### Returns

Series/DataFrame with changed index.

#### See Also

DataFrame.set\_index : Set row labels.

DataFrame.reset\_index : Remove row labels or move them to new columns.

DataFrame.reindex\_like : Change to same indices as other DataFrame.

#### Examples

``DataFrame.reindex`` supports two calling conventions

```
* ``(index=index_labels, columns=column_labels, ...)``
* ``(labels, axis={'index', 'columns'}, ...)``
```

We *highly* recommend using keyword arguments to clarify your intent.

Create a dataframe with some fictional data.

```
>>> index = ['Firefox', 'Chrome', 'Safari', 'IE10', 'Konqueror']
>>> df = pd.DataFrame({'http_status': [200, 200, 404, 404, 301],
... 'response_time': [0.04, 0.02, 0.07, 0.08, 1.0]},
... index=index)
>>> df
```

|           | http_status | response_time |
|-----------|-------------|---------------|
| Firefox   | 200         | 0.04          |
| Chrome    | 200         | 0.02          |
| Safari    | 404         | 0.07          |
| IE10      | 404         | 0.08          |
| Konqueror | 301         | 1.00          |

Create a new index and reindex the dataframe. By default values in the new index that do not have corresponding records in the dataframe are assigned ``NaN``.

```
>>> new_index = ['Safari', 'Iceweasel', 'Comodo Dragon', 'IE10',
... 'Chrome']
>>> df.reindex(new_index)
```

|               | http_status | response_time |
|---------------|-------------|---------------|
| Safari        | 404.0       | 0.07          |
| Iceweasel     | NaN         | NaN           |
| Comodo Dragon | NaN         | NaN           |
| IE10          | 404.0       | 0.08          |
| Chrome        | 200.0       | 0.02          |

We can fill in the missing values by passing a value to the keyword ``fill\_value``. Because the index is not monotonically increasing or decreasing, we cannot use arguments to the keyword ``method`` to fill the ``NaN`` values.

```
>>> df.reindex(new_index, fill_value=0)
```

|               | http_status | response_time |
|---------------|-------------|---------------|
| Safari        | 404         | 0.07          |
| Iceweasel     | 0           | 0.00          |
| Comodo Dragon | 0           | 0.00          |
| IE10          | 404         | 0.08          |
| Chrome        | 200         | 0.02          |

```
>>> df.reindex(new_index, fill_value='missing')
```

|               | http_status | response_time |
|---------------|-------------|---------------|
| Safari        | 404         | 0.07          |
| Iceweasel     | missing     | missing       |
| Comodo Dragon | missing     | missing       |
| IE10          | 404         | 0.08          |
| Chrome        | 200         | 0.02          |

We can also reindex the columns.

```
>>> df.reindex(columns=['http_status', 'user_agent'])
```

|           | http_status | user_agent |
|-----------|-------------|------------|
| Firefox   | 200         | NaN        |
| Chrome    | 200         | NaN        |
| Safari    | 404         | NaN        |
| IE10      | 404         | NaN        |
| Konqueror | 301         | NaN        |

Or we can use "axis-style" keyword arguments

```
>>> df.reindex(['http_status', 'user_agent'], axis="columns")
 http_status user_agent
Firefox 200 NaN
Chrome 200 NaN
Safari 404 NaN
IE10 404 NaN
Konqueror 301 NaN
```

To further illustrate the filling functionality in ``reindex``, we will create a dataframe with a monotonically increasing index (for example, a sequence of dates).

```
>>> date_index = pd.date_range('1/1/2010', periods=6, freq='D')
>>> df2 = pd.DataFrame({"prices": [100, 101, np.nan, 100, 89, 88]},
... index=date_index)
>>> df2
 prices
2010-01-01 100.0
2010-01-02 101.0
2010-01-03 NaN
2010-01-04 100.0
2010-01-05 89.0
2010-01-06 88.0
```

Suppose we decide to expand the dataframe to cover a wider date range.

```
>>> date_index2 = pd.date_range('12/29/2009', periods=10, freq='D')
>>> df2.reindex(date_index2)
 prices
2009-12-29 NaN
2009-12-30 NaN
2009-12-31 NaN
2010-01-01 100.0
2010-01-02 101.0
2010-01-03 NaN
2010-01-04 100.0
2010-01-05 89.0
2010-01-06 88.0
2010-01-07 NaN
```

The index entries that did not have a value in the original data frame (for example, '2009-12-29') are by default filled with ``NaN``. If desired, we can fill in the missing values using one of several options.

For example, to back-propagate the last valid value to fill the ``NaN`` values, pass ``bfill`` as an argument to the ``method`` keyword.

```
>>> df2.reindex(date_index2, method='bfill')
 prices
2009-12-29 100.0
2009-12-30 100.0
2009-12-31 100.0
2010-01-01 100.0
2010-01-02 101.0
2010-01-03 NaN
2010-01-04 100.0
2010-01-05 89.0
2010-01-06 88.0
2010-01-07 NaN
```

Please note that the ``NaN`` value present in the original dataframe (at index value 2010-01-03) will not be filled by any of the value propagation schemes. This is because filling while reindexing does not look at dataframe values, but only compares the original and desired indexes. If you do want to fill in the ``NaN`` values present in the original dataframe, use the ``fillna()`` method.

See the :ref:`user guide <basics.reindexing>` for more.

```
rename(self, mapper=None, index=None, columns=None, axis=None, copy=True, inplace=False, level=None, errors='ignore')
```

Alter axes labels.

Function / dict values must be unique (1-to-1). Labels not contained in a dict / Series will be left as-is. Extra labels listed don't throw an error.

See the :ref:`user guide <basics.rename>` for more.

#### Parameters

**mapper** : dict-like or function

Dict-like or function transformations to apply to that axis' values. Use either ``mapper`` and ``axis`` to specify the axis to target with ``mapper``, or ``index`` and ``columns``.

**index** : dict-like or function

Alternative to specifying axis (``mapper, axis=0`` is equivalent to ``index=mapper``).

**columns** : dict-like or function

Alternative to specifying axis (``mapper, axis=1`` is equivalent to ``columns=mapper``).

**axis** : {0 or 'index', 1 or 'columns'}, default 0

Axis to target with ``mapper``. Can be either the axis name ('index', 'columns') or number (0, 1). The default is 'index'.

**copy** : bool, default True

Also copy underlying data.

**inplace** : bool, default False

Whether to return a new DataFrame. If True then value of copy is ignored.

**level** : int or level name, default None

In case of a MultiIndex, only rename labels in the specified level.

**errors** : {'ignore', 'raise'}, default 'ignore'

If 'raise', raise a `KeyError` when a dict-like `mapper`, `index`, or `columns` contains labels that are not present in the Index being transformed.

If 'ignore', existing keys will be renamed and extra keys will be ignored.

#### Returns

DataFrame or None

DataFrame with the renamed axis labels or None if ``inplace=True``.

#### Raises

KeyError

If any of the labels is not found in the selected axis and "errors='raise'".

#### See Also

DataFrame.rename\_axis : Set the name of the axis.



## Examples

`DataFrame.rename` supports two calling conventions

```
* `(index=index_mapper, columns=columns_mapper, ...)`
* `(mapper, axis={'index', 'columns'}, ...)`
```

We *highly* recommend using keyword arguments to clarify your intent.

Rename columns using a mapping:

```
>>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
>>> df.rename(columns={"A": "a", "B": "c"})
 a c
0 1 4
1 2 5
2 3 6
```

Rename index using a mapping:

```
>>> df.rename(index={0: "x", 1: "y", 2: "z"})
 A B
x 1 4
y 2 5
z 3 6
```

Cast index labels to a different type:

```
>>> df.index
RangeIndex(start=0, stop=3, step=1)
>>> df.rename(index=str).index
Index(['0', '1', '2'], dtype='object')
```

```
>>> df.rename(columns={"A": "a", "B": "b", "C": "c"}, errors="raise")
Traceback (most recent call last):
KeyError: ['C'] not found in axis
```

Using axis-style parameters:

```
>>> df.rename(str.lower, axis='columns')
 a b
0 1 4
1 2 5
2 3 6

>>> df.rename({1: 2, 2: 4}, axis='index')
 A B
0 1 4
2 2 5
4 3 6
```

`reorder_levels(self, order, axis=0) -> 'DataFrame'`  
Rearrange index levels using input order. May not drop or duplicate level

ls.

## Parameters

**order** : list of int or list of str  
List representing new level order. Reference level by number (position) or by key (label).

**axis** : {0 or 'index', 1 or 'columns'}, default 0  
Where to reorder levels.

## Returns

## ----- DataFrame

```
replace(self, to_replace=None, value=None, inplace=False, limit=None, regex=False, method='pad')
```

Replace values given in `to\_replace` with `value`.

Values of the DataFrame are replaced with other values dynamically. This differs from updating with ``.loc`` or ``.iloc``, which require you to specify a location to update with some value.

## Parameters

`to_replace` : str, regex, list, dict, Series, int, float, or None  
How to find the values that will be replaced.

### \* numeric, str or regex:

- numeric: numeric values equal to `to\_replace` will be replaced with `value`
- str: string exactly matching `to\_replace` will be replaced with `value`
- regex: regexs matching `to\_replace` will be replaced with `value`

### \* list of str, regex, or numeric:

- First, if `to\_replace` and `value` are both lists, they **must** be the same length.
- Second, if ``regex=True`` then all of the strings in **both** lists will be interpreted as regexs otherwise they will match directly. This doesn't matter much for `value` since there are only a few possible substitution regexes you can use.
- str, regex and numeric rules apply as above.

### \* dict:

- Dicts can be used to specify different replacement values for different existing values. For example, ``{'a': 'b', 'y': 'z'}`` replaces the value 'a' with 'b' and 'y' with 'z'. To use a dict in this way the `value` parameter should be `None`.
- For a DataFrame a dict can specify that different values should be replaced in different columns. For example, ``{'a': 1, 'b': 'z'}`` looks for the value 1 in column 'a' and the value 'z' in column 'b' and replaces these values with whatever is specified in `value`. The `value` parameter should not be `None` in this case. You can treat this as a special case of passing two lists except that you are specifying the column to search in.
- For a DataFrame nested dictionaries, e.g., ``{'a': {'b': np.nan}}``, are read as follows: look in column 'a' for the value 'b' and replace it with NaN. The `value` parameter should be `None` to use a nested dict in this way. You can nest regular expressions as well. Note that column names (the top-level dictionary keys in a nested dictionary) **cannot** be regular expressions.

### \* None:

- This means that the `regex` argument must be a string, compiled regular expression, or list, dict, ndarray or Series of such elements. If `value` is also `None` then this **must** be a nested dictionary or Series.

See the examples section for examples of each of these.

`value` : scalar, dict, list, str, regex, default None  
 Value to replace any values matching ``to_replace`` with.  
 For a DataFrame a dict of values can be used to specify which value to use for each column (columns not in the dict will not be filled). Regular expressions, strings and lists or dicts of such objects are also allowed.

`inplace` : bool, default False  
 If True, in place. Note: this will modify any other views on this object (e.g. a column from a DataFrame). Returns the caller if this is True.

`limit` : int or None, default None  
 Maximum size gap to forward or backward fill.

`regex` : bool or same types as ``to_replace``, default False  
 Whether to interpret ``to_replace`` and/or ``value`` as regular expressions. If this is ``True`` then ``to_replace`` *must* be a string. Alternatively, this could be a regular expression or a list, dict, or array of regular expressions in which case ``to_replace`` must be ``None``.

`method` : {'pad', 'ffill', 'bfill', 'None'}  
 The method to use when for replacement, when ``to_replace`` is a scalar, list or tuple and ``value`` is ``None``.

### Returns

-----  
 DataFrame or None

Object after replacement or None if ``inplace=True``.

### Raises

-----  
 AssertionError

\* If ``regex`` is not a ``bool`` and ``to_replace`` is not ``None``.

TypeError

\* If ``to_replace`` is not a scalar, array-like, ``dict``, or ``None``  
 \* If ``to_replace`` is a ``dict`` and ``value`` is not a ``list``, ``dict``, ``ndarray``, or ``Series``  
 \* If ``to_replace`` is ``None`` and ``regex`` is not compilable into a regular expression or is a list, dict, ndarray, or Series.  
 \* When replacing multiple ``bool`` or ``datetime64`` objects and the arguments to ``to_replace`` does not match the type of the value being replaced

ValueError

\* If a ``list`` or an ``ndarray`` is passed to ``to_replace`` and ``value`` but they are not the same length.

### See Also

-----  
 DataFrame.fillna : Fill NA values.

DataFrame.where : Replace values based on boolean condition.

Series.str.replace : Simple string replacement.

### Notes

-----  
 \* Regex substitution is performed under the hood with ``re.sub``. The rules for substitution for ``re.sub`` are the same.  
 \* Regular expressions will only substitute on strings, meaning you cannot provide, for example, a regular expression matching floating point numbers and expect the columns in your frame that have a numeric dtype to be matched. However, if those floating point numbers *are* strings, then you can do this.  
 \* This method has *a lot* of options. You are encouraged to experiment

and play with this method to gain intuition about how it works.  
 \* When dict is used as the `to\_replace` value, it is like  
 key(s) in the dict are the to\_replace part and  
 value(s) in the dict are the value parameter.

### Examples

**\*\*Scalar `to\_replace` and `value`\*\***

```
>>> s = pd.Series([0, 1, 2, 3, 4])
>>> s.replace(0, 5)
0 5
1 1
2 2
3 3
4 4
dtype: int64

>>> df = pd.DataFrame({'A': [0, 1, 2, 3, 4],
... 'B': [5, 6, 7, 8, 9],
... 'C': ['a', 'b', 'c', 'd', 'e']})
>>> df.replace(0, 5)
 A B C
0 5 5 a
1 1 6 b
2 2 7 c
3 3 8 d
4 4 9 e
```

**\*\*List-like `to\_replace`\*\***

```
>>> df.replace([0, 1, 2, 3], 4)
 A B C
0 4 5 a
1 4 6 b
2 4 7 c
3 4 8 d
4 4 9 e

>>> df.replace([0, 1, 2, 3], [4, 3, 2, 1])
 A B C
0 4 5 a
1 3 6 b
2 2 7 c
3 1 8 d
4 4 9 e
```

```
>>> s.replace([1, 2], method='bfill')
0 0
1 3
2 3
3 3
4 4
dtype: int64
```

**\*\*dict-like `to\_replace`\*\***

```
>>> df.replace({0: 10, 1: 100})
 A B C
0 10 5 a
1 100 6 b
2 2 7 c
3 3 8 d
4 4 9 e
```

```

>>> df.replace({'A': 0, 'B': 5}, 100)
 A B C
0 100 100 a
1 1 6 b
2 2 7 c
3 3 8 d
4 4 9 e

>>> df.replace({'A': {0: 100, 4: 400}})
 A B C
0 100 5 a
1 1 6 b
2 2 7 c
3 3 8 d
4 400 9 e

Regular expression `to_replace`

>>> df = pd.DataFrame({'A': ['bat', 'foo', 'bait'],
... 'B': ['abc', 'bar', 'xyz']})
>>> df.replace(to_replace=r'^ba.$', value='new', regex=True)
 A B
0 new abc
1 foo new
2 bait xyz

>>> df.replace({'A': r'^ba.$'}, {'A': 'new'}, regex=True)
 A B
0 new abc
1 foo bar
2 bait xyz

>>> df.replace(regex=r'^ba.$', value='new')
 A B
0 new abc
1 foo new
2 bait xyz

>>> df.replace(regex={r'^ba.$': 'new', 'foo': 'xyz'})
 A B
0 new abc
1 xyz new
2 bait xyz

>>> df.replace(regex=[r'^ba.$', 'foo'], value='new')
 A B
0 new abc
1 new new
2 bait xyz

Compare the behavior of ``s.replace({'a': None})`` and
``s.replace('a', None)`` to understand the peculiarities
of the `to_replace` parameter:

>>> s = pd.Series([10, 'a', 'a', 'b', 'a'])

When one uses a dict as the `to_replace` value, it is like the
value(s) in the dict are equal to the `value` parameter.
``s.replace({'a': None})`` is equivalent to
``s.replace(to_replace={'a': None}, value=None, method=None)``:

>>> s.replace({'a': None})
0 10
1 None

```

```

2 None
3 b
4 None
dtype: object

```

When ``value=None`` and `to\_replace` is a scalar, list or tuple, `replace` uses the method parameter (default 'pad') to do the replacement. So this is why the 'a' values are being replaced by 10 in rows 1 and 2 and 'b' in row 4 in this case.

The command ``s.replace('a', None)`` is actually equivalent to ``s.replace(to\_replace='a', value=None, method='pad')``:

```

>>> s.replace('a', None)
0 10
1 10
2 10
3 b
4 b
dtype: object

```

`reset_index(self, level: 'Optional[Union[Hashable, Sequence[Hashable]]]' = None, drop: 'bool' = False, inplace: 'bool' = False, col_level: 'Hashable' = 0, col_fill: 'Label' = '') -> 'Optional[DataFrame]'`

Reset the index, or a level of it.

Reset the index of the DataFrame, and use the default one instead. If the DataFrame has a MultiIndex, this method can remove one or more levels.

#### Parameters

`level` : int, str, tuple, or list, default None  
Only remove the given levels from the index. Removes all levels by default.

`drop` : bool, default False  
Do not try to insert index into dataframe columns. This resets the index to the default integer index.

`inplace` : bool, default False  
Modify the DataFrame in place (do not create a new object).

`col_level` : int or str, default 0  
If the columns have multiple levels, determines which level the labels are inserted into. By default it is inserted into the first level.

`col_fill` : object, default ''  
If the columns have multiple levels, determines how the other levels are named. If None then the index name is repeated.

#### Returns

DataFrame or None  
DataFrame with the new index or None if ``inplace=True``.

#### See Also

`DataFrame.set_index` : Opposite of `reset_index`.  
`DataFrame.reindex` : Change to new indices or expand indices.  
`DataFrame.reindex_like` : Change to same indices as other DataFrame.

#### Examples

```

>>> df = pd.DataFrame([('bird', 389.0),
... ('bird', 24.0),
... ('mammal', 80.5),
... ('mammal', np.nan)],
... index=['falcon', 'parrot', 'lion', 'monkey'],

```

```

... columns=('class', 'max_speed'))
>>> df
 class max_speed
falcon bird 389.0
parrot bird 24.0
lion mammal 80.5
monkey mammal NaN

```

When we reset the index, the old index is added as a column, and a new sequential index is used:

```

>>> df.reset_index()
 index class max_speed
0 falcon bird 389.0
1 parrot bird 24.0
2 lion mammal 80.5
3 monkey mammal NaN

```

We can use the `drop` parameter to avoid the old index being added as a column:

```

>>> df.reset_index(drop=True)
 class max_speed
0 bird 389.0
1 bird 24.0
2 mammal 80.5
3 mammal NaN

```

You can also use `reset\_index` with `MultiIndex`.

```

>>> index = pd.MultiIndex.from_tuples([('bird', 'falcon'),
... ('bird', 'parrot'),
... ('mammal', 'lion'),
... ('mammal', 'monkey')],
... names=['class', 'name'])
>>> columns = pd.MultiIndex.from_tuples([('speed', 'max'),
... ('species', 'type')])
>>> df = pd.DataFrame([(389.0, 'fly'),
... (24.0, 'fly'),
... (80.5, 'run'),
... (np.nan, 'jump')],
... index=index,
... columns=columns)
>>> df

```

```

 speed species
 max type
class name
bird falcon 389.0 fly
 parrot 24.0 fly
mammal lion 80.5 run
 monkey NaN jump

```

If the index has multiple levels, we can reset a subset of them:

```

>>> df.reset_index(level='class')
 class speed species
 max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump

```

If we are not dropping the index, by default, it is placed in the top level. We can place it in another level:

```
>>> df.reset_index(level='class', col_level=1)
 speed species
 class max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

When the index is inserted under another level, we can specify under which one with the parameter `col\_fill`:

```
>>> df.reset_index(level='class', col_level=1, col_fill='species')
 species speed species
 class max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

If we specify a nonexistent level for `col\_fill`, it is created:

```
>>> df.reset_index(level='class', col_level=1, col_fill='genus')
 genus speed species
 class max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

`rfloordiv(self, other, axis='columns', level=None, fill_value=None)`  
Get Integer division of dataframe and other, element-wise (binary operation `rfloordiv`).

Equivalent to ``other // dataframe``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, `floordiv`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

#### Parameters

`other` : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.

`axis` : {0 or 'index', 1 or 'columns'}  
Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

`level` : int or label  
Broadcast across a level, matching Index values on the passed MultiIndex level.

`fill_value` : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

DataFrame  
Result of the arithmetic operation.



## See Also

-----  
 DataFrame.add : Add DataFrames.  
 DataFrame.sub : Subtract DataFrames.  
 DataFrame.mul : Multiply DataFrames.  
 DataFrame.div : Divide DataFrames (float division).  
 DataFrame.truediv : Divide DataFrames (float division).  
 DataFrame.floordiv : Divide DataFrames (integer division).  
 DataFrame.mod : Calculate modulo (remainder after division).  
 DataFrame.pow : Calculate exponential power.

## Notes

-----  
 Mismatched indices will be unioned together.

## Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub([1, 2], axis='columns')
```

|  | angles | degrees |
|--|--------|---------|
|--|--------|---------|

```

circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
e')),
... axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359

Multiply a DataFrame of different shape with operator version.

>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4

>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
0}},
... 'degrees': [360, 180, 360, 360, 540, 72
e',
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
n']]))
... 'square', 'pentagon', 'hexagon'])

>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

rmod(self, other, axis='columns', level=None, fill_value=None)
Get Modulo of dataframe and other, element-wise (binary operator `rmod
`).
```

```

ll_value Equivalent to ``other % dataframe``, but with support to substitute a fi
for missing data in one of the inputs. With reverse version, ``mod``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to
arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same
results.

>>> df + 1
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

>>> df.add(1)

```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub([1, 2], axis='columns')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
axis='index')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 359     |
| triangle  | 2      | 179     |
| rectangle | 3      | 359     |

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
index=['circle', 'triangle', 'rectangle'])
```

```
>>> other
```

|           | angles |
|-----------|--------|
| circle    | 0      |
| triangle  | 3      |
| rectangle | 4      |

```
>>> df * other
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | NaN     |
| triangle  | 9      | NaN     |
| rectangle | 16     | NaN     |

```
>>> df.mul(other, fill_value=0)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 0.0     |
| triangle  | 9      | 0.0     |
| rectangle | 16     | 0.0     |

Divide by a MultiIndex by level.

```

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720],
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']})

```

```

>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

```

```

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

```

`rmul(self, other, axis='columns', level=None, fill_value=None)`

Get Multiplication of dataframe and other, element-wise (binary operator ``rmul``).

Equivalent to ``other * dataframe``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, ``mul``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

#### Parameters

`other` : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

`axis` : {0 or 'index', 1 or 'columns'}

Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

`level` : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

`fill_value` : float or None, default None

Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

DataFrame

Result of the arithmetic operation.

#### See Also

`DataFrame.add` : Add DataFrames.

`DataFrame.sub` : Subtract DataFrames.

`DataFrame.mul` : Multiply DataFrames.

`DataFrame.div` : Divide DataFrames (float division).

`DataFrame.truediv` : Divide DataFrames (float division).

DataFrame.floordiv : Divide DataFrames (integer division).  
 DataFrame.mod : Calculate modulo (remainder after division).  
 DataFrame.pow : Calculate exponential power.

## Notes

-----  
 Mismatched indices will be unioned together.

## Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
```

```
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360
```

Add a scalar with operator version which return the same results.

```
>>> df + 1
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361
```

```
>>> df.add(1)
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361
```

Divide by constant with reverse version.

```
>>> df.div(10)
 angles degrees
circle 0.0 36.0
triangle 0.3 18.0
rectangle 0.4 36.0
```

```
>>> df.rdiv(10)
 angles degrees
circle inf 0.027778
triangle 3.333333 0.055556
rectangle 2.500000 0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
... axis='index'))
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 359     |
| triangle  | 2      | 179     |
| rectangle | 3      | 359     |

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

|           | angles |
|-----------|--------|
| circle    | 0      |
| triangle  | 3      |
| rectangle | 4      |

```
>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN
```

```
>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
```

```
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0
```

`round(self, decimals=0, *args, **kwargs) -> 'DataFrame'`  
 Round a DataFrame to a variable number of decimal places.

Parameters

`decimals : int, dict, Series`

Number of decimal places to round each column to. If an int is given, round each column to the same number of places. Otherwise dict and Series round to variable numbers of places. Column names should be in the keys if `decimals` is a

dict-like, or in the index if `decimals` is a Series. Any columns not included in `decimals` will be left as is. Elements of `decimals` which are not columns of the input will be ignored.

#### \*args

Additional keywords have no effect but might be accepted for compatibility with numpy.

#### \*\*kwargs

Additional keywords have no effect but might be accepted for compatibility with numpy.

#### Returns

##### DataFrame

A DataFrame with the affected columns rounded to the specified number of decimal places.

#### See Also

`numpy.around` : Round a numpy array to the given number of decimals.  
`Series.round` : Round a Series to the given number of decimals.

#### Examples

```
>>> df = pd.DataFrame([(0.21, 0.32), (0.01, 0.67), (0.66, 0.03), (0.21, 0.18)],
... columns=['dogs', 'cats'])
>>> df
 dogs cats
0 0.21 0.32
1 0.01 0.67
2 0.66 0.03
3 0.21 0.18
```

By providing an integer each column is rounded to the same number of decimal places

```
>>> df.round(1)
 dogs cats
0 0.2 0.3
1 0.0 0.7
2 0.7 0.0
3 0.2 0.2
```

With a dict, the number of places for specific columns can be specified with the column names as key and the number of decimal places as value

```
>>> df.round({'dogs': 1, 'cats': 0})
 dogs cats
0 0.2 0.0
1 0.0 1.0
2 0.7 0.0
3 0.2 0.0
```

Using a Series, the number of places for specific columns can be specified with the column names as index and the number of decimal places as value

```
>>> decimals = pd.Series([0, 1], index=['cats', 'dogs'])
>>> df.round(decimals)
 dogs cats
0 0.2 0.0
1 0.0 1.0
2 0.7 0.0
3 0.2 0.0
```



```

 rpow(self, other, axis='columns', level=None, fill_value=None)
 Get Exponential power of dataframe and other, element-wise (binary opera
tor `rpow`).

 Equivalent to ``other ** dataframe``, but with support to substitute a f
ill_value
 for missing data in one of the inputs. With reverse version, `pow`.

 Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
 arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

 Parameters

 other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
 axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
 level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
 fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

 Returns

 DataFrame
 Result of the arithmetic operation.

 See Also

 DataFrame.add : Add DataFrames.
 DataFrame.sub : Subtract DataFrames.
 DataFrame.mul : Multiply DataFrames.
 DataFrame.div : Divide DataFrames (float division).
 DataFrame.truediv : Divide DataFrames (float division).
 DataFrame.floordiv : Divide DataFrames (integer division).
 DataFrame.mod : Calculate modulo (remainder after division).
 DataFrame.pow : Calculate exponential power.

 Notes

 Mismatched indices will be unioned together.

 Examples

 >>> df = pd.DataFrame({'angles': [0, 3, 4],
 ... 'degrees': [360, 180, 360]},
 ... index=['circle', 'triangle', 'rectangle'])
 >>> df
 angles degrees
 circle 0 360
 triangle 3 180
 rectangle 4 360

 Add a scalar with operator version which return the same
 results.

 >>> df + 1
 angles degrees
 circle 1 361

```

|           |   |     |
|-----------|---|-----|
| triangle  | 4 | 181 |
| rectangle | 5 | 361 |

```
>>> df.add(1)
```

|           |        |         |
|-----------|--------|---------|
|           | angles | degrees |
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           |        |         |
|-----------|--------|---------|
|           | angles | degrees |
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           |          |          |
|-----------|----------|----------|
|           | angles   | degrees  |
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|           |        |         |
|-----------|--------|---------|
|           | angles | degrees |
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub([1, 2], axis='columns')
```

|           |        |         |
|-----------|--------|---------|
|           | angles | degrees |
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
...
... axis='index'))
```

|           |        |         |
|-----------|--------|---------|
|           | angles | degrees |
| circle    | -1     | 359     |
| triangle  | 2      | 179     |
| rectangle | 3      | 359     |

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

|           |        |
|-----------|--------|
|           | angles |
| circle    | 0      |
| triangle  | 3      |
| rectangle | 4      |

```
>>> df * other
```

|           |        |         |
|-----------|--------|---------|
|           | angles | degrees |
| circle    | 0      | NaN     |
| triangle  | 9      | NaN     |
| rectangle | 16     | NaN     |

```
>>> df.mul(other, fill_value=0)
```

|          |        |         |
|----------|--------|---------|
|          | angles | degrees |
| circle   | 0      | 0.0     |
| triangle | 9      | 0.0     |

```

rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])

>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

rsub(self, other, axis='columns', level=None, fill_value=None)
Get Subtraction of dataframe and other, element-wise (binary operator `rsub`).

Equivalent to ``other - dataframe``, but with support to substitute a fill_value
for missing data in one of the inputs. With reverse version, `sub`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.

```

DataFrame.sub : Subtract DataFrames.  
 DataFrame.mul : Multiply DataFrames.  
 DataFrame.div : Divide DataFrames (float division).  
 DataFrame.truediv : Divide DataFrames (float division).  
 DataFrame.floordiv : Divide DataFrames (integer division).  
 DataFrame.mod : Calculate modulo (remainder after division).  
 DataFrame.pow : Calculate exponential power.

#### Notes

-----  
 Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
```

```
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360
```

Add a scalar with operator version which return the same results.

```
>>> df + 1
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361
```

```
>>> df.add(1)
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361
```

Divide by constant with reverse version.

```
>>> df.div(10)
 angles degrees
circle 0.0 36.0
triangle 0.3 18.0
rectangle 0.4 36.0
```

```
>>> df.rdiv(10)
 angles degrees
circle inf 0.027778
triangle 3.333333 0.055556
rectangle 2.500000 0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```

e]],
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
... axis='index'))
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359

Multiply a DataFrame of different shape with operator version.

>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4

>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df_multindex.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

rtruediv(self, other, axis='columns', level=None, fill_value=None)
Get Floating division of dataframe and other, element-wise (binary operator `rtruediv`).

Equivalent to ``other / dataframe``, but with support to substitute a fill_value

```

for missing data in one of the inputs. With reverse version, `truediv`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

#### Parameters

**other** : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.

**axis** : {0 or 'index', 1 or 'columns'}  
Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

**level** : int or label  
Broadcast across a level, matching Index values on the passed MultiIndex level.

**fill\_value** : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

**DataFrame**  
Result of the arithmetic operation.

#### See Also

**DataFrame.add** : Add DataFrames.  
**DataFrame.sub** : Subtract DataFrames.  
**DataFrame.mul** : Multiply DataFrames.  
**DataFrame.div** : Divide DataFrames (float division).  
**DataFrame.truediv** : Divide DataFrames (float division).  
**DataFrame.floordiv** : Divide DataFrames (integer division).  
**DataFrame.mod** : Calculate modulo (remainder after division).  
**DataFrame.pow** : Calculate exponential power.

#### Notes

Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|          | angles | degrees |
|----------|--------|---------|
| circle   | 1      | 361     |
| triangle | 4      | 181     |

```
rectangle 5 361
```

Divide by constant with reverse version.

```
>>> df.div(10)
 angles degrees
circle 0.0 36.0
triangle 0.3 18.0
rectangle 0.4 36.0
```

```
>>> df.rdiv(10)
 angles degrees
circle inf 0.027778
triangle 3.333333 0.055556
rectangle 2.500000 0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
e']),
... axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4
```

```
>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN
```

```
>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 72
0]}),
```

```
e',
n']]
...
index=[['A', 'A', 'A', 'B', 'B', 'B'],
 ['circle', 'triangle', 'rectangle',
 'square', 'pentagon', 'hexagon']]
...

```

```
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0
```

`select_dtypes(self, include=None, exclude=None) -> 'DataFrame'`  
 Return a subset of the DataFrame's columns based on the column dtypes.

#### Parameters

`include, exclude` : scalar or list-like  
 A selection of dtypes or strings to be included/excluded. At least one of these parameters must be supplied.

#### Returns

##### DataFrame

The subset of the frame including the dtypes in ``include`` and excluding the dtypes in ``exclude``.

#### Raises

##### ValueError

- \* If both of ``include`` and ``exclude`` are empty
- \* If ``include`` and ``exclude`` have overlapping elements
- \* If any kind of string dtype is passed in.

#### See Also

`DataFrame.dtypes`: Return Series with the data type of each column.

#### Notes

- \* To select all *numeric* types, use ``np.number`` or ``'number'``
- \* To select strings you must use the ``object`` dtype, but note that this will return *all* object dtype columns
- \* See the *numpy* dtype hierarchy  
<https://numpy.org/doc/stable/reference/arrays.scalars.html>
- \* To select datetimes, use ``np.datetime64``, ``'datetime'`` or ``'datetime64'``
- \* To select timedeltas, use ``np.timedelta64``, ``'timedelta'`` or ``'timedelta64'``
- \* To select Pandas categorical dtypes, use ``'category'``
- \* To select Pandas datetimetz dtypes, use ``'datetime64[ns, tz]'`` (new in 0.20.0) or ``'datetime64[ns, tz]'``

#### Examples



```

>>> df = pd.DataFrame({'a': [1, 2] * 3,
... 'b': [True, False] * 3,
... 'c': [1.0, 2.0] * 3})
>>> df
 a b c
0 1 True 1.0
1 2 False 2.0
2 1 True 1.0
3 2 False 2.0
4 1 True 1.0
5 2 False 2.0

>>> df.select_dtypes(include='bool')
 b
0 True
1 False
2 True
3 False
4 True
5 False

>>> df.select_dtypes(include=['float64'])
 c
0 1.0
1 2.0
2 1.0
3 2.0
4 1.0
5 2.0

>>> df.select_dtypes(exclude=['int64'])
 b c
0 True 1.0
1 False 2.0
2 True 1.0
3 False 2.0
4 True 1.0
5 False 2.0

```

sem(self, axis=None, skipna=None, level=None, ddof=1, numeric\_only=None, \*\*kwargs)

Return unbiased standard error of the mean over requested axis.

Normalized by N-1 by default. This can be changed using the ddof argument

#### Parameters

axis : {index (0), columns (1)}

skipna : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

ddof : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is N - ddof,

where N represents the number of elements.

numeric\_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use

everything, then use only numeric data. Not implemented for Series.

## Returns

Series or DataFrame (if level specified)

## Notes

To have the same behaviour as ``numpy.std``, use ``ddof=0`` (instead of the default ``ddof=1``)

```
set_axis(self, labels, axis: 'Axis' = 0, inplace: 'bool' = False)
```

Assign desired index to given axis.

Indexes for column or row labels can be changed by assigning a list-like or Index.

## Parameters

`labels` : list-like, Index  
The values for the new index.

`axis` : {0 or 'index', 1 or 'columns'}, default 0  
The axis to update. The value 0 identifies the rows, and 1 identifies the columns.

`inplace` : bool, default False  
Whether to return a new DataFrame instance.

## Returns

`renamed` : DataFrame or None  
An object of type DataFrame or None if ``inplace=True``.

## See Also

`DataFrame.rename_axis` : Alter the name of the index or columns.

## Examples

```
>>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
```

Change the row labels.

```
>>> df.set_axis(['a', 'b', 'c'], axis='index')
```

```
 A B
a 1 4
b 2 5
c 3 6
```

Change the column labels.

```
>>> df.set_axis(['I', 'II'], axis='columns')
```

```
 I II
0 1 4
1 2 5
2 3 6
```

Now, update the labels inplace.

```
>>> df.set_axis(['i', 'ii'], axis='columns', inplace=True)
```

```
>>> df
 i ii
0 1 4
1 2 5
2 3 6
```

```
| set_index(self, keys, drop=True, append=False, inplace=False, verify_integrity=False)
```

Set the DataFrame index using existing columns.

Set the DataFrame index (row labels) using one or more existing columns or arrays (of the correct length). The index can replace the existing index or expand on it.

#### Parameters

**keys** : label or array-like or list of labels/arrays

This parameter can be either a single column key, a single array of the same length as the calling DataFrame, or a list containing an arbitrary combination of column keys and arrays. Here, "array" encompasses `:class:`Series``, `:class:`Index``, ``np.ndarray``, and instances of `:class:`~collections.abc.Iterator``.

**drop** : bool, default True

Delete columns to be used as the new index.

**append** : bool, default False

Whether to append columns to existing index.

**inplace** : bool, default False

If True, modifies the DataFrame in place (do not create a new object).

**verify\_integrity** : bool, default False

Check the new index for duplicates. Otherwise defer the check until necessary. Setting to False will improve the performance of this method.

#### Returns

DataFrame or None

Changed row labels or None if ``inplace=True``.

#### See Also

**DataFrame.reset\_index** : Opposite of `set_index`.

**DataFrame.reindex** : Change to new indices or expand indices.

**DataFrame.reindex\_like** : Change to same indices as other DataFrame.

#### Examples

```
>>> df = pd.DataFrame({'month': [1, 4, 7, 10],
... 'year': [2012, 2014, 2013, 2014],
... 'sale': [55, 40, 84, 31]})
```

```
>>> df
 month year sale
0 1 2012 55
1 4 2014 40
2 7 2013 84
3 10 2014 31
```

Set the index to become the 'month' column:

```
>>> df.set_index('month')
 year sale
month
1 2012 55
4 2014 40
7 2013 84
10 2014 31
```

Create a MultiIndex using columns 'year' and 'month':

```
>>> df.set_index(['year', 'month'])
 sale
year month
2012 1 55
2014 4 40
2013 7 84
2014 10 31
```

| year | month | sale |
|------|-------|------|
| 2012 | 1     | 55   |
| 2014 | 4     | 40   |
| 2013 | 7     | 84   |
| 2014 | 10    | 31   |

Create a MultiIndex using an Index and a column:

```
>>> df.set_index([pd.Index([1, 2, 3, 4]), 'year'])
```

|   | year | month | sale |
|---|------|-------|------|
| 1 | 2012 | 1     | 55   |
| 2 | 2014 | 4     | 40   |
| 3 | 2013 | 7     | 84   |
| 4 | 2014 | 10    | 31   |

Create a MultiIndex using two Series:

```
>>> s = pd.Series([1, 2, 3, 4])
>>> df.set_index([s, s**2])
```

|      | month | year | sale |
|------|-------|------|------|
| 1 1  | 1     | 2012 | 55   |
| 2 4  | 4     | 2014 | 40   |
| 3 9  | 7     | 2013 | 84   |
| 4 16 | 10    | 2014 | 31   |

```
shift(self, periods=1, freq=None, axis=0, fill_value=<object object at 0x152b0a4c52f0>) -> 'DataFrame'
```

Shift index by desired number of periods with an optional time `freq`.

When `freq` is not passed, shift the index without realigning the data. If `freq` is passed (in this case, the index must be date or datetime, or it will raise a `NotImplementedError`), the index will be increased using the periods and the `freq`. `freq` can be inferred when specified as "infer" as long as either freq or inferred\_freq attribute is set in the index.

#### Parameters

**periods** : int

Number of periods to shift. Can be positive or negative.

**freq** : DateOffset, tseries.offsets, timedelta, or str, optional

Offset to use from the tseries module or time rule (e.g. 'EOM').

If `freq` is specified then the index values are shifted but the data is not realigned. That is, use `freq` if you would like to extend the index when shifting and preserve the original data.

If `freq` is specified as "infer" then it will be inferred from the freq or inferred\_freq attributes of the index. If neither of those attributes exist, a ValueError is thrown.

**axis** : {0 or 'index', 1 or 'columns', None}, default None

Shift direction.

**fill\_value** : object, optional

The scalar value to use for newly introduced missing values. the default depends on the dtype of `self`.

For numeric data, ``np.nan`` is used.

For datetime, timedelta, or period data, etc. :attr:`NaT` is used.

For extension dtypes, ``self.dtype.na\_value`` is used.

.. versionchanged:: 1.1.0

#### Returns

**DataFrame**

Copy of input object, shifted.

## See Also

-----

Index.shift : Shift values of Index.

DatetimeIndex.shift : Shift values of DatetimeIndex.

PeriodIndex.shift : Shift values of PeriodIndex.

tshift : Shift the time index, using the index's frequency if available.

## Examples

-----

```
>>> df = pd.DataFrame({"Col1": [10, 20, 15, 30, 45],
... "Col2": [13, 23, 18, 33, 48],
... "Col3": [17, 27, 22, 37, 52]},
... index=pd.date_range("2020-01-01", "2020-01-05"))
>>> df
```

|            | Col1 | Col2 | Col3 |
|------------|------|------|------|
| 2020-01-01 | 10   | 13   | 17   |
| 2020-01-02 | 20   | 23   | 27   |
| 2020-01-03 | 15   | 18   | 22   |
| 2020-01-04 | 30   | 33   | 37   |
| 2020-01-05 | 45   | 48   | 52   |

```
>>> df.shift(periods=3)
 Col1 Col2 Col3
2020-01-01 NaN NaN NaN
2020-01-02 NaN NaN NaN
2020-01-03 NaN NaN NaN
2020-01-04 10.0 13.0 17.0
2020-01-05 20.0 23.0 27.0
```

```
>>> df.shift(periods=1, axis="columns")
 Col1 Col2 Col3
2020-01-01 NaN 10 13
2020-01-02 NaN 20 23
2020-01-03 NaN 15 18
2020-01-04 NaN 30 33
2020-01-05 NaN 45 48
```

```
>>> df.shift(periods=3, fill_value=0)
 Col1 Col2 Col3
2020-01-01 0 0 0
2020-01-02 0 0 0
2020-01-03 0 0 0
2020-01-04 10 13 17
2020-01-05 20 23 27
```

```
>>> df.shift(periods=3, freq="D")
 Col1 Col2 Col3
2020-01-04 10 13 17
2020-01-05 20 23 27
2020-01-06 15 18 22
2020-01-07 30 33 37
2020-01-08 45 48 52
```

```
>>> df.shift(periods=3, freq="infer")
 Col1 Col2 Col3
2020-01-04 10 13 17
2020-01-05 20 23 27
2020-01-06 15 18 22
2020-01-07 30 33 37
2020-01-08 45 48 52
```

```
skew(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
Return unbiased skew over requested axis.
```

Normalized by N-1.

#### Parameters

`axis` : {index (0), columns (1)}  
 Axis for the function to be applied on.

`skipna` : bool, default True  
 Exclude NA/null values when computing the result.

`level` : int or level name, default None  
 If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

`numeric_only` : bool, default None  
 Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

`**kwargs`  
 Additional keyword arguments to be passed to the function.

#### Returns

Series or DataFrame (if level specified)

`sort_index(self, axis=0, level=None, ascending: 'Union[Union[bool, int], Sequence[Union[bool, int]]]' = True, inplace: 'bool' = False, kind: 'str' = 'quicksort', na_position: 'str' = 'last', sort_remaining: 'bool' = True, ignore_index: 'bool' = False, key: 'IndexKeyFunc' = None)`  
 Sort object by labels (along an axis).

Returns a new DataFrame sorted by label if ``inplace`` argument is ``False``, otherwise updates the original DataFrame and returns None.

#### Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0  
 The axis along which to sort. The value 0 identifies the rows, and 1 identifies the columns.

`level` : int or level name or list of ints or list of level names  
 If not None, sort on values in specified index level(s).

`ascending` : bool or list-like of bools, default True  
 Sort ascending vs. descending. When the index is a MultiIndex the sort direction can be controlled for each level individually.

`inplace` : bool, default False  
 If True, perform operation in-place.

`kind` : {'quicksort', 'mergesort', 'heapsort'}, default 'quicksort'  
 Choice of sorting algorithm. See also `ndarray.sort` for more information. ``mergesort`` is the only stable algorithm. For DataFrames, this option is only applied when sorting on a single column or label.

`na_position` : {'first', 'last'}, default 'last'  
 Puts NaNs at the beginning if ``first``; ``last`` puts NaNs at the end. Not implemented for MultiIndex.

`sort_remaining` : bool, default True  
 If True and sorting by level and index is multilevel, sort by other levels too (in order) after sorting by specified level.

`ignore_index` : bool, default False  
 If True, the resulting axis will be labeled 0, 1, ..., n - 1.

`.. versionadded:: 1.0.0`

`key` : callable, optional  
 If not None, apply the key function to the index values before sorting. This is similar to the ``key`` argument in the builtin `:meth:`sorted`` function, with the notable difference that this ``key`` function should be `*vectorized*`. It should expect an ``Index`` and return an ``Index`` of the same shape. For MultiIndex

inputs, the key is applied \*per level\*.

.. versionadded:: 1.1.0

#### Returns

DataFrame or None

The original DataFrame sorted by the labels or None if ``inplace=True

#### See Also

Series.sort\_index : Sort Series by the index.

DataFrame.sort\_values : Sort DataFrame by the value.

Series.sort\_values : Sort Series by the value.

#### Examples

```
>>> df = pd.DataFrame([1, 2, 3, 4, 5], index=[100, 29, 234, 1, 150],
... columns=['A'])
```

```
>>> df.sort_index()
```

```
 A
1 4
29 2
100 1
150 5
234 3
```

By default, it sorts in ascending order, to sort in descending order, use ``ascending=False``

```
>>> df.sort_index(ascending=False)
```

```
 A
234 3
150 5
100 1
29 2
1 4
```

A key function can be specified which is applied to the index before sorting. For a ``MultiIndex`` this is applied to each level separately.

```
>>> df = pd.DataFrame({"a": [1, 2, 3, 4]}, index=['A', 'b', 'C', 'd'])
```

```
>>> df.sort_index(key=lambda x: x.str.lower())
```

```
 a
A 1
b 2
C 3
d 4
```

sort\_values(self, by, axis=0, ascending=True, inplace=False, kind='quicksort', na\_position='last', ignore\_index=False, key: 'ValueKeyFunc' = None)  
Sort by the values along either axis.

#### Parameters

by : str or list of str

Name or list of names to sort by.

- if ``axis`` is 0 or ``'index'`` then ``by`` may contain index levels and/or column labels.

- if ``axis`` is 1 or ``'columns'`` then ``by`` may contain column levels and/or index labels.

axis : {0 or 'index', 1 or 'columns'}, default 0

Axis to be sorted.

ascending : bool or list of bool, default True  
Sort ascending vs. descending. Specify list for multiple sort orders. If this is a list of bools, must match the length of the by.

inplace : bool, default False  
If True, perform operation in-place.

kind : {'quicksort', 'mergesort', 'heapsort'}, default 'quicksort'  
Choice of sorting algorithm. See also ndarray.np.sort for more information. `mergesort` is the only stable algorithm. For DataFrames, this option is only applied when sorting on a single column or label.

na\_position : {'first', 'last'}, default 'last'  
Puts NaNs at the beginning if `first`; `last` puts NaNs at the end.

ignore\_index : bool, default False  
If True, the resulting axis will be labeled 0, 1, ..., n - 1.

.. versionadded:: 1.0.0

key : callable, optional  
Apply the key function to the values before sorting. This is similar to the `key` argument in the builtin :meth:`sorted` function, with the notable difference that this `key` function should be *\*vectorized\**. It should expect a ``Series`` and return a Series with the same shape as the input. It will be applied to each column in `by` independently.

.. versionadded:: 1.1.0

#### Returns

DataFrame or None

DataFrame with sorted values or None if ``inplace=True``.

#### See Also

DataFrame.sort\_index : Sort a DataFrame by the index.  
Series.sort\_values : Similar method for a Series.

#### Examples

```
>>> df = pd.DataFrame({
... 'col1': ['A', 'A', 'B', np.nan, 'D', 'C'],
... 'col2': [2, 1, 9, 8, 7, 4],
... 'col3': [0, 1, 9, 4, 2, 3],
... 'col4': ['a', 'B', 'c', 'D', 'e', 'F']
... })
>>> df
```

|   | col1 | col2 | col3 | col4 |
|---|------|------|------|------|
| 0 | A    | 2    | 0    | a    |
| 1 | A    | 1    | 1    | B    |
| 2 | B    | 9    | 9    | c    |
| 3 | NaN  | 8    | 4    | D    |
| 4 | D    | 7    | 2    | e    |
| 5 | C    | 4    | 3    | F    |

Sort by col1

```
>>> df.sort_values(by=['col1'])
```

|   | col1 | col2 | col3 | col4 |
|---|------|------|------|------|
| 0 | A    | 2    | 0    | a    |
| 1 | A    | 1    | 1    | B    |
| 2 | B    | 9    | 9    | c    |
| 5 | C    | 4    | 3    | F    |
| 4 | D    | 7    | 2    | e    |



```
3 NaN 8 4 D
```

Sort by multiple columns

```
>>> df.sort_values(by=['col1', 'col2'])
```

```
 col1 col2 col3 col4
1 A 1 1 B
0 A 2 0 a
2 B 9 9 c
5 C 4 3 F
4 D 7 2 e
3 NaN 8 4 D
```

Sort Descending

```
>>> df.sort_values(by='col1', ascending=False)
```

```
 col1 col2 col3 col4
4 D 7 2 e
5 C 4 3 F
2 B 9 9 c
0 A 2 0 a
1 A 1 1 B
3 NaN 8 4 D
```

Putting NAs first

```
>>> df.sort_values(by='col1', ascending=False, na_position='first')
```

```
 col1 col2 col3 col4
3 NaN 8 4 D
4 D 7 2 e
5 C 4 3 F
2 B 9 9 c
0 A 2 0 a
1 A 1 1 B
```

Sorting with a key function

```
>>> df.sort_values(by='col4', key=lambda col: col.str.lower())
```

```
 col1 col2 col3 col4
0 A 2 0 a
1 A 1 1 B
2 B 9 9 c
3 NaN 8 4 D
4 D 7 2 e
5 C 4 3 F
```

Natural sort with the key argument,  
using the `natsort` <<https://github.com/SethMMorton/natsort>> package.

```
>>> df = pd.DataFrame({
... "time": ['0hr', '128hr', '72hr', '48hr', '96hr'],
... "value": [10, 20, 30, 40, 50]
... })
>>> df
 time value
0 0hr 10
1 128hr 20
2 72hr 30
3 48hr 40
4 96hr 50
>>> from natsort import index_natsorted
>>> df.sort_values(
... by="time",
... key=lambda x: np.argsort(index_natsorted(df["time"]))
...)
```

|   | time  | value |
|---|-------|-------|
| 0 | 0hr   | 10    |
| 3 | 48hr  | 40    |
| 2 | 72hr  | 30    |
| 4 | 96hr  | 50    |
| 1 | 128hr | 20    |

```
stack(self, level=-1, dropna=True)
```

Stack the prescribed level(s) from columns to index.

Return a reshaped DataFrame or Series having a multi-level index with one or more new inner-most levels compared to the current DataFrame. The new inner-most levels are created by pivoting the columns of the current dataframe:

- if the columns have a single level, the output is a Series;
- if the columns have multiple levels, the new index level(s) is (are) taken from the prescribed level(s) and the output is a DataFrame.

#### Parameters

**level** : int, str, list, default -1

Level(s) to stack from the column axis onto the index axis, defined as one index or label, or a list of indices or labels.

**dropna** : bool, default True

Whether to drop rows in the resulting Frame/Series with missing values. Stacking a column level onto the index axis can create combinations of index and column values that are missing from the original dataframe. See Examples section.

#### Returns

DataFrame or Series

Stacked dataframe or series.

#### See Also

**DataFrame.unstack** : Unstack prescribed level(s) from index axis onto column axis.

**DataFrame.pivot** : Reshape dataframe from long format to wide format.

**DataFrame.pivot\_table** : Create a spreadsheet-style pivot table as a DataFrame.

#### Notes

The function is named by analogy with a collection of books being reorganized from being side by side on a horizontal position (the columns of the dataframe) to being stacked vertically on top of each other (in the index of the dataframe).

#### Examples

**\*\*Single level columns\*\***

```
>>> df_single_level_cols = pd.DataFrame([[0, 1], [2, 3]],
... index=['cat', 'dog'],
... columns=['weight', 'height'])
```

Stacking a dataframe with a single level column axis returns a Series:

```
>>> df_single_level_cols
 weight height
cat 0 1
dog 2 3
>>> df_single_level_cols.stack()
cat weight 0
 height 1
dog weight 2
 height 3
dtype: int64

Multi level columns: simple case

>>> multicol1 = pd.MultiIndex.from_tuples([('weight', 'kg'),
... ('weight', 'pounds')])
>>> df_multi_level_cols1 = pd.DataFrame([[1, 2], [2, 4]],
... index=['cat', 'dog'],
... columns=multicol1)
```

Stacking a dataframe with a multi-level column axis:

```
>>> df_multi_level_cols1
 weight
 kg pounds
cat 1 2
dog 2 4
>>> df_multi_level_cols1.stack()
cat weight
kg 1
pounds 2
dog kg 2
 pounds 4

Missing values

>>> multicol2 = pd.MultiIndex.from_tuples([('weight', 'kg'),
... ('height', 'm')])
>>> df_multi_level_cols2 = pd.DataFrame([[1.0, 2.0], [3.0, 4.0]],
... index=['cat', 'dog'],
... columns=multicol2)
```

It is common to have missing values when stacking a dataframe with multi-level columns, as the stacked dataframe typically has more values than the original dataframe. Missing values are filled with NaNs:

```
>>> df_multi_level_cols2
 weight height
 kg m
cat 1.0 2.0
dog 3.0 4.0
>>> df_multi_level_cols2.stack()
cat height weight
kg NaN 1.0
m 2.0 NaN
dog kg 3.0
 m 4.0 NaN
```

**\*\*Prescribing the level(s) to be stacked\*\***

The first parameter controls which level or levels are stacked:

```
>>> df_multi_level_cols2.stack(0)
 kg m
cat height NaN 2.0
```

```

weight 1.0 NaN
dog height NaN 4.0
weight 3.0 NaN
>>> df_multi_level_cols2.stack([0, 1])
cat height m 2.0
weight kg 1.0
dog height m 4.0
weight kg 3.0
dtype: float64

Dropping missing values

>>> df_multi_level_cols3 = pd.DataFrame([[None, 1.0], [2.0, 3.0]],
... index=['cat', 'dog'],
... columns=multicol2)

```

Note that rows where all values are missing are dropped by default but this behaviour can be controlled via the `dropna` keyword parameter:

```

>>> df_multi_level_cols3
weight height
kg m
cat NaN 1.0
dog 2.0 3.0
>>> df_multi_level_cols3.stack(dropna=False)
height weight
cat kg NaN NaN
m 1.0 NaN
dog kg NaN 2.0
m 3.0 NaN
>>> df_multi_level_cols3.stack(dropna=True)
height weight
cat m 1.0 NaN
dog kg NaN 2.0
m 3.0 NaN

```

`std(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None, **kwargs)`

Return sample standard deviation over requested axis.

Normalized by N-1 by default. This can be changed using the `ddof` argument.

#### Parameters

`axis` : {index (0), columns (1)}

`skipna` : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

`level` : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

`ddof` : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is  $N - ddof$

where  $N$  represents the number of elements.

`numeric_only` : bool, default None

Include only float, int, boolean columns. If None, will attempt to use

everything, then use only numeric data. Not implemented for Series.

#### Returns

Series or DataFrame (if level specified)

## Notes

-----  
To have the same behaviour as ``numpy.std``, use ``ddof=0`` (instead of the default ``ddof=1``)

`sub(self, other, axis='columns', level=None, fill_value=None)`  
Get Subtraction of dataframe and other, element-wise (binary operator ``sub``).

Equivalent to ``dataframe - other``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, ``rsub``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

## Parameters

-----  
`other` : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.  
`axis` : {0 or 'index', 1 or 'columns'}  
Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.  
`level` : int or label  
Broadcast across a level, matching Index values on the passed MultiIndex level.  
`fill_value` : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

## Returns

-----  
DataFrame  
Result of the arithmetic operation.

## See Also

-----  
`DataFrame.add` : Add DataFrames.  
`DataFrame.sub` : Subtract DataFrames.  
`DataFrame.mul` : Multiply DataFrames.  
`DataFrame.div` : Divide DataFrames (float division).  
`DataFrame.truediv` : Divide DataFrames (float division).  
`DataFrame.floordiv` : Divide DataFrames (integer division).  
`DataFrame.mod` : Calculate modulo (remainder after division).  
`DataFrame.pow` : Calculate exponential power.

## Notes

-----  
Mismatched indices will be unioned together.

## Examples

-----  

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | 360     |
| triangle  | 3      | 180     |
| rectangle | 4      | 360     |

Add a scalar with operator version which return the same

results.

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub([1, 2], axis='columns')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
...
... axis='index'))
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 359     |
| triangle  | 2      | 179     |
| rectangle | 3      | 359     |

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

|           | angles |
|-----------|--------|
| circle    | 0      |
| triangle  | 3      |
| rectangle | 4      |

```
>>> df * other
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | NaN     |
| triangle  | 9      | NaN     |
| rectangle | 16     | NaN     |

```

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

subtract = sub(self, other, axis='columns', level=None, fill_value=None)

sum(self, axis=None, skipna=None, level=None, numeric_only=None, min_count=
0, **kwargs)
 Return the sum of the values over the requested axis.

 This is equivalent to the method ``numpy.sum``.

 Parameters

 axis : {index (0), columns (1)}
 Axis for the function to be applied on.
 skipna : bool, default True
 Exclude NA/null values when computing the result.
 level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a Series.
 numeric_only : bool, default None
 Include only float, int, boolean columns. If None, will attempt to u
se
 everything, then use only numeric data. Not implemented for Series.
 min_count : int, default 0
 The required number of valid values to perform the operation. If few
er than
 ``min_count`` non-NA values are present the result will be NA.
 **kwargs
 Additional keyword arguments to be passed to the function.

 Returns

```

Series or DataFrame (if level specified)

See Also

Series.sum : Return the sum.  
 Series.min : Return the minimum.  
 Series.max : Return the maximum.  
 Series.idxmin : Return the index of the minimum.  
 Series.idxmax : Return the index of the maximum.  
 DataFrame.sum : Return the sum over the requested axis.  
 DataFrame.min : Return the minimum over the requested axis.  
 DataFrame.max : Return the maximum over the requested axis.  
 DataFrame.idxmin : Return the index of the minimum over the requested axis.  
 DataFrame.idxmax : Return the index of the maximum over the requested axis.

Examples

```
>>> idx = pd.MultiIndex.from_arrays([
... ['warm', 'warm', 'cold', 'cold'],
... ['dog', 'falcon', 'fish', 'spider']],
... names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
>>> s
```

| blooded | animal |   |
|---------|--------|---|
| warm    | dog    | 4 |
|         | falcon | 2 |
| cold    | fish   | 0 |
|         | spider | 8 |

```
Name: legs, dtype: int64

>>> s.sum()
14
```

Sum using level names, as well as indices.

```
>>> s.sum(level='blooded')
blooded
warm 6
cold 8
Name: legs, dtype: int64

>>> s.sum(level=0)
blooded
warm 6
cold 8
Name: legs, dtype: int64
```

By default, the sum of an empty or all-NA Series is ``0``.

```
>>> pd.Series([]).sum() # min_count=0 is the default
0.0
```

This can be controlled with the ``min\_count`` parameter. For example, if you'd like the sum of an empty series to be NaN, pass ``min\_count=1``.

```
>>> pd.Series([]).sum(min_count=1)
nan
```

Thanks to the ``skipna`` parameter, ``min\_count`` handles all-NA and empty series identically.

```
>>> pd.Series([np.nan]).sum()
0.0
```



```
>>> pd.Series([np.nan]).sum(min_count=1)
nan
```

`swaplevel(self, i=-2, j=-1, axis=0) -> 'DataFrame'`  
 Swap levels i and j in a MultiIndex on a particular axis.

#### Parameters

`i, j : int or str`  
 Levels of the indices to be swapped. Can pass level name as string.  
`axis : {0 or 'index', 1 or 'columns'}, default 0`  
 The axis to swap levels on. 0 or 'index' for row-wise, 1 or 'columns' for column-wise.

#### Returns

`DataFrame`

`to_dict(self, orient='dict', into=<class 'dict'>)`  
 Convert the DataFrame to a dictionary.

The type of the key-value pairs can be customized with the parameters (see below).

#### Parameters

`orient : str {'dict', 'list', 'series', 'split', 'records', 'index'}`  
 Determines the type of the values of the dictionary.

- 'dict' (default) : dict like {column -> {index -> value}}
- 'list' : dict like {column -> [values]}
- 'series' : dict like {column -> Series(values)}
- 'split' : dict like  
 {'index' -> [index], 'columns' -> [columns], 'data' -> [values]}
- 'records' : list like  
 [{column -> value}, ... , {column -> value}]
- 'index' : dict like {index -> {column -> value}}

Abbreviations are allowed. `s` indicates `series` and `sp` indicates `split`.

`into : class, default dict`  
 The collections.abc.Mapping subclass used for all Mappings in the return value. Can be the actual class or an empty instance of the mapping type you want. If you want a collections.defaultdict, you must pass it initialized.

#### Returns

`dict, list or collections.abc.Mapping`  
 Return a collections.abc.Mapping object representing the DataFrame. The resulting transformation depends on the `orient` parameter.

#### See Also

`DataFrame.from_dict`: Create a DataFrame from a dictionary.  
`DataFrame.to_json`: Convert a DataFrame to JSON format.

#### Examples

```
>>> df = pd.DataFrame({'col1': [1, 2],
... 'col2': [0.5, 0.75]},
... index=['row1', 'row2'])
>>> df
```

```

 col1 col2
row1 1 0.50
row2 2 0.75
>>> df.to_dict()
{'col1': {'row1': 1, 'row2': 2}, 'col2': {'row1': 0.5, 'row2': 0.75}}

```

You can specify the return orientation.

```

>>> df.to_dict('series')
{'col1': row1 1
 row2 2
Name: col1, dtype: int64,
 'col2': row1 0.50
 row2 0.75
Name: col2, dtype: float64}

>>> df.to_dict('split')
{'index': ['row1', 'row2'], 'columns': ['col1', 'col2'],
 'data': [[1, 0.5], [2, 0.75]]}

>>> df.to_dict('records')
[{'col1': 1, 'col2': 0.5}, {'col1': 2, 'col2': 0.75}]

>>> df.to_dict('index')
{'row1': {'col1': 1, 'col2': 0.5}, 'row2': {'col1': 2, 'col2': 0.75}}

```

You can also specify the mapping type.

```

>>> from collections import OrderedDict, defaultdict
>>> df.to_dict(into=OrderedDict)
OrderedDict([('col1', OrderedDict([('row1', 1), ('row2', 2)])),
 ('col2', OrderedDict([('row1', 0.5), ('row2', 0.75)]))])

```

If you want a `defaultdict`, you need to initialize it:

```

>>> dd = defaultdict(list)
>>> df.to_dict('records', into=dd)
[defaultdict(<class 'list'>, {'col1': 1, 'col2': 0.5}),
 defaultdict(<class 'list'>, {'col1': 2, 'col2': 0.75})]

```

`to_feather(self, path: 'FilePathOrBuffer[AnyStr]', **kwargs) -> 'None'`  
 Write a DataFrame to the binary Feather format.

#### Parameters

`path` : str or file-like object

If a string, it will be used as Root Directory path.

`**kwargs` :

Additional keywords passed to `:func:`pyarrow.feather.write_feather``.  
 Starting with pyarrow 0.17, this includes the ``compression``,  
``compression_level``, ``chunksize`` and ``version`` keywords.

.. versionadded:: 1.1.0

`to_gbq(self, destination_table, project_id=None, chunksize=None, reauth=False, if_exists='fail', auth_local_webserver=False, table_schema=None, location=None, progress_bar=True, credentials=None) -> 'None'`

Write a DataFrame to a Google BigQuery table.

This function requires the ``pandas-gbq`` package

<https://pandas-gbq.readthedocs.io>>`\_\_.

See the ``How to authenticate with Google BigQuery``

<https://pandas-gbq.readthedocs.io/en/latest/howto/authentication.html>>`

guide for authentication instructions.

### Parameters

```

destination_table : str
 Name of table to be written, in the form ``dataset.tablename``.
project_id : str, optional
 Google BigQuery Account project ID. Optional when available from
 the environment.
chunksize : int, optional
 Number of rows to be inserted in each chunk from the dataframe.
 Set to ``None`` to load the whole dataframe at once.
reauth : bool, default False
 Force Google BigQuery to re-authenticate the user. This is useful
 if multiple accounts are used.
if_exists : str, default 'fail'
 Behavior when the destination table exists. Value can be one of:

 ``'fail'``
 If table exists raise pandas_gbq.gbq.TableCreationError.
 ``'replace'``
 If table exists, drop it, recreate it, and insert data.
 ``'append'``
 If table exists, insert data. Create if does not exist.
auth_local_webserver : bool, default False
 Use the `local webserver flow`_ instead of the `console flow`_
 when getting user credentials.

 .. _local webserver flow:
 https://google-auth-oauthlib.readthedocs.io/en/latest/reference/
google_auth_oauthlib.flow.html#google_auth_oauthlib.flow.InstalledAppFlow.run_lo
cal_server
 .. _console flow:
 https://google-auth-oauthlib.readthedocs.io/en/latest/reference/
google_auth_oauthlib.flow.html#google_auth_oauthlib.flow.InstalledAppFlow.run_co
nsole

 New in version 0.2.0 of pandas-gbq.
table_schema : list of dicts, optional
 List of BigQuery table fields to which according DataFrame
 columns conform to, e.g. ``[{'name': 'col1', 'type':
 'STRING'},...]``. If schema is not provided, it will be
 generated according to dtypes of DataFrame columns. See
 BigQuery API documentation on available names of a field.

 New in version 0.3.1 of pandas-gbq.
location : str, optional
 Location where the load job should run. See the `BigQuery locations
 documentation
 <https://cloud.google.com/bigquery/docs/dataset-locations>`_ for a
 list of available locations. The location must match that of the
 target dataset.

 New in version 0.5.0 of pandas-gbq.
progress_bar : bool, default True
 Use the library `tqdm` to show the progress bar for the upload,
 chunk by chunk.

 New in version 0.5.0 of pandas-gbq.
credentials : google.auth.credentials.Credentials, optional
 Credentials for accessing Google APIs. Use this parameter to
 override default credentials, such as to use Compute Engine
 :class:`google.auth.compute_engine.Credentials` or Service
 Account :class:`google.oauth2.service_account.Credentials`
 directly.

```

\*New in version 0.8.0 of pandas-gbq\*.

.. versionadded:: 0.24.0

See Also

-----  
pandas\_gbq.to\_gbq : This function in the pandas-gbq library.  
read\_gbq : Read a DataFrame from Google BigQuery.

to\_html(self, buf=None, columns=None, col\_space=None, header=True, index=True, na\_rep='NaN', formatters=None, float\_format=None, sparsify=None, index\_names=True, justify=None, max\_rows=None, max\_cols=None, show\_dimensions=False, decimal='.', bold\_rows=True, classes=None, escape=True, notebook=False, border=None, table\_id=None, render\_links=False, encoding=None)

Render a DataFrame as an HTML table.

Parameters

-----  
buf : str, Path or StringIO-like, optional, default None  
Buffer to write to. If None, the output is returned as a string.  
columns : sequence, optional, default None  
The subset of columns to write. Writes all columns by default.  
col\_space : str or int, list or dict of int or str, optional  
The minimum width of each column in CSS length units. An int is assumed to be px units.

.. versionadded:: 0.25.0

Ability to use str.

header : bool, optional

Whether to print column labels, default True.

index : bool, optional, default True

Whether to print index (row) labels.

na\_rep : str, optional, default 'NaN'

String representation of ``NaN`` to use.

formatters : list, tuple or dict of one-param. functions, optional

Formatter functions to apply to columns' elements by position or name.

The result of each function must be a unicode string.

List/tuple must be of length equal to the number of columns.

float\_format : one-parameter function, optional, default None

Formatter function to apply to columns' elements if they are floats. This function must return a unicode string and will be applied only to the non-``NaN`` elements, with ``NaN`` being handled by ``na\_rep``.

.. versionchanged:: 1.2.0

sparsify : bool, optional, default True

Set to False for a DataFrame with a hierarchical index to print every multiindex key at each row.

index\_names : bool, optional, default True

Prints the names of the indexes.

justify : str, default None

How to justify the column labels. If None uses the option from the print configuration (controlled by set\_option), 'right' out of the box. Valid values are

- \* left
- \* right
- \* center
- \* justify
- \* justify-all
- \* start
- \* end

```

 * inherit
 * match-parent
 * initial
 * unset.
max_rows : int, optional
 Maximum number of rows to display in the console.
min_rows : int, optional
 The number of rows to display in the console in a truncated repr
 (when number of rows is above `max_rows`).
max_cols : int, optional
 Maximum number of columns to display in the console.
show_dimensions : bool, default False
 Display DataFrame dimensions (number of rows by number of columns).
decimal : str, default '.'
 Character recognized as decimal separator, e.g. ',' in Europe.

bold_rows : bool, default True
 Make the row labels bold in the output.
classes : str or list or tuple, default None
 CSS class(es) to apply to the resulting html table.
escape : bool, default True
 Convert the characters <, >, and & to HTML-safe sequences.
notebook : {True, False}, default False
 Whether the generated HTML is for IPython Notebook.
border : int
 A ``border=border`` attribute is included in the opening
 `<table>` tag. Default ``pd.options.display.html.border``.
encoding : str, default "utf-8"
 Set character encoding.

```

```
.. versionadded:: 1.0
```

```

table_id : str, optional
 A css id is included in the opening `<table>` tag if specified.
render_links : bool, default False
 Convert URLs to HTML links.

```

```
.. versionadded:: 0.24.0
```

#### Returns

```

str or None
 If buf is None, returns the result as a string. Otherwise returns
 None.

```

#### See Also

```
to_string : Convert DataFrame to a string.
```

```

to_markdown(self, buf: 'Optional[Union[IO[str], str]]' = None, mode: 'str' =
'wt', index: 'bool' = True, storage_options: 'StorageOptions' = None, **kwargs)
-> 'Optional[str]'

```

```
Print DataFrame in Markdown-friendly format.
```

```
.. versionadded:: 1.0.0
```

#### Parameters

```

buf : str, Path or StringIO-like, optional, default None
 Buffer to write to. If None, the output is returned as a string.
mode : str, optional
 Mode in which file is opened, "wt" by default.
index : bool, optional, default True
 Add index (row) labels.

```

e.g. `.. versionadded:: 1.1.0`  
`storage_options : dict, optional`  
 Extra options that make sense for a particular storage connection,  
 host, port, username, password, etc., if using a URL that will  
 be parsed by ```fsspec```, e.g., starting `"s3://"`, `"gcs://"`. An error  
 will be raised if providing this argument with a non-`fsspec` URL.  
 See the `fsspec` and backend storage implementation docs for the set o  
 f allowed keys and values.

`.. versionadded:: 1.2.0`  
`**kwargs`  
 These parameters will be passed to ``tabulate`` `<http`  
`s://pypi.org/project/tabulate>`_`.`

#### Returns

`str`

DataFrame in Markdown-friendly format.

#### Notes

Requires the ``tabulate`` `<https://pypi.org/project/tabulate>`_`` package.

#### Examples

```
>>> s = pd.Series(["elk", "pig", "dog", "quetzal"], name="animal")
>>> print(s.to_markdown())
```

```
	animal
0	elk
1	pig
2	dog
3	quetzal
```

Output markdown with a `tabulate` option.

```
>>> print(s.to_markdown(tablefmt="grid"))
```

```
+-----+
| | animal |
+====+=====+
| 0 | elk |
+-----+-----+
| 1 | pig |
+-----+-----+
| 2 | dog |
+-----+-----+
| 3 | quetzal|
+-----+-----+
```

`to_numpy(self, dtype=None, copy: 'bool' = False, na_value=<object object at 0x152b0a4c52f0>) -> 'np.ndarray'`  
 Convert the DataFrame to a NumPy array.

`.. versionadded:: 0.24.0`

By default, the dtype of the returned array will be the common NumPy dtype of all types in the DataFrame. For example, if the dtypes are ```float16``` and ```float32```, the results dtype will be ```float32```. This may require copying data and coercing values, which may be expensive.

#### Parameters

```

dtype : str or numpy.dtype, optional
 The dtype to pass to :meth:`numpy.asarray`.
copy : bool, default False
 Whether to ensure that the returned value is not a view on
 another array. Note that ``copy=False`` does not *ensure* that
 ``to_numpy()`` is no-copy. Rather, ``copy=True`` ensure that
 a copy is made, even if not strictly necessary.
na_value : Any, optional
 The value to use for missing values. The default value depends
 on `dtype` and the dtypes of the DataFrame columns.

.. versionadded:: 1.1.0

```

Returns

```
numpy.ndarray
```

See Also

Series.to\_numpy : Similar method for Series.

Examples

```

>>> pd.DataFrame({"A": [1, 2], "B": [3, 4]}).to_numpy()
array([[1, 3],
 [2, 4]])

```

With heterogeneous data, the lowest common type will have to be used.

```

>>> df = pd.DataFrame({"A": [1, 2], "B": [3.0, 4.5]})
>>> df.to_numpy()
array([[1. , 3.],
 [2. , 4.5]])

```

For a mix of numeric and non-numeric types, the output array will have object dtype.

```

>>> df['C'] = pd.date_range('2000', periods=2)
>>> df.to_numpy()
array([[1, 3.0, Timestamp('2000-01-01 00:00:00')],
 [2, 4.5, Timestamp('2000-01-02 00:00:00')]], dtype=object)

```

```

to_parquet(self, path: 'Optional[FilePathOrBuffer]' = None, engine: 'str' =
'auto', compression: 'Optional[str]' = 'snappy', index: 'Optional[bool]' = None,
partition_cols: 'Optional[List[str]]' = None, storage_options: 'StorageOptions'
= None, **kwargs) -> 'Optional[bytes]'

```

Write a DataFrame to the binary parquet format.

This function writes the dataframe as a `parquet` file [<https://parquet.apache.org/>](https://parquet.apache.org/). You can choose different parquet backends, and have the option of compression. See [:ref:`the user guide <io.parquet>](#) for more details.

Parameters

```

path : str or file-like object, default None
 If a string, it will be used as Root Directory path
 when writing a partitioned dataset. By file-like object,
 we refer to objects with a write() method, such as a file handle
 (e.g. via builtin open function) or io.BytesIO. The engine
 fastparquet does not accept file-like objects. If path is None,
 a bytes object is returned.

```

```
.. versionchanged:: 1.2.0
```

Previously this was "fname"

```
engine : {'auto', 'pyarrow', 'fastparquet'}, default 'auto'
 Parquet library to use. If 'auto', then the option
 ``io.parquet.engine`` is used. The default ``io.parquet.engine``
 behavior is to try 'pyarrow', falling back to 'fastparquet' if
 'pyarrow' is unavailable.
compression : {'snappy', 'gzip', 'brotli', None}, default 'snappy'
 Name of the compression to use. Use ``None`` for no compression.
index : bool, default None
 If ``True``, include the dataframe's index(es) in the file output.
 If ``False``, they will not be written to the file.
 If ``None``, similar to ``True`` the dataframe's index(es)
 will be saved. However, instead of being saved as values,
 the RangeIndex will be stored as a range in the metadata so it
 doesn't require much space and is faster. Other indexes will
 be included as columns in the file output.
```

```
.. versionadded:: 0.24.0
```

```
partition_cols : list, optional, default None
 Column names by which to partition the dataset.
 Columns are partitioned in the order they are given.
 Must be None if path is not a string.
```

```
.. versionadded:: 0.24.0
```

```
storage_options : dict, optional
 Extra options that make sense for a particular storage connection,
 e.g.
 host, port, username, password, etc., if using a URL that will
 be parsed by ``fsspec``, e.g., starting "s3://", "gcs://". An error
 will be raised if providing this argument with a non-fsspec URL.
 See the fsspec and backend storage implementation docs for the set o
 f
 allowed keys and values.
```

```
.. versionadded:: 1.2.0
```

```
**kwargs
 Additional arguments passed to the parquet library. See
 :ref:`pandas io <io.parquet>` for more details.
```

Returns

bytes if no path argument is provided else None

See Also

```
read_parquet : Read a parquet file.
DataFrame.to_csv : Write a csv file.
DataFrame.to_sql : Write to a sql table.
DataFrame.to_hdf : Write to hdf.
```

Notes

This function requires either the `fastparquet` <https://pypi.org/project/fastparquet> or `pyarrow` <https://arrow.apache.org/docs/python/> library.

Examples

```
>>> df = pd.DataFrame(data={'col1': [1, 2], 'col2': [3, 4]})
```



```
>>> df.to_parquet('df.parquet.gzip',
... compression='gzip') # doctest: +SKIP
>>> pd.read_parquet('df.parquet.gzip') # doctest: +SKIP
 col1 col2
0 1 3
1 2 4
```

sIO

files.

If you want to get a buffer to the parquet content you can use a `io.BytesIO` object, as long as you don't use `partition_cols`, which creates multiple

```
>>> import io
>>> f = io.BytesIO()
>>> df.to_parquet(f)
>>> f.seek(0)
0
>>> content = f.read()
```

ame'

`to_period(self, freq=None, axis: 'Axis' = 0, copy: 'bool' = True) -> 'DataFr`

Convert DataFrame from DatetimeIndex to PeriodIndex.

Convert DataFrame from DatetimeIndex to PeriodIndex with desired frequency (inferred from index if not passed).

Parameters

```
freq : str, default
 Frequency of the PeriodIndex.
axis : {0 or 'index', 1 or 'columns'}, default 0
 The axis to convert (the index by default).
copy : bool, default True
 If False then underlying input data is not copied.
```

Returns

DataFrame with PeriodIndex

`to_records(self, index=True, column_dtypes=None, index_dtypes=None) -> 'np.r`  
ecarray'

Convert DataFrame to a NumPy record array.

Index will be included as the first field of the record array if requested.

Parameters

```
index : bool, default True
 Include index in resulting record array, stored in 'index'
 field or using the index label, if set.
column_dtypes : str, type, dict, default None
 .. versionadded:: 0.24.0

 If a string or type, the data type to store all columns. If
 a dictionary, a mapping of column names and indices (zero-indexed)
 to specific data types.
index_dtypes : str, type, dict, default None
 .. versionadded:: 0.24.0
```

If a string or type, the data type to store all index levels. If a dictionary, a mapping of index level names and indices (zero-indexed) to specific data types.

This mapping is applied only if ``index=True``.

## Returns

### numpy.recarray

NumPy ndarray with the DataFrame labels as fields and each row of the DataFrame as entries.

## See Also

`DataFrame.from_records`: Convert structured or record ndarray to DataFrame.

`numpy.recarray`: An ndarray that allows field access using attributes, analogous to typed columns in a spreadsheet.

## Examples

```
>>> df = pd.DataFrame({'A': [1, 2], 'B': [0.5, 0.75]},
... index=['a', 'b'])
>>> df
 A B
a 1 0.50
b 2 0.75
>>> df.to_records()
rec.array([('a', 1, 0.5), ('b', 2, 0.75)],
 dtype=[('index', 'O'), ('A', '<i8'), ('B', '<f8')])
```

If the DataFrame index has no label then the recarray field name is set to 'index'. If the index has a label then this is used as the field name:

```
>>> df.index = df.index.rename("I")
>>> df.to_records()
rec.array([('a', 1, 0.5), ('b', 2, 0.75)],
 dtype=[('I', 'O'), ('A', '<i8'), ('B', '<f8')])
```

The index can be excluded from the record array:

```
>>> df.to_records(index=False)
rec.array([(1, 0.5), (2, 0.75)],
 dtype=[('A', '<i8'), ('B', '<f8')])
```

Data types can be specified for the columns:

```
>>> df.to_records(column_dtypes={"A": "int32"})
rec.array([('a', 1, 0.5), ('b', 2, 0.75)],
 dtype=[('I', 'O'), ('A', '<i4'), ('B', '<f8')])
```

As well as for the index:

```
>>> df.to_records(index_dtypes="<S2")
rec.array([(b'a', 1, 0.5), (b'b', 2, 0.75)],
 dtype=[('I', 'S2'), ('A', '<i8'), ('B', '<f8')])

>>> index_dtypes = f"<S{df.index.str.len().max()}"
>>> df.to_records(index_dtypes=index_dtypes)
rec.array([(b'a', 1, 0.5), (b'b', 2, 0.75)],
 dtype=[('I', 'S1'), ('A', '<i8'), ('B', '<f8')])
```

```
to_stata(self, path: 'FilePathOrBuffer', convert_dates: 'Optional[Dict[Label, str]]' = None, write_index: 'bool' = True, byteorder: 'Optional[str]' = None,
time_stamp: 'Optional[datetime.datetime]' = None, data_label: 'Optional[str]' = None, variable_labels: 'Optional[Dict[Label, str]]' = None, version: 'Optional[int]' = 114, convert_strl: 'Optional[Sequence[Label]]' = None, compression: 'CompressionOptions' = 'infer', storage_options: 'StorageOptions' = None) -> 'None'
```

Export DataFrame object to Stata dta format.

Writes the DataFrame to a Stata dataset file.

"dta" files contain a Stata dataset.

#### Parameters

`path` : str, buffer or path object

String, path object (pathlib.Path or py.\_path.local.LocalPath) or object implementing a binary write() function. If using a buffer then the buffer will not be automatically closed after the file data has been written.

.. versionchanged:: 1.0.0

Previously this was "fname"

`convert_dates` : dict

Dictionary mapping columns containing datetime types to stata internal format to use when writing the dates. Options are 'tc', 'td', 'tm', 'tw', 'th', 'tq', 'ty'. Column can be either an integer or a name. Datetime columns that do not have a conversion type specified will be converted to 'tc'. Raises NotImplementedError if a datetime column has timezone information.

`write_index` : bool

Write the index to Stata dataset.

`byteorder` : str

Can be ">", "<", "little", or "big". default is `sys.byteorder`.

`time_stamp` : datetime

A datetime to use as file creation date. Default is the current time.

`data_label` : str, optional

A label for the data set. Must be 80 characters or smaller.

`variable_labels` : dict

Dictionary containing columns as keys and variable labels as values. Each label must be 80 characters or smaller.

`version` : {114, 117, 118, 119, None}, default 114

Version to use in the output dta file. Set to None to let pandas decide between 118 or 119 formats depending on the number of columns in the frame. Version 114 can be read by Stata 10 and later. Version 117 can be read by Stata 13 or later. Version 118 is supported in Stata 14 and later. Version 119 is supported in Stata 15 and later. Version 114 limits string variables to 244 characters or fewer while versions 117 and later allow strings with lengths up to 2,000,000 characters. Versions 118 and 119 support Unicode characters, and version 119 supports more than 32,767 variables.

Version 119 should usually only be used when the number of variables exceeds the capacity of dta format 118. Exporting smaller datasets in format 119 may have unintended consequences, and, as of November 2020, Stata SE cannot read version 119 files.

.. versionchanged:: 1.0.0

Added support for formats 118 and 119.

`convert_strl` : list, optional

List of column names to convert to string columns to Stata StrL format. Only available if version is 117. Storing strings in the StrL format can produce smaller dta files if strings have more than 8 characters and values are repeated.

`compression` : str or dict, default 'infer'

For on-the-fly compression of the output dta. If string, specifies compression mode. If dict, value at key 'method' specifies

compression mode. Compression mode must be one of {'infer', 'gzip', 'bz2', 'zip', 'xz', None}. If compression mode is 'infer' and 'fname' is path-like, then detect compression from the following extensions: '.gz', '.bz2', '.zip', or '.xz' (otherwise no compression). If dict and compression mode is one of {'zip', 'gzip', 'bz2'}, or inferred as one of the above, other entries passed as additional compression options.

.. versionadded:: 1.1.0

storage\_options : dict, optional

Extra options that make sense for a particular storage connection,

e.g.

host, port, username, password, etc., if using a URL that will be parsed by ``fsspec``, e.g., starting "s3://", "gcs://". An error will be raised if providing this argument with a non-fsspec URL. See the fsspec and backend storage implementation docs for the set o

f

allowed keys and values.

.. versionadded:: 1.2.0

Raises

-----  
NotImplementedError

- \* If datetimes contain timezone information
- \* Column dtype is not representable in Stata

ValueError

- \* Columns listed in convert\_dates are neither datetime64[ns] or datetime.datetime
- \* Column listed in convert\_dates is not in DataFrame
- \* Categorical label contains more than 32,000 characters

See Also

-----  
read\_stata : Import Stata data files.

io.stata.StataWriter : Low-level writer for Stata data files.

io.stata.StataWriter117 : Low-level writer for version 117 files.

Examples

```
>>> df = pd.DataFrame({'animal': ['falcon', 'parrot', 'falcon',
... 'parrot'],
... 'speed': [350, 18, 361, 15]})
>>> df.to_stata('animals.dta') # doctest: +SKIP
```

to\_string(self, buf: 'Optional[FilePathOrBuffer[str]]' = None, columns: 'Optional[Sequence[str]]' = None, col\_space: 'Optional[int]' = None, header: 'Union[bool, Sequence[str]]' = True, index: 'bool' = True, na\_rep: 'str' = 'NaN', formatters: 'Optional[fmt.FormatterType]' = None, float\_format: 'Optional[fmt.FloatFormatType]' = None, sparsify: 'Optional[bool]' = None, index\_names: 'bool' = True, justify: 'Optional[str]' = None, max\_rows: 'Optional[int]' = None, min\_rows: 'Optional[int]' = None, max\_cols: 'Optional[int]' = None, show\_dimensions: 'bool' = False, decimal: 'str' = '.', line\_width: 'Optional[int]' = None, max\_colwidth: 'Optional[int]' = None, encoding: 'Optional[str]' = None) -> 'Optional[str]'  
Render a DataFrame to a console-friendly tabular output.

Parameters

-----  
buf : str, Path or StringIO-like, optional, default None

Buffer to write to. If None, the output is returned as a string.

columns : sequence, optional, default None

The subset of columns to write. Writes all columns by default.

col\_space : int, list or dict of int, optional

The minimum width of each column.

```

header : bool or sequence, optional
 Write out the column names. If a list of strings is given, it is assumed to be aliases for the column names.
index : bool, optional, default True
 Whether to print index (row) labels.
na_rep : str, optional, default 'NaN'
 String representation of ``NaN`` to use.
formatters : list, tuple or dict of one-param. functions, optional
 Formatter functions to apply to columns' elements by position or name.
 The result of each function must be a unicode string.
 List/tuple must be of length equal to the number of columns.
float_format : one-parameter function, optional, default None
 Formatter function to apply to columns' elements if they are floats. This function must return a unicode string and will be applied only to the non-``NaN`` elements, with ``NaN`` being handled by ``na_rep``.

 .. versionchanged:: 1.2.0

sparsify : bool, optional, default True
 Set to False for a DataFrame with a hierarchical index to print every multiindex key at each row.
index_names : bool, optional, default True
 Prints the names of the indexes.
justify : str, default None
 How to justify the column labels. If None uses the option from the print configuration (controlled by set_option), 'right' out of the box. Valid values are

 * left
 * right
 * center
 * justify
 * justify-all
 * start
 * end
 * inherit
 * match-parent
 * initial
 * unset.
max_rows : int, optional
 Maximum number of rows to display in the console.
min_rows : int, optional
 The number of rows to display in the console in a truncated repr (when number of rows is above `max_rows`).
max_cols : int, optional
 Maximum number of columns to display in the console.
show_dimensions : bool, default False
 Display DataFrame dimensions (number of rows by number of columns).
decimal : str, default '.'
 Character recognized as decimal separator, e.g. ',' in Europe.

line_width : int, optional
 Width to wrap a line in characters.
max_colwidth : int, optional
 Max width to truncate each column in characters. By default, no limit.

 .. versionadded:: 1.0.0
encoding : str, default "utf-8"
 Set character encoding.

 .. versionadded:: 1.0

```

## Returns

str or None

If buf is None, returns the result as a string. Otherwise returns None.

## See Also

to\_html : Convert DataFrame to HTML.

## Examples

```
>>> d = {'col1': [1, 2, 3], 'col2': [4, 5, 6]}
>>> df = pd.DataFrame(d)
>>> print(df.to_string())
 col1 col2
0 1 4
1 2 5
2 3 6
```

```
to_timestamp(self, freq=None, how: 'str' = 'start', axis: 'Axis' = 0, copy:
'bool' = True) -> 'DataFrame'
```

Cast to DatetimeIndex of timestamps, at \*beginning\* of period.

## Parameters

freq : str, default frequency of PeriodIndex

Desired frequency.

how : {'s', 'e', 'start', 'end'}

Convention for converting period to timestamp; start of period vs. end.

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to convert (the index by default).

copy : bool, default True

If False then underlying input data is not copied.

## Returns

DataFrame with DatetimeIndex

```
transform(self, func: 'AggFuncType', axis: 'Axis' = 0, *args, **kwargs) ->
'DataFrame'
```

Call ``func`` on self producing a DataFrame with transformed values.

Produced DataFrame will have same axis length as self.

## Parameters

func : function, str, list-like or dict-like

Function to use for transforming the data. If a function, must either work when passed a DataFrame or when passed to DataFrame.apply. If func is both list-like and dict-like, dict-like behavior takes precedence.

Accepted combinations are:

- function
- string function name
- list-like of functions and/or function names, e.g. ``[np.exp, 'sqrt']``
- dict-like of axis labels -> functions, function names or list-like of such.

axis : {0 or 'index', 1 or 'columns'}, default 0

If 0 or 'index': apply function to each column.  
 If 1 or 'columns': apply function to each row.

\*args  
 Positional arguments to pass to `func`.  
 \*\*kwargs  
 Keyword arguments to pass to `func`.

Returns

-----  
 DataFrame

A DataFrame that must have the same length as self.

Raises

-----  
 ValueError : If the returned DataFrame has a different length than self.

See Also

-----  
 DataFrame.agg : Only perform aggregating type operations.  
 DataFrame.apply : Invoke function on a DataFrame.

Examples

```
>>> df = pd.DataFrame({'A': range(3), 'B': range(1, 4)})
>>> df
 A B
0 0 1
1 1 2
2 2 3
>>> df.transform(lambda x: x + 1)
 A B
0 1 2
1 2 3
2 3 4
```

Even though the resulting DataFrame must have the same length as the input DataFrame, it is possible to provide several input functions:

```
>>> s = pd.Series(range(3))
>>> s
0 0
1 1
2 2
dtype: int64
>>> s.transform([np.sqrt, np.exp])
 sqrt exp
0 0.000000 1.000000
1 1.000000 2.718282
2 1.414214 7.389056
```

You can call transform on a GroupBy object:

```
>>> df = pd.DataFrame({
... "Date": [
... "2015-05-08", "2015-05-07", "2015-05-06", "2015-05-05",
... "2015-05-08", "2015-05-07", "2015-05-06", "2015-05-05"],
... "Data": [5, 8, 6, 1, 50, 100, 60, 120],
... })
>>> df
 Date Data
0 2015-05-08 5
1 2015-05-07 8
2 2015-05-06 6
3 2015-05-05 1
4 2015-05-08 50
```

```

5 2015-05-07 100
6 2015-05-06 60
7 2015-05-05 120
>>> df.groupby('Date')['Data'].transform('sum')
0 55
1 108
2 66
3 121
4 55
5 108
6 66
7 121
Name: Data, dtype: int64

>>> df = pd.DataFrame({
... "c": [1, 1, 1, 2, 2, 2, 2],
... "type": ["m", "n", "o", "m", "m", "n", "n"]
... })
>>> df
 c type
0 1 m
1 1 n
2 1 o
3 2 m
4 2 m
5 2 n
6 2 n
>>> df['size'] = df.groupby('c')['type'].transform(len)
>>> df
 c type size
0 1 m 3
1 1 n 3
2 1 o 3
3 2 m 4
4 2 m 4
5 2 n 4
6 2 n 4

```

`transpose(self, *args, copy: 'bool' = False) -> 'DataFrame'`  
 Transpose index and columns.

Reflect the DataFrame over its main diagonal by writing rows as columns and vice-versa. The property `:attr:'.T'` is an accessor to the method `:meth:'.transpose'`.

#### Parameters

`*args` : tuple, optional  
 Accepted for compatibility with NumPy.

`copy` : bool, default False  
 Whether to copy the data after transposing, even for DataFrames with a single dtype.

Note that a copy is always required for mixed dtype DataFrames, or for DataFrames with any extension types.

#### Returns

##### DataFrame

The transposed DataFrame.

#### See Also

`numpy.transpose` : Permute the dimensions of a given array.



## Notes

Transposing a DataFrame with mixed dtypes will result in a homogeneous DataFrame with the `object` dtype. In such a case, a copy of the data is always made.

## Examples

**\*\*Square DataFrame with homogeneous dtype\*\***

```
>>> d1 = {'col1': [1, 2], 'col2': [3, 4]}
>>> df1 = pd.DataFrame(data=d1)
>>> df1
 col1 col2
0 1 3
1 2 4

>>> df1_transposed = df1.T # or df1.transpose()
>>> df1_transposed
 0 1
col1 1 2
col2 3 4
```

When the dtype is homogeneous in the original DataFrame, we get a transposed DataFrame with the same dtype:

```
>>> df1.dtypes
col1 int64
col2 int64
dtype: object
>>> df1_transposed.dtypes
0 int64
1 int64
dtype: object
```

**\*\*Non-square DataFrame with mixed dtypes\*\***

```
>>> d2 = {'name': ['Alice', 'Bob'],
... 'score': [9.5, 8],
... 'employed': [False, True],
... 'kids': [0, 0]}
>>> df2 = pd.DataFrame(data=d2)
>>> df2
 name score employed kids
0 Alice 9.5 False 0
1 Bob 8.0 True 0

>>> df2_transposed = df2.T # or df2.transpose()
>>> df2_transposed
 0 1
name Alice Bob
score 9.5 8.0
employed False True
kids 0 0
```

When the DataFrame has mixed dtypes, we get a transposed DataFrame with the `object` dtype:

```
>>> df2.dtypes
name object
score float64
employed bool
kids int64
dtype: object
>>> df2_transposed.dtypes
```

```

0 object
1 object
dtype: object

truediv(self, other, axis='columns', level=None, fill_value=None)
 Get Floating division of dataframe and other, element-wise (binary opera
tor `truediv`).

 Equivalent to ``dataframe / other``, but with support to substitute a fi
ll_value
 for missing data in one of the inputs. With reverse version, `rtruediv`.

 Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
 arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same
results.

```

```
>>> df + 1
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

```
>>> df.add(1)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 1      | 361     |
| triangle  | 4      | 181     |
| rectangle | 5      | 361     |

Divide by constant with reverse version.

```
>>> df.div(10)
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0.0    | 36.0    |
| triangle  | 0.3    | 18.0    |
| rectangle | 0.4    | 36.0    |

```
>>> df.rdiv(10)
```

|           | angles   | degrees  |
|-----------|----------|----------|
| circle    | inf      | 0.027778 |
| triangle  | 3.333333 | 0.055556 |
| rectangle | 2.500000 | 0.027778 |

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub([1, 2], axis='columns')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 358     |
| triangle  | 2      | 178     |
| rectangle | 3      | 358     |

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle'],
e')),
```

```
... axis='index')
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | -1     | 359     |
| triangle  | 2      | 179     |
| rectangle | 3      | 359     |

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
```

```
>>> other
```

|           | angles |
|-----------|--------|
| circle    | 0      |
| triangle  | 3      |
| rectangle | 4      |

```
>>> df * other
```

|           | angles | degrees |
|-----------|--------|---------|
| circle    | 0      | NaN     |
| triangle  | 9      | NaN     |
| rectangle | 16     | NaN     |

```
>>> df.mul(other, fill_value=0)
```

```

 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0

Divide by a MultiIndex by level.

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]
0}},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
e',
... ['circle', 'triangle', 'rectangl
n']]
... 'square', 'pentagon', 'hexago

>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

unstack(self, level=-1, fill_value=None)
Pivot a level of the (necessarily hierarchical) index labels.

Returns a DataFrame having a new level of column labels whose inner-most
level consists of the pivoted index labels.

If the index is not a MultiIndex, the output will be a Series
(the analogue of stack when the columns are not a MultiIndex).

Parameters

level : int, str, or list of these, default -1 (last level)
 Level(s) of index to unstack, can pass level name.
fill_value : int, str or dict
 Replace NaN with this value if the unstack produces missing values.

Returns

Series or DataFrame

See Also

DataFrame.pivot : Pivot a table based on column values.
DataFrame.stack : Pivot a level of the column labels (inverse operation
 from `unstack`).

Examples

>>> index = pd.MultiIndex.from_tuples([('one', 'a'), ('one', 'b'),
... ('two', 'a'), ('two', 'b')])
>>> s = pd.Series(np.arange(1.0, 5.0), index=index)

```

```

>>> s
one a 1.0
 b 2.0
two a 3.0
 b 4.0
dtype: float64

>>> s.unstack(level=-1)
 a b
one 1.0 2.0
two 3.0 4.0

>>> s.unstack(level=0)
 one two
a 1.0 3.0
b 2.0 4.0

>>> df = s.unstack(level=0)
>>> df.unstack()
one a 1.0
 b 2.0
two a 3.0
 b 4.0
dtype: float64

```

`update(self, other, join='left', overwrite=True, filter_func=None, errors='ignore') -> 'None'`

Modify in place using non-NA values from another DataFrame.

Aligns on indices. There is no return value.

#### Parameters

`other` : DataFrame, or object coercible into a DataFrame  
Should have at least one matching index/column label with the original DataFrame. If a Series is passed, its name attribute must be set, and that will be used as the column name to align with the original DataFrame.

`join` : {'left'}, default 'left'  
Only left join is implemented, keeping the index and columns of the original object.

`overwrite` : bool, default True  
How to handle non-NA values for overlapping keys:

- \* True: overwrite original DataFrame's values with values from ``other``.
- \* False: only update values that are NA in the original DataFrame.

`filter_func` : callable(1d-array) -> bool 1d-array, optional  
Can choose to replace values other than NA. Return True for values that should be updated.

`errors` : {'raise', 'ignore'}, default 'ignore'  
If 'raise', will raise a ValueError if the DataFrame and ``other`` both contain non-NA data in the same place.

.. versionchanged:: 0.24.0  
Changed from ``raise_conflict=False|True`` to ``errors='ignore'|'raise'``.

#### Returns

None : method directly changes calling object

#### Raises

```

ValueError
 * When `errors='raise'` and there's overlapping non-NA data.
 * When `errors` is not either `ignore` or `raise`
NotImplementedError
 * If `join != 'left'`

```

See Also

```

dict.update : Similar method for dictionaries.
DataFrame.merge : For column(s)-on-column(s) operations.

```

Examples

```

>>> df = pd.DataFrame({'A': [1, 2, 3],
... 'B': [400, 500, 600]})
>>> new_df = pd.DataFrame({'B': [4, 5, 6],
... 'C': [7, 8, 9]})
>>> df.update(new_df)
>>> df
 A B
0 1 4
1 2 5
2 3 6

```

The DataFrame's length does not increase as a result of the update, only values at matching index/column labels are updated.

```

>>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
... 'B': ['x', 'y', 'z']})
>>> new_df = pd.DataFrame({'B': ['d', 'e', 'f', 'g', 'h', 'i']})
>>> df.update(new_df)
>>> df
 A B
0 a d
1 b e
2 c f

```

For Series, its name attribute must be set.

```

>>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
... 'B': ['x', 'y', 'z']})
>>> new_column = pd.Series(['d', 'e'], name='B', index=[0, 2])
>>> df.update(new_column)
>>> df
 A B
0 a d
1 b y
2 c e
>>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
... 'B': ['x', 'y', 'z']})
>>> new_df = pd.DataFrame({'B': ['d', 'e'], index=[1, 2]})
>>> df.update(new_df)
>>> df
 A B
0 a x
1 b d
2 c e

```

If `other` contains NaNs the corresponding values are not updated in the original dataframe.

```

>>> df = pd.DataFrame({'A': [1, 2, 3],
... 'B': [400, 500, 600]})
>>> new_df = pd.DataFrame({'B': [4, np.nan, 6]})

```

```

>>> df.update(new_df)
>>> df
 A B
0 1 4.0
1 2 500.0
2 3 6.0

```

value\_counts(self, subset: 'Optional[Sequence[Label]]' = None, normalize: 'bool' = False, sort: 'bool' = True, ascending: 'bool' = False)

Return a Series containing counts of unique rows in the DataFrame.

.. versionadded:: 1.1.0

Parameters

-----

subset : list-like, optional  
Columns to use when counting unique combinations.

normalize : bool, default False  
Return proportions rather than frequencies.

sort : bool, default True  
Sort by frequencies.

ascending : bool, default False  
Sort in ascending order.

Returns

-----

Series

See Also

-----

Series.value\_counts: Equivalent method on Series.

Notes

-----

The returned Series will have a MultiIndex with one level per input column. By default, rows that contain any NA values are omitted from the result. By default, the resulting Series will be in descending order so that the first element is the most frequently-occurring row.

Examples

-----

```

>>> df = pd.DataFrame({'num_legs': [2, 4, 4, 6],
... 'num_wings': [2, 0, 0, 0]},
... index=['falcon', 'dog', 'cat', 'ant'])
>>> df
 num_legs num_wings
falcon 2 2
dog 4 0
cat 4 0
ant 6 0

>>> df.value_counts()
num_legs num_wings
4 0 2
2 2 1
6 0 1
dtype: int64

>>> df.value_counts(sort=False)
num_legs num_wings
2 2 1
4 0 2
6 0 1
dtype: int64

```

```
>>> df.value_counts(ascending=True)
num_legs num_wings
2 2 1
6 0 1
4 0 2
dtype: int64
```

```
>>> df.value_counts(normalize=True)
num_legs num_wings
4 0 0.50
2 2 0.25
6 0 0.25
dtype: float64
```

`var(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None, **kwargs)`

Return unbiased variance over requested axis.

Normalized by N-1 by default. This can be changed using the ddof argument

#### Parameters

`axis` : {index (0), columns (1)}

`skipna` : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

`level` : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

`ddof` : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is  $N - ddof$

where  $N$  represents the number of elements.

`numeric_only` : bool, default None

Include only float, int, boolean columns. If None, will attempt to use

everything, then use only numeric data. Not implemented for Series.

#### Returns

Series or DataFrame (if level specified)

#### Notes

To have the same behaviour as ``numpy.std``, use ``ddof=0`` (instead of the default ``ddof=1``)

---

Class methods defined here:

`from_dict(data, orient='columns', dtype=None, columns=None) -> 'DataFrame'` from builtins.type

Construct DataFrame from dict of array-like or dicts.

Creates DataFrame object from dictionary by columns or by index allowing dtype specification.

#### Parameters

`data` : dict

Of the form `{field : array-like}` or `{field : dict}`.

`orient` : {'columns', 'index'}, default 'columns'

The "orientation" of the data. If the keys of the passed dict should be the columns of the resulting DataFrame, pass 'columns'



(default). Otherwise if the keys should be rows, pass 'index'.  
 dtype : dtype, default None  
 Data type to force, otherwise infer.  
 columns : list, default None  
 Column labels to use when ``orient='index'``. Raises a ValueError  
 if used with ``orient='columns'``.

Returns

-----  
 DataFrame

See Also

-----  
 DataFrame.from\_records : DataFrame from structured ndarray, sequence  
 of tuples or dicts, or DataFrame.

DataFrame : DataFrame object creation using constructor.

Examples

-----  
 By default the keys of the dict become the DataFrame columns:

```
>>> data = {'col_1': [3, 2, 1, 0], 'col_2': ['a', 'b', 'c', 'd']}
>>> pd.DataFrame.from_dict(data)
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d
```

Specify ``orient='index'`` to create the DataFrame using dictionary  
 keys as rows:

```
>>> data = {'row_1': [3, 2, 1, 0], 'row_2': ['a', 'b', 'c', 'd']}
>>> pd.DataFrame.from_dict(data, orient='index')
 0 1 2 3
row_1 3 2 1 0
row_2 a b c d
```

When using the 'index' orientation, the column names can be  
 specified manually:

```
>>> pd.DataFrame.from_dict(data, orient='index',
... columns=['A', 'B', 'C', 'D'])
 A B C D
row_1 3 2 1 0
row_2 a b c d
```

from\_records(data, index=None, exclude=None, columns=None, coerce\_float=False,  
 e, nrows=None) -> 'DataFrame' from builtins.type  
 Convert structured or record ndarray to DataFrame.

Creates a DataFrame object from a structured ndarray, sequence of  
 tuples or dicts, or DataFrame.

Parameters

-----  
 data : structured ndarray, sequence of tuples or dicts, or DataFrame  
 Structured input data.  
 index : str, list of fields, array-like  
 Field of array to use as the index, alternately a specific set of  
 input labels to use.  
 exclude : sequence, default None  
 Columns or fields to exclude.  
 columns : sequence, default None  
 Column names to use. If the passed data do not have names

associated with them, this argument provides names for the columns. Otherwise this argument indicates the order of the columns in the result (any names not found in the data will become all-NA columns).

`coerce_float` : bool, default False

Attempt to convert values of non-string, non-numeric objects (like decimal.Decimal) to floating point, useful for SQL result sets.

`nrows` : int, default None

Number of rows to read if data is an iterator.

Returns

-----  
DataFrame

See Also

-----  
DataFrame.from\_dict : DataFrame from dict of array-like or dicts.

DataFrame : DataFrame object creation using constructor.

Examples

-----  
Data can be provided as a structured ndarray:

```
>>> data = np.array([(3, 'a'), (2, 'b'), (1, 'c'), (0, 'd')],
... dtype=[('col_1', 'i4'), ('col_2', 'U1')])
>>> pd.DataFrame.from_records(data)
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d
```

Data can be provided as a list of dicts:

```
>>> data = [{'col_1': 3, 'col_2': 'a'},
... {'col_1': 2, 'col_2': 'b'},
... {'col_1': 1, 'col_2': 'c'},
... {'col_1': 0, 'col_2': 'd'}]
>>> pd.DataFrame.from_records(data)
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d
```

Data can be provided as a list of tuples with corresponding columns:

```
>>> data = [(3, 'a'), (2, 'b'), (1, 'c'), (0, 'd')]
>>> pd.DataFrame.from_records(data, columns=['col_1', 'col_2'])
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d
```

-----  
Readonly properties defined here:

T

axes

Return a list representing the axes of the DataFrame.

It has the row axis labels and column axis labels as the only members. They are returned in that order.

### Examples

```
>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.axes
[RangeIndex(start=0, stop=2, step=1), Index(['col1', 'col2'],
dtype='object')]
```

### shape

Return a tuple representing the dimensionality of the DataFrame.

### See Also

`ndarray.shape` : Tuple of array dimensions.

### Examples

```
>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.shape
(2, 2)

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4],
... 'col3': [5, 6]})
>>> df.shape
(2, 3)
```

### style

Returns a Styler object.

Contains methods for building a styled HTML representation of the DataFrame.

ame.

### See Also

`io.formats.style.Styler` : Helps style a DataFrame or Series according to the data with HTML and CSS.

---

Data descriptors defined here:

### columns

The column labels of the DataFrame.

### index

The index (row labels) of the DataFrame.

---

Data and other attributes defined here:

`__annotations__` = {'\_AXIS\_TO\_AXIS\_NUMBER': 'Dict[Axis, int]', '\_access...

`plot` = <class 'pandas.plotting.\_core.PlotAccessor'>

Make plots of Series or DataFrame.

Uses the backend specified by the option ```plotting.backend```. By default, `matplotlib` is used.

### Parameters

`data` : Series or DataFrame

The object for which the method is called.

`x` : label or position, default None

Only used if data is a DataFrame.

`y` : label, position or list of label, positions, default None

```

 Allows plotting of one column versus another. Only used if data is a
 DataFrame.
kind : str
 The kind of plot to produce:

 - 'line' : line plot (default)
 - 'bar' : vertical bar plot
 - 'barh' : horizontal bar plot
 - 'hist' : histogram
 - 'box' : boxplot
 - 'kde' : Kernel Density Estimation plot
 - 'density' : same as 'kde'
 - 'area' : area plot
 - 'pie' : pie plot
 - 'scatter' : scatter plot
 - 'hexbin' : hexbin plot.
ax : matplotlib axes object, default None
 An axes of the current figure.
subplots : bool, default False
 Make separate subplots for each column.
sharex : bool, default True if ax is None else False
 In case ``subplots=True``, share x axis and set some x axis labels
 to invisible; defaults to True if ax is None otherwise False if
 an ax is passed in; Be aware, that passing in both an ax and
 ``sharex=True`` will alter all x axis labels for all axis in a figure.
e.
sharey : bool, default False
 In case ``subplots=True``, share y axis and set some y axis labels to
o invisible.
layout : tuple, optional
 (rows, columns) for the layout of subplots.
figsize : a tuple (width, height) in inches
 Size of a figure object.
use_index : bool, default True
 Use index as ticks for x axis.
title : str or list
 Title to use for the plot. If a string is passed, print the string
 at the top of the figure. If a list is passed and `subplots` is
 True, print each item in the list above the corresponding subplot.
grid : bool, default None (matlab style default)
 Axis grid lines.
legend : bool or {'reverse'}
 Place legend on axis subplots.
style : list or dict
 The matplotlib line style per column.
logx : bool or 'sym', default False
 Use log scaling or symlog scaling on x axis.
 .. versionchanged:: 0.25.0

logy : bool or 'sym' default False
 Use log scaling or symlog scaling on y axis.
 .. versionchanged:: 0.25.0

loglog : bool or 'sym', default False
 Use log scaling or symlog scaling on both x and y axes.
 .. versionchanged:: 0.25.0

xticks : sequence
 Values to use for the xticks.
yticks : sequence
 Values to use for the yticks.
xlim : 2-tuple/list
 Set the x limits of the current axes.
ylim : 2-tuple/list
 Set the y limits of the current axes.

```

```

xlabel : label, optional
 Name to use for the xlabel on x-axis. Default uses index name as xla
 bel, or the
 x-column name for planar plots.

 .. versionadded:: 1.1.0

 .. versionchanged:: 1.2.0

 Now applicable to planar plots (`scatter`, `hexbin`).

ylabel : label, optional
 Name to use for the ylabel on y-axis. Default will show no ylabel, o
 r the
 y-column name for planar plots.

 .. versionadded:: 1.1.0

 .. versionchanged:: 1.2.0

 Now applicable to planar plots (`scatter`, `hexbin`).

rot : int, default None
 Rotation for ticks (xticks for vertical, yticks for horizontal
 plots).

fontsize : int, default None
 Font size for xticks and yticks.

colormap : str or matplotlib colormap object, default None
 Colormap to select colors from. If string, load colormap with that
 name from matplotlib.

colorbar : bool, optional
 If True, plot colorbar (only relevant for 'scatter' and 'hexbin'
 plots).

position : float
 Specify relative alignments for bar plot layout.
 From 0 (left/bottom-end) to 1 (right/top-end). Default is 0.5
 (center).

table : bool, Series or DataFrame, default False
 If True, draw a table using the data in the DataFrame and the data
 will be transposed to meet matplotlib's default layout.
 If a Series or DataFrame is passed, use passed data to draw a
 table.

yerr : DataFrame, Series, array-like, dict and str
 See :ref:`Plotting with Error Bars <visualization.errorbars>` for
 detail.

xerr : DataFrame, Series, array-like, dict and str
 Equivalent to yerr.

stacked : bool, default False in line and bar plots, and True in area pl
 ot
 If True, create stacked plot.

sort_columns : bool, default False
 Sort column names to determine plot ordering.

secondary_y : bool or sequence, default False
 Whether to plot on the secondary y-axis if a list/tuple, which
 columns to plot on secondary y-axis.

mark_right : bool, default True
 When using a secondary_y axis, automatically mark the column
 labels with "(right)" in the legend.

include_bool : bool, default is False
 If True, boolean values can be plotted.

backend : str, default None
 Backend to use instead of the backend specified in the option
 ``plotting.backend``. For instance, 'matplotlib'. Alternatively, to
 specify the ``plotting.backend`` for the whole session, set
 ``pd.options.plotting.backend``.

```

```

 .. versionadded:: 1.0.0

 **kwargs
 Options to pass to matplotlib plotting method.

 Returns

 :class:`matplotlib.axes.Axes` or numpy.ndarray of them
 If the backend is not the default matplotlib one, the return value
 will be the object returned by the backend.

 Notes

 - See matplotlib documentation online for more on this subject
 - If `kind` = 'bar' or 'barh', you can specify relative alignments
 for bar plot layout by `position` keyword.
 From 0 (left/bottom-end) to 1 (right/top-end). Default is 0.5
 (center)

 sparse = <class 'pandas.core.arrays.sparse.accessor.SparseFrameAccesso...
 DataFrame accessor for sparse data.

 .. versionadded:: 0.25.0

 Methods inherited from pandas.core.generic.NDFrame:

 __abs__(self: 'FrameOrSeries') -> 'FrameOrSeries'

 __array__(self, dtype=None) -> 'np.ndarray'

 __array_ufunc__(self, ufunc: 'Callable', method: 'str', *inputs: 'Any', **kw
args: 'Any')

 __array_wrap__(self, result: 'np.ndarray', context: 'Optional[Tuple[Callabl
e, Tuple[Any, ...], int]]' = None)
 Gets called after a ufunc and other functions.

 Parameters

 result: np.ndarray
 The result of the ufunc or other function called on the NumPy array
 returned by __array__
 context: tuple of (func, tuple, int)
 This parameter is returned by ufuncs as a 3-element tuple: (name of
the
 ufunc, arguments of the ufunc, domain of the ufunc), but is not set
by
 other numpy functions.q

 Notes

 Series implements __array_ufunc__ so this not called for ufunc on Series.

 __bool__ = __nonzero__(self)

 __contains__(self, key) -> 'bool_t'
 True if the key is in the info axis

 __copy__(self: 'FrameOrSeries', deep: 'bool_t' = True) -> 'FrameOrSeries'

 __deepcopy__(self: 'FrameOrSeries', memo=None) -> 'FrameOrSeries'

 Parameters

```

```

memo, default None
 Standard signature. Unused

__delitem__(self, key) -> 'None'
 Delete item

__finalize__(self: 'FrameOrSeries', other, method: 'Optional[str]' = None, *
*kwargs) -> 'FrameOrSeries'
 Propagate metadata from other to self.

 Parameters

 other : the object from which to get the attributes that we are going
 to propagate
 method : str, optional
 A passed method name providing context on where ``__finalize__``
 was called.

 .. warning::

 The value passed as `method` are not currently considered
 stable across pandas releases.

__getattr__(self, name: 'str')
 After regular attribute access, try looking up the name
 This allows simpler access to columns for interactive use.

__getstate__(self) -> 'Dict[str, Any]'

__hash__(self) -> 'int'
 Return hash(self).

__iadd__(self, other)

__iand__(self, other)

__ifloordiv__(self, other)

__imod__(self, other)

__imul__(self, other)

__invert__(self)

__ior__(self, other)

__ipow__(self, other)

__isub__(self, other)

__iter__(self)
 Iterate over info axis.

 Returns

 iterator
 Info axis as iterator.

__itruediv__(self, other)

__ixor__(self, other)

__neg__(self)

__nonzero__(self)

```

```

__pos__(self)

__round__(self: 'FrameOrSeries', decimals: 'int' = 0) -> 'FrameOrSeries'

__setattr__(self, name: 'str', value) -> 'None'
 After regular attribute access, try setting the name
 This allows simpler access to columns for interactive use.

__setstate__(self, state)

abs(self: 'FrameOrSeries') -> 'FrameOrSeries'
 Return a Series/DataFrame with absolute numeric value of each element.

 This function only applies to elements that are all numeric.

Returns

abs
 Series/DataFrame containing the absolute value of each element.

See Also

numpy.absolute : Calculate the absolute value element-wise.

Notes

For ``complex`` inputs, ``1.2 + 1j``, the absolute value is
:math:`\sqrt{a^2 + b^2}``.

Examples

Absolute numeric values in a Series.

>>> s = pd.Series([-1.10, 2, -3.33, 4])
>>> s.abs()
0 1.10
1 2.00
2 3.33
3 4.00
dtype: float64

Absolute numeric values in a Series with complex numbers.

>>> s = pd.Series([1.2 + 1j])
>>> s.abs()
0 1.56205
dtype: float64

Absolute numeric values in a Series with a Timedelta element.

>>> s = pd.Series([pd.Timedelta('1 days')])
>>> s.abs()
0 1 days
dtype: timedelta64[ns]

Select rows with data closest to certain value using argsort (from
`StackOverflow <https://stackoverflow.com/a/17758115>`).

>>> df = pd.DataFrame({
... 'a': [4, 5, 6, 7],
... 'b': [10, 20, 30, 40],
... 'c': [100, 50, -30, -50]
... })
>>> df

```



```

 a b c
0 4 10 100
1 5 20 50
2 6 30 -30
3 7 40 -50
>>> df.loc[(df.c - 43).abs().argsort()]
 a b c
1 5 20 50
0 4 10 100
2 6 30 -30
3 7 40 -50

```

`add_prefix(self: 'FrameOrSeries', prefix: 'str') -> 'FrameOrSeries'`  
 Prefix labels with string `prefix`.

For Series, the row labels are prefixed.  
 For DataFrame, the column labels are prefixed.

#### Parameters

`prefix` : str  
 The string to add before each label.

#### Returns

Series or DataFrame  
 New Series or DataFrame with updated labels.

#### See Also

`Series.add_suffix`: Suffix row labels with string `suffix`.  
`DataFrame.add_suffix`: Suffix column labels with string `suffix`.

#### Examples

```

>>> s = pd.Series([1, 2, 3, 4])
>>> s
0 1
1 2
2 3
3 4
dtype: int64

>>> s.add_prefix('item_')
item_0 1
item_1 2
item_2 3
item_3 4
dtype: int64

>>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
>>> df
 A B
0 1 3
1 2 4
2 3 5
3 4 6

>>> df.add_prefix('col_')
 col_A col_B
0 1 3
1 2 4
2 3 5
3 4 6

```

```
add_suffix(self: 'FrameOrSeries', suffix: 'str') -> 'FrameOrSeries'
 Suffix labels with string `suffix`.
```

For Series, the row labels are suffixed.  
For DataFrame, the column labels are suffixed.

Parameters

`suffix` : str  
The string to add after each label.

Returns

Series or DataFrame  
New Series or DataFrame with updated labels.

See Also

`Series.add_prefix`: Prefix row labels with string `prefix`.  
`DataFrame.add_prefix`: Prefix column labels with string `prefix`.

Examples

```
>>> s = pd.Series([1, 2, 3, 4])
```

```
>>> s
0 1
1 2
2 3
3 4
dtype: int64
```

```
>>> s.add_suffix('_item')
```

```
0_item 1
1_item 2
2_item 3
3_item 4
dtype: int64
```

```
>>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
```

```
>>> df
 A B
0 1 3
1 2 4
2 3 5
3 4 6
```

```
>>> df.add_suffix('_col')
```

```
 A_col B_col
0 1 3
1 2 4
2 3 5
3 4 6
```

```
asfreq(self: 'FrameOrSeries', freq, method=None, how: 'Optional[str]' = None,
normalize: 'bool_t' = False, fill_value=None) -> 'FrameOrSeries'
 Convert TimeSeries to specified frequency.
```

Optionally provide filling method to pad/backfill missing values.

Returns the original data conformed to a new index with the specified frequency. ``resample`` is more appropriate if an operation, such as summarization, is necessary to represent the data at the new frequency.

Parameters

```

freq : DateOffset or str
 Frequency DateOffset or string.
method : {'backfill'/'bfill', 'pad'/'ffill'}, default None
 Method to use for filling holes in reindexed Series (note this
 does not fill NaNs that already were present):

 * 'pad' / 'ffill': propagate last valid observation forward to next
 valid
 * 'backfill' / 'bfill': use NEXT valid observation to fill.
how : {'start', 'end'}, default end
 For PeriodIndex only (see PeriodIndex.asfreq).
normalize : bool, default False
 Whether to reset output index to midnight.
fill_value : scalar, optional
 Value to use for missing values, applied during upsampling (note
 this does not fill NaNs that already were present).

```

#### Returns

Same type as caller  
Object converted to the specified frequency.

#### See Also

reindex : Conform DataFrame to new index with optional filling logic.

#### Notes

To learn more about the frequency strings, please see [this link](https://pandas.pydata.org/pandas-docs/stable/user_guide/timeseries.html)  
<[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/timeseries.html](https://pandas.pydata.org/pandas-docs/stable/user_guide/timeseries.html)

#offset-aliases>`\_.

#### Examples

Start by creating a series with 4 one minute timestamps.

```

>>> index = pd.date_range('1/1/2000', periods=4, freq='T')
>>> series = pd.Series([0.0, None, 2.0, 3.0], index=index)
>>> df = pd.DataFrame({'s':series})
>>> df

```

```

 s
2000-01-01 00:00:00 0.0
2000-01-01 00:01:00 NaN
2000-01-01 00:02:00 2.0
2000-01-01 00:03:00 3.0

```

Upsample the series into 30 second bins.

```

>>> df.asfreq(freq='30S')
 s
2000-01-01 00:00:00 0.0
2000-01-01 00:00:30 NaN
2000-01-01 00:01:00 NaN
2000-01-01 00:01:30 NaN
2000-01-01 00:02:00 2.0
2000-01-01 00:02:30 NaN
2000-01-01 00:03:00 3.0

```

Upsample again, providing a ``fill value``.

```

>>> df.asfreq(freq='30S', fill_value=9.0)
 s
2000-01-01 00:00:00 0.0
2000-01-01 00:00:30 9.0
2000-01-01 00:01:00 NaN

```

```
2000-01-01 00:01:30 9.0
2000-01-01 00:02:00 2.0
2000-01-01 00:02:30 9.0
2000-01-01 00:03:00 3.0
```

Upsample again, providing a ``method``.

```
>>> df.asfreq(freq='30S', method='bfill')
```

```

s
2000-01-01 00:00:00 0.0
2000-01-01 00:00:30 NaN
2000-01-01 00:01:00 NaN
2000-01-01 00:01:30 2.0
2000-01-01 00:02:00 2.0
2000-01-01 00:02:30 3.0
2000-01-01 00:03:00 3.0
```

`asof(self, where, subset=None)`

Return the last row(s) without any NaNs before `where`.

The last row (for each element in `where`, if list) without any NaN is taken.

In case of a `:class:`~pandas.DataFrame``, the last row without NaN considering only the subset of columns (if not `None`)

If there is no good value, NaN is returned for a Series or a Series of NaN values for a DataFrame

Parameters

`where` : date or array-like of dates

Date(s) before which the last row(s) are returned.

`subset` : str or array-like of str, default `None`

For DataFrame, if not `None`, only use these columns to check for NaNs.

Returns

scalar, Series, or DataFrame

The return can be:

- \* scalar : when `self` is a Series and `where` is a scalar
- \* Series: when `self` is a Series and `where` is an array-like, or when `self` is a DataFrame and `where` is a scalar
- \* DataFrame : when `self` is a DataFrame and `where` is an array-like

Return scalar, Series, or DataFrame.

See Also

`merge_asof` : Perform an asof merge. Similar to left join.

Notes

Dates are assumed to be sorted. Raises if this is not the case.

Examples

A Series and a scalar `where`.

```
>>> s = pd.Series([1, 2, np.nan, 4], index=[10, 20, 30, 40])
>>> s
10 1.0
```

```
20 2.0
30 NaN
40 4.0
dtype: float64
```

```
>>> s.asof(20)
2.0
```

For a sequence `where`, a Series is returned. The first value is NaN, because the first element of `where` is before the first index value.

```
>>> s.asof([5, 20])
5 NaN
20 2.0
dtype: float64
```

Missing values are not considered. The following is ``2.0``, not NaN, even though NaN is at the index location for ``30``.

```
>>> s.asof(30)
2.0
```

Take all columns into consideration

```
>>> df = pd.DataFrame({'a': [10, 20, 30, 40, 50],
... 'b': [None, None, None, None, 500]},
... index=pd.DatetimeIndex(['2018-02-27 09:01:00',
... '2018-02-27 09:02:00',
... '2018-02-27 09:03:00',
... '2018-02-27 09:04:00',
... '2018-02-27 09:05:00']))
>>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
... '2018-02-27 09:04:30']))
 a b
2018-02-27 09:03:30 NaN NaN
2018-02-27 09:04:30 NaN NaN
```

Take a single column into consideration

```
>>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
... '2018-02-27 09:04:30']),
... subset=['a'])
 a b
2018-02-27 09:03:30 30.0 NaN
2018-02-27 09:04:30 40.0 NaN
```

```
astype(self: 'FrameOrSeries', dtype, copy: 'bool_t' = True, errors: 'str' =
'raise') -> 'FrameOrSeries'
```

Cast a pandas object to a specified dtype ``dtype``.

Parameters

```
dtype : data type, or dict of column name -> data type
 Use a numpy.dtype or Python type to cast entire pandas object to
 the same type. Alternatively, use {col: dtype, ...}, where col is a
 column label and dtype is a numpy.dtype or Python type to cast one
 or more of the DataFrame's columns to column-specific types.
copy : bool, default True
 Return a copy when ``copy=True`` (be very careful setting
 ``copy=False`` as changes to values then may propagate to other
 pandas objects).
errors : {'raise', 'ignore'}, default 'raise'
 Control raising of exceptions on invalid data for provided dtype.
```

- ``raise`` : allow exceptions to be raised
- ``ignore`` : suppress exceptions. On error return original object.

### Returns

-----  
 casted : same type as caller

### See Also

-----  
 to\_datetime : Convert argument to datetime.  
 to\_timedelta : Convert argument to timedelta.  
 to\_numeric : Convert argument to a numeric type.  
 numpy.ndarray.astype : Cast a numpy array to a specified type.

### Examples

-----  
 Create a DataFrame:

```
>>> d = {'col1': [1, 2], 'col2': [3, 4]}
>>> df = pd.DataFrame(data=d)
>>> df.dtypes
col1 int64
col2 int64
dtype: object
```

Cast all columns to int32:

```
>>> df.astype('int32').dtypes
col1 int32
col2 int32
dtype: object
```

Cast col1 to int32 using a dictionary:

```
>>> df.astype({'col1': 'int32'}).dtypes
col1 int32
col2 int64
dtype: object
```

Create a series:

```
>>> ser = pd.Series([1, 2], dtype='int32')
>>> ser
0 1
1 2
dtype: int32
>>> ser.astype('int64')
0 1
1 2
dtype: int64
```

Convert to categorical type:

```
>>> ser.astype('category')
0 1
1 2
dtype: category
Categories (2, int64): [1, 2]
```

Convert to ordered categorical type with custom ordering:

```
>>> cat_dtype = pd.api.types.CategoricalDtype(
... categories=[2, 1], ordered=True)
>>> ser.astype(cat_dtype)
0 1
```

```
1 2
dtype: category
Categories (2, int64): [2 < 1]
```

Note that using ``copy=False`` and changing data on a new pandas object may propagate changes:

```
>>> s1 = pd.Series([1, 2])
>>> s2 = s1.astype('int64', copy=False)
>>> s2[0] = 10
>>> s1 # note that s1[0] has changed too
0 10
1 2
dtype: int64
```

Create a series of dates:

```
>>> ser_date = pd.Series(pd.date_range('20200101', periods=3))
>>> ser_date
0 2020-01-01
1 2020-01-02
2 2020-01-03
dtype: datetime64[ns]
```

Datetimes are localized to UTC first before converting to the specified timezone:

```
>>> ser_date.astype('datetime64[ns, US/Eastern]')
0 2019-12-31 19:00:00-05:00
1 2020-01-01 19:00:00-05:00
2 2020-01-02 19:00:00-05:00
dtype: datetime64[ns, US/Eastern]
```

`at_time(self: 'FrameOrSeries', time, asof: 'bool_t' = False, axis=None) -> 'FrameOrSeries'`

Select values at particular time of day (e.g., 9:30AM).

Parameters

time : datetime.time or str  
axis : {0 or 'index', 1 or 'columns'}, default 0

.. versionadded:: 0.24.0

Returns

Series or DataFrame

Raises

TypeError

If the index is not a :class:`DatetimeIndex`

See Also

between\_time : Select values between particular times of the day.  
first : Select initial periods of time series based on a date offset.  
last : Select final periods of time series based on a date offset.  
DatetimeIndex.indexer\_at\_time : Get just the index locations for values at particular time of the day.

Examples

```
>>> i = pd.date_range('2018-04-09', periods=4, freq='12H')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
```

```

>>> ts
 A
2018-04-09 00:00:00 1
2018-04-09 12:00:00 2
2018-04-10 00:00:00 3
2018-04-10 12:00:00 4

>>> ts.at_time('12:00')
 A
2018-04-09 12:00:00 2
2018-04-10 12:00:00 4

```

`backfill = bfill(self: 'FrameOrSeries', axis=None, inplace: 'bool_t' = False, limit=None, downcast=None) -> 'Optional[FrameOrSeries]'`

`between_time(self: 'FrameOrSeries', start_time, end_time, include_start: 'bool_t' = True, include_end: 'bool_t' = True, axis=None) -> 'FrameOrSeries'`

Select values between particular times of the day (e.g., 9:00–9:30 AM).

By setting ``start\_time`` to be later than ``end\_time``, you can get the times that are *not* between the two times.

Parameters

-----

`start_time` : `datetime.time` or `str`  
Initial time as a time filter limit.

`end_time` : `datetime.time` or `str`  
End time as a time filter limit.

`include_start` : `bool`, default `True`  
Whether the start time needs to be included in the result.

`include_end` : `bool`, default `True`  
Whether the end time needs to be included in the result.

`axis` : {0 or 'index', 1 or 'columns'}, default 0  
Determine range time on index or columns value.

.. versionadded:: 0.24.0

Returns

-----

Series or DataFrame  
Data from the original object filtered to the specified dates range.

Raises

-----

`TypeError`  
If the index is not a `:class:`DatetimeIndex``

See Also

-----

`at_time` : Select values at a particular time of the day.  
`first` : Select initial periods of time series based on a date offset.  
`last` : Select final periods of time series based on a date offset.  
`DatetimeIndex.indexer_between_time` : Get just the index locations for values between particular times of the day.

Examples

-----

```

>>> i = pd.date_range('2018-04-09', periods=4, freq='1D20min')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
>>> ts
 A
2018-04-09 00:00:00 1
2018-04-10 00:20:00 2
2018-04-11 00:40:00 3
2018-04-12 01:00:00 4

```



```
>>> ts.between_time('0:15', '0:45')
```

```
A
2018-04-10 00:20:00 2
2018-04-11 00:40:00 3
```

You get the times that are *\*not\** between two times by setting ``start\_time`` later than ``end\_time``:

```
>>> ts.between_time('0:45', '0:15')
```

```
A
2018-04-09 00:00:00 1
2018-04-12 01:00:00 4
```

```
bfill(self: 'FrameOrSeries', axis=None, inplace: 'bool_t' = False, limit=None,
downcast=None) -> 'Optional[FrameOrSeries]'
 Synonym for :meth:`DataFrame.fillna` with ``method='bfill'``.
```

Returns

Series/DataFrame or None

Object with missing values filled or None if ``inplace=True``.

bool(self)

Return the bool of a single element Series or DataFrame.

a

r that

ion).

Returns

bool

The value in the Series or DataFrame.

See Also

Series.astype : Change the data type of a Series, including to boolean.  
 DataFrame.astype : Change the data type of a DataFrame, including to boolean.

lean.

numpy.bool\_ : NumPy boolean data type, used by pandas for boolean values.

s.

Examples

The method will only work for single element objects with a boolean value:

e:

```
>>> pd.Series([True]).bool()
```

```
True
```

```
>>> pd.Series([False]).bool()
```

```
False
```

```
>>> pd.DataFrame({'col': [True]}).bool()
```

```
True
```

```
>>> pd.DataFrame({'col': [False]}).bool()
```

```
False
```

```
clip(self: 'FrameOrSeries', lower=None, upper=None, axis=None, inplace: 'bool_t' = False, *args, **kwargs) -> 'FrameOrSeries'
 Trim values at input threshold(s).
```

Assigns values outside boundary to boundary values. Thresholds can be singular values or array like, and in the latter case the clipping is performed element-wise in the specified axis.

#### Parameters

`lower` : float or array\_like, default None  
Minimum threshold value. All values below this threshold will be set to it.

`upper` : float or array\_like, default None  
Maximum threshold value. All values above this threshold will be set to it.

`axis` : int or str axis name, optional  
Align object with lower and upper along the given axis.

`inplace` : bool, default False  
Whether to perform the operation in place on the data.

`*args, **kwargs`  
Additional keywords have no effect but might be accepted for compatibility with numpy.

#### Returns

Series or DataFrame or None  
Same type as calling object with the values outside the clip boundaries replaced or None if ``inplace=True``.

#### See Also

`Series.clip` : Trim values at input threshold in series.  
`DataFrame.clip` : Trim values at input threshold in dataframe.  
`numpy.clip` : Clip (limit) the values in an array.

#### Examples

```
>>> data = {'col_0': [9, -3, 0, -1, 5], 'col_1': [-2, -7, 6, 8, -5]}
>>> df = pd.DataFrame(data)
>>> df
```

|   | col_0 | col_1 |
|---|-------|-------|
| 0 | 9     | -2    |
| 1 | -3    | -7    |
| 2 | 0     | 6     |
| 3 | -1    | 8     |
| 4 | 5     | -5    |

Clips per column using lower and upper thresholds:

```
>>> df.clip(-4, 6)
```

|   | col_0 | col_1 |
|---|-------|-------|
| 0 | 6     | -2    |
| 1 | -3    | -4    |
| 2 | 0     | 6     |
| 3 | -1    | 6     |
| 4 | 5     | -4    |

Clips using specific lower and upper thresholds per column element:

```
>>> t = pd.Series([2, -4, -1, 6, 3])
>>> t
```

|   | t  |
|---|----|
| 0 | 2  |
| 1 | -4 |
| 2 | -1 |
| 3 | 6  |
| 4 | 3  |

dtype: int64

```
>>> df.clip(t, t + 4, axis=0)
 col_0 col_1
0 6 2
1 -3 -4
2 0 3
3 6 8
4 5 3
```

```
convert_dtypes(self: 'FrameOrSeries', infer_objects: 'bool_t' = True, convert_string: 'bool_t' = True, convert_integer: 'bool_t' = True, convert_boolean: 'bool_t' = True, convert_floating: 'bool_t' = True) -> 'FrameOrSeries'
```

```
Convert columns to best possible dtypes using dtypes supporting ``pd.NA``
```

```
.. versionadded:: 1.0.0
```

#### Parameters

```
infer_objects : bool, default True
```

```
Whether object dtypes should be converted to the best possible type
```

```
convert_string : bool, default True
```

```
Whether object dtypes should be converted to ``StringDtype``.
```

```
convert_integer : bool, default True
```

```
Whether, if possible, conversion can be done to integer extension types.
```

```
convert_boolean : bool, defaults True
```

```
Whether object dtypes should be converted to ``BooleanDtype``.
```

```
convert_floating : bool, defaults True
```

```
Whether, if possible, conversion can be done to floating extension types.
```

```
If ``convert_integer`` is also True, preference will be given to integer dtypes if the floats can be faithfully casted to integers.
```

```
.. versionadded:: 1.2.0
```

#### Returns

```
Series or DataFrame
```

```
Copy of input object with new dtype.
```

#### See Also

```
infer_objects : Infer dtypes of objects.
```

```
to_datetime : Convert argument to datetime.
```

```
to_timedelta : Convert argument to timedelta.
```

```
to_numeric : Convert argument to a numeric type.
```

#### Notes

By default, ``convert\_dtypes`` will attempt to convert a Series (or each Series in a DataFrame) to dtypes that support ``pd.NA``. By using the options

``convert\_string``, ``convert\_integer``, ``convert\_boolean`` and ``convert\_floating``, it is possible to turn off individual conversions to ``StringDtype``, the integer extension types, ``BooleanDtype`` or floating extension types, respectively.

For object-dtyped columns, if ``infer\_objects`` is ``True``, use the inference rules as during normal Series/DataFrame construction. Then, if possible, convert to ``StringDtype``, ``BooleanDtype`` or an appropriate integer or floating extension type, otherwise leave as ``object``.

e.

If the dtype is integer, convert to an appropriate integer extension typ

If the dtype is numeric, and consists of all integers, convert to an appropriate integer extension type. Otherwise, convert to an appropriate floating extension type.

.. versionchanged:: 1.2

Starting with pandas 1.2, this method also converts float columns to the nullable floating extension type.

ts

In the future, as new dtypes are added that support ``pd.NA``, the result of this method will change to support those new dtypes.

Examples

```
>>> df = pd.DataFrame(
... {
... "a": pd.Series([1, 2, 3], dtype=np.dtype("int32")),
... "b": pd.Series(["x", "y", "z"], dtype=np.dtype("O")),
... "c": pd.Series([True, False, np.nan], dtype=np.dtype("O")),
... "d": pd.Series(["h", "i", np.nan], dtype=np.dtype("O")),
... "e": pd.Series([10, np.nan, 20], dtype=np.dtype("float")),
... "f": pd.Series([np.nan, 100.5, 200], dtype=np.dtype("floo
t"))),
... }
...)
```

Start with a DataFrame with default dtypes.

```
>>> df
 a b c d e f
0 1 x True h 10.0 NaN
1 2 y False i NaN 100.5
2 3 z NaN NaN 20.0 200.0
```

```
>>> df.dtypes
a int32
b object
c object
d object
e float64
f float64
dtype: object
```

Convert the DataFrame to use best possible dtypes.

```
>>> dfn = df.convert_dtypes()
>>> dfn
 a b c d e f
0 1 x True h 10 <NA>
1 2 y False i <NA> 100.5
2 3 z <NA> <NA> 20 200.0

>>> dfn.dtypes
a Int32
b string
c boolean
d string
e Int64
f Float64
dtype: object
```

Start with a Series of strings and missing data represented by ``np.nan`

```
>>> s = pd.Series(["a", "b", np.nan])
>>> s
0 a
1 b
2 NaN
dtype: object
```

Obtain a Series with dtype ``StringDtype``.

```
>>> s.convert_dtypes()
0 a
1 b
2 <NA>
dtype: string
```

`copy(self: 'FrameOrSeries', deep: 'bool_t' = True) -> 'FrameOrSeries'`  
Make a copy of this object's indices and data.

When ``deep=True`` (default), a new object will be created with a copy of the calling object's data and indices. Modifications to the data or indices of the copy will not be reflected in the original object (see notes below).

When ``deep=False``, a new object will be created without copying the calling object's data or index (only references to the data and index are copied). Any changes to the data of the original will be reflected in the shallow copy (and vice versa).

#### Parameters

`deep` : bool, default True  
Make a deep copy, including a copy of the data and the indices.  
With ``deep=False`` neither the indices nor the data are copied.

#### Returns

`copy` : Series or DataFrame  
Object type matches caller.

#### Notes

When ``deep=True``, data is copied but actual Python objects will not be copied recursively, only the reference to the object. This is in contrast to `copy.deepcopy` in the Standard Library, which recursively copies object data (see examples below).

While ``Index`` objects are copied when ``deep=True``, the underlying numpy array is not copied for performance reasons. Since ``Index`` is immutable, the underlying data can be safely shared and a copy is not needed.

#### Examples

```
>>> s = pd.Series([1, 2], index=["a", "b"])
>>> s
a 1
b 2
dtype: int64

>>> s_copy = s.copy()
>>> s_copy
a 1
```

```
b 2
dtype: int64
```

**\*\*Shallow copy versus default (deep) copy:\*\***

```
>>> s = pd.Series([1, 2], index=["a", "b"])
>>> deep = s.copy()
>>> shallow = s.copy(deep=False)
```

Shallow copy shares data and index with original.

```
>>> s is shallow
False
>>> s.values is shallow.values and s.index is shallow.index
True
```

Deep copy has own copy of data and index.

```
>>> s is deep
False
>>> s.values is deep.values or s.index is deep.index
False
```

Updates to the data shared by shallow copy and original is reflected in both; deep copy remains unchanged.

```
>>> s[0] = 3
>>> shallow[1] = 4
>>> s
a 3
b 4
dtype: int64
>>> shallow
a 3
b 4
dtype: int64
>>> deep
a 1
b 2
dtype: int64
```

Note that when copying an object containing Python objects, a deep copy will copy the data, but will not do so recursively. Updating a nested data object will be reflected in the deep copy.

```
>>> s = pd.Series([[1, 2], [3, 4]])
>>> deep = s.copy()
>>> s[0][0] = 10
>>> s
0 [10, 2]
1 [3, 4]
dtype: object
>>> deep
0 [10, 2]
1 [3, 4]
dtype: object
```

```
describe(self: 'FrameOrSeries', percentiles=None, include=None, exclude=None,
e, datetime_is_numeric=False) -> 'FrameOrSeries'
Generate descriptive statistics.
```

Descriptive statistics include those that summarize the central tendency, dispersion and shape of a dataset's distribution, excluding ``NaN`` values.

Analyzes both numeric and object series, as well as ``DataFrame`` column sets of mixed data types. The output will vary depending on what is provided. Refer to the notes below for more detail.

### Parameters

**percentiles** : list-like of numbers, optional  
The percentiles to include in the output. All should fall between 0 and 1. The default is ``[.25, .5, .75]``, which returns the 25th, 50th, and 75th percentiles.

**include** : 'all', list-like of dtypes or None (default), optional  
A white list of data types to include in the result. Ignored for ``Series``. Here are the options:

- 'all' : All columns of the input will be included in the output.
- A list-like of dtypes : Limits the results to the provided data types.  
To limit the result to numeric types submit ``numpy.number``. To limit it instead to object columns submit the ``numpy.object`` data type. Strings can also be used in the style of ``select\_dtypes`` (e.g. ``df.describe(include=['O'])``). To select pandas categorical columns, use ``'category'``
- None (default) : The result will include all numeric columns.

**exclude** : list-like of dtypes or None (default), optional  
A black list of data types to omit from the result. Ignored for ``Series``. Here are the options:

- A list-like of dtypes : Excludes the provided data types from the result. To exclude numeric types submit ``numpy.number``. To exclude object columns submit the data type ``numpy.object``. Strings can also be used in the style of ``select\_dtypes`` (e.g. ``df.describe(include=['O'])``). To exclude pandas categorical columns, use ``'category'``
- None (default) : The result will exclude nothing.

**datetime\_is\_numeric** : bool, default False  
Whether to treat datetime dtypes as numeric. This affects statistics calculated for the column. For DataFrame input, this also controls whether datetime columns are included by default.

.. versionadded:: 1.1.0

### Returns

Series or DataFrame

Summary statistics of the Series or Dataframe provided.

### See Also

**DataFrame.count**: Count number of non-NA/null observations.  
**DataFrame.max**: Maximum of the values in the object.  
**DataFrame.min**: Minimum of the values in the object.  
**DataFrame.mean**: Mean of the values.  
**DataFrame.std**: Standard deviation of the observations.  
**DataFrame.select\_dtypes**: Subset of a DataFrame including/excluding columns based on their dtype.

### Notes

For numeric data, the result's index will include ``count``, ``mean``, ``std``, ``min``, ``max`` as well as lower, ``50`` and upper percentiles. By default the lower percentile is ``25`` and the upper percentile is ``75``. The ``50`` percentile is the

same as the median.

For object data (e.g. strings or timestamps), the result's index will include ``count``, ``unique``, ``top``, and ``freq``. The ``top`` is the most common value. The ``freq`` is the most common value's frequency. Timestamps also include the ``first`` and ``last`` items.

If multiple object values have the highest count, then the ``count`` and ``top`` results will be arbitrarily chosen from among those with the highest count.

For mixed data types provided via a ``DataFrame``, the default is to return only an analysis of numeric columns. If the dataframe consists only of object and categorical data without any numeric columns, the default is to return an analysis of both the object and categorical columns. If ``include='all'`` is provided as an option, the result will include a union of attributes of each type.

The ``include`` and ``exclude`` parameters can be used to limit which columns in a ``DataFrame`` are analyzed for the output. The parameters are ignored when analyzing a ``Series``.

### Examples

Describing a numeric ``Series``.

```
>>> s = pd.Series([1, 2, 3])
>>> s.describe()
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
dtype: float64
```

Describing a categorical ``Series``.

```
>>> s = pd.Series(['a', 'a', 'b', 'c'])
>>> s.describe()
count 4
unique 3
top a
freq 2
dtype: object
```

Describing a timestamp ``Series``.

```
>>> s = pd.Series([
... np.datetime64("2000-01-01"),
... np.datetime64("2010-01-01"),
... np.datetime64("2010-01-01")
...])
>>> s.describe(datetime_is_numeric=True)
count 3
mean 2006-09-01 08:00:00
min 2000-01-01 00:00:00
25% 2004-12-31 12:00:00
50% 2010-01-01 00:00:00
75% 2010-01-01 00:00:00
max 2010-01-01 00:00:00
dtype: object
```



Describing a ``DataFrame``. By default only numeric fields are returned.

```
>>> df = pd.DataFrame({'categorical': pd.Categorical(['d','e','f']),
... 'numeric': [1, 2, 3],
... 'object': ['a', 'b', 'c']}
... })
>>> df.describe()
 numeric
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
```

Describing all columns of a ``DataFrame`` regardless of data type.

```
>>> df.describe(include='all') # doctest: +SKIP
 categorical numeric object
count 3 3.0 3
unique 3 NaN 3
top f NaN a
freq 1 NaN 1
mean NaN 2.0 NaN
std NaN 1.0 NaN
min NaN 1.0 NaN
25% NaN 1.5 NaN
50% NaN 2.0 NaN
75% NaN 2.5 NaN
max NaN 3.0 NaN
```

Describing a column from a ``DataFrame`` by accessing it as an attribute.

```
>>> df.numeric.describe()
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
Name: numeric, dtype: float64
```

Including only numeric columns in a ``DataFrame`` description.

```
>>> df.describe(include=[np.number])
 numeric
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
```

Including only string columns in a ``DataFrame`` description.

```
>>> df.describe(include=[object]) # doctest: +SKIP
 object
```

```
count 3
unique 3
top a
freq 1
```

Including only categorical columns from a ``DataFrame`` description.

```
>>> df.describe(include=['category'])
 categorical
count 3
unique 3
top d
freq 1
```

Excluding numeric columns from a ``DataFrame`` description.

```
>>> df.describe(exclude=[np.number]) # doctest: +SKIP
 categorical object
count 3 3
unique 3 3
top f a
freq 1 1
```

Excluding object columns from a ``DataFrame`` description.

```
>>> df.describe(exclude=[object]) # doctest: +SKIP
 categorical numeric
count 3 3.0
unique 3 NaN
top f NaN
freq 1 NaN
mean NaN 2.0
std NaN 1.0
min NaN 1.0
25% NaN 1.5
50% NaN 2.0
75% NaN 2.5
max NaN 3.0
```

`droplevel(self: 'FrameOrSeries', level, axis=0) -> 'FrameOrSeries'`  
Return DataFrame with requested index / column level(s) removed.

.. versionadded:: 0.24.0

Parameters

`level` : int, str, or list-like

If a string is given, must be the name of a level

If list-like, elements must be names or positional indexes of levels.

`axis` : {0 or 'index', 1 or 'columns'}, default 0

Axis along which the level(s) is removed:

\* 0 or 'index': remove level(s) in column.

\* 1 or 'columns': remove level(s) in row.

Returns

DataFrame

DataFrame with requested index / column level(s) removed.

Examples

```
>>> df = pd.DataFrame([
```

```

... [1, 2, 3, 4],
... [5, 6, 7, 8],
... [9, 10, 11, 12]
...]).set_index([0, 1]).rename_axis(['a', 'b'])

>>> df.columns = pd.MultiIndex.from_tuples([
... ('c', 'e'), ('d', 'f')
...], names=['level_1', 'level_2'])

>>> df
level_1 c d
level_2 e f
a b
1 2 3 4
5 6 7 8
9 10 11 12

>>> df.droplevel('a')
level_1 c d
level_2 e f
b
2 3 4
6 7 8
10 11 12

>>> df.droplevel('level_2', axis=1)
level_1 c d
a b
1 2 3 4
5 6 7 8
9 10 11 12

```

`equals(self, other: 'object') -> 'bool'`

Test whether two objects contain the same elements.

This function allows two Series or DataFrames to be compared against each other to see if they have the same shape and elements. NaNs in the same location are considered equal.

The row/column index do not need to have the same type, as long as the values are considered equal. Corresponding columns must be of the same dtype.

#### Parameters

`other` : Series or DataFrame

The other Series or DataFrame to be compared with the first.

#### Returns

`bool`

True if all elements are the same in both objects, False otherwise.

#### See Also

`Series.eq` : Compare two Series objects of the same length and return a Series where each element is True if the element in each Series is equal, False otherwise.

`DataFrame.eq` : Compare two DataFrame objects of the same shape and return a DataFrame where each element is True if the respective element in each DataFrame is equal, False otherwise.

`testing.assert_series_equal` : Raises an `AssertionError` if left and right are not equal. Provides an easy interface to ignore inequality in dtypes, indexes and precision among others.

testing.assert\_frame\_equal : Like assert\_series\_equal, but targets DataFrames.

numpy.array\_equal : Return True if two arrays have the same shape and elements, False otherwise.

#### Examples

```
>>> df = pd.DataFrame({1: [10], 2: [20]})
```

```
>>> df
 1 2
0 10 20
```

DataFrames df and exactly\_equal have the same types and values for their elements and column labels, which will return True.

```
>>> exactly_equal = pd.DataFrame({1: [10], 2: [20]})
```

```
>>> exactly_equal
```

```
 1 2
0 10 20
```

```
>>> df.equals(exactly_equal)
```

```
True
```

DataFrames df and different\_column\_type have the same element types and values, but have different types for the column labels, which will still return True.

```
>>> different_column_type = pd.DataFrame({1.0: [10], 2.0: [20]})
```

```
>>> different_column_type
```

```
 1.0 2.0
0 10 20
```

```
>>> df.equals(different_column_type)
```

```
True
```

DataFrames df and different\_data\_type have different types for the same values for their elements, and will return False even though their column labels are the same values and types.

```
>>> different_data_type = pd.DataFrame({1: [10.0], 2: [20.0]})
```

```
>>> different_data_type
```

```
 1 2
0 10.0 20.0
```

```
>>> df.equals(different_data_type)
```

```
False
```

```
ewm(self, com: 'Optional[float]' = None, span: 'Optional[float]' = None, halflife: 'Optional[Union[float, TimedeltaConvertibleTypes]]' = None, alpha: 'Optional[float]' = None, min_periods: 'int' = 0, adjust: 'bool_t' = True, ignore_na: 'bool_t' = False, axis: 'Axis' = 0, times: 'Optional[Union[str, np.ndarray, FrameOrSeries]]' = None) -> 'ExponentialMovingWindow'
```

Provide exponential weighted (EW) functions.

```
cov()``. Available EW functions: ``mean()`, ``var()`, ``std()`, ``corr()`, ``
```

t be

Exactly one parameter: ``com``, ``span``, ``halflife``, or ``alpha`` must be provided.

#### Parameters

com : float, optional

Specify decay in terms of center of mass,

:math:\alpha = 1 / (1 + com), for :math:com \geq 0.

span : float, optional

Specify decay in terms of span,

```

 :math:\alpha = 2 / (\text{span} + 1), for $\text{span} \geq 1$.
halflife : float, str, timedelta, optional
 Specify decay in terms of half-life,
 :math:\alpha = 1 - \exp(-\ln(2) / \text{halflife}), for
 :math:\text{halflife} > 0.

 If ``times`` is specified, the time unit (str or timedelta) over which
 observation decays to half its value. Only applicable to ``mean()``
 and halflife value will not apply to the other functions.

 .. versionadded:: 1.1.0

alpha : float, optional
 Specify smoothing factor α directly,
 :math:0 < \alpha \leq 1.
min_periods : int, default 0
 Minimum number of observations in window required to have a value
 (otherwise result is NA).
adjust : bool, default True
 Divide by decaying adjustment factor in beginning periods to account
 for imbalance in relative weightings (viewing EWMA as a moving average).

 - When ``adjust=True`` (default), the EW function is calculated using
 weights
 :math:w_i = (1 - \alpha)^i. For example, the EW moving average of
 the series
 $[x_0, x_1, \dots, x_t]$ would be:

 .. math::
 y_t = \frac{x_t + (1 - \alpha)x_{t-1} + (1 - \alpha)^2 x_{t-2} + \dots + (1 - \alpha)^t x_0}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots + (1 - \alpha)^t}

 - When ``adjust=False``, the exponentially weighted function is calculated
 recursively:

 .. math::
 \begin{split}
 y_0 &= x_0 \\
 y_t &= (1 - \alpha) y_{t-1} + \alpha x_t,
 \end{split}

 ignore_na : bool, default False
 Ignore missing values when calculating weights; specify ``True`` to
 reproduce pre-0.15.0 behavior.

 - When ``ignore_na=False`` (default), weights are based on absolute
 positions.
 For example, the weights of x_0 and x_2 used in calculating
 the final weighted average of $[x_0, \text{None}, x_2]$ are
 :math:(1 - \alpha)^2 and 1 if ``adjust=True``, and
 :math:(1 - \alpha)^2 and α if ``adjust=False``.

 - When ``ignore_na=True`` (reproducing pre-0.15.0 behavior), weights
 are based
 on relative positions. For example, the weights of x_0 and
 x_2
 used in calculating the final weighted average of
 $[x_0, \text{None}, x_2]$ are $1 - \alpha$ and 1
 if

```

```

 ``adjust=True``, and :math:`1-\alpha` and :math:`\alpha` if ``adjust=False``.
 axis : {0, 1}, default 0
 The axis to use. The value 0 identifies the rows, and 1 identifies the columns.
 times : str, np.ndarray, Series, default None
 .. versionadded:: 1.1.0
 Times corresponding to the observations. Must be monotonically increasing and
 ``datetime64[ns]`` dtype.
 If str, the name of the column in the DataFrame representing the times.
 If 1-D array like, a sequence with the same shape as the observations.

 Only applicable to ``mean()``.

Returns

DataFrame
 A Window sub-classed for the particular operation.

See Also

rolling : Provides rolling window calculations.
expanding : Provides expanding transformations.

Notes

More details can be found at:
:ref:`Exponentially weighted windows <window.exponentially_weighted>`.

Examples

>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
>>> df
 B
0 0.0
1 1.0
2 2.0
3 NaN
4 4.0

>>> df.ewm(com=0.5).mean()
 B
0 0.000000
1 0.750000
2 1.615385
3 1.615385
4 3.670213

Specifying ``times`` with a timedelta ``halflife`` when computing mean.
>>> times = ['2020-01-01', '2020-01-03', '2020-01-10', '2020-01-15', '2020-01-17']
>>> df.ewm(halflife='4 days', times=pd.DatetimeIndex(times)).mean()
 B
0 0.000000
1 0.585786
2 1.523889

```

```

3 1.523889
4 3.233686
expanding(self, min_periods: 'int' = 1, center: 'Optional[bool_t]' = None, a
xis: 'Axis' = 0) -> 'Expanding'
 Provide expanding transformations.

 Parameters

 min_periods : int, default 1
 Minimum number of observations in window required to have a value
 (otherwise result is NA).
 center : bool, default False
 Set the labels at the center of the window.
 axis : int or str, default 0

 Returns

 a Window sub-classed for the particular operation

 See Also

 rolling : Provides rolling window calculations.
 ewm : Provides exponential weighted functions.

 Notes

 By default, the result is set to the right edge of the window. This can
 be changed to the center of the window by setting ``center=True``.

 Examples

 >>> df = pd.DataFrame({"B": [0, 1, 2, np.nan, 4]})
 >>> df
 B
 0 0.0
 1 1.0
 2 2.0
 3 NaN
 4 4.0

 >>> df.expanding(2).sum()
 B
 0 NaN
 1 1.0
 2 3.0
 3 3.0
 4 7.0

 ffill(self: 'FrameOrSeries', axis=None, inplace: 'bool_t' = False, limit=None,
e, downcast=None) -> 'Optional[FrameOrSeries]'
 Synonym for :meth:`DataFrame.fillna` with ``method='ffill'``.

 Returns

 Series/DataFrame or None
 Object with missing values filled or None if ``inplace=True``.

 filter(self: 'FrameOrSeries', items=None, like: 'Optional[str]' = None, rege
x: 'Optional[str]' = None, axis=None) -> 'FrameOrSeries'
 Subset the dataframe rows or columns according to the specified index la
bels.

 Note that this routine does not filter a dataframe on its

```

contents. The filter is applied to the labels of the index.

#### Parameters

items : list-like  
Keep labels from axis which are in items.

like : str  
Keep labels from axis for which "like in label == True".

regex : str (regular expression)  
Keep labels from axis for which re.search(regex, label) == True.

axis : {0 or 'index', 1 or 'columns', None}, default None  
The axis to filter on, expressed either as an index (int) or axis name (str). By default this is the info axis, 'index' for Series, 'columns' for DataFrame.

#### Returns

same type as input object

#### See Also

DataFrame.loc : Access a group of rows and columns by label(s) or a boolean array.

#### Notes

The ``items``, ``like``, and ``regex`` parameters are enforced to be mutually exclusive.

``axis`` defaults to the info axis that is used when indexing with ``[]``.

#### Examples

```
>>> df = pd.DataFrame(np.array([[1, 2, 3], [4, 5, 6]]),
... index=['mouse', 'rabbit'],
... columns=['one', 'two', 'three'])
>>> df
```

```
 one two three
mouse 1 2 3
rabbit 4 5 6
```

```
>>> # select columns by name
>>> df.filter(items=['one', 'three'])
```

```
 one three
mouse 1 3
rabbit 4 6
```

```
>>> # select columns by regular expression
>>> df.filter(regex='e$', axis=1)
```

```
 one three
mouse 1 3
rabbit 4 6
```

```
>>> # select rows containing 'bbi'
>>> df.filter(like='bbi', axis=0)
```

```
 one two three
rabbit 4 5 6
```

first(self: 'FrameOrSeries', offset) -> 'FrameOrSeries'

Select initial periods of time series data based on a date offset.

When having a DataFrame with dates as index, this function can select the first few rows based on a date offset.



onth.

## Parameters

offset : str, DateOffset or dateutil.relativedelta

The offset length of the data that will be selected. For instance, '1M' will display all the rows having their index within the first m

## Returns

Series or DataFrame

A subset of the caller.

## Raises

TypeError

If the index is not a :class:`DatetimeIndex`

## See Also

last : Select final periods of time series based on a date offset.

at\_time : Select values at a particular time of the day.

between\_time : Select values between particular times of the day.

## Examples

```
>>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
```

```
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
```

```
>>> ts
```

|            | A |
|------------|---|
| 2018-04-09 | 1 |
| 2018-04-11 | 2 |
| 2018-04-13 | 3 |
| 2018-04-15 | 4 |

Get the rows for the first 3 days:

```
>>> ts.first('3D')
```

|            | A |
|------------|---|
| 2018-04-09 | 1 |
| 2018-04-11 | 2 |

Notice the data for 3 first calendar days were returned, not the first 3 days observed in the dataset, and therefore data for 2018-04-13 was not returned.

first\_valid\_index(self)

Return index for first non-NA/null value.

## Returns

scalar : type of index

## Notes

If all elements are non-NA/null, returns None.

Also returns None for empty Series/DataFrame.

get(self, key, default=None)

Get item from object for given key (ex: DataFrame column).

Returns default value if not found.

## Parameters

key : object

## Returns

value : same type as items contained in object

head(self: 'FrameOrSeries', n: 'int' = 5) -> 'FrameOrSeries'  
Return the first `n` rows.

This function returns the first `n` rows for the object based on position. It is useful for quickly testing if your object has the right type of data in it.

For negative values of `n`, this function returns all rows except the last `n` rows, equivalent to ``df[:-n]``.

## Parameters

n : int, default 5  
Number of rows to select.

## Returns

same type as caller  
The first `n` rows of the caller object.

## See Also

DataFrame.tail: Returns the last `n` rows.

## Examples

```
>>> df = pd.DataFrame({'animal': ['alligator', 'bee', 'falcon', 'lion',
... 'monkey', 'parrot', 'shark', 'whale', 'zebra']})
>>> df
```

```
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot
6 shark
7 whale
8 zebra
```

Viewing the first 5 lines

```
>>> df.head()
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
```

Viewing the first `n` lines (three in this case)

```
>>> df.head(3)
 animal
0 alligator
1 bee
2 falcon
```

For negative values of `n`

```
>>> df.head(-3)
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot
```

`infer_objects(self: 'FrameOrSeries') -> 'FrameOrSeries'`

Attempt to infer better dtypes for object columns.

Attempts soft conversion of object-dtyped columns, leaving non-object and unconvertible columns unchanged. The inference rules are the same as during normal Series/DataFrame construction.

Returns

converted : same type as input object

See Also

`to_datetime` : Convert argument to datetime.

`to_timedelta` : Convert argument to timedelta.

`to_numeric` : Convert argument to numeric type.

`convert_dtypes` : Convert argument to best possible dtype.

Examples

```
>>> df = pd.DataFrame({"A": ["a", 1, 2, 3]})
```

```
>>> df = df.iloc[1:]
```

```
>>> df
```

```
 A
1 1
2 2
3 3
```

```
>>> df.dtypes
```

```
A object
```

```
dtype: object
```

```
>>> df.infer_objects().dtypes
```

```
A int64
```

```
dtype: object
```

`interpolate(self: 'FrameOrSeries', method: 'str' = 'linear', axis: 'Axis' = 0, limit: 'Optional[int]' = None, inplace: 'bool_t' = False, limit_direction: 'Optional[str]' = None, limit_area: 'Optional[str]' = None, downcast: 'Optional[str]' = None, **kwargs) -> 'Optional[FrameOrSeries]'`

Fill NaN values using an interpolation method.

Please note that only ``method='linear'`` is supported for DataFrame/Series with a MultiIndex.

Parameters

`method` : str, default 'linear'

Interpolation technique to use. One of:

- \* 'linear': Ignore the index and treat the values as equally spaced. This is the only method supported on MultiIndexes.
- \* 'time': Works on daily and higher resolution data to interpolate given length of interval.
- \* 'index', 'values': use the actual numerical values of the index.

```

* 'pad': Fill in NaNs using existing values.
* 'nearest', 'zero', 'slinear', 'quadratic', 'cubic', 'spline',
 'barycentric', 'polynomial': Passed to
 `scipy.interpolate.interpld`. These methods use the numerical
 values of the index. Both 'polynomial' and 'spline' require that
 you also specify an `order` (int), e.g.
 ``df.interpolate(method='polynomial', order=5)``.
* 'krogh', 'piecewise_polynomial', 'spline', 'pchip', 'akima',
 'cubicspline': Wrappers around the SciPy interpolation methods of
 similar names. See `Notes`.
* 'from_derivatives': Refers to
 `scipy.interpolate.BPoly.from_derivatives` which
 replaces 'piecewise_polynomial' interpolation method in
 scipy 0.18.

axis : {{0 or 'index', 1 or 'columns', None}}, default None
 Axis to interpolate along.
limit : int, optional
 Maximum number of consecutive NaNs to fill. Must be greater than
 0.
inplace : bool, default False
 Update the data in place if possible.
limit_direction : {{'forward', 'backward', 'both'}}, Optional
 Consecutive NaNs will be filled in this direction.

 If limit is specified:
 * If 'method' is 'pad' or 'ffill', 'limit_direction' must be 'fo
 rward'.
 * If 'method' is 'backfill' or 'bfill', 'limit_direction' must b
 e
 'backwards'.

 If 'limit' is not specified:
 * If 'method' is 'backfill' or 'bfill', the default is 'backwar
 d'.
 * else the default is 'forward'

.. versionchanged:: 1.1.0
 raises ValueError if `limit_direction` is 'forward' or 'both' an
 d
 method is 'backfill' or 'bfill'.
 raises ValueError if `limit_direction` is 'backward' or 'both' a
 nd
 method is 'pad' or 'ffill'.

limit_area : {{`None`, 'inside', 'outside'}}, default None
 If limit is specified, consecutive NaNs will be filled with this
 restriction.

 * ``None``: No fill restriction.
 * 'inside': Only fill NaNs surrounded by valid values
 (interpolate).
 * 'outside': Only fill NaNs outside valid values (extrapolate).

downcast : optional, 'infer' or None, defaults to None
 Downcast dtypes if possible.
`**kwargs` : optional
 Keyword arguments to pass on to the interpolating function.

Returns

Series or DataFrame or None
 Returns the same object type as the caller, interpolated at
 some or all ``NaN`` values or None if ``inplace=True``.

```

## See Also

-----  
 fillna : Fill missing values using different methods.  
 scipy.interpolate.Akima1DInterpolator : Piecewise cubic polynomials (Akima interpolator).  
 scipy.interpolate.BPoly.from\_derivatives : Piecewise polynomial in the Bernstein basis.  
 scipy.interpolate.interp1d : Interpolate a 1-D function.  
 scipy.interpolate.KroghInterpolator : Interpolate polynomial (Krogh interpolator).  
 scipy.interpolate.PchipInterpolator : PCHIP 1-d monotonic cubic interpolation.  
 scipy.interpolate.CubicSpline : Cubic spline data interpolator.

## Notes

-----  
 The 'krogh', 'piecewise\_polynomial', 'spline', 'pchip' and 'akima' methods are wrappers around the respective SciPy implementations of similar names. These use the actual numerical values of the index. For more information on their behavior, see the `SciPy documentation` <<https://docs.scipy.org/doc/scipy/reference/interpolate.html#univariate-interpolation>>` and `SciPy tutorial` <<https://docs.scipy.org/doc/scipy/reference/tutorial/interpolate.html>>`\_

## Examples

-----  
 Filling in ``NaN`` in a :class:`~pandas.Series` via linear interpolation.

```
>>> s = pd.Series([0, 1, np.nan, 3])
>>> s
0 0.0
1 1.0
2 NaN
3 3.0
dtype: float64
>>> s.interpolate()
0 0.0
1 1.0
2 2.0
3 3.0
dtype: float64
```

Filling in ``NaN`` in a Series by padding, but filling at most two consecutive ``NaN`` at a time.

```
>>> s = pd.Series([np.nan, "single_one", np.nan,
... "fill_two_more", np.nan, np.nan, np.nan,
... 4.71, np.nan])
>>> s
0 NaN
1 single_one
2 NaN
3 fill_two_more
4 NaN
5 NaN
6 NaN
7 4.71
8 NaN
dtype: object
>>> s.interpolate(method='pad', limit=2)
0 NaN
```

```

1 single_one
2 single_one
3 fill_two_more
4 fill_two_more
5 fill_two_more
6 NaN
7 4.71
8 4.71
dtype: object

```

Filling in ``NaN`` in a Series via polynomial interpolation or splines: Both 'polynomial' and 'spline' methods require that you also specify an ``order`` (int).

```

>>> s = pd.Series([0, 2, np.nan, 8])
>>> s.interpolate(method='polynomial', order=2)
0 0.000000
1 2.000000
2 4.666667
3 8.000000
dtype: float64

```

Fill the DataFrame forward (that is, going down) along each column using linear interpolation.

Note how the last entry in column 'a' is interpolated differently, because there is no entry after it to use for interpolation. Note how the first entry in column 'b' remains ``NaN``, because there is no entry before it to use for interpolation.

```

>>> df = pd.DataFrame([(0.0, np.nan, -1.0, 1.0),
... (np.nan, 2.0, np.nan, np.nan),
... (2.0, 3.0, np.nan, 9.0),
... (np.nan, 4.0, -4.0, 16.0)],
... columns=list('abcd'))
>>> df
 a b c d
0 0.0 NaN -1.0 1.0
1 NaN 2.0 NaN NaN
2 2.0 3.0 NaN 9.0
3 NaN 4.0 -4.0 16.0
>>> df.interpolate(method='linear', limit_direction='forward', axis=0)
 a b c d
0 0.0 NaN -1.0 1.0
1 1.0 2.0 -2.0 5.0
2 2.0 3.0 -3.0 9.0
3 2.0 4.0 -4.0 16.0

```

Using polynomial interpolation.

```

>>> df['d'].interpolate(method='polynomial', order=2)
0 1.0
1 4.0
2 9.0
3 16.0
Name: d, dtype: float64

```

keys(self)

Get the 'info axis' (see Indexing for more).

This is index for Series, columns for DataFrame.

Returns

-----  
Index

Info axis.

```
last(self: 'FrameOrSeries', offset) -> 'FrameOrSeries'
```

Select final periods of time series data based on a date offset.

When having a DataFrame with dates as index, this function can select the last few rows based on a date offset.

Parameters

`offset` : str, DateOffset, dateutil.relativedelta  
The offset length of the data that will be selected. For instance, '3D' will display all the rows having their index within the last 3

days.

Returns

Series or DataFrame  
A subset of the caller.

Raises

TypeError  
If the index is not a :class:`DatetimeIndex`

See Also

`first` : Select initial periods of time series based on a date offset.  
`at_time` : Select values at a particular time of the day.  
`between_time` : Select values between particular times of the day.

Examples

```
>>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
>>> ts
```

```
 A
2018-04-09 1
2018-04-11 2
2018-04-13 3
2018-04-15 4
```

Get the rows for the last 3 days:

```
>>> ts.last('3D')
 A
2018-04-13 3
2018-04-15 4
```

Notice the data for 3 last calendar days were returned, not the last 3 observed days in the dataset, and therefore data for 2018-04-11 was not returned.

```
last_valid_index(self)
```

Return index for last non-NA/null value.

Returns

scalar : type of index

Notes

If all elements are non-NA/null, returns None.  
Also returns None for empty Series/DataFrame.

```
| mask(self, cond, other=nan, inplace=False, axis=None, level=None, errors='raise', try_cast=False)
```

```
| Replace values where the condition is True.
```

#### Parameters

```
| cond : bool Series/DataFrame, array-like, or callable
| Where `cond` is False, keep the original value. Where
| True, replace with corresponding value from `other`.
| If `cond` is callable, it is computed on the Series/DataFrame and
| should return boolean Series/DataFrame or array. The callable must
| not change input Series/DataFrame (though pandas doesn't check it).
| other : scalar, Series/DataFrame, or callable
| Entries where `cond` is True are replaced with
| corresponding value from `other`.
| If other is callable, it is computed on the Series/DataFrame and
| should return scalar or Series/DataFrame. The callable must not
| change input Series/DataFrame (though pandas doesn't check it).
| inplace : bool, default False
| Whether to perform the operation in place on the data.
| axis : int, default None
| Alignment axis if needed.
| level : int, default None
| Alignment level if needed.
| errors : str, {'raise', 'ignore'}, default 'raise'
| Note that currently this parameter won't affect
| the results and will always coerce to a suitable dtype.
|
| - 'raise' : allow exceptions to be raised.
| - 'ignore' : suppress exceptions. On error return original object.
```

```
| try_cast : bool, default False
| Try to cast the result back to the input type (if possible).
```

#### Returns

```
| Same type as caller or None if ``inplace=True``.
```

#### See Also

```
| :func:`DataFrame.where` : Return an object of same shape as
| self.
```

#### Notes

```
| The mask method is an application of the if-then idiom. For each
| element in the calling DataFrame, if ``cond`` is ``False`` the
| element is used; otherwise the corresponding element from the DataFrame
| ``other`` is used.
```

```
| The signature for :func:`DataFrame.where` differs from
| :func:`numpy.where`. Roughly ``df1.where(m, df2)`` is equivalent to
| ``np.where(m, df1, df2)``.
```

```
| For further details and examples see the ``mask`` documentation in
| :ref:`indexing <indexing.where_mask>`.
```

#### Examples

```
| >>> s = pd.Series(range(5))
| >>> s.where(s > 0)
| 0 NaN
| 1 1.0
| 2 2.0
| 3 3.0
```



```

4 4.0
dtype: float64
>>> s.mask(s > 0)
0 0.0
1 NaN
2 NaN
3 NaN
4 NaN
dtype: float64

>>> s.where(s > 1, 10)
0 10
1 10
2 2
3 3
4 4
dtype: int64
>>> s.mask(s > 1, 10)
0 0
1 1
2 10
3 10
4 10
dtype: int64

>>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A', 'B'])
>>> df
 A B
0 0 1
1 2 3
2 4 5
3 6 7
4 8 9
>>> m = df % 3 == 0
>>> df.where(m, -df)
 A B
0 0 -1
1 -2 3
2 -4 -5
3 6 -7
4 -8 9
>>> df.where(m, -df) == np.where(m, df, -df)
 A B
0 True True
1 True True
2 True True
3 True True
4 True True
>>> df.where(m, -df) == df.mask(~m, -df)
 A B
0 True True
1 True True
2 True True
3 True True
4 True True

```

```

pad = ffill(self: 'FrameOrSeries', axis=None, inplace: 'bool_t' = False, limit=None, downcast=None) -> 'Optional[FrameOrSeries]'

```

```

pct_change(self: 'FrameOrSeries', periods=1, fill_method='pad', limit=None, freq=None, **kwargs) -> 'FrameOrSeries'

```

Percentage change between the current and a prior element.

Computes the percentage change from the immediately previous row by default. This is useful in comparing the percentage of change in a time

series of elements.

#### Parameters

`periods` : int, default 1  
Periods to shift for forming percent change.

`fill_method` : str, default 'pad'  
How to handle NAs before computing percent changes.

`limit` : int, default None  
The number of consecutive NAs to fill before stopping.

`freq` : DateOffset, timedelta, or str, optional  
Increment to use from time series API (e.g. 'M' or BDay()).

`**kwargs`  
Additional keyword arguments are passed into `'DataFrame.shift'` or `'Series.shift'`.

#### Returns

`chg` : Series or DataFrame  
The same type as the calling object.

#### See Also

`Series.diff` : Compute the difference of two elements in a Series.

`DataFrame.diff` : Compute the difference of two elements in a DataFrame.

`Series.shift` : Shift the index by some number of periods.

`DataFrame.shift` : Shift the index by some number of periods.

#### Examples

**\*\*Series\*\***

```
>>> s = pd.Series([90, 91, 85])
>>> s
0 90
1 91
2 85
dtype: int64
```

```
>>> s.pct_change()
0 NaN
1 0.011111
2 -0.065934
dtype: float64
```

```
>>> s.pct_change(periods=2)
0 NaN
1 NaN
2 -0.055556
dtype: float64
```

See the percentage change in a Series where filling NAs with last valid observation forward to next valid.

```
>>> s = pd.Series([90, 91, None, 85])
>>> s
0 90.0
1 91.0
2 NaN
3 85.0
dtype: float64

>>> s.pct_change(fill_method='ffill')
0 NaN
1 0.011111
```

```
2 0.000000
3 -0.065934
dtype: float64
```

```
DataFrame
```

Percentage change in French franc, Deutsche Mark, and Italian lira from 1980-01-01 to 1980-03-01.

```
>>> df = pd.DataFrame({
... 'FR': [4.0405, 4.0963, 4.3149],
... 'GR': [1.7246, 1.7482, 1.8519],
... 'IT': [804.74, 810.01, 860.13]},
... index=['1980-01-01', '1980-02-01', '1980-03-01'])
>>> df
```

|            | FR     | GR     | IT     |
|------------|--------|--------|--------|
| 1980-01-01 | 4.0405 | 1.7246 | 804.74 |
| 1980-02-01 | 4.0963 | 1.7482 | 810.01 |
| 1980-03-01 | 4.3149 | 1.8519 | 860.13 |

```
>>> df.pct_change()
 FR GR IT
1980-01-01 NaN NaN NaN
1980-02-01 0.013810 0.013684 0.006549
1980-03-01 0.053365 0.059318 0.061876
```

Percentage of change in GOOG and APPL stock volume. Shows computing the percentage change between columns.

```
>>> df = pd.DataFrame({
... '2016': [1769950, 30586265],
... '2015': [1500923, 40912316],
... '2014': [1371819, 41403351]},
... index=['GOOG', 'APPL'])
>>> df
```

|      | 2016     | 2015     | 2014     |
|------|----------|----------|----------|
| GOOG | 1769950  | 1500923  | 1371819  |
| APPL | 30586265 | 40912316 | 41403351 |

```
>>> df.pct_change(axis='columns')
 2016 2015 2014
GOOG NaN -0.151997 -0.086016
APPL NaN 0.337604 0.012002
```

```
pipe(self, func, *args, **kwargs)
Apply func(self, *args, **kwargs).
```

Parameters

```
func : function
 Function to apply to the Series/DataFrame.
 ``args``, and ``kwargs`` are passed into ``func``.
 Alternatively a ``(callable, data_keyword)`` tuple where
 ``data_keyword`` is a string indicating the keyword of
 ``callable`` that expects the Series/DataFrame.
args : iterable, optional
 Positional arguments passed into ``func``.
kwargs : mapping, optional
 A dictionary of keyword arguments passed into ``func``.
```

Returns

```
object : the return type of ``func``.
```

See Also

```

DataFrame.apply : Apply a function along input axis of DataFrame.
DataFrame.applymap : Apply a function elementwise on a whole DataFrame.
Series.map : Apply a mapping correspondence on a
 :class:`~pandas.Series`.

```

## Notes

Use ``.pipe`` when chaining together functions that expect Series, DataFrames or GroupBy objects. Instead of writing

```
>>> func(g(h(df), arg1=a), arg2=b, arg3=c) # doctest: +SKIP
```

You can write

```
>>> (df.pipe(h)
... .pipe(g, arg1=a)
... .pipe(func, arg2=b, arg3=c)
...) # doctest: +SKIP
```

If you have a function that takes the data as (say) the second argument, pass a tuple indicating which keyword expects the data. For example, suppose ``f`` takes its data as ``arg2``:

```
>>> (df.pipe(h)
... .pipe(g, arg1=a)
... .pipe((func, 'arg2'), arg1=a, arg3=c)
...) # doctest: +SKIP
```

```
rank(self: 'FrameOrSeries', axis=0, method: 'str' = 'average', numeric_only:
Optional[bool_t] = None, na_option: 'str' = 'keep', ascending: 'bool_t' = Tru
e, pct: 'bool_t' = False) -> 'FrameOrSeries'
Compute numerical data ranks (1 through n) along axis.
```

By default, equal values are assigned a rank that is the average of the ranks of those values.

## Parameters

```
axis : {0 or 'index', 1 or 'columns'}, default 0
 Index to direct ranking.
method : {'average', 'min', 'max', 'first', 'dense'}, default 'average'
 How to rank the group of records that have the same value (i.e. tie
```

s):

- \* average: average rank of the group
- \* min: lowest rank in the group
- \* max: highest rank in the group
- \* first: ranks assigned in order they appear in the array
- \* dense: like 'min', but rank always increases by 1 between groups.

```
numeric_only : bool, optional
 For DataFrame objects, rank only numeric columns if set to True.
na_option : {'keep', 'top', 'bottom'}, default 'keep'
 How to rank NaN values:
```

- \* keep: assign NaN rank to NaN values
- \* top: assign smallest rank to NaN values if ascending
- \* bottom: assign highest rank to NaN values if ascending.

```
ascending : bool, default True
 Whether or not the elements should be ranked in ascending order.
pct : bool, default False
 Whether or not to display the returned rankings in percentile
 form.
```

## Returns

same type as caller

Return a Series or DataFrame with data ranks as values.

## See Also

`core.groupby.GroupBy.rank` : Rank of values within each group.

## Examples

```
>>> df = pd.DataFrame(data={'Animal': ['cat', 'penguin', 'dog',
... 'spider', 'snake'],
... 'Number_legs': [4, 2, 4, 8, np.nan]})
```

```
>>> df
 Animal Number_legs
0 cat 4.0
1 penguin 2.0
2 dog 4.0
3 spider 8.0
4 snake NaN
```

The following example shows how the method behaves with the above parameters:

- \* `default_rank`: this is the default behaviour obtained without using any parameter.
- \* `max_rank`: setting `method = 'max'` the records that have the same values are ranked using the highest rank (e.g.: since 'cat' and 'dog' are both in the 2nd and 3rd position, rank 3 is assigned.)
- \* `NA_bottom`: choosing `na_option = 'bottom'`, if there are records with NaN values they are placed at the bottom of the ranking.
- \* `pct_rank`: when setting `pct = True`, the ranking is expressed as percentile rank.

```
>>> df['default_rank'] = df['Number_legs'].rank()
>>> df['max_rank'] = df['Number_legs'].rank(method='max')
>>> df['NA_bottom'] = df['Number_legs'].rank(na_option='bottom')
>>> df['pct_rank'] = df['Number_legs'].rank(pct=True)
>>> df
 Animal Number_legs default_rank max_rank NA_bottom pct_rank
0 cat 4.0 2.5 3.0 2.5 0.625
1 penguin 2.0 1.0 1.0 1.0 0.250
2 dog 4.0 2.5 3.0 2.5 0.625
3 spider 8.0 4.0 4.0 4.0 1.000
4 snake NaN NaN NaN 5.0 NaN
```

```
reindex_like(self: 'FrameOrSeries', other, method: 'Optional[str]' = None, copy: 'bool_t' = True, limit=None, tolerance=None) -> 'FrameOrSeries'
```

Return an object with matching indices as other object.

Conform the object to the same index on all axes. Optional filling logic, placing NaN in locations having no value in the previous index. A new object is produced unless the new index is equivalent to the current one and `copy=False`.

## Parameters

`other` : Object of the same data type

Its row and column indices are used to define the new indices of this object.

`method` : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}

Method to use for filling holes in reindexed DataFrame.

Please note: this is only applicable to DataFrames/Series with a

monotonically increasing/decreasing index.

- \* None (default): don't fill gaps
- \* pad / ffill: propagate last valid observation forward to next valid
- \* backfill / bfill: use next valid observation to fill gap
- \* nearest: use nearest valid observations to fill gap.

copy : bool, default True

Return a new object, even if the passed indexes are the same.

limit : int, default None

Maximum number of consecutive labels to fill for inexact matches.

tolerance : optional

Maximum distance between original and new labels for inexact matches. The values of the index at the matching locations must satisfy the equation ``abs(index[indexer] - target) <= tolerance``.

Tolerance may be a scalar value, which applies the same tolerance to all values, or list-like, which applies variable tolerance per element. List-like includes list, tuple, array, Series, and must be the same size as the index and its dtype must exactly match the index's type.

## Returns

Series or DataFrame

Same type as caller, but with changed indices on each axis.

## See Also

DataFrame.set\_index : Set row labels.

DataFrame.reset\_index : Remove row labels or move them to new columns.

DataFrame.reindex : Change to new indices or expand indices.

## Notes

Same as calling

```.reindex(index=other.index, columns=other.columns,...)``.

Examples

```
>>> df1 = pd.DataFrame([[24.3, 75.7, 'high'],
...                     [31, 87.8, 'high'],
...                     [22, 71.6, 'medium'],
...                     [35, 95, 'medium']],
...                     columns=['temp_celsius', 'temp_fahrenheit',
...                               'windspeed'],
...                     index=pd.date_range(start='2014-02-12',
...                                           end='2014-02-15', freq='D'))
```

```
>>> df1
              temp_celsius  temp_fahrenheit  windspeed
2014-02-12             24.3              75.7        high
2014-02-13             31.0              87.8        high
2014-02-14             22.0              71.6      medium
2014-02-15             35.0              95.0      medium
```

```
>>> df2 = pd.DataFrame([[28, 'low'],
...                     [30, 'low'],
...                     [35.1, 'medium']],
...                     columns=['temp_celsius', 'windspeed'],
...                     index=pd.DatetimeIndex(['2014-02-12', '2014-02-1
...                                             '2014-02-15']))
```

```
>>> df2
      temp_celsius windspeed
2014-02-12      28.0      low
2014-02-13      30.0      low
2014-02-15      35.1    medium

>>> df2.reindex_like(df1)
      temp_celsius  temp_fahrenheit windspeed
2014-02-12      28.0             NaN      low
2014-02-13      30.0             NaN      low
2014-02-14       NaN             NaN      NaN
2014-02-15      35.1             NaN    medium
```

`rename_axis(self, mapper=None, index=None, columns=None, axis=None, copy=True, inplace=False)`

Set the name of the axis for the index or columns.

Parameters

`mapper` : scalar, list-like, optional
Value to set the axis name attribute.

`index, columns` : scalar, list-like, dict-like or function, optional
A scalar, list-like, dict-like or functions transformations to apply to that axis' values.
Note that the ``columns`` parameter is not allowed if the object is a Series. This parameter only apply for DataFrame type objects.

Use either ``mapper`` and ``axis`` to specify the axis to target with ``mapper``, or ``index`` and/or ``columns``.

.. versionchanged:: 0.24.0

`axis` : {0 or 'index', 1 or 'columns'}, default 0
The axis to rename.

`copy` : bool, default True
Also copy underlying data.

`inplace` : bool, default False
Modifies the object directly, instead of creating a new Series or DataFrame.

Returns

Series, DataFrame, or None
The same type as the caller or None if ``inplace=True``.

See Also

`Series.rename` : Alter Series index labels or name.
`DataFrame.rename` : Alter DataFrame index labels or name.
`Index.rename` : Set new names on index.

Notes

``DataFrame.rename_axis`` supports two calling conventions

```
* ``(index=index_mapper, columns=columns_mapper, ...)``
* ``(mapper, axis={'index', 'columns'}, ...)``
```

The first calling convention will only modify the names of the index and/or the names of the Index object that is the columns. In this case, the parameter ``copy`` is ignored.

The second calling convention will modify the names of the

corresponding index if mapper is a list or a scalar.
However, if mapper is dict-like or a function, it will use the deprecated behavior of modifying the axis `*labels*`.

We *highly* recommend using keyword arguments to clarify your intent.

Examples

****Series****

```
>>> s = pd.Series(["dog", "cat", "monkey"])
>>> s
0      dog
1      cat
2    monkey
dtype: object
>>> s.rename_axis("animal")
animal
0      dog
1      cat
2    monkey
dtype: object
```

****DataFrame****

```
>>> df = pd.DataFrame({"num_legs": [4, 4, 2],
...                    "num_arms": [0, 0, 2]},
...                    ["dog", "cat", "monkey"])
>>> df
   num_legs  num_arms
dog         4         0
cat         4         0
monkey      2         2
>>> df = df.rename_axis("animal")
>>> df
   num_legs  num_arms
animal
dog         4         0
cat         4         0
monkey      2         2
>>> df = df.rename_axis("limbs", axis="columns")
>>> df
limbs  num_legs  num_arms
animal
dog         4         0
cat         4         0
monkey      2         2
```

****MultiIndex****

```
>>> df.index = pd.MultiIndex.from_product(['mammal'],
...                                       ['dog', 'cat', 'monkey']],
...                               names=['type', 'name'])
>>> df
limbs  num_legs  num_arms
type  name
mammal dog         4         0
      cat         4         0
      monkey      2         2
>>> df.rename_axis(index={'type': 'class'})
limbs  num_legs  num_arms
class  name
mammal dog         4         0
```


| | | |
|--------|---|---|
| cat | 4 | 0 |
| monkey | 2 | 2 |

```
>>> df.rename_axis(columns=str.upper)
LIMBS      num_legs  num_arms
type  name
mammal dog          4         0
      cat          4         0
      monkey       2         2
```

```
resample(self, rule, axis=0, closed: 'Optional[str]' = None, label: 'Optional[str]' = None, convention: 'str' = 'start', kind: 'Optional[str]' = None, loffset=None, base: 'Optional[int]' = None, on=None, level=None, origin: 'Union[str, TimestampConvertibleTypes]' = 'start_day', offset: 'Optional[TimedeltaConvertibleTypes]' = None) -> 'Resampler'
```

Resample time-series data.

Convenience method for frequency conversion and resampling of time series. Object must have a datetime-like index (`'DatetimeIndex'`, `'PeriodIndex'`, or `'TimedeltaIndex'`), or pass datetime-like values to the `'on'` or `'level'` keyword.

Parameters

`rule` : DateOffset, Timedelta or str

The offset string or object representing target conversion.

`axis` : {0 or 'index', 1 or 'columns'}, default 0

Which axis to use for up- or down-sampling. For `'Series'` this will default to 0, i.e. along the rows. Must be `'DatetimeIndex'`, `'TimedeltaIndex'` or `'PeriodIndex'`.

`closed` : {'right', 'left'}, default None

Which side of bin interval is closed. The default is 'left' for all frequency offsets except for 'M', 'A', 'Q', 'BM', 'BA', 'BQ', and 'W' which all have a default of 'right'.

`label` : {'right', 'left'}, default None

Which bin edge label to label bucket with. The default is 'left' for all frequency offsets except for 'M', 'A', 'Q', 'BM', 'BA', 'BQ', and 'W' which all have a default of 'right'.

`convention` : {'start', 'end', 's', 'e'}, default 'start'

For `'PeriodIndex'` only, controls whether to use the start or end of `'rule'`.

`kind` : {'timestamp', 'period'}, optional, default None

Pass 'timestamp' to convert the resulting index to a `'DateTimeIndex'` or 'period' to convert it to a `'PeriodIndex'`. By default the input representation is retained.

`loffset` : timedelta, default None

Adjust the resampled time labels.

.. deprecated:: 1.1.0

You should add the loffset to the `'df.index'` after the resample. See below.

`base` : int, default 0

For frequencies that evenly subdivide 1 day, the "origin" of the aggregated intervals. For example, for '5min' frequency, base could range from 0 through 4. Defaults to 0.

.. deprecated:: 1.1.0

The new arguments that you should use are 'offset' or 'origin'.

`on` : str, optional

For a DataFrame, column to use instead of index for resampling. Column must be datetime-like.

`level` : str or int, optional

For a MultiIndex, level (name or number) to use for

```

    resampling. `level` must be datetime-like.
    origin : {'epoch', 'start', 'start_day'}, Timestamp or str, default 'start_day'
        The timestamp on which to adjust the grouping. The timezone of origin
        must match the timezone of the index.
        If a timestamp is not used, these values are also supported:
        - 'epoch': `origin` is 1970-01-01
        - 'start': `origin` is the first value of the timeseries
        - 'start_day': `origin` is the first day at midnight of the timeseries

    .. versionadded:: 1.1.0

    offset : Timedelta or str, default is None
        An offset timedelta added to the origin.

    .. versionadded:: 1.1.0

    Returns
    -----
    Resampler object

    See Also
    -----
    groupby : Group by mapping, function, label, or list of labels.
    Series.resample : Resample a Series.
    DataFrame.resample: Resample a DataFrame.

    Notes
    -----
    See the `user guide
    <https://pandas.pydata.org/pandas-docs/stable/user\_guide/timeseries.html
    #resampling>`_
    for more.

    To learn more about the offset strings, please see `this link
    <https://pandas.pydata.org/pandas-docs/stable/user\_guide/timeseries.html
    #dateoffset-objects>`_.

    Examples
    -----
    Start by creating a series with 9 one minute timestamps.

    >>> index = pd.date_range('1/1/2000', periods=9, freq='T')
    >>> series = pd.Series(range(9), index=index)
    >>> series
    2000-01-01 00:00:00    0
    2000-01-01 00:01:00    1
    2000-01-01 00:02:00    2
    2000-01-01 00:03:00    3
    2000-01-01 00:04:00    4
    2000-01-01 00:05:00    5
    2000-01-01 00:06:00    6
    2000-01-01 00:07:00    7
    2000-01-01 00:08:00    8
    Freq: T, dtype: int64

    Downsample the series into 3 minute bins and sum the values
    of the timestamps falling into a bin.

    >>> series.resample('3T').sum()
    2000-01-01 00:00:00    3
    2000-01-01 00:03:00   12

```

```
2000-01-01 00:06:00    21
Freq: 3T, dtype: int64
```

Downsample the series into 3 minute bins as above, but label each bin using the right edge instead of the left. Please note that the value in the bucket used as the label is not included in the bucket, which it labels. For example, in the original series the bucket ``2000-01-01 00:03:00`` contains the value 3, but the summed value in the resampled bucket with the label ``2000-01-01 00:03:00`` does not include 3 (if it did, the summed value would be 6, not 3). To include this value close the right side of the bin interval as illustrated in the example below this one.

```
>>> series.resample('3T', label='right').sum()
2000-01-01 00:03:00     3
2000-01-01 00:06:00    12
2000-01-01 00:09:00    21
Freq: 3T, dtype: int64
```

Downsample the series into 3 minute bins as above, but close the right side of the bin interval.

```
>>> series.resample('3T', label='right', closed='right').sum()
2000-01-01 00:00:00     0
2000-01-01 00:03:00     6
2000-01-01 00:06:00    15
2000-01-01 00:09:00    15
Freq: 3T, dtype: int64
```

Upsample the series into 30 second bins.

```
>>> series.resample('30S').asfreq()[0:5]    # Select first 5 rows
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30    NaN
2000-01-01 00:01:00    1.0
2000-01-01 00:01:30    NaN
2000-01-01 00:02:00    2.0
Freq: 30S, dtype: float64
```

Upsample the series into 30 second bins and fill the ``NaN`` values using the ``pad`` method.

```
>>> series.resample('30S').pad()[0:5]
2000-01-01 00:00:00     0
2000-01-01 00:00:30     0
2000-01-01 00:01:00     1
2000-01-01 00:01:30     1
2000-01-01 00:02:00     2
Freq: 30S, dtype: int64
```

Upsample the series into 30 second bins and fill the ``NaN`` values using the ``bfill`` method.

```
>>> series.resample('30S').bfill()[0:5]
2000-01-01 00:00:00     0
2000-01-01 00:00:30     1
2000-01-01 00:01:00     1
2000-01-01 00:01:30     2
2000-01-01 00:02:00     2
Freq: 30S, dtype: int64
```

Pass a custom function via ``apply``

```
>>> def custom_resampler(array_like):
...     return np.sum(array_like) + 5
```

```
...
>>> series.resample('3T').apply(custom_resampler)
2000-01-01 00:00:00      8
2000-01-01 00:03:00     17
2000-01-01 00:06:00     26
Freq: 3T, dtype: int64
```

For a Series with a PeriodIndex, the keyword `'convention'` can be used to control whether to use the start or end of `'rule'`.

Resample a year by quarter using `'start'` `'convention'`. Values are assigned to the first quarter of the period.

```
>>> s = pd.Series([1, 2], index=pd.period_range('2012-01-01',
...                                             freq='A',
...                                             periods=2))
>>> s
2012      1
2013      2
Freq: A-DEC, dtype: int64
>>> s.resample('Q', convention='start').asfreq()
2012Q1      1.0
2012Q2      NaN
2012Q3      NaN
2012Q4      NaN
2013Q1      2.0
2013Q2      NaN
2013Q3      NaN
2013Q4      NaN
Freq: Q-DEC, dtype: float64
```

Resample quarters by month using `'end'` `'convention'`. Values are assigned to the last month of the period.

```
>>> q = pd.Series([1, 2, 3, 4], index=pd.period_range('2018-01-01',
...                                             freq='Q',
...                                             periods=4))
>>> q
2018Q1      1
2018Q2      2
2018Q3      3
2018Q4      4
Freq: Q-DEC, dtype: int64
>>> q.resample('M', convention='end').asfreq()
2018-03      1.0
2018-04      NaN
2018-05      NaN
2018-06      2.0
2018-07      NaN
2018-08      NaN
2018-09      3.0
2018-10      NaN
2018-11      NaN
2018-12      4.0
Freq: M, dtype: float64
```

For DataFrame objects, the keyword `'on'` can be used to specify the column instead of the index for resampling.

```
>>> d = {'price': [10, 11, 9, 13, 14, 18, 17, 19],
...      'volume': [50, 60, 40, 100, 50, 100, 40, 50]}
>>> df = pd.DataFrame(d)
>>> df['week_starting'] = pd.date_range('01/01/2018',
...                                     periods=8,
...                                     freq='W')
```

```
>>> df
   price  volume week_starting
0     10     50   2018-01-07
1     11     60   2018-01-14
2      9     40   2018-01-21
3     13    100   2018-01-28
4     14     50   2018-02-04
5     18    100   2018-02-11
6     17     40   2018-02-18
7     19     50   2018-02-25
>>> df.resample('M', on='week_starting').mean()
           price  volume
week_starting
2018-01-31    10.75    62.5
2018-02-28    17.00    60.0
```

For a DataFrame with MultiIndex, the keyword `level` can be used to specify on which level the resampling needs to take place.

```
>>> days = pd.date_range('1/1/2000', periods=4, freq='D')
>>> d2 = {'price': [10, 11, 9, 13, 14, 18, 17, 19],
...       'volume': [50, 60, 40, 100, 50, 100, 40, 50]}
>>> df2 = pd.DataFrame(d2,
...                     index=pd.MultiIndex.from_product([days,
...                                                         ['morning',
...                                                         'afternoon']])
>>> df2
           price  volume
2000-01-01 morning     10     50
           afternoon    11     60
2000-01-02 morning      9     40
           afternoon    13    100
2000-01-03 morning     14     50
           afternoon    18    100
2000-01-04 morning     17     40
           afternoon    19     50
>>> df2.resample('D', level=0).sum()
           price  volume
2000-01-01     21    110
2000-01-02     22    140
2000-01-03     32    150
2000-01-04     36     90
```

If you want to adjust the start of the bins based on a fixed timestamp:

```
>>> start, end = '2000-10-01 23:30:00', '2000-10-02 00:30:00'
>>> rng = pd.date_range(start, end, freq='7min')
>>> ts = pd.Series(np.arange(len(rng)) * 3, index=rng)
>>> ts
2000-10-01 23:30:00    0
2000-10-01 23:37:00    3
2000-10-01 23:44:00    6
2000-10-01 23:51:00    9
2000-10-01 23:58:00   12
2000-10-02 00:05:00   15
2000-10-02 00:12:00   18
2000-10-02 00:19:00   21
2000-10-02 00:26:00   24
Freq: 7T, dtype: int64

>>> ts.resample('17min').sum()
2000-10-01 23:14:00    0
2000-10-01 23:31:00    9
2000-10-01 23:48:00   21
```

```
2000-10-02 00:05:00    54
2000-10-02 00:22:00    24
Freq: 17T, dtype: int64
```

```
>>> ts.resample('17min', origin='epoch').sum()
2000-10-01 23:18:00     0
2000-10-01 23:35:00    18
2000-10-01 23:52:00    27
2000-10-02 00:09:00    39
2000-10-02 00:26:00    24
Freq: 17T, dtype: int64

>>> ts.resample('17min', origin='2000-01-01').sum()
2000-10-01 23:24:00     3
2000-10-01 23:41:00    15
2000-10-01 23:58:00    45
2000-10-02 00:15:00    45
Freq: 17T, dtype: int64
```

If you want to adjust the start of the bins with an `offset` Timedelta, the two following lines are equivalent:

```
>>> ts.resample('17min', origin='start').sum()
2000-10-01 23:30:00     9
2000-10-01 23:47:00    21
2000-10-02 00:04:00    54
2000-10-02 00:21:00    24
Freq: 17T, dtype: int64

>>> ts.resample('17min', offset='23h30min').sum()
2000-10-01 23:30:00     9
2000-10-01 23:47:00    21
2000-10-02 00:04:00    54
2000-10-02 00:21:00    24
Freq: 17T, dtype: int64
```

To replace the use of the deprecated `base` argument, you can now use `offset`, in this example it is equivalent to have `base=2`:

```
>>> ts.resample('17min', offset='2min').sum()
2000-10-01 23:16:00     0
2000-10-01 23:33:00     9
2000-10-01 23:50:00    36
2000-10-02 00:07:00    39
2000-10-02 00:24:00    24
Freq: 17T, dtype: int64
```

To replace the use of the deprecated `loffset` argument:

```
>>> from pandas.tseries.frequencies import to_offset
>>> loffset = '19min'
>>> ts_out = ts.resample('17min').sum()
>>> ts_out.index = ts_out.index + to_offset(loffset)
>>> ts_out
2000-10-01 23:33:00     0
2000-10-01 23:50:00     9
2000-10-02 00:07:00    21
2000-10-02 00:24:00    54
2000-10-02 00:41:00    24
Freq: 17T, dtype: int64
```

rolling(self, window: 'Union[int, timedelta, BaseOffset, BaseIndexer]', min_periods: 'Optional[int]' = None, center: 'bool_t' = False, win_type: 'Optional[s

```

tr]' = None, on: 'Optional[str]' = None, axis: 'Axis' = 0, closed: 'Optional[str]' = None)
    Provide rolling window calculations.

    Parameters
    -----
    window : int, offset, or BaseIndexer subclass
        Size of the moving window. This is the number of observations used for
        calculating the statistic. Each window will be a fixed size.

        If its an offset then this will be the time period of each window. Each
        window will be a variable sized based on the observations included in
        the time-period. This is only valid for datetimelike indexes.

        If a BaseIndexer subclass is passed, calculates the window boundaries
        based on the defined ``get_window_bounds`` method. Additional rolling
        keyword arguments, namely ``min_periods``, ``center``, and
        ``closed`` will be passed to ``get_window_bounds``.
    min_periods : int, default None
        Minimum number of observations in window required to have a value
        (otherwise result is NA). For a window that is specified by an offset,
        ``min_periods`` will default to 1. Otherwise, ``min_periods`` will default
        to the size of the window.
    center : bool, default False
        Set the labels at the center of the window.
    win_type : str, default None
        Provide a window type. If ``None``, all points are evenly weighted.
        See the notes below for further information.
    on : str, optional
        For a DataFrame, a datetime-like column or MultiIndex level on which
        to calculate the rolling window, rather than the DataFrame's index.
        Provided integer column is ignored and excluded from result since
        an integer index is not used to calculate the rolling window.
    axis : int or str, default 0
    closed : str, default None
        Make the interval closed on the 'right', 'left', 'both' or
        'neither' endpoints. Defaults to 'right'.

    .. versionchanged:: 1.2.0

        The closed parameter with fixed windows is now supported.

    Returns
    -----
    a Window or Rolling sub-classed for the particular operation

    See Also
    -----
    expanding : Provides expanding transformations.
    ewm : Provides exponential weighted functions.

    Notes
    -----
    By default, the result is set to the right edge of the window. This can
    be changed to the center of the window by setting ``center=True``.

    To learn more about the offsets & frequency strings, please see `this link

```

```
nk
|   <https://pandas.pydata.org/pandas-docs/stable/user\_guide/timeseries.html
#offset-aliases>`__`.
|
|   If ``win_type=None``, all points are evenly weighted; otherwise, ``win_t
ype``
|   can accept a string of any `scipy.signal` window function
|   <https://docs.scipy.org/doc/scipy/reference/signal.windows.html#module-s
cipy.signal.windows>`__`.
|
|   Certain Scipy window types require additional parameters to be passed
|   in the aggregation function. The additional parameters must match
|   the keywords specified in the Scipy window type method signature.
|   Please see the third example below on how to add the additional paramete
rs.
|
Examples
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
>>> df
B
0  0.0
1  1.0
2  2.0
3  NaN
4  4.0
Rolling sum with a window length of 2, using the 'triang'
window type.
>>> df.rolling(2, win_type='triang').sum()
B
0  NaN
1  0.5
2  1.5
3  NaN
4  NaN
Rolling sum with a window length of 2, using the 'gaussian'
window type (note how we need to specify std).
>>> df.rolling(2, win_type='gaussian').sum(std=3)
B
0  NaN
1  0.986207
2  2.958621
3  NaN
4  NaN
Rolling sum with a window length of 2, min_periods defaults
to the window length.
>>> df.rolling(2).sum()
B
0  NaN
1  1.0
2  3.0
3  NaN
4  NaN
Same as above, but explicitly set the min_periods
>>> df.rolling(2, min_periods=1).sum()
B
0  0.0
```



```
1  1.0
2  3.0
3  2.0
4  4.0
```

Same as above, but with forward-looking windows

```
>>> indexer = pd.api.indexers.FixedForwardWindowIndexer(window_size=2)
>>> df.rolling(window=indexer, min_periods=1).sum()
```

```
      B
0  1.0
1  3.0
2  2.0
3  4.0
4  4.0
```

A ragged (meaning not-a-regular frequency), time-indexed DataFrame

```
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]},
...                    index = [pd.Timestamp('20130101 09:00:00'),
...                             pd.Timestamp('20130101 09:00:02'),
...                             pd.Timestamp('20130101 09:00:03'),
...                             pd.Timestamp('20130101 09:00:05'),
...                             pd.Timestamp('20130101 09:00:06')])
```

```
>>> df
              B
2013-01-01 09:00:00  0.0
2013-01-01 09:00:02  1.0
2013-01-01 09:00:03  2.0
2013-01-01 09:00:05  NaN
2013-01-01 09:00:06  4.0
```

Contrasting to an integer rolling window, this will roll a variable length window corresponding to the time period. The default for min_periods is 1.

```
>>> df.rolling('2s').sum()
              B
2013-01-01 09:00:00  0.0
2013-01-01 09:00:02  1.0
2013-01-01 09:00:03  3.0
2013-01-01 09:00:05  NaN
2013-01-01 09:00:06  4.0
```

```
sample(self: 'FrameOrSeries', n=None, frac=None, replace=False, weights=None,
random_state=None, axis=None) -> 'FrameOrSeries'
Return a random sample of items from an axis of object.
```

You can use `random_state` for reproducibility.

Parameters

```
n : int, optional
    Number of items from axis to return. Cannot be used with `frac`.
    Default = 1 if `frac` = None.
frac : float, optional
    Fraction of axis items to return. Cannot be used with `n`.
replace : bool, default False
    Allow or disallow sampling of the same row more than once.
weights : str or ndarray-like, optional
    Default 'None' results in equal probability weighting.
    If passed a Series, will align with target object on index. Index
    values in weights not found in sampled object will be ignored and
    index values in sampled object not in weights will be assigned
```

weights of zero.

If called on a DataFrame, will accept the name of a column when axis = 0.

Unless weights are a Series, weights must be same length as axis being sampled.

If weights do not sum to 1, they will be normalized to sum to 1.

Missing values in the weights column will be treated as zero.

Infinite values not allowed.

random_state : int, array-like, BitGenerator, np.random.RandomState, opt

If int, array-like, or BitGenerator (NumPy>=1.17), seed for random number generator

If np.random.RandomState, use as numpy RandomState object.

.. versionchanged:: 1.1.0

array-like and BitGenerator (for NumPy>=1.17) object now passed

np.random.RandomState() as seed

axis : {0 or 'index', 1 or 'columns', None}, default None

Axis to sample. Accepts axis number or name. Default is stat axis for given data type (0 for Series and DataFrames).

Returns

Series or DataFrame

A new object of same type as caller containing `n` items randomly sampled from the caller object.

See Also

DataFrameGroupBy.sample: Generates random samples from each group of a DataFrame object.

SeriesGroupBy.sample: Generates random samples from each group of a Series object.

numpy.random.choice: Generates a random sample from a given 1-D numpy array.

Notes

If `frac` > 1, `replacement` should be set to `True`.

Examples

```
>>> df = pd.DataFrame({'num_legs': [2, 4, 8, 0],
...                     'num_wings': [2, 0, 0, 0],
...                     'num_specimen_seen': [10, 2, 1, 8]},
...                     index=['falcon', 'dog', 'spider', 'fish'])
>>> df
```

| | num_legs | num_wings | num_specimen_seen |
|--------|----------|-----------|-------------------|
| falcon | 2 | 2 | 10 |
| dog | 4 | 0 | 2 |
| spider | 8 | 0 | 1 |
| fish | 0 | 0 | 8 |

Extract 3 random elements from the ``Series`` ``df['num_legs']``:
Note that we use `random_state` to ensure the reproducibility of the examples.

```
>>> df['num_legs'].sample(n=3, random_state=1)
fish      0
spider    8
falcon    2
Name: num_legs, dtype: int64
```

A random 50% sample of the ``DataFrame`` with replacement:

```
>>> df.sample(frac=0.5, replace=True, random_state=1)
   num_legs  num_wings  num_specimen_seen
dog         4         0                   2
fish        0         0                   8
```

An upsample sample of the ``DataFrame`` with replacement:

Note that `replace` parameter has to be `True` for `frac` parameter > 1.

```
>>> df.sample(frac=2, replace=True, random_state=1)
   num_legs  num_wings  num_specimen_seen
dog         4         0                   2
fish        0         0                   8
falcon       2         2                  10
falcon       2         2                  10
fish         0         0                   8
dog         4         0                   2
fish         0         0                   8
dog         4         0                   2
```

Using a DataFrame column as weights. Rows with larger value in the `num_specimen_seen` column are more likely to be sampled.

```
>>> df.sample(n=2, weights='num_specimen_seen', random_state=1)
   num_legs  num_wings  num_specimen_seen
falcon       2         2                  10
fish         0         0                   8
```

```
set_flags(self: 'FrameOrSeries', *, copy: 'bool' = False, allows_duplicate_labels: 'Optional[bool]' = None) -> 'FrameOrSeries'
Return a new object with updated flags.
```

Parameters

`allows_duplicate_labels` : bool, optional
Whether the returned object allows duplicate labels.

Returns

Series or DataFrame
The same type as the caller.

See Also

`DataFrame.attrs` : Global metadata applying to this dataset.
`DataFrame.flags` : Global flags applying to this object.

Notes

This method returns a new object that's a view on the same data as the input. Mutating the input or the output values will be reflected in the other.

This method is intended to be used in method chains.

"Flags" differ from "metadata". Flags reflect properties of the pandas object (the Series or DataFrame). Metadata refer to properties of the dataset, and should be stored in `:attr:'DataFrame.attrs'`.

Examples

```
>>> df = pd.DataFrame({"A": [1, 2]})
>>> df.flags.allows_duplicate_labels
```

```

True
>>> df2 = df.set_flags(allows_duplicate_labels=False)
>>> df2.flags.allows_duplicate_labels
False

```

`slice_shift(self: 'FrameOrSeries', periods: 'int' = 1, axis=0) -> 'FrameOrSeries'`

Equivalent to ``shift`` without copying data.
The shifted data will not include the dropped periods and the shifted axis will be smaller than the original.

.. deprecated:: 1.2.0
 `slice_shift` is deprecated,
 use `DataFrame/Series.shift` instead.

Parameters

`periods` : int
 Number of periods to move, can be positive or negative.

Returns

`shifted` : same type as caller

Notes

While the ``slice_shift`` is faster than ``shift``, you may pay for it later during alignment.

`squeeze(self, axis=None)`
 Squeeze 1 dimensional axis objects into scalars.

Series or DataFrames with a single element are squeezed to a scalar.
DataFrames with a single column or a single row are squeezed to a Series. Otherwise the object is unchanged.

This method is most useful when you don't know if your object is a Series or DataFrame, but you do know it has just a single column. In that case you can safely call ``squeeze`` to ensure you have a Series.

Parameters

`axis` : {0 or 'index', 1 or 'columns', None}, default None
 A specific axis to squeeze. By default, all length-1 axes are squeezed.

Returns

DataFrame, Series, or scalar
 The projection after squeezing ``axis`` or all the axes.

See Also

`Series.iloc` : Integer-location based indexing for selecting scalars.
`DataFrame.iloc` : Integer-location based indexing for selecting Series.
`Series.to_frame` : Inverse of `DataFrame.squeeze` for a single-column DataFrame.

Examples

```

>>> primes = pd.Series([2, 3, 5, 7])

```

Slicing might produce a Series with a single value:

```
>>> even_primes = primes[primes % 2 == 0]
>>> even_primes
0    2
dtype: int64
```

```
>>> even_primes.squeeze()
2
```

Squeezing objects with more than one value in every axis does nothing:

```
>>> odd_primes = primes[primes % 2 == 1]
>>> odd_primes
1    3
2    5
3    7
dtype: int64
```

```
>>> odd_primes.squeeze()
1    3
2    5
3    7
dtype: int64
```

Squeezing is even more effective when used with DataFrames.

```
>>> df = pd.DataFrame([[1, 2], [3, 4]], columns=['a', 'b'])
>>> df
   a  b
0  1  2
1  3  4
```

Slicing a single column will produce a DataFrame with the columns having only one value:

```
>>> df_a = df[['a']]
>>> df_a
   a
0  1
1  3
```

So the columns can be squeezed down, resulting in a Series:

```
>>> df_a.squeeze('columns')
0    1
1    3
Name: a, dtype: int64
```

Slicing a single row from a single column will produce a single scalar DataFrame:

```
>>> df_0a = df.loc[df.index < 1, ['a']]
>>> df_0a
   a
0  1
```

Squeezing the rows produces a single scalar Series:

```
>>> df_0a.squeeze('rows')
a    1
Name: 0, dtype: int64
```

Squeezing all axes will project directly into a scalar:

```
>>> df_0a.squeeze()
1
```

```
swapaxes(self: 'FrameOrSeries', axis1, axis2, copy=True) -> 'FrameOrSeries'
Interchange axes and swap values axes appropriately.
```

Returns

y : same as input

```
tail(self: 'FrameOrSeries', n: 'int' = 5) -> 'FrameOrSeries'
Return the last `n` rows.
```

This function returns last `n` rows from the object based on position. It is useful for quickly verifying data, for example, after sorting or appending rows.

For negative values of `n`, this function returns all rows except the first `n` rows, equivalent to `df[n:]`.

Parameters

n : int, default 5
Number of rows to select.

Returns

type of caller
The last `n` rows of the caller object.

See Also

DataFrame.head : The first `n` rows of the caller object.

Examples

```
>>> df = pd.DataFrame({'animal': ['alligator', 'bee', 'falcon', 'lion',
...                               'monkey', 'parrot', 'shark', 'whale', 'zebra']})
>>> df
```

```
   animal
0  alligator
1      bee
2   falcon
3     lion
4   monkey
5   parrot
6    shark
7    whale
8    zebra
```

Viewing the last 5 lines

```
>>> df.tail()
   animal
4  monkey
5  parrot
6   shark
7   whale
8   zebra
```

Viewing the last `n` lines (three in this case)

```
>>> df.tail(3)
   animal
6   shark
7   whale
8   zebra
```

For negative values of `n`

```
>>> df.tail(-3)
      animal
3      lion
4    monkey
5    parrot
6     shark
7     whale
8     zebra
```

`take(self: 'FrameOrSeries', indices, axis=0, is_copy: 'Optional[bool_t]' = None, **kwargs) -> 'FrameOrSeries'`

Return the elements in the given *positional* indices along an axis.

This means that we are not indexing according to actual values in the index attribute of the object. We are indexing according to the actual position of the element in the object.

Parameters

`indices` : array-like

An array of ints indicating which positions to take.

`axis` : {0 or 'index', 1 or 'columns', None}, default 0

The axis on which to select elements. ``0`` means that we are selecting rows, ``1`` means that we are selecting columns.

`is_copy` : bool

Before pandas 1.0, ``is_copy=False`` can be specified to ensure that the return value is an actual copy. Starting with pandas 1.0, ``take`` always returns a copy, and the keyword is therefore deprecated.

.. deprecated:: 1.0.0

`**kwargs`

For compatibility with :meth:`numpy.take`. Has no effect on the output.

Returns

`taken` : same type as caller

An array-like containing the elements taken from the object.

See Also

`DataFrame.loc` : Select a subset of a DataFrame by labels.

`DataFrame.iloc` : Select a subset of a DataFrame by positions.

`numpy.take` : Take elements from an array along an axis.

Examples

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0),
...                    ('parrot', 'bird', 24.0),
...                    ('lion', 'mammal', 80.5),
...                    ('monkey', 'mammal', np.nan)],
...                    columns=['name', 'class', 'max_speed'],
...                    index=[0, 2, 3, 1])
>>> df
   name  class  max_speed
0  falcon   bird    389.0
2  parrot   bird     24.0
3   lion  mammal     80.5
1  monkey  mammal      NaN
```

Take elements at positions 0 and 3 along the axis 0 (default).

Note how the actual indices selected (0 and 1) do not correspond to our selected indices 0 and 3. That's because we are selecting the 0th and 3rd rows, not rows whose indices equal 0 and 3.

```
>>> df.take([0, 3])
      name  class  max_speed
0  falcon   bird    389.0
1  monkey  mammal      NaN
```

Take elements at indices 1 and 2 along the axis 1 (column selection).

```
>>> df.take([1, 2], axis=1)
      class  max_speed
0     bird    389.0
2     bird    24.0
3  mammal    80.5
1  mammal      NaN
```

We may take elements using negative integers for positive indices, starting from the end of the object, just like with Python lists.

```
>>> df.take([-1, -2])
      name  class  max_speed
1  monkey  mammal      NaN
3    lion  mammal    80.5
```

`to_clipboard(self, excel: 'bool_t' = True, sep: 'Optional[str]' = None, **kw
args) -> 'None'`

Copy object to the system clipboard.

Write a text representation of object to the system clipboard.
This can be pasted into Excel, for example.

Parameters

`excel` : bool, default True

Produce output in a csv format for easy pasting into excel.

- True, use the provided separator for csv pasting.
- False, write a string representation of the object to the clipboard.

`sep` : str, default ```\t```
Field delimiter.

`**kwargs`

These parameters will be passed to `DataFrame.to_csv`.

See Also

`DataFrame.to_csv` : Write a DataFrame to a comma-separated values (csv) file.

`read_clipboard` : Read text from clipboard and pass to `read_table`.

Notes

Requirements for your platform.

- Linux : ``xclip``, or ``xsel`` (with ``PyQt4`` modules)
- Windows : none
- OS X : none

Examples

Copy the contents of a DataFrame to the clipboard.


```
>>> df = pd.DataFrame([[1, 2, 3], [4, 5, 6]], columns=['A', 'B', 'C'])

>>> df.to_clipboard(sep=',') # doctest: +SKIP
... # Wrote the following to the system clipboard:
... # ,A,B,C
... # 0,1,2,3
... # 1,4,5,6
```

We can omit the index by passing the keyword `index` and setting it to false.

```
>>> df.to_clipboard(sep=',', index=False) # doctest: +SKIP
... # Wrote the following to the system clipboard:
... # A,B,C
... # 1,2,3
... # 4,5,6
```

```
to_csv(self, path_or_buf: 'Optional[FilePathOrBuffer]' = None, sep: 'str' =
',', na_rep: 'str' = '', float_format: 'Optional[str]' = None, columns: 'Optional[Sequence[Label]]' = None, header: 'Union[bool_t, List[str]]' = True, index: 'bool_t' = True, index_label: 'Optional[IndexLabel]' = None, mode: 'str' = 'w', encoding: 'Optional[str]' = None, compression: 'CompressionOptions' = 'infer', quoting: 'Optional[int]' = None, quotechar: 'str' = "'", line_terminator: 'Optional[str]' = None, chunksize: 'Optional[int]' = None, date_format: 'Optional[str]' = None, doublequote: 'bool_t' = True, escapechar: 'Optional[str]' = None, decimal: 'str' = '.', errors: 'str' = 'strict', storage_options: 'StorageOptions' = None)
-> 'Optional[str]'
```

Write object to a comma-separated values (csv) file.

.. versionchanged:: 0.24.0

The order of arguments for Series was changed.

Parameters

path_or_buf : str or file handle, default None

File path or object, if None is provided the result is returned as a string. If a non-binary file object is passed, it should be opened

with `newline=''`, disabling universal newlines. If a binary file object is passed, `mode` might need to contain a `b`.

.. versionchanged:: 0.24.0

Was previously named "path" for Series.

.. versionchanged:: 1.2.0

Support for binary file objects was introduced.

sep : str, default ','

String of length 1. Field delimiter for the output file.

na_rep : str, default ''

Missing data representation.

float_format : str, default None

Format string for floating point numbers.

columns : sequence, optional

Columns to write.

header : bool or list of str, default True

Write out the column names. If a list of strings is given it is assumed to be aliases for the column names.

.. versionchanged:: 0.24.0

Previously defaulted to False for Series.

```

index : bool, default True
    Write row names (index).
index_label : str or sequence, or False, default None
    Column label for index column(s) if desired. If None is given, and
    `header` and `index` are True, then the index names are used. A
    sequence should be given if the object uses MultiIndex. If
    False do not print fields for index names. Use index_label=False
    for easier importing in R.
mode : str
    Python write mode, default 'w'.
encoding : str, optional
    A string representing the encoding to use in the output file,
    defaults to 'utf-8'. `encoding` is not supported if `path_or_buf`
    is a non-binary file object.
compression : str or dict, default 'infer'
    If str, represents compression mode. If dict, value at 'method' is
    the compression mode. Compression mode may be any of the following
    possible values: {'infer', 'gzip', 'bz2', 'zip', 'xz', None}. If
    compression mode is 'infer' and `path_or_buf` is path-like, then
    detect compression mode from the following extensions: '.gz',
    '.bz2', '.zip' or '.xz'. (otherwise no compression). If dict given
    and mode is one of {'zip', 'gzip', 'bz2'}, or inferred as
    one of the above, other entries passed as
    additional compression options.

.. versionchanged:: 1.0.0

    May now be a dict with key 'method' as compression mode
    and other entries as additional compression options if
    compression mode is 'zip'.

.. versionchanged:: 1.1.0

    Passing compression options as keys in dict is
    supported for compression modes 'gzip' and 'bz2'
    as well as 'zip'.

.. versionchanged:: 1.2.0

    Compression is supported for binary file objects.

.. versionchanged:: 1.2.0

    Previous versions forwarded dict entries for 'gzip' to
    `gzip.open` instead of `gzip.GzipFile` which prevented
    setting `mtime`.

quoting : optional constant from csv module
    Defaults to csv.QUOTE_MINIMAL. If you have set a `float_format`
    then floats are converted to strings and thus csv.QUOTE_NONNUMERIC
    will treat them as non-numeric.
quotechar : str, default '"'
    String of length 1. Character used to quote fields.
line_terminator : str, optional
    The newline character or character sequence to use in the output
    file. Defaults to `os.linesep`, which depends on the OS in which
    this method is called ('\n' for linux, '\r\n' for Windows, i.e.).

.. versionchanged:: 0.24.0
chunksize : int or None
    Rows to write at a time.
date_format : str, default None
    Format string for datetime objects.
doublequote : bool, default True

```

```

Control quoting of `quotechar` inside a field.
escapechar : str, default None
    String of length 1. Character used to escape `sep` and `quotechar`
    when appropriate.
decimal : str, default '.'
    Character recognized as decimal separator. E.g. use ',' for
    European data.
errors : str, default 'strict'
    Specifies how encoding and decoding errors are to be handled.
    See the errors argument for :func:`open` for a full list
    of options.

    .. versionadded:: 1.1.0

storage_options : dict, optional
    Extra options that make sense for a particular storage connection,
    e.g.
    host, port, username, password, etc., if using a URL that will
    be parsed by ``fsspec``, e.g., starting "s3://", "gcs://". An error
    will be raised if providing this argument with a non-fsspec URL.
    See the fsspec and backend storage implementation docs for the set o
    f
    allowed keys and values.

    .. versionadded:: 1.2.0

```

Returns

None or str

If path_or_buf is None, returns the resulting csv format as a string. Otherwise returns None.

See Also

read_csv : Load a CSV file into a DataFrame.
to_excel : Write DataFrame to an Excel file.

Examples

```

>>> df = pd.DataFrame({'name': ['Raphael', 'Donatello'],
...                     'mask': ['red', 'purple'],
...                     'weapon': ['sai', 'bo staff']})
>>> df.to_csv(index=False)
'name,mask,weapon\nRaphael,red,sai\nDonatello,purple,bo staff\n'

```

Create 'out.zip' containing 'out.csv'

```

>>> compression_opts = dict(method='zip',
...                           archive_name='out.csv') # doctest: +SKIP
>>> df.to_csv('out.zip', index=False,
...           compression=compression_opts) # doctest: +SKIP

```

```

to_excel(self, excel_writer, sheet_name: 'str' = 'Sheet1', na_rep: 'str' =
'', float_format: 'Optional[str]' = None, columns=None, header=True, index=True,
index_label=None, startrow=0, startcol=0, engine=None, merge_cells=True, encodin
g=None, inf_rep='inf', verbose=True, freeze_panes=None, storage_options: 'Storag
eOptions' = None) -> 'None'

```

Write object to an Excel sheet.

To write a single object to an Excel .xlsx file it is only necessary to specify a target file name. To write to multiple sheets it is necessary

create an `ExcelWriter` object with a target file name, and specify a sheet in the file to write to.

Multiple sheets may be written to by specifying unique `sheet_name`. With all data written to the file it is necessary to save the changes. Note that creating an `ExcelWriter` object with a file name that already exists will result in the contents of the existing file being erased.

Parameters

`excel_writer` : path-like, file-like, or `ExcelWriter` object
File path or existing `ExcelWriter`.

`sheet_name` : str, default 'Sheet1'
Name of sheet which will contain `DataFrame`.

`na_rep` : str, default ''
Missing data representation.

`float_format` : str, optional
Format string for floating point numbers. For example
`float_format="%.2f"` will format 0.1234 to 0.12.

`columns` : sequence or list of str, optional
Columns to write.

`header` : bool or list of str, default True
Write out the column names. If a list of string is given it is assumed to be aliases for the column names.

`index` : bool, default True
Write row names (index).

`index_label` : str or sequence, optional
Column label for index column(s) if desired. If not specified, and `header` and `index` are True, then the index names are used. A sequence should be given if the `DataFrame` uses `MultiIndex`.

`startrow` : int, default 0
Upper left cell row to dump data frame.

`startcol` : int, default 0
Upper left cell column to dump data frame.

`engine` : str, optional
Write engine to use, 'openpyxl' or 'xlsxwriter'. You can also set this via the options ``io.excel.xlsx.writer``, ``io.excel.xls.writer``, and ``io.excel.xlsm.writer``.

.. deprecated:: 1.2.0

As the `xlwt` <<https://pypi.org/project/xlwt/>> package is no longer maintained, the ``xlwt`` engine will be removed in a future version of pandas.

`merge_cells` : bool, default True
Write `MultiIndex` and Hierarchical Rows as merged cells.

`encoding` : str, optional
Encoding of the resulting excel file. Only necessary for `xlwt`, other writers support unicode natively.

`inf_rep` : str, default 'inf'
Representation for infinity (there is no native representation for infinity in Excel).

`verbose` : bool, default True
Display more information in the error logs.

`freeze_panes` : tuple of int (length 2), optional
Specifies the one-based bottommost row and rightmost column that is to be frozen.

`storage_options` : dict, optional
Extra options that make sense for a particular storage connection, e.g. host, port, username, password, etc., if using a URL that will be parsed by ``fsspec``, e.g., starting "s3://", "gcs://". An error

will be raised if providing this argument with a non-fsspec URL.
See the fsspec and backend storage implementation docs for the set o

allowed keys and values.

.. versionadded:: 1.2.0

See Also

`to_csv` : Write DataFrame to a comma-separated values (csv) file.
`ExcelWriter` : Class for writing DataFrame objects into excel sheets.
`read_excel` : Read an Excel file into a pandas DataFrame.
`read_csv` : Read a comma-separated values (csv) file into DataFrame.

Notes

For compatibility with `:meth:`~DataFrame.to_csv``,
`to_excel` serializes lists and dicts to strings before writing.

Once a workbook has been saved it is not possible write further data
without rewriting the whole workbook.

Examples

Create, write to and save a workbook:

```
>>> df1 = pd.DataFrame([[ 'a', 'b'], [ 'c', 'd']],
...                     index=[ 'row 1', 'row 2'],
...                     columns=[ 'col 1', 'col 2'])
>>> df1.to_excel("output.xlsx") # doctest: +SKIP
```

To specify the sheet name:

```
>>> df1.to_excel("output.xlsx",
...              sheet_name='Sheet_name_1') # doctest: +SKIP
```

If you wish to write to more than one sheet in the workbook, it is
necessary to specify an `ExcelWriter` object:

```
>>> df2 = df1.copy()
>>> with pd.ExcelWriter('output.xlsx') as writer: # doctest: +SKIP
...     df1.to_excel(writer, sheet_name='Sheet_name_1')
...     df2.to_excel(writer, sheet_name='Sheet_name_2')
```

`ExcelWriter` can also be used to append to an existing Excel file:

```
>>> with pd.ExcelWriter('output.xlsx',
...                     mode='a') as writer: # doctest: +SKIP
...     df.to_excel(writer, sheet_name='Sheet_name_3')
```

To set the library that is used to write the Excel file,
you can pass the `'engine'` keyword (the default engine is
automatically chosen depending on the file extension):

```
>>> df1.to_excel('output1.xlsx', engine='xlsxwriter') # doctest: +SKIP
```

```
to_hdf(self, path_or_buf, key: 'str', mode: 'str' = 'a', complevel: 'Optiona
l[int]' = None, complib: 'Optional[str]' = None, append: 'bool_t' = False, forma
t: 'Optional[str]' = None, index: 'bool_t' = True, min_itemsize: 'Optional[Union
[int, Dict[str, int]]]' = None, nan_rep=None, dropna: 'Optional[bool_t]' = None,
data_columns: 'Optional[Union[bool_t, List[str]]]' = None, errors: 'str' = 'stri
ct', encoding: 'str' = 'UTF-8') -> 'None'
```

Write the contained data to an HDF5 file using `HDFStore`.

Hierarchical Data Format (HDF) is self-describing, allowing an application to interpret the structure and contents of a file with no outside information. One HDF file can hold a mix of related objects which can be accessed as a group or as individual objects.

In order to add another DataFrame or Series to an existing HDF file please use append mode and a different a key.

For more information see the :ref:`user guide <io.hdf5>`.

Parameters

```

path_or_buf : str or pandas.HDFStore
    File path or HDFStore object.
key : str
    Identifier for the group in the store.
mode : {'a', 'w', 'r+'}, default 'a'
    Mode to open file:

    - 'w': write, a new file is created (an existing file with
      the same name would be deleted).
    - 'a': append, an existing file is opened for reading and
      writing, and if the file does not exist it is created.
    - 'r+': similar to 'a', but the file must already exist.
complevel : {0-9}, optional
    Specifies a compression level for data.
    A value of 0 disables compression.
complib : {'zlib', 'lzo', 'bzip2', 'blosc'}, default 'zlib'
    Specifies the compression library to be used.
    As of v0.20.2 these additional compressors for Blosc are supported
    (default if no compressor specified: 'blosc:blosclz'):
    {'blosc:blosclz', 'blosc:lz4', 'blosc:lz4hc', 'blosc:snappy',
     'blosc:zlib', 'blosc:zstd'}.
    Specifying a compression library which is not available issues
    a ValueError.
append : bool, default False
    For Table formats, append the input data to the existing.
format : {'fixed', 'table', None}, default 'fixed'
    Possible values:

    - 'fixed': Fixed format. Fast writing/reading. Not-appendable,
      nor searchable.
    - 'table': Table format. Write as a PyTables Table structure
      which may perform worse but allow more flexible operations
      like searching / selecting subsets of the data.
    - If None, pd.get_option('io.hdf.default_format') is checked,
      followed by fallback to "fixed"
errors : str, default 'strict'
    Specifies how encoding and decoding errors are to be handled.
    See the errors argument for :func:`open` for a full list
    of options.
encoding : str, default "UTF-8"
min_itemsize : dict or int, optional
    Map column names to minimum string sizes for columns.
nan_rep : Any, optional
    How to represent null values as str.
    Not allowed with append=True.
data_columns : list of columns or True, optional
    List of columns to create as indexed data columns for on-disk
    queries, or True to use all columns. By default only the axes
    of the object are indexed. See :ref:`io.hdf5-query-data-columns`.
    Applicable only to format='table'.

```

See Also

```
DataFrame.read_hdf : Read from HDF file.
DataFrame.to_parquet : Write a DataFrame to the binary parquet format.
DataFrame.to_sql : Write to a sql table.
DataFrame.to_feather : Write out feather-format for DataFrames.
DataFrame.to_csv : Write out to a csv file.
```

Examples

```
>>> df = pd.DataFrame({'A': [1, 2, 3], 'B': [4, 5, 6]},
...                    index=['a', 'b', 'c'])
>>> df.to_hdf('data.h5', key='df', mode='w')
```

We can add another object to the same file:

```
>>> s = pd.Series([1, 2, 3, 4])
>>> s.to_hdf('data.h5', key='s')
```

Reading from HDF file:

```
>>> pd.read_hdf('data.h5', 'df')
A  B
a  1  4
b  2  5
c  3  6
>>> pd.read_hdf('data.h5', 's')
0    1
1    2
2    3
3    4
dtype: int64
```

Deleting file with data:

```
>>> import os
>>> os.remove('data.h5')
```

```
to_json(self, path_or_buf: 'Optional[FilePathOrBuffer]' = None, orient: 'Optional[str]' = None, date_format: 'Optional[str]' = None, double_precision: 'int' = 10, force_ascii: 'bool_t' = True, date_unit: 'str' = 'ms', default_handler: 'Optional[Callable[[Any], JSONSerializable]]' = None, lines: 'bool_t' = False, compression: 'CompressionOptions' = 'infer', index: 'bool_t' = True, indent: 'Optional[int]' = None, storage_options: 'StorageOptions' = None) -> 'Optional[str]'
```

Convert the object to a JSON string.

Note NaN's and None will be converted to null and datetime objects will be converted to UNIX timestamps.

Parameters

path_or_buf : str or file handle, optional
File path or object. If not specified, the result is returned as a string.

orient : str
Indication of expected JSON string format.

* Series:

- default is 'index'
- allowed values are: {'split', 'records', 'index', 'table'}.

* DataFrame:

- default is 'columns'
- allowed values are: {'split', 'records', 'index', 'columns', 'values', 'table'}.

```

s],
ue}]
rds``.
    * The format of the JSON string:
        - 'split' : dict like {'index' -> [index], 'columns' -> [column
            'data' -> [values]]
        - 'records' : list like [{column -> value}, ... , {column -> val
            - 'index' : dict like {index -> {column -> value}}
            - 'columns' : dict like {column -> {index -> value}}
            - 'values' : just the values array
            - 'table' : dict like {'schema': {schema}, 'data': {data}}

        Describing the data, where data component is like ``orient='reco

```

```

date_format : {None, 'epoch', 'iso'}
    Type of date conversion. 'epoch' = epoch milliseconds,
    'iso' = ISO8601. The default depends on the `orient`. For
    ``orient='table'``, the default is 'iso'. For all other orients,
    the default is 'epoch'.
double_precision : int, default 10
    The number of decimal places to use when encoding
    floating point values.
force_ascii : bool, default True
    Force encoded string to be ASCII.
date_unit : str, default 'ms' (milliseconds)
    The time unit to encode to, governs timestamp and ISO8601
    precision. One of 's', 'ms', 'us', 'ns' for second, millisecond,
    microsecond, and nanosecond respectively.
default_handler : callable, default None
    Handler to call if object cannot otherwise be converted to a
    suitable format for JSON. Should receive a single argument which is
    the object to convert and return a serialisable object.
lines : bool, default False
    If 'orient' is 'records' write out line delimited json format. Will
    throw ValueError if incorrect 'orient' since others are not list
    like.

```

```

compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}

```

A string representing the compression to use in the output file, only used when the first argument is a filename. By default, the compression is inferred from the filename.

```

.. versionchanged:: 0.24.0
    'infer' option added and set to default

```

```

index : bool, default True
    Whether to include the index values in the JSON string. Not
    including the index (`index=False`) is only supported when
    orient is 'split' or 'table'.

```

```

indent : int, optional
    Length of whitespace used to indent each record.

```

```

.. versionadded:: 1.0.0

```

```

storage_options : dict, optional
    Extra options that make sense for a particular storage connection,
    e.g.
        host, port, username, password, etc., if using a URL that will
        be parsed by ``fsspec``, e.g., starting "s3://", "gcs://". An error
        will be raised if providing this argument with a non-fsspec URL.
        See the fsspec and backend storage implementation docs for the set o
        allowed keys and values.

```



```
.. versionadded:: 1.2.0
```

Returns

None or str

If `path_or_buf` is None, returns the resulting json format as a string. Otherwise returns None.

See Also

`read_json` : Convert a JSON string to pandas object.

Notes

The behavior of `indent=0` varies from the `stdlib`, which does not indent the output but does insert newlines. Currently, `indent=0` and the default `indent=None` are equivalent in pandas, though this may change in a future release.

`orient='table'` contains a `'pandas_version'` field under `'schema'`. This stores the version of `'pandas'` used in the latest revision of the schema.

Examples

```
>>> import json
>>> df = pd.DataFrame(
...     [{"a", "b"}, {"c", "d"}],
...     index=["row 1", "row 2"],
...     columns=["col 1", "col 2"],
... )

>>> result = df.to_json(orient="split")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
    "columns": [
        "col 1",
        "col 2"
    ],
    "index": [
        "row 1",
        "row 2"
    ],
    "data": [
        [
            "a",
            "b"
        ],
        [
            "c",
            "d"
        ]
    ]
}
```

Encoding/decoding a Dataframe using `'records'` formatted JSON. Note that index labels are not preserved with this encoding.

```
>>> result = df.to_json(orient="records")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
[
    {

```

```

        "col 1": "a",
        "col 2": "b"
    },
    {
        "col 1": "c",
        "col 2": "d"
    }
]

```

Encoding/decoding a Dataframe using ``'index'`` formatted JSON:

```

>>> result = df.to_json(orient="index")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
    "row 1": {
        "col 1": "a",
        "col 2": "b"
    },
    "row 2": {
        "col 1": "c",
        "col 2": "d"
    }
}

```

Encoding/decoding a Dataframe using ``'columns'`` formatted JSON:

```

>>> result = df.to_json(orient="columns")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
    "col 1": {
        "row 1": "a",
        "row 2": "c"
    },
    "col 2": {
        "row 1": "b",
        "row 2": "d"
    }
}

```

Encoding/decoding a Dataframe using ``'values'`` formatted JSON:

```

>>> result = df.to_json(orient="values")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
[
    [
        "a",
        "b"
    ],
    [
        "c",
        "d"
    ]
]

```

Encoding with Table Schema:

```

>>> result = df.to_json(orient="table")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
    "schema": {
        "fields": [

```

```

        {
            "name": "index",
            "type": "string"
        },
        {
            "name": "col 1",
            "type": "string"
        },
        {
            "name": "col 2",
            "type": "string"
        }
    ],
    "primaryKey": [
        "index"
    ],
    "pandas_version": "0.20.0"
},
"data": [
    {
        "index": "row 1",
        "col 1": "a",
        "col 2": "b"
    },
    {
        "index": "row 2",
        "col 1": "c",
        "col 2": "d"
    }
]
}

```

`to_latex(self, buf=None, columns=None, col_space=None, header=True, index=True, na_rep='NaN', formatters=None, float_format=None, sparsify=None, index_names=None, bold_rows=False, column_format=None, longtable=None, escape=None, encoding=None, decimal='.', multicolumn=None, multicolumn_format=None, multirow=None, caption=None, label=None, position=None)`

Render object to a LaTeX tabular, longtable, or nested table/tabular.

Requires `\\usepackage{booktabs}`. The output can be copy/pasted into a main LaTeX document or read from an external file with `\\input{table.tex}`.

.. versionchanged:: 1.0.0
Added caption and label arguments.

.. versionchanged:: 1.2.0
Added position argument, changed meaning of caption argument.

Parameters

`buf` : str, Path or StringIO-like, optional, default None
Buffer to write to. If None, the output is returned as a string.

`columns` : list of label, optional
The subset of columns to write. Writes all columns by default.

`col_space` : int, optional
The minimum width of each column.

`header` : bool or list of str, default True
Write out the column names. If a list of strings is given, it is assumed to be aliases for the column names.

`index` : bool, default True
Write row names (index).

`na_rep` : str, default 'NaN'
Missing data representation.

`formatters` : list of functions or dict of {str: function}, optional

```

    Formatter functions to apply to columns' elements by position or
    name. The result of each function must be a unicode string.
    List must be of length equal to the number of columns.
float_format : one-parameter function or str, optional, default None
    Formatter for floating point numbers. For example
    ``float_format="%0.2f"`` and ``float_format="{:0.2f}".format`` will
    both result in 0.1234 being formatted as 0.12.
sparsify : bool, optional
    Set to False for a DataFrame with a hierarchical index to print
    every multiindex key at each row. By default, the value will be
    read from the config module.
index_names : bool, default True
    Prints the names of the indexes.
bold_rows : bool, default False
    Make the row labels bold in the output.
column_format : str, optional
    The columns format as specified in `LaTeX table format
    <https://en.wikibooks.org/wiki/LaTeX/Tables>`__ e.g. 'rcl' for 3
    columns. By default, 'l' will be used for all columns except
    columns of numbers, which default to 'r'.
longtable : bool, optional
    By default, the value will be read from the pandas config
    module. Use a longtable environment instead of tabular. Requires
    adding a \usepackage{longtable} to your LaTeX preamble.
escape : bool, optional
    By default, the value will be read from the pandas config
    module. When set to False prevents from escaping latex special
    characters in column names.
encoding : str, optional
    A string representing the encoding to use in the output file,
    defaults to 'utf-8'.
decimal : str, default '.'
    Character recognized as decimal separator, e.g. ',' in Europe.
multicolumn : bool, default True
    Use \multicolumn to enhance MultiIndex columns.
    The default will be read from the config module.
multicolumn_format : str, default 'l'
    The alignment for multicolumns, similar to `column_format`
    The default will be read from the config module.
multirow : bool, default False
    Use \multirow to enhance MultiIndex rows. Requires adding a
    \usepackage{multirow} to your LaTeX preamble. Will print
    centered labels (instead of top-aligned) across the contained
    rows, separating groups via clines. The default will be read
    from the pandas config module.
caption : str or tuple, optional
    Tuple (full_caption, short_caption),
    which results in ``\caption[short_caption]{full_caption}``;
    if a single string is passed, no short caption will be set.

    .. versionadded:: 1.0.0

    .. versionchanged:: 1.2.0
        Optionally allow caption to be a tuple ``(full_caption, short_cap
tion)``.

label : str, optional
    The LaTeX label to be placed inside ``\label{}`` in the output.
    This is used with ``\ref{}`` in the main ``.tex`` file.

    .. versionadded:: 1.0.0
position : str, optional
    The LaTeX positional argument for tables, to be placed after
    ``\begin{}`` in the output.

```

.. versionadded:: 1.2.0

Returns

str or None

If buf is None, returns the result as a string. Otherwise returns None.

See Also

DataFrame.to_string : Render a DataFrame to a console-friendly tabular output.

DataFrame.to_html : Render a DataFrame as an HTML table.

Examples

```
>>> df = pd.DataFrame(dict(name=['Raphael', 'Donatello'],
...                          mask=['red', 'purple'],
...                          weapon=['sai', 'bo staff']))
>>> print(df.to_latex(index=False)) # doctest: +NORMALIZE_WHITESPACE
\begin{tabular}{lll}
\toprule
    name & mask & weapon \\
\midrule
    Raphael & red & sai \\
    Donatello & purple & bo staff \\
\bottomrule
\end{tabular}
```

to_pickle(self, path, compression: 'CompressionOptions' = 'infer', protocol: 'int' = 5, storage_options: 'StorageOptions' = None) -> 'None'

Pickle (serialize) object to file.

Parameters

path : str

File path where the pickled object will be stored.

compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}, default

to 'infer'

A string representing the compression to use in the output file. By default, infers from the file extension in specified path.

Compression mode may be any of the following possible values: {'infer', 'gzip', 'bz2', 'zip', 'xz', None}. If compression mode is 'infer' and path_or_buf is path-like, then detect compression mode from the following extensions: '.gz', '.bz2', '.zip' or '.xz'. (otherwise no compression).

If dict given and mode is 'zip' or inferred as 'zip', other entries passed as additional compression options.

protocol : int

Int which indicates which protocol should be used by the pickler, default HIGHEST_PROTOCOL (see [1]_ paragraph 12.1.2). The possible values are 0, 1, 2, 3, 4, 5. A negative value for the protocol parameter is equivalent to setting its value to HIGHEST_PROTOCOL.

.. [1] <https://docs.python.org/3/library/pickle.html>.

storage_options : dict, optional

Extra options that make sense for a particular storage connection,

e.g.

host, port, username, password, etc., if using a URL that will be parsed by ``fsspec``, e.g., starting "s3://", "gcs://". An error will be raised if providing this argument with a non-fsspec URL.

See the fsspec and backend storage implementation docs for the set o

f

allowed keys and values.

.. versionadded:: 1.2.0

See Also

 read_pickle : Load pickled pandas object (or any object) from file.
 DataFrame.to_hdf : Write DataFrame to an HDF5 file.
 DataFrame.to_sql : Write DataFrame to a SQL database.
 DataFrame.to_parquet : Write a DataFrame to the binary parquet format.

Examples

 >>> original_df = pd.DataFrame({"foo": range(5), "bar": range(5, 10)})
 >>> original_df
 foo bar
 0 0 5
 1 1 6
 2 2 7
 3 3 8
 4 4 9
 >>> original_df.to_pickle("./dummy.pkl")

 >>> unpickled_df = pd.read_pickle("./dummy.pkl")
 >>> unpickled_df
 foo bar
 0 0 5
 1 1 6
 2 2 7
 3 3 8
 4 4 9

 >>> import os
 >>> os.remove("./dummy.pkl")

to_sql(self, name: 'str', con, schema=None, if_exists: 'str' = 'fail', index_label=None, chunksize=None, dtype=None, method=None) -> 'None'

Write records stored in a DataFrame to a SQL database.

Databases supported by SQLAlchemy [1]_ are supported. Tables can be newly created, appended to, or overwritten.

Parameters

 name : str
 Name of SQL table.
 con : sqlalchemy.engine.(Engine or Connection) or sqlite3.Connection
 Using SQLAlchemy makes it possible to use any DB supported by that library. Legacy support is provided for sqlite3.Connection objects.

The user

is responsible for engine disposal and connection closure for the SQLAlchemy

connectable

See `here`_

<<https://docs.sqlalchemy.org/en/13/core/connections.html>>`_.

schema : str, optional
 Specify the schema (if database flavor supports this). If None, use default schema.
 if_exists : {'fail', 'replace', 'append'}, default 'fail'
 How to behave if the table already exists.

 * fail: Raise a ValueError.
 * replace: Drop the table before inserting new values.
 * append: Insert new values to the existing table.

```

index : bool, default True
    Write DataFrame index as a column. Uses `index_label` as the column
    name in the table.
index_label : str or sequence, default None
    Column label for index column(s). If None is given (default) and
    `index` is True, then the index names are used.
    A sequence should be given if the DataFrame uses MultiIndex.
chunksize : int, optional
    Specify the number of rows in each batch to be written at a time.
    By default, all rows will be written at once.
dtype : dict or scalar, optional
    Specifying the datatype for columns. If a dictionary is used, the
    keys should be the column names and the values should be the
    SQLAlchemy types or strings for the sqlite3 legacy mode. If a
    scalar is provided, it will be applied to all columns.
method : {None, 'multi', callable}, optional
    Controls the SQL insertion clause used:

    * None : Uses standard SQL ``INSERT`` clause (one per row).
    * 'multi': Pass multiple values in a single ``INSERT`` clause.
    * callable with signature ``(pd_table, conn, keys, data_iter)``.

    Details and a sample callable implementation can be found in the
    section :ref:`insert method <io.sql.method>`.

.. versionadded:: 0.24.0

```

Raises

ValueError

When the table already exists and `if_exists` is 'fail' (the default).

See Also

`read_sql` : Read a DataFrame from a table.

Notes

Timezone aware datetime columns will be written as ``Timestamp with timezone`` type with SQLAlchemy if supported by the database. Otherwise, the datetimes will be stored as timezone unaware timestamps local to the original timezone.

.. versionadded:: 0.24.0

References

- .. [1] <https://docs.sqlalchemy.org>
- .. [2] <https://www.python.org/dev/peps/pep-0249/>

Examples

Create an in-memory SQLite database.

```

>>> from sqlalchemy import create_engine
>>> engine = create_engine('sqlite://', echo=False)

```

Create a table from scratch with 3 rows.

```

>>> df = pd.DataFrame({'name' : ['User 1', 'User 2', 'User 3']})
>>> df
   name
0  User 1

```

```
1 User 2
2 User 3
```

```
>>> df.to_sql('users', con=engine)
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 1'), (1, 'User 2'), (2, 'User 3')]
```

An `sqlalchemy.engine.Connection` can also be passed to `con`:

```
>>> with engine.begin() as connection:
...     df1 = pd.DataFrame({'name' : ['User 4', 'User 5']})
...     df1.to_sql('users', con=connection, if_exists='append')
```

This is allowed to support operations that require that the same DBAPI connection is used for the entire operation.

```
>>> df2 = pd.DataFrame({'name' : ['User 6', 'User 7']})
>>> df2.to_sql('users', con=engine, if_exists='append')
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 1'), (1, 'User 2'), (2, 'User 3'),
 (0, 'User 4'), (1, 'User 5'), (0, 'User 6'),
 (1, 'User 7')]
```

Overwrite the table with just `df2`.

```
>>> df2.to_sql('users', con=engine, if_exists='replace',
...           index_label='id')
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 6'), (1, 'User 7')]
```

Specify the dtype (especially useful for integers with missing values). Notice that while pandas is forced to store the data as floating point, the database supports nullable integers. When fetching the data with Python, we get back integer scalars.

```
>>> df = pd.DataFrame({"A": [1, None, 2]})
>>> df
   A
0  1.0
1  NaN
2  2.0
```

```
>>> from sqlalchemy.types import Integer
>>> df.to_sql('integers', con=engine, index=False,
...          dtype={"A": Integer()})
>>> engine.execute("SELECT * FROM integers").fetchall()
[(1,), (None,), (2,)]
```

`to_xarray(self)`

Return an xarray object from the pandas object.

Returns

xarray.DataArray or xarray.Dataset

Data in the pandas structure converted to Dataset if the object is a DataFrame, or a DataArray if the object is a Series.

See Also

`DataFrame.to_hdf` : Write DataFrame to an HDF5 file.

`DataFrame.to_parquet` : Write a DataFrame to the binary parquet format.

Notes

See the `xarray docs <<https://xarray.pydata.org/en/stable/>>`__

Examples

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0, 2),
...                    ('parrot', 'bird', 24.0, 2),
...                    ('lion', 'mammal', 80.5, 4),
...                    ('monkey', 'mammal', np.nan, 4)],
...                    columns=['name', 'class', 'max_speed',
...                              'num_legs'])
```

```
>>> df
   name  class  max_speed  num_legs
0  falcon   bird     389.0         2
1  parrot   bird     24.0         2
2   lion  mammal     80.5         4
3  monkey  mammal      NaN         4
```

```
>>> df.to_xarray()
<xarray.Dataset>
Dimensions:  (index: 4)
Coordinates:
  * index    (index) int64 0 1 2 3
Data variables:
  name      (index) object 'falcon' 'parrot' 'lion' 'monkey'
  class     (index) object 'bird' 'bird' 'mammal' 'mammal'
  max_speed (index) float64 389.0 24.0 80.5 nan
  num_legs  (index) int64 2 2 4 4
```

```
>>> df['max_speed'].to_xarray()
<xarray.DataArray 'max_speed' (index: 4)>
array([389. , 24. , 80.5,  nan])
Coordinates:
  * index    (index) int64 0 1 2 3
```

```
>>> dates = pd.to_datetime(['2018-01-01', '2018-01-01',
...                          '2018-01-02', '2018-01-02'])
>>> df_multiindex = pd.DataFrame({'date': dates,
...                                'animal': ['falcon', 'parrot',
...                                           'falcon', 'parrot'],
...                                'speed': [350, 18, 361, 15]})
>>> df_multiindex = df_multiindex.set_index(['date', 'animal'])
```

```
>>> df_multiindex
              speed
date  animal
2018-01-01 falcon    350
           parrot    18
2018-01-02 falcon    361
           parrot    15
```

```
>>> df_multiindex.to_xarray()
<xarray.Dataset>
Dimensions:  (animal: 2, date: 2)
Coordinates:
  * date      (date) datetime64[ns] 2018-01-01 2018-01-02
  * animal    (animal) object 'falcon' 'parrot'
Data variables:
  speed      (date, animal) int64 350 18 361 15
```

```
truncate(self: 'FrameOrSeries', before=None, after=None, axis=None, copy: 'bool_t' = True) -> 'FrameOrSeries'
```

Truncate a Series or DataFrame before and after some index value.

This is a useful shorthand for boolean indexing based on index values above or below certain thresholds.

Parameters

before : date, str, int
Truncate all rows before this index value.
after : date, str, int
Truncate all rows after this index value.
axis : {0 or 'index', 1 or 'columns'}, optional
Axis to truncate. Truncates the index (rows) by default.
copy : bool, default is True,
Return a copy of the truncated section.

Returns

type of caller
The truncated Series or DataFrame.

See Also

DataFrame.loc : Select a subset of a DataFrame by label.
DataFrame.iloc : Select a subset of a DataFrame by position.

Notes

If the index being truncated contains only datetime values, 'before' and 'after' may be specified as strings instead of Timestamps.

Examples

```
>>> df = pd.DataFrame({'A': ['a', 'b', 'c', 'd', 'e'],
...                    'B': ['f', 'g', 'h', 'i', 'j'],
...                    'C': ['k', 'l', 'm', 'n', 'o']},
...                    index=[1, 2, 3, 4, 5])
```

```
>>> df
   A  B  C
1  a  f  k
2  b  g  l
3  c  h  m
4  d  i  n
5  e  j  o
```

```
>>> df.truncate(before=2, after=4)
```

```
   A  B  C
2  b  g  l
3  c  h  m
4  d  i  n
```

The columns of a DataFrame can be truncated.

```
>>> df.truncate(before="A", after="B", axis="columns")
```

```
   A  B
1  a  f
2  b  g
3  c  h
4  d  i
5  e  j
```

For Series, only rows can be truncated.

```
>>> df['A'].truncate(before=2, after=4)
```

```
2    b
3    c
4    d
Name: A, dtype: object
```

The index values in ``truncate`` can be datetimes or string dates.

```
>>> dates = pd.date_range('2016-01-01', '2016-02-01', freq='s')
>>> df = pd.DataFrame(index=dates, data={'A': 1})
>>> df.tail()
```

```

      A
2016-01-31 23:59:56  1
2016-01-31 23:59:57  1
2016-01-31 23:59:58  1
2016-01-31 23:59:59  1
2016-02-01 00:00:00  1
```

```
>>> df.truncate(before=pd.Timestamp('2016-01-05'),
...              after=pd.Timestamp('2016-01-10')).tail()
```

```

      A
2016-01-09 23:59:56  1
2016-01-09 23:59:57  1
2016-01-09 23:59:58  1
2016-01-09 23:59:59  1
2016-01-10 00:00:00  1
```

Because the index is a DatetimeIndex containing only dates, we can specify `before` and `after` as strings. They will be coerced to Timestamps before truncation.

```
>>> df.truncate('2016-01-05', '2016-01-10').tail()
```

```

      A
2016-01-09 23:59:56  1
2016-01-09 23:59:57  1
2016-01-09 23:59:58  1
2016-01-09 23:59:59  1
2016-01-10 00:00:00  1
```

Note that ``truncate`` assumes a 0 value for any unspecified time component (midnight). This differs from partial string slicing, which returns any partially matching dates.

```
>>> df.loc['2016-01-05':'2016-01-10', :].tail()
```

```

      A
2016-01-10 23:59:55  1
2016-01-10 23:59:56  1
2016-01-10 23:59:57  1
2016-01-10 23:59:58  1
2016-01-10 23:59:59  1
```

```
0) tshift(self: 'FrameOrSeries', periods: 'int' = 1, freq=None, axis: 'Axis' =
-> 'FrameOrSeries'
    Shift the time index, using the index's frequency if available.
```

```
.. deprecated:: 1.1.0
   Use `shift` instead.
```

Parameters

```

periods : int
    Number of periods to move, can be positive or negative.
freq : DateOffset, timedelta, or str, default None
    Increment to use from the tseries module
    or time rule expressed as a string (e.g. 'EOM').
axis : {0 or 'index', 1 or 'columns', None}, default 0
    Corresponds to the axis that contains the Index.
```

Returns

shifted : Series/DataFrame

Notes

If freq is not specified then tries to use the freq or inferred_freq attributes of the index. If neither of those attributes exist, a ValueError is thrown

tz_convert(self: 'FrameOrSeries', tz, axis=0, level=None, copy: 'bool_t' = True) -> 'FrameOrSeries'

Convert tz-aware axis to target time zone.

Parameters

tz : str or tzinfo object

axis : the axis to convert

level : int, str, default None

If axis is a MultiIndex, convert a specific level. Otherwise must be None.

copy : bool, default True

Also make a copy of the underlying data.

Returns

{klass}

Object with time zone converted axis.

Raises

TypeError

If the axis is tz-naive.

tz_localize(self: 'FrameOrSeries', tz, axis=0, level=None, copy: 'bool_t' = True, ambiguous='raise', nonexistent: 'str' = 'raise') -> 'FrameOrSeries'

Localize tz-naive index of a Series or DataFrame to target time zone.

This operation localizes the Index. To localize the values in a timezone-naive Series, use :meth:`Series.dt.tz_localize`.

Parameters

tz : str or tzinfo

axis : the axis to localize

level : int, str, default None

If axis is a MultiIndex, localize a specific level. Otherwise must be None.

copy : bool, default True

Also make a copy of the underlying data.

ambiguous : 'infer', bool-ndarray, 'NaT', default 'raise'

When clocks moved backward due to DST, ambiguous times may arise. For example in Central European Time (UTC+01), when going from 03:00 DST to 02:00 non-DST, 02:30:00 local time occurs both at 00:30:00 UTC and at 01:30:00 UTC. In such a situation, the `ambiguous` parameter dictates how ambiguous times should be handled.

- 'infer' will attempt to infer fall dst-transition hours based on order
- bool-ndarray where True signifies a DST time, False designates a non-DST time (note that this flag is only applicable for ambiguous times)
- 'NaT' will return NaT where there are ambiguous times
- 'raise' will raise an AmbiguousTimeError if there are ambiguous times.

```

nonexistent : str, default 'raise'
    A nonexistent time does not exist in a particular timezone
    where clocks moved forward due to DST. Valid values are:

    - 'shift_forward' will shift the nonexistent time forward to the
      closest existing time
    - 'shift_backward' will shift the nonexistent time backward to the
      closest existing time
    - 'NaT' will return NaT where there are nonexistent times
    - timedelta objects will shift nonexistent times by the timedelta
    - 'raise' will raise a NonExistentTimeError if there are
      nonexistent times.

    .. versionadded:: 0.24.0

```

Returns

```

Series or DataFrame
    Same type as the input.

```

Raises

```

TypeError
    If the TimeSeries is tz-aware and tz is not None.

```

Examples

Localize local times:

```

>>> s = pd.Series([1],
...                 index=pd.DatetimeIndex(['2018-09-15 01:30:00']))
>>> s.tz_localize('CET')
2018-09-15 01:30:00+02:00    1
dtype: int64

```

Be careful with DST changes. When there is sequential data, pandas can infer the DST time:

```

>>> s = pd.Series(range(7),
...                 index=pd.DatetimeIndex(['2018-10-28 01:30:00',
...                                         '2018-10-28 02:00:00',
...                                         '2018-10-28 02:30:00',
...                                         '2018-10-28 02:00:00',
...                                         '2018-10-28 02:30:00',
...                                         '2018-10-28 03:00:00',
...                                         '2018-10-28 03:30:00']))
>>> s.tz_localize('CET', ambiguous='infer')
2018-10-28 01:30:00+02:00    0
2018-10-28 02:00:00+02:00    1
2018-10-28 02:30:00+02:00    2
2018-10-28 02:00:00+01:00    3
2018-10-28 02:30:00+01:00    4
2018-10-28 03:00:00+01:00    5
2018-10-28 03:30:00+01:00    6
dtype: int64

```

In some cases, inferring the DST is impossible. In such cases, you can pass an ndarray to the ambiguous parameter to set the DST explicitly

```

>>> s = pd.Series(range(3),
...                 index=pd.DatetimeIndex(['2018-10-28 01:20:00',
...                                         '2018-10-28 02:36:00',
...                                         '2018-10-28 03:46:00']))
>>> s.tz_localize('CET', ambiguous=np.array([True, True, False]))
2018-10-28 01:20:00+02:00    0

```

```
2018-10-28 02:36:00+02:00    1
2018-10-28 03:46:00+01:00    2
dtype: int64
```

If the DST transition causes nonexistent times, you can shift these dates forward or backward with a `timedelta` object or `'shift_forward'` or `'shift_backward'`.

```
>>> s = pd.Series(range(2),
...               index=pd.DatetimeIndex(['2015-03-29 02:30:00',
...                                       '2015-03-29 03:30:00']))
>>> s.tz_localize('Europe/Warsaw', nonexistent='shift_forward')
2015-03-29 03:00:00+02:00    0
2015-03-29 03:30:00+02:00    1
dtype: int64
>>> s.tz_localize('Europe/Warsaw', nonexistent='shift_backward')
2015-03-29 01:59:59.999999999+01:00    0
2015-03-29 03:30:00+02:00    1
dtype: int64
>>> s.tz_localize('Europe/Warsaw', nonexistent=pd.Timedelta('1H'))
2015-03-29 03:30:00+02:00    0
2015-03-29 03:30:00+02:00    1
dtype: int64
```

```
where(self, cond, other=nan, inplace=False, axis=None, level=None, errors='raise', try_cast=False)
```

Replace values where the condition is False.

Parameters

cond : bool Series/DataFrame, array-like, or callable
Where `'cond'` is True, keep the original value. Where False, replace with corresponding value from `'other'`. If `'cond'` is callable, it is computed on the Series/DataFrame and should return boolean Series/DataFrame or array. The callable must not change input Series/DataFrame (though pandas doesn't check it).

other : scalar, Series/DataFrame, or callable
Entries where `'cond'` is False are replaced with corresponding value from `'other'`. If `other` is callable, it is computed on the Series/DataFrame and should return scalar or Series/DataFrame. The callable must not change input Series/DataFrame (though pandas doesn't check it).

inplace : bool, default False
Whether to perform the operation in place on the data.

axis : int, default None
Alignment axis if needed.

level : int, default None
Alignment level if needed.

errors : str, {'raise', 'ignore'}, default 'raise'
Note that currently this parameter won't affect the results and will always coerce to a suitable dtype.

- 'raise' : allow exceptions to be raised.
- 'ignore' : suppress exceptions. On error return original object.

try_cast : bool, default False
Try to cast the result back to the input type (if possible).

Returns

Same type as caller or None if `'inplace=True'`.

See Also

`:func:`DataFrame.mask`` : Return an object of same shape as

self.

Notes

The where method is an application of the if-then idiom. For each element in the calling DataFrame, if ``cond`` is ``True`` the element is used; otherwise the corresponding element from the DataFrame ``other`` is used.

The signature for :func:`DataFrame.where` differs from :func:`numpy.where`. Roughly ``df1.where(m, df2)`` is equivalent to ``np.where(m, df1, df2)``.

For further details and examples see the ``where`` documentation in :ref:`indexing <indexing.where_mask>`.

Examples

```
>>> s = pd.Series(range(5))
>>> s.where(s > 0)
0    NaN
1    1.0
2    2.0
3    3.0
4    4.0
dtype: float64
>>> s.mask(s > 0)
0    0.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64

>>> s.where(s > 1, 10)
0    10
1    10
2     2
3     3
4     4
dtype: int64
>>> s.mask(s > 1, 10)
0     0
1     1
2    10
3    10
4    10
dtype: int64

>>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A', 'B'])
>>> df
   A  B
0  0  1
1  2  3
2  4  5
3  6  7
4  8  9
>>> m = df % 3 == 0
>>> df.where(m, -df)
   A  B
0  0 -1
1 -2  3
2 -4 -5
3  6 -7
4 -8  9
```

```
>>> df.where(m, -df) == np.where(m, df, -df)
      A      B
0  True  True
1  True  True
2  True  True
3  True  True
4  True  True
>>> df.where(m, -df) == df.mask(~m, -df)
      A      B
0  True  True
1  True  True
2  True  True
3  True  True
4  True  True
```

`xs(self, key, axis=0, level=None, drop_level: 'bool_t' = True)`
Return cross-section from the Series/DataFrame.

This method takes a `key` argument to select data at a particular level of a MultiIndex.

Parameters

`key` : label or tuple of label
Label contained in the index, or partially in a MultiIndex.
`axis` : {0 or 'index', 1 or 'columns'}, default 0
Axis to retrieve cross-section on.
`level` : object, defaults to first n levels (n=1 or len(key))
In case of a key partially contained in a MultiIndex, indicate which levels are used. Levels can be referred by label or position.
`drop_level` : bool, default True
If False, returns object with same levels as self.

Returns

Series or DataFrame
Cross-section from the original Series or DataFrame corresponding to the selected index levels.

See Also

`DataFrame.loc` : Access a group of rows and columns by label(s) or a boolean array.
`DataFrame.iloc` : Purely integer-location based indexing for selection by position.

Notes

`xs` can not be used to set values.

MultiIndex Slicers is a generic way to get/set values on any level or levels.

It is a superset of `xs` functionality, see :ref:`MultiIndex Slicers <advanced.mi_slicers>`.

Examples

```
>>> d = {'num_legs': [4, 4, 2, 2],
...      'num_wings': [0, 0, 2, 2],
...      'class': ['mammal', 'mammal', 'mammal', 'bird'],
...      'animal': ['cat', 'dog', 'bat', 'penguin'],
...      'locomotion': ['walks', 'walks', 'flies', 'walks']}
>>> df = pd.DataFrame(data=d)
>>> df = df.set_index(['class', 'animal', 'locomotion'])
>>> df
```


| | | | num_legs | num_wings |
|--------|---------|------------|----------|-----------|
| class | animal | locomotion | | |
| mammal | cat | walks | 4 | 0 |
| | dog | walks | 4 | 0 |
| | bat | flies | 2 | 2 |
| bird | penguin | walks | 2 | 2 |

Get values at specified index

```
>>> df.xs('mammal')
          num_legs  num_wings
animal locomotion
cat     walks      4          0
dog     walks      4          0
bat     flies      2          2
```

Get values at several indexes

```
>>> df.xs(('mammal', 'dog'))
          num_legs  num_wings
locomotion
walks           4          0
```

Get values at specified index and level

```
>>> df.xs('cat', level=1)
          num_legs  num_wings
class locomotion
mammal walks      4          0
```

Get values at several indexes and levels

```
>>> df.xs(('bird', 'walks'),
...       level=[0, 'locomotion'])
          num_legs  num_wings
animal
penguin           2          2
```

Get values at specified column and axis

```
>>> df.xs('num_wings', axis=1)
class animal locomotion
mammal cat     walks      0
        dog     walks      0
        bat     flies      2
bird   penguin walks      2
Name: num_wings, dtype: int64
```

Readonly properties inherited from `pandas.core.generic.NDFrame`:

`dtypes`

Return the dtypes in the DataFrame.

This returns a Series with the data type of each column. The result's index is the original DataFrame's columns. Columns with mixed types are stored with the ``object`` dtype. See :ref:`the User Guide <basics.dtypes>` for more.

Returns

`pandas.Series`

The data type of each column.

Examples

```

>>> df = pd.DataFrame({'float': [1.0],
...                     'int': [1],
...                     'datetime': [pd.Timestamp('20180310')],
...                     'string': ['foo']})
>>> df.dtypes
float          float64
int            int64
datetime      datetime64[ns]
string         object
dtype: object

```

empty

Indicator whether DataFrame is empty.

True if DataFrame is entirely empty (no items), meaning any of the axes are of length 0.

Returns

bool

If DataFrame is empty, return True, if not return False.

See Also

`Series.dropna` : Return series without null values.

`DataFrame.dropna` : Return DataFrame with labels on given axis omitted where (all or any) data are missing.

Notes

If DataFrame contains only NaNs, it is still not considered empty. See the example below.

Examples

An example of an actual empty DataFrame. Notice the index is empty:

```

>>> df_empty = pd.DataFrame({'A' : []})
>>> df_empty
Empty DataFrame
Columns: [A]
Index: []
>>> df_empty.empty
True

```

If we only have NaNs in our DataFrame, it is not considered empty! We will need to drop the NaNs to make the DataFrame empty:

```

>>> df = pd.DataFrame({'A' : [np.nan]})
>>> df
   A
0 NaN
>>> df.empty
False
>>> df.dropna().empty
True

```

flags

Get the properties associated with this pandas object.

The available flags are

* :attr:`Flags.allows_duplicate_labels`

See Also

Flags : Flags that apply to pandas objects.
 DataFrame.attrs : Global metadata applying to this dataset.

Notes

"Flags" differ from "metadata". Flags reflect properties of the pandas object (the Series or DataFrame). Metadata refer to properties of the dataset, and should be stored in :attr:`DataFrame.attrs`.

Examples

```
>>> df = pd.DataFrame({"A": [1, 2]})
>>> df.flags
<Flags(allows_duplicate_labels=True)>
```

Flags can be get or set using ``.``

```
>>> df.flags.allows_duplicate_labels
True
>>> df.flags.allows_duplicate_labels = False
```

Or by slicing with a key

```
>>> df.flags["allows_duplicate_labels"]
False
>>> df.flags["allows_duplicate_labels"] = True
```

ndim

Return an int representing the number of axes / array dimensions.

Return 1 if Series. Otherwise return 2 if DataFrame.

See Also

ndarray.ndim : Number of array dimensions.

Examples

```
>>> s = pd.Series({'a': 1, 'b': 2, 'c': 3})
>>> s.ndim
1

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.ndim
2
```

size

Return an int representing the number of elements in this object.

Return the number of rows if Series. Otherwise return the number of rows times number of columns if DataFrame.

See Also

ndarray.size : Number of elements in the array.

Examples

```
>>> s = pd.Series({'a': 1, 'b': 2, 'c': 3})
>>> s.size
3

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
```

```
>>> df.size
4
```

values

Return a Numpy representation of the DataFrame.

.. warning::

We recommend using :meth:`DataFrame.to_numpy` instead.

Only the values in the DataFrame will be returned, the axes labels will be removed.

Returns

numpy.ndarray
The values of the DataFrame.

See Also

DataFrame.to_numpy : Recommended alternative to this method.

DataFrame.index : Retrieve the index labels.

DataFrame.columns : Retrieving the column names.

Notes

The dtype will be a lower-common-denominator dtype (implicit upcasting); that is to say if the dtypes (even of numeric types) are mixed, the one that accommodates all will be chosen. Use this with care if you are not dealing with the blocks.

e.g. If the dtypes are float16 and float32, dtype will be upcast to float32. If dtypes are int32 and uint8, dtype will be upcast to int32. By :func:`numpy.find_common_type` convention, mixing int64 and uint64 will result in a float64 dtype.

Examples

A DataFrame where all columns are the same type (e.g., int64) results in an array of the same type.

```
>>> df = pd.DataFrame({'age': [3, 29],
...                     'height': [94, 170],
...                     'weight': [31, 115]})
>>> df
   age  height  weight
0    3     94     31
1   29    170    115
>>> df.dtypes
age          int64
height       int64
weight       int64
dtype: object
>>> df.values
array([[ 3, 94, 31],
       [29, 170, 115]])
```

A DataFrame with mixed type columns (e.g., str/object, int64, float32) results in an ndarray of the broadest type that accommodates these mixed types (e.g., object).

```
>>> df2 = pd.DataFrame([('parrot', 24.0, 'second'),
...                     ('lion', 80.5, 1),
...                     ('monkey', np.nan, None)],
...                     columns=('name', 'max_speed', 'rank'))
```

```
>>> df2.dtypes
name          object
max_speed     float64
rank          object
dtype: object
>>> df2.values
array([[ 'parrot', 24.0, 'second'],
       [ 'lion', 80.5, 1],
       [ 'monkey', nan, None]], dtype=object)
```

Data descriptors inherited from pandas.core.generic.NDFrame:

attrs

Dictionary of global attributes of this dataset.

.. warning::

attrs is experimental and may change without warning.

See Also

DataFrame.flags : Global flags applying to this object.

Data and other attributes inherited from pandas.core.generic.NDFrame:

__array_priority__ = 1000

Methods inherited from pandas.core.base.PandasObject:

__sizeof__(self)

Generates the total memory usage for an object that returns either a value or Series of values

Methods inherited from pandas.core.accessor.DirNamesMixin:

__dir__(self) -> List[str]

Provide method name lookup and completion.

Notes

Only provide 'public' methods.

Data descriptors inherited from pandas.core.accessor.DirNamesMixin:

__dict__

dictionary for instance variables (if defined)

__weakref__

list of weak references to the object (if defined)

Readonly properties inherited from pandas.core.indexing.IndexingMixin:

at

Access a single value for a row/column label pair.

Similar to ``loc``, in that both provide label-based lookups. Use ``at`` if you only need to get or set a single value in a DataFrame or Series.

Raises

KeyError

If 'label' does not exist in DataFrame.

See Also

DataFrame.iat : Access a single value for a row/column pair by integer position.

DataFrame.loc : Access a group of rows and columns by label(s).

Series.at : Access a single value using a label.

Examples

```
>>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
...                     index=[4, 5, 6], columns=['A', 'B', 'C'])
```

```
>>> df
   A  B  C
4  0  2  3
5  0  4  1
6 10 20 30
```

Get value at specified row/column pair

```
>>> df.at[4, 'B']
2
```

Set value at specified row/column pair

```
>>> df.at[4, 'B'] = 10
>>> df.at[4, 'B']
10
```

Get value within a Series

```
>>> df.loc[5].at['B']
4
```

iat

Access a single value for a row/column pair by integer position.

Similar to ``iloc``, in that both provide integer-based lookups. Use ``iat`` if you only need to get or set a single value in a DataFrame or Series.

Raises

IndexError

When integer position is out of bounds.

See Also

DataFrame.at : Access a single value for a row/column label pair.

DataFrame.loc : Access a group of rows and columns by label(s).

DataFrame.iloc : Access a group of rows and columns by integer position

(s).

Examples

```
>>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
...                     columns=['A', 'B', 'C'])
```

```
>>> df
   A  B  C
0  0  2  3
1  0  4  1
```

```
2 10 20 30
```

Get value at specified row/column pair

```
>>> df.iat[1, 2]
1
```

Set value at specified row/column pair

```
>>> df.iat[1, 2] = 10
>>> df.iat[1, 2]
10
```

Get value within a series

```
>>> df.loc[0].iat[1]
2
```

iloc

Purely integer-location based indexing for selection by position.

`df.iloc[]` is primarily integer position based (from `0` to `length-1` of the axis), but may also be used with a boolean array.

Allowed inputs are:

- An integer, e.g. `5`.
- A list or array of integers, e.g. `[4, 3, 0]`.
- A slice object with ints, e.g. `1:7`.
- A boolean array.
- A callable function with one argument (the calling Series or DataFrame) and that returns valid output for indexing (one of the above).

This is useful in method chains, when you don't have a reference to the calling object, but would like to base your selection on some value.

`df.iloc` will raise `IndexError` if a requested indexer is out-of-bounds, except `*slice*` indexers which allow out-of-bounds indexing (this conforms with python/numpy `*slice*` semantics).

See more at :ref:`Selection by Position <indexing.integer>`.

See Also

`DataFrame.iat` : Fast integer location scalar accessor.

`DataFrame.loc` : Purely label-location based indexer for selection by label.

`Series.iloc` : Purely integer-location based indexing for selection by position.

Examples

```
>>> mydict = [{'a': 1, 'b': 2, 'c': 3, 'd': 4},
...           {'a': 100, 'b': 200, 'c': 300, 'd': 400},
...           {'a': 1000, 'b': 2000, 'c': 3000, 'd': 4000}]
>>> df = pd.DataFrame(mydict)
>>> df
```

| | a | b | c | d |
|---|------|------|------|------|
| 0 | 1 | 2 | 3 | 4 |
| 1 | 100 | 200 | 300 | 400 |
| 2 | 1000 | 2000 | 3000 | 4000 |

****Indexing just the rows****

With a scalar integer.

```
>>> type(df.iloc[0])
<class 'pandas.core.series.Series'>
>>> df.iloc[0]
a    1
b    2
c    3
d    4
Name: 0, dtype: int64
```

With a list of integers.

```
>>> df.iloc[[0]]
   a  b  c  d
0  1  2  3  4
>>> type(df.iloc[[0]])
<class 'pandas.core.frame.DataFrame'>

>>> df.iloc[[0, 1]]
   a    b    c    d
0   1    2    3    4
1  100  200  300  400
```

With a `slice` object.

```
>>> df.iloc[:3]
   a    b    c    d
0   1    2    3    4
1  100  200  300  400
2 1000 2000 3000 4000
```

With a boolean mask the same length as the index.

```
>>> df.iloc[[True, False, True]]
   a    b    c    d
0   1    2    3    4
2 1000 2000 3000 4000
```

With a callable, useful in method chains. The `x` passed to the ``lambda`` is the DataFrame being sliced. This selects the rows whose index label even.

```
>>> df.iloc[lambda x: x.index % 2 == 0]
   a    b    c    d
0   1    2    3    4
2 1000 2000 3000 4000
```

****Indexing both axes****

You can mix the indexer types for the index and columns. Use ``:``` to select the entire axis.

With scalar integers.

```
>>> df.iloc[0, 1]
2
```

With lists of integers.

```
>>> df.iloc[[0, 2], [1, 3]]
   b    d
0   2    4
2 2000 4000
```


With `slice` objects.

```
>>> df.iloc[1:3, 0:3]
      a      b      c
1    100    200    300
2   1000   2000   3000
```

With a boolean array whose length matches the columns.

```
>>> df.iloc[:, [True, False, True, False]]
      a      c
0      1      3
1    100    300
2   1000   3000
```

With a callable function that expects the Series or DataFrame.

```
>>> df.iloc[:, lambda df: [0, 2]]
      a      c
0      1      3
1    100    300
2   1000   3000
```

loc

Access a group of rows and columns by label(s) or a boolean array.

``.loc[]`` is primarily label based, but may also be used with a boolean array.

Allowed inputs are:

- A single label, e.g. ``5`` or ``'a'``, (note that ``5`` is interpreted as a *label* of the index, and *never* as an integer position along the index).
- A list or array of labels, e.g. ``['a', 'b', 'c']``.
- A slice object with labels, e.g. ``'a':'f'``.

.. warning:: Note that contrary to usual python slices, *both* the start and the stop are included

- A boolean array of the same length as the axis being sliced, e.g. ``[True, False, True]``.
- An alignable boolean Series. The index of the key will be aligned before masking.
- An alignable Index. The Index of the returned selection will be the index of the original array.
- A ``callable`` function with one argument (the calling Series or DataFrame) and that returns valid output for indexing (one of the above types).

See more at :ref:`Selection by Label <indexing.label>`.

Raises

KeyError

If any items are not found.

IndexingError

If an indexed key is passed and its index is unalignable to the frame

Index.

See Also

DataFrame.at : Access a single value for a row/column label pair.

```

| DataFrame.iloc : Access group of rows and columns by integer position
(s).
| DataFrame.xs : Returns a cross-section (row(s) or column(s)) from the
| Series/DataFrame.
| Series.loc : Access group of values using labels.
|
Examples
**Getting values**
>>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
...     index=['cobra', 'viper', 'sidewinder'],
...     columns=['max_speed', 'shield'])
>>> df
max_speed  shield
cobra              1      2
viper              4      5
sidewinder         7      8
Single label. Note this returns the row as a Series.
>>> df.loc['viper']
max_speed      4
shield         5
Name: viper, dtype: int64
List of labels. Note using ``[]`` returns a DataFrame.
>>> df.loc[['viper', 'sidewinder']]
max_speed  shield
viper              4      5
sidewinder         7      8
Single label for row and column
>>> df.loc['cobra', 'shield']
2
Slice with labels for row and single label for column. As mentioned
above, note that both the start and stop of the slice are included.
>>> df.loc['cobra':'viper', 'max_speed']
cobra      1
viper      4
Name: max_speed, dtype: int64
Boolean list with the same length as the row axis
>>> df.loc[[False, False, True]]
max_speed  shield
sidewinder         7      8
Alignable boolean Series:
>>> df.loc[pd.Series([False, True, False],
...     index=['viper', 'sidewinder', 'cobra'])]
max_speed  shield
sidewinder         7      8
Index (same behavior as ``df.reindex``)
>>> df.loc[pd.Index(["cobra", "viper"], name="foo")]
max_speed  shield
foo
cobra              1      2

```

```
viper          4      5
```

Conditional that returns a boolean Series

```
>>> df.loc[df['shield'] > 6]
           max_speed  shield
sidewinder          7      8
```

Conditional that returns a boolean Series with column labels specified

```
>>> df.loc[df['shield'] > 6, ['max_speed']]
           max_speed
sidewinder          7
```

Callable that returns a boolean Series

```
>>> df.loc[lambda df: df['shield'] == 8]
           max_speed  shield
sidewinder          7      8
```

****Setting values****

Set value for all items matching the list of labels

```
>>> df.loc[['viper', 'sidewinder'], ['shield']] = 50
>>> df
           max_speed  shield
cobra              1      2
viper              4     50
sidewinder         7     50
```

Set value for an entire row

```
>>> df.loc['cobra'] = 10
>>> df
           max_speed  shield
cobra           10     10
viper           4     50
sidewinder       7     50
```

Set value for an entire column

```
>>> df.loc[:, 'max_speed'] = 30
>>> df
           max_speed  shield
cobra             30     10
viper             30     50
sidewinder        30     50
```

Set value for rows matching callable condition

```
>>> df.loc[df['shield'] > 35] = 0
>>> df
           max_speed  shield
cobra             30     10
viper              0      0
sidewinder         0      0
```

****Getting values on a DataFrame with an index that has integer labels****

Another example using integers for the index

```
>>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
...                    index=[7, 8, 9], columns=['max_speed', 'shield'])
>>> df
```

```

max_speed  shield
7          1      2
8          4      5
9          7      8

```

Slice with integer labels for rows. As mentioned above, note that both the start and stop of the slice are included.

```

>>> df.loc[7:9]
max_speed  shield
7          1      2
8          4      5
9          7      8

```

****Getting values with a MultiIndex****

A number of examples using a DataFrame with a MultiIndex

```

>>> tuples = [
...     ('cobra', 'mark i'), ('cobra', 'mark ii'),
...     ('sidewinder', 'mark i'), ('sidewinder', 'mark ii'),
...     ('viper', 'mark ii'), ('viper', 'mark iii')
... ]
>>> index = pd.MultiIndex.from_tuples(tuples)
>>> values = [[12, 2], [0, 4], [10, 20],
...           [1, 4], [7, 1], [16, 36]]
>>> df = pd.DataFrame(values, columns=['max_speed', 'shield'], index=index)

```

ex)

```

>>> df
max_speed  shield
cobra      mark i      12      2
           mark ii      0      4
sidewinder mark i      10     20
           mark ii      1      4
viper      mark ii      7      1
           mark iii     16     36

```

Single label. Note this returns a DataFrame with a single index.

```

>>> df.loc['cobra']
max_speed  shield
mark i      12      2
mark ii     0      4

```

Single index tuple. Note this returns a Series.

```

>>> df.loc[('cobra', 'mark ii')]
max_speed    0
shield        4
Name: (cobra, mark ii), dtype: int64

```

Single label for row and column. Similar to passing in a tuple, this returns a Series.

```

>>> df.loc['cobra', 'mark i']
max_speed    12
shield        2
Name: (cobra, mark i), dtype: int64

```

Single tuple. Note using ``[]`` returns a DataFrame.

```

>>> df.loc[['cobra', 'mark ii']]
max_speed  shield
cobra mark ii      0      4

```

Single tuple for the index with a single label for the column

```
>>> df.loc[('cobra', 'mark i'), 'shield']
2
```

Slice from index tuple to single label

```
>>> df.loc[('cobra', 'mark i'):'viper']
      max_speed  shield
cobra  mark i      12      2
      mark ii      0      4
sidewinder mark i     10     20
      mark ii      1      4
viper   mark ii      7      1
      mark iii     16     36
```

Slice from index tuple to index tuple

```
>>> df.loc[('cobra', 'mark i'):(('viper', 'mark ii'))]
      max_speed  shield
cobra  mark i      12      2
      mark ii      0      4
sidewinder mark i     10     20
      mark ii      1      4
viper   mark ii      7      1
```

Methods inherited from pandas.core.arraylike.OpsMixin:

```
__add__(self, other)
__and__(self, other)
__eq__(self, other)
    Return self==value.
__floordiv__(self, other)
__ge__(self, other)
    Return self>=value.
__gt__(self, other)
    Return self>value.
__le__(self, other)
    Return self<=value.
__lt__(self, other)
    Return self<value.
__mod__(self, other)
__mul__(self, other)
__ne__(self, other)
    Return self!=value.
__or__(self, other)
__pow__(self, other)
__radd__(self, other)
__rand__(self, other)
```

```

|  __rfloordiv__(self, other)
|  __rmod__(self, other)
|  __rmul__(self, other)
|  __ror__(self, other)
|  __rpow__(self, other)
|  __rsub__(self, other)
|  __rtruediv__(self, other)
|  __rxor__(self, other)
|  __sub__(self, other)
|  __truediv__(self, other)
|  __xor__(self, other)

```

Let's review summary statistics for the nobel laureates dataframe - df: use `describe()` .

`df.describe()`

`df.describe()` does not really tell us much. The `info()` method provides more information as you can check by running the code cell below.

In [75]:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    1000 non-null  int64
1   firstname            1000 non-null  object
2   surname              968 non-null   object
3   born                 999 non-null   object
4   died                 1000 non-null  object
5   bornCountry          969 non-null   object
6   bornCountryCode      969 non-null   object
7   bornCity             966 non-null   object
8   diedCountry          653 non-null   object
9   diedCountryCode      653 non-null   object
10  diedCity             647 non-null   object
11  gender               1000 non-null  object
12  year                 1000 non-null  int64
13  category             1000 non-null  object
14  overallMotivation    23 non-null    object
15  share                1000 non-null  int64
16  motivation           1000 non-null  object
17  name                 736 non-null   object
18  city                 731 non-null   object
19  country              733 non-null   object
dtypes: int64(3), object(17)
memory usage: 156.4+ KB

```

We can use the dataframe `head()` or `tail()` method to view some actual entries. Without a numeric parameter, both methods return 5 entries!

Run the next 2 code cells below to see the first 10 and the last 10 entries

In [76]:

```
df.head(10)
```

Out[76]:

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|---|----|----------------|-----------|------------|------------|-----------------------------|-----------------|-----------------------------|-----------------|
| 0 | 1 | Wilhelm Conrad | Röntgen | 1845-03-27 | 1923-02-10 | Prussia (now Germany) | DE | Lennepe (now Remscheid) | Germany |
| 1 | 2 | Hendrik A. | Lorentz | 1853-07-18 | 1928-02-04 | the Netherlands | NL | Arnhem | the Netherlands |
| 2 | 3 | Pieter | Zeeman | 1865-05-25 | 1943-10-09 | the Netherlands | NL | Zonnemaire | the Netherlands |
| 3 | 4 | Henri | Becquerel | 1852-12-15 | 1908-08-25 | France | FR | Paris | France |
| 4 | 5 | Pierre | Curie | 1859-05-15 | 1906-04-19 | France | FR | Paris | France |
| 5 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw | France |
| 6 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw | France |
| 7 | 8 | Lord | Rayleigh | 1842-11-12 | 1919-06-30 | United Kingdom | GB | Langford Grove Maldon Essex | United Kingdom |
| 8 | 9 | Philipp | Lenard | 1862-06-07 | 1947-05-20 | Hungary (now Slovakia) | SK | Pressburg (now Bratislava) | Germany |
| 9 | 10 | J.J. | Thomson | 1856-12-18 | 1940-08-30 | United Kingdom | GB | Cheetham Hill | United Kingdom |

In [77]:

```
df.tail(10)
```


Out[77]:

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | died |
|------------|------|-----------|------------|------------|------------|--|-----------------|--------------|------|
| 990 | 1025 | Drew | Weissman | 1959-09-07 | 0000-00-00 | USA | US | Lexington MA | |
| 991 | 1026 | Pierre | Agostini | 1941-07-23 | 0000-00-00 | French protectorate of Tunisia (now Tunisia) | TN | Tunis | |
| 992 | 1027 | Ferenc | Krausz | 1962-05-17 | 0000-00-00 | Hungary | HU | Mór | |
| 993 | 1028 | Anne | L'Huillier | 1958-08-16 | 0000-00-00 | France | FR | Paris | |
| 994 | 1029 | Moungi | Bawendi | 1961-00-00 | 0000-00-00 | France | FR | Paris | |
| 995 | 1030 | Louis | Brus | 1943-00-00 | 0000-00-00 | USA | US | Cleveland OH | |
| 996 | 1031 | Aleksey | Yekimov | 1945-00-00 | 0000-00-00 | USSR (now Russia) | RU | NaN | |
| 997 | 1032 | Jon | Fosse | 1959-09-29 | 0000-00-00 | Norway | NO | Haugesund | |
| 998 | 1033 | Narges | Mohammadi | 1972-04-21 | 0000-00-00 | Iran | IR | Zanjan | |
| 999 | 1034 | Claudia | Goldin | 1946-00-00 | 0000-00-00 | USA | US | New York NY | |

Both `head()` and `tail()` and indeed select statements will return the full set of attributes for the entries. How can we project to just a select set of attributes? This will be the equivalent of an SQL project operator listing the output columns.

Run the code cell below to show only 'firstname','surname','city' and 'country' columns for the first 10 entries.

In [78]: `df.head(10)[['firstname','surname','city','country']]`

Out[78]:

| | firstname | surname | city | country |
|---|----------------|-----------|-----------|-----------------|
| 0 | Wilhelm Conrad | Röntgen | Munich | Germany |
| 1 | Hendrik A. | Lorentz | Leiden | the Netherlands |
| 2 | Pieter | Zeeman | Amsterdam | the Netherlands |
| 3 | Henri | Becquerel | Paris | France |
| 4 | Pierre | Curie | Paris | France |
| 5 | Marie | Curie | NaN | NaN |
| 6 | Marie | Curie | Paris | France |
| 7 | Lord | Rayleigh | London | United Kingdom |
| 8 | Philipp | Lenard | Kiel | Germany |
| 9 | J.J. | Thomson | Cambridge | United Kingdom |

Following the example of the last code cell, write a code cell below to return the last 10 records of nobel laureates but only showing their firstname, surname, year of birth, year of death and country

In [79]: `# Add your code below this line`

Querying & Locating Data in the DataFrame

One of the most useful tasks in pandas is locating data that satisfies desired criteria. For example, we can locate a Nobel laureate with a particular surname. Let's look at the record of Caltech's most beloved figure, physicist Richard Feynman (pronounced "FINE-men"). In addition to his groundbreaking work in theoretical physics (especially quantum electrodynamics and its associated Feynman diagrams), Feynman is known for The Feynman Lectures on Physics, which covers the elementary physics curriculum (mechanics, thermal physics, electrodynamics, etc.) in an unusually entertaining and insightful way. Let's use square brackets and a boolean criterion on the "surname" column to find Feynman's record in the laureates da.

What can you conclude from the format of this query below and results produce?

In [80]: `df[df['surname'] == 'Feynman']`

Out[80]:

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|----|----|------------|---------|------------|------------|-------------|-----------------|-------------|-------------|
| 86 | 86 | Richard P. | Feynman | 1918-05-11 | 1988-02-15 | USA | US | New York NY | USA |

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|--|----|-----------|---------|------|------|-------------|-----------------|----------|-------------|
|--|----|-----------|---------|------|------|-------------|-----------------|----------|-------------|

The DataFrame object can take a Boolean condition on columns in it's index and will return records that meet that condition!. The inner part of the syntax for query above returns a Series consisting of boolean values for every laureate, with True if the surname is equal to "Feynman" and False otherwise.

By using the correct index (i.e., 86), we can confirm that the value in that case is True.

Run the code cell below to confirm only this index returns True.

```
In [81]: (df["surname"] == "Feynman")[86]
```

```
Out[81]: True
```

The `loc` attribute can be used in place of brackets in many places and is generally a more flexible way to pull out data items of interest. Let's use the `loc` attribute to retrieve the year when Feynman won.

Run the code cell below for this result.

```
In [82]: df.loc[df["surname"] == "Feynman", "year"]
```

```
Out[82]: 86    1965
Name: year, dtype: int64
```

Use the `loc` attribute illustrated above to code a query on this dataset to:

Find all Nobel laureates named named 'Curie'

Tip: Use the Boolean condition: `df["surname"].str.contains("Curie", na=False)`

```
In [83]: # Code and test your scriptlet to find all 'Curies'
# Assigned the results to a variable named curies.
df.loc[df['surname'] == 'Curie']
```

```
Out[83]:
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|--|----|-----------|---------|------|------|-------------|-----------------|----------|-------------|
|--|----|-----------|---------|------|------|-------------|-----------------|----------|-------------|

| | | | | | | | | | |
|---|---|--------|-------|------------|------------|-----------------------------|----|--------|--------|
| 4 | 5 | Pierre | Curie | 1859-05-15 | 1906-04-19 | France | FR | Paris | France |
| 5 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw | France |

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|---|----|-----------|---------|------------|------------|-----------------------------|-----------------|----------|-------------|
| 6 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw | France |

Let's **find all the winners of multiple nobel prizes** using the `groupby` method

```
In [84]: laureates = df.groupby(["id", "firstname", "surname"])
        sizes = laureates.size()
        sizes[sizes > 1]      # result should show 5 winners
```

```
Out[84]: id  firstname  surname
6      Marie      Curie      2
66     John      Bardeen     2
217   Linus      Pauling     2
222   Frederick  Sanger      2
743   Barry      Sharpless   2
dtype: int64
```

Selecting Dates & Time information Pandas provides good support for `datetime`. Let's search for laureates by exact birthday as a string - the default pandas storage format for dates.

Run the code cell below to get the result entry for Einstein born 3/14/1879 fondly known as Pi Day.

```
In [85]: df.loc[df['born'] == '1879-03-14']
```

```
Out[85]:
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|----|----|-----------|----------|------------|------------|-------------|-----------------|----------|-------------|
| 25 | 26 | Albert | Einstein | 1879-03-14 | 1955-04-18 | Germany | DE | Ulm | USA |

Like the query above let's see if there are any laureates born 6/28 known as Tau Day

Run the code cell below to confirm!

```
In [86]: df.loc[df['born'].str.contains('06-08', na=False)]
```

```
Out[86]:
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | died |
|-----|-----|------------|---------|------------|------------|----------------|-----------------|-------------|------|
| 120 | 121 | Kenneth G. | Wilson | 1936-06-08 | 2013-06-15 | USA | US | Waltham MA | |
| 370 | 372 | Francis | Crick | 1916-06-08 | 2004-07-28 | United Kingdom | GB | Northampton | |

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | died |
|-----|-----|-----------|-----------|------------|------------|-------------|-----------------|-----------------------|------|
| 452 | 454 | Eric F. | Wieschaus | 1947-06-08 | 0000-00-00 | USA | US | South Bend IN | |
| 667 | 683 | Wassily | Leontief | 1906-08-05 | 1999-02-05 | Russia | RU | St. Petersburg | |
| 781 | 799 | Robert J. | Aumann | 1930-06-08 | 0000-00-00 | Germany | DE | Frankfurt-on-the-Main | |
| 918 | 944 | Jacques | Dubochet | 1942-06-08 | 0000-00-00 | Switzerland | CH | Aigle | |

Rewrite the query above to filter the result to just laureates in Physics born on Tau Day ...

Tip: use (df['category'] == "physics") as the other part of your Boolean function!

In [87]:

Your code for the Physics category Tau Day Nobel Laureates below
df[(df['category'] == 'physics') & df['born'].str.contains('06-08', na=False)]

Out[87]:

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry |
|-----|-----|------------|---------|------------|------------|-------------|-----------------|------------|-------------|
| 120 | 121 | Kenneth G. | Wilson | 1936-06-08 | 2013-06-15 | USA | US | Waltham MA | US |